

# EECS 370 – Lecture 1

Introduction

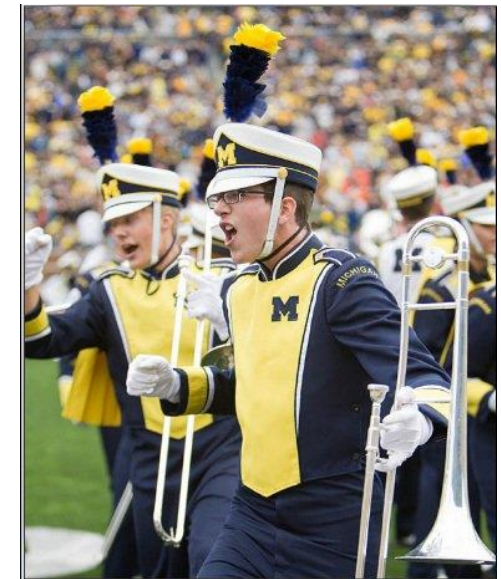


Live Poll + Q&A: [slido.com #eecs370](https://slido.com/#eecs370)

*Poll and Q&A Link*

# Who am I?

- Bachelors (CE), Masters (CSE), and PhD (CSE) all from Michigan!
- Research in computer architecture
  - How to make devices like Graphics Processing Units (GPUs) better at non-graphics stuff?



# Who am I?



- When not doing school stuff, you can find me:
  - Triathlon training
  - Cooking
  - Designing retro games / emulators
- Call me:
  - Jon
  - Dr / Prof Beaumont
  - “Yo Teach!”



# Other Faculty Instructors

- Mark Brehob, [brehob@umich.edu](mailto:brehob@umich.edu)
  - Teaching section 1, 3, and 4
- Krisztian Flautner, <mailto:manowar@umich.edu>
  - Teaching section 5

# Who are You?

Poll: Who are you!



I'll make use of polls during class – these are optional

I have you log in with your name because it helps me learn names

(you can also ask questions asynchronously on Slido)

# Class resources

- Course homepage: [eecs370.github.io/](https://eecs370.github.io/)
  - All assignments will be posted here.
  - Also links for administrative requests (SSD, Medical emergencies, etc.)
- Piazza: [piazza.com](https://piazza.com)
  - Use for general questions on lectures, projects and homework assignments. Can discuss with your classmates.
- Gradescope: <https://www.gradescope.com>
  - Turn in homework assignments.
  - Contact us if you don't have access

For other issues, submit form on website



# Course Logistics

- In line with the general CoE and University policies, default mode of instruction is **in-person**
- Lectures (T/Th)
  - Attendance is not required
  - Feel free to ask questions at anytime!
    - Interaction is highly encouraged (and critical!)
- Discussions (Fr,M)
  - Attendance not required, but strongly encouraged
  - Cover tips for working through projects
- Office hours will have some remote options, but mostly in-person
  - See website – start this weekend

# Your work in 370

- Programming assignments (40%)
- Homeworks (12% total)
  - Total of 7 homeworks, drop lowest
  - HW 0 out today, due **WEDNESDAY**
- One midterm and a final exam (48% total)
  - **In-person**
  - Midterm - Thursday March 9<sup>th</sup>, 7-9 pm
  - Final – Thursday April 20<sup>th</sup>, 10:30-12:30am





# Programming assignments

- 4 programming assignments simulating the execution of a simple microprocessor
  - Assembler / functional processor simulation
  - Linker
  - Pipeline simulation
  - Cache simulation
- First programming assignment will be posted soon
- Using C to program, C is a subset of C++ without several features like classes
- The challenge is to understand computer organization enough that you can build a complete computer emulator

# Auto-grading assignments

- We use a program to grade your assignments
  - Program submitted using [autograder.io](https://autograder.io)
- **Assignments due at 11:59 pm on due date.**  
You may use up to 4 late days just for projects over the course of the semester
- Make sure code runs correctly on CAEN machines
- Due today: **make sure to have a CAEN account**
- Help on C available from staff

# Academic integrity

- We want you to collaborate as you learn!
  - But **DON'T SHARE CODE**

Suspected code  
copying reported to  
Honor Council

– usual penalty is 0 on  
assignment and 1/3  
letter grade reduction  
in course

## DO

Share high level strategies

Help someone debug

Explain compiler errors to  
someone

Discuss test strategies

## DO NOT

Share code

Debug for someone

Fix someone's compiler  
error

Share test cases

# How we assign course grades

- Grade on a straight scale
- We may adjust this in your favor if exams are more difficult than expected
- Average is usually ~B-B+
- 1/3 As, 1/3 Bs, 1/3 Cs
  - Not a guarantee, just how it usually works out

Total Weighted Score	Letter Grade
0 - 56%	E
57 - 59%	D
60 - 64%	D+
65 - 69%	C-
70 - 74%	C
75 - 78%	C+
79 - 81%	B-
82 - 85%	B
86 - 89%	B+
90 - 92%	A-
93 - 96%	A
97 - 100%	A+

# Course textbooks

## Computer Organization and Design ARM Edition

by Patterson and Hennessy

- Not required, but it's good



# What is 370 about?

- You're used to writing programs like this

```
void daxpy(int n, double a,  
           double *x, double *y)  
{  
    for (int i = 0; i < n; ++i)  
        y[i] = a*x[i] + y[i];  
}
```

- But what's actually happening *inside* the computer when we compile / run this?
  - Come to think it... what does "compiling" even mean??
- 370 in a nutshell: **How do computer's execute programs?**

# You will need to know this stuff if...

- You work in designing processors at Intel, ARM, NVIDIA, etc
- You write optimized library code
- You work on designing operating systems or compilers
- You work in computer security

- Remember this?



- You work in designing embedded systems (IOT, etc)

# You might need to know this stuff if...

- Even if you just write software for the rest of your life
  - Important to know what your computer is doing when it executes your code!
  - It can make a big difference in your performance
- My favorite example comes from the #1 StackOverflow question of all time...



# Example

- Consider 1 version of code that loops through an array of random values and adds all elements larger than 128

```
for (unsigned c = 0; c < arraySize; ++c)
    data[c] = std::rand() % 256;

// Test
clock_t start = clock();
long long sum = 0;
// Primary loop
for (unsigned c = 0; c < arraySize; ++c)
{
    if (data[c] >= 128)
        sum += data[c];
}

double elapsedTime =
    static_cast<double>(clock() - start);
```

<https://stackoverflow.com/questions/11227809/>

# Example

- What will happen to the execution time of the loop if we sort the array beforehand?

```
for (unsigned c = 0; c < arraySize; ++c)
    data[c] = std::rand() % 256;
std::sort(data, data + arraySize);

// Test
clock_t start = clock();
long long sum = 0;
// Primary loop
for (unsigned c = 0; c < arraySize; ++c)
{
    if (data[c] >= 128)
        sum += data[c];
}

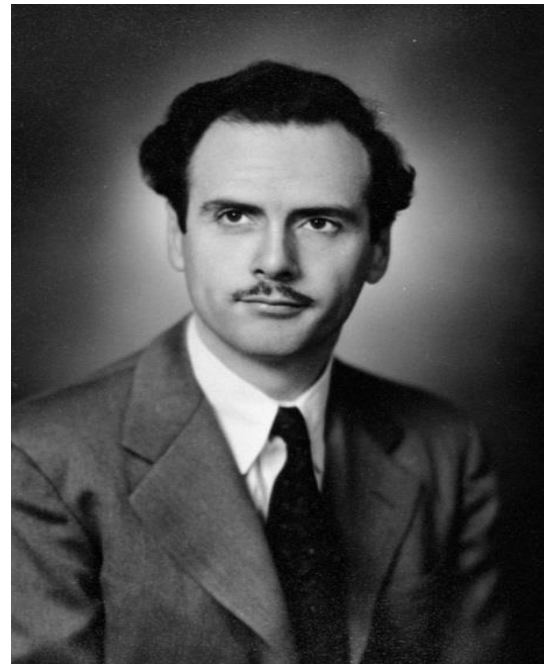
double elapsedTime =
    static_cast<double>(clock() - start);
```

**Poll: What do you think will happen?**

- A. Sorted array sums much faster
- B. Sorted array sums much slower
- C. No significant difference between sorted and unsorted arrays
- D. I have no idea!

# Why take 370?

- Inherent value in understanding how the tools we use work



*“We shape our tools and thereafter  
our tools shape us”*

- Marshal McLuhan

# Let's take a break

- In the meantime...
- What processor are you using?
  - Windows: Control Panel > System and Security > System
  - Mac: Click top-left Apple icon > About this Mac
  - Linux: lscpu
- How many cores does it have?
  - Windows: Ctrl-Shift-Esc > Performance

**Poll: What hardware do you have with you today?**



# Course overview

- Introduction (this lecture)
- ISAs and Assembly (~6 lectures)
- Processor implementation (~9 lectures)
- Memory (~7 lectures)
- Exams/catch-up (3 lectures)

# How programs work in a nutshell

- We're used to seeing programs like this
- But computer hardware is limited and can't understand code this complicated
  - Tons of different keywords and variable names...
  - Matching up different parentheses
  - Do we have to hardwire all this in logical circuits???

```
void daxpy(int n, double a,  
           double *x, double *y)  
{  
    for (int i = 0; i < n; ++i)  
        y[i] = a*x[i] + y[i];  
}
```

# How programs work in a nutshell

- High level languages are intended to be easy for humans to understand / write
  - **NOT** for hardware to execute
- To be executed, must **compile** this complicated code into a set of **very simple** and easy to understand instructions that hardware can execute

```
void daxpy(int n, double a,  
           double *x, double *y)  
{  
    for (int i = 0; i < n; ++i)  
        y[i] = a*x[i] + y[i];  
}
```

# How programs work in a nutshell

`$xxd -b [file]`  
will show you the  
binary contents... not  
very useful

- THIS is what computers can understand and execute

Address	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
00000510	48	83	ec	08	48	8b	05	cd	0a	20	00	48	85	c0	74	02
00000520	ff	d0	48	83	c4	08	c3	00	00	00	00	00	00	00	00	00
00000530	ff	35	8a	0a	20	00	ff	25	8c	0a	20	00	0f	1f	40	00
00000540	ff	25	8a	0a	20	00	68	00	00	00	00	e9	e0	ff	ff	ff
00000550	ff	25	a2	0a	20	00	66	90	00	00	00	00	00	00	00	00

- This is called "machine code": just a bunch of 0s and 1s (easier for us to read if we convert it into hex digits)
- It's what's produced when you type:  
`g++ example.c -o example.exe`



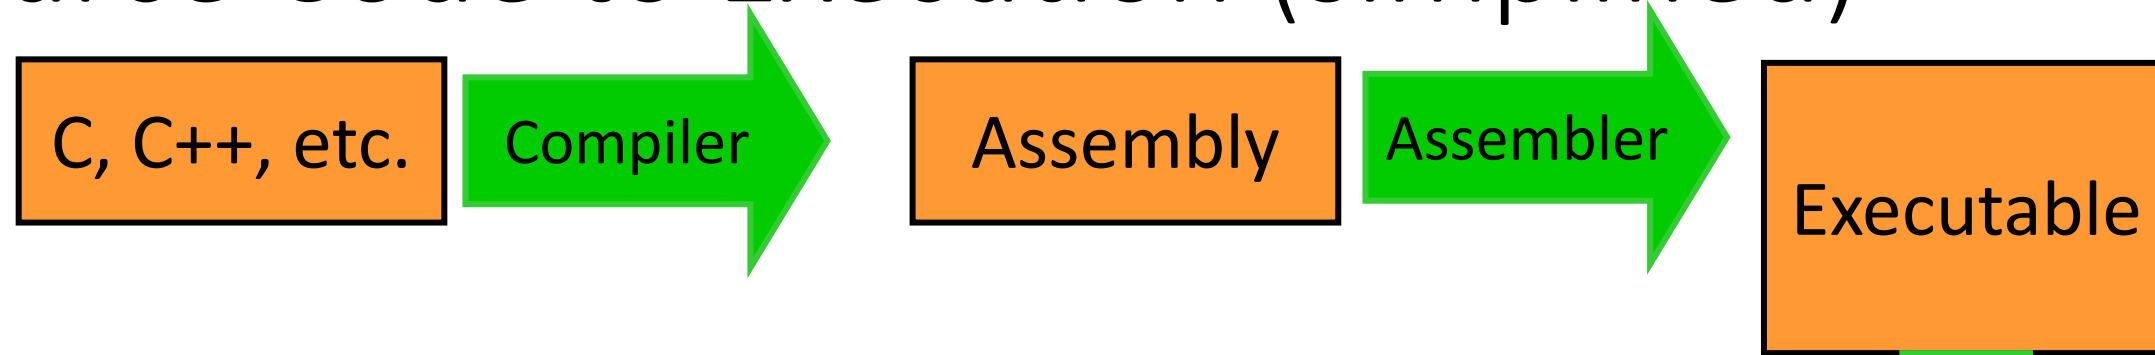
# How programs work in a nutshell

- Humans often work at an intermediate level by writing **assembly code**
- Usually has a 1-1 correspondence with machine code instructions
  - Gives the programmer fine control over the final executable
- But it's (relatively) easy to read
- You can view generated assembly code with -s flag in g++:

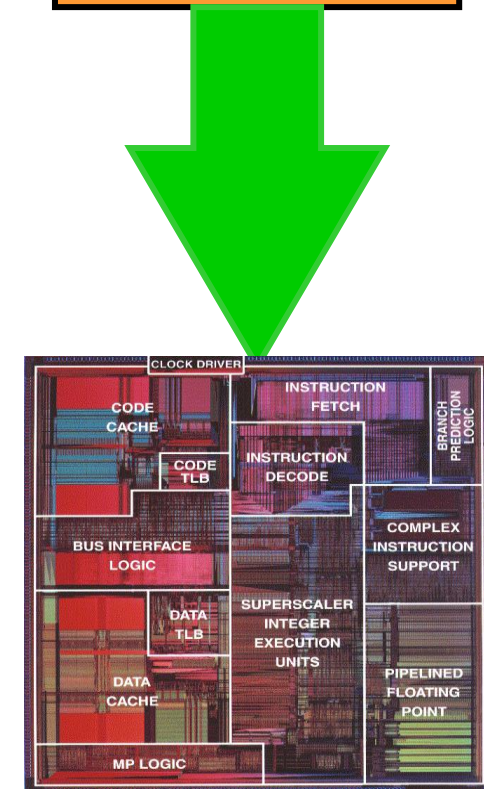
g++ -s example.c

```
5  daxpy:
6  .LFB0:
7          .cfi_startproc
8          pushq   %rbp
9          .cfi_def_cfa_offset 16
10         .cfi_offset 6, -16
11         movq    %rsp, %rbp
12         .cfi_def_cfa_register 6
13         movl    %edi, -20(%rbp)
14         movsd   %xmm0, -32(%rbp)
15         movq    %rsi, -40(%rbp)
16         movq    %rdx, -48(%rbp)
17         movl    $0, -4(%rbp)
18         jmp     .L2
```

# Source Code to Execution (simplified)



- There are some things missing here... we'll fill those in later
- First 2 weeks of the course will cover this process
- Remainder of the class will be... how do you build this thing?? (and memory)



# Architectures

- Not just one type of machine code produced for all types of computers
- Just like how there are several different programming languages (C/C++, Java, Python, etc)...
  - there are also many different types of **architectures** that code can be compiled to run on
- Popular architectures:
  - x86, ARM, RISC-V
- Code compiled for one architecture will not run on another

# x86

- Designed by Intel (AMD designed 64-bit version)
- Beefy, complex, fast, power-hungry
- Used in:
  - Desktops
  - Most laptops
  - Servers
  - PlayStation 4/5, Xbox One



# ARM

- Designed by... Arm
- Versatile: can be used for higher performance or low-power usage
- Used in:
  - Most smartphones
  - Recent Macbooks
  - Recent supercomputer clusters
  - Nintendo Switch



# RISC-V

- Open source
- Very popular in academia
  - Don't need to pay super-expensive licensing fees
- Starting to make its way into actual products



# Architectures Discussed in this Class

- We primarily focus on:
  - A subset of ARM called "LEG" (hardy-har-har)
  - A made-up ISA we call LC2K (Little Computer 2000)
    - Extremely simple, lets us focus on the concepts
    - Not practical for real applications





# The Trend of Computing

**Moore's Law:** The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World  
in Data

- Moore's Law

Transistor count

50,000,000,000

10,000,000,000

5,000,000,000

1,000,000,000

500,000,000

100,000,000

50,000,000

10,000,000

5,000,000

1,000,000

500,000

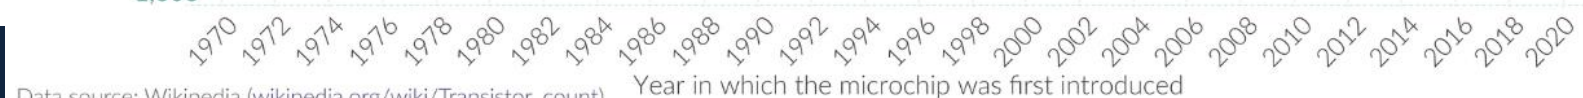
100,000

50,000

10,000

5,000

1,000



Data source: Wikipedia ([wikipedia.org/wiki/Transistor\\_count](https://wikipedia.org/wiki/Transistor_count))

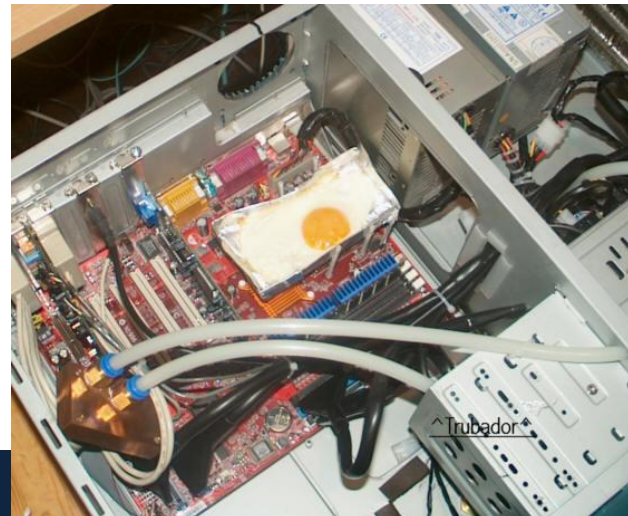
OurWorldinData.org – Research and data to make progress against the world's largest problems.

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser.



# The End of Moore's Law?: Dennard Scaling

- Dennard Scaling: as transistors get smaller their power density stays constant
- Translation: as the number of transistors on a chip grows (Moore's Law), the power stays roughly constant
- Mid-2000's Dennard Scaling broke. Why? Transistors got so small that they began to leak a lot of power. Leaking lots of power caused a chip heat up a lot.
- Conclusion: you can put lots of transistors on a chip, but you can't use them all at full power at the same time.
  - You'll melt the processor!



# Reminder

- Make sure you have a CAEN account! (today)
- Homework 0 posted; due Wednesday
- Discussion 1 this week
  - learn about C programming, debugging methods and tools, and more

# Next Time

- Introduce Instruction Set Architectures (ISAs)
- Lingering questions / feedback? I'll include an anonymous form at the end of every lecture: <https://bit.ly/3oXr4Ah>

