

Poll: Why is multiplying numbers by 100 easier than multiplying by 128?

EECS 370 - Lecture 2

ISA and Binary



Live Poll + Q&A: [slido.com #eecs370](https://slido.com/#eecs370)

Poll and Q&A Link

Unresolved questions from last time

- *"Is there a reason that most programming languages are open source but architectures IP?"*
 - Designing hardware is expensive!
 - Anyone can write software: incentive to develop toolchains and share them with people
 - Actually implementing a processor design takes millions of dollars. Mostly only specialized companies develop toolchains, and sell them to other companies for big \$\$\$
 - But that's gradually changing!
- Remember, you can post questions on Slido, or fill out this "end of lecture" form:



Announcements

- HW 0
 - Posted on website, due Wednesday
- P1
 - Released soon
- Setup clinic
 - Friday
 - Need help getting your debugger setup?
- OH
 - Going on now: see website->Google Calendar->Office Hours

My Office Hours

- 2 types:
- Group:
 - In-person
 - T/Th - 30 minutes right after class (3901 BBB)
 - Prioritize group questions over individual debugging
 - Starting this Thursday
- Individual:
 - Wednesday – 4:30-5:30
 - Over Zoom (for now) – see website for link
 - One-on-one: any questions welcome

Extra Resources

- Want more examples on binary? Two's complement?
 - See "resources tab" on website
 - Extra videos, review sheets

The screenshot displays the 'Course Resources' page for EECS 370. On the left is a dark sidebar with a menu containing: 'EECS 370', 'Calendar', 'Lecture', 'Discussion', 'Assignments', 'Exams', 'Admin Requests', 'Schedule', and 'Course Resources' (which is highlighted in teal). The main content area is titled 'Course Resources' and is organized into four columns, each with a header and a list of resource buttons:

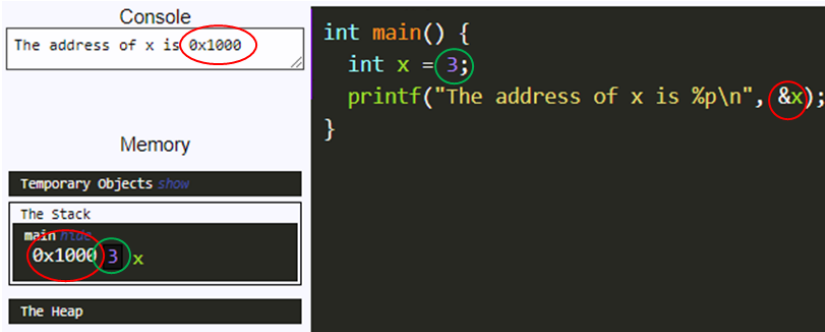
- Review Content**
 - EECS 370 Youtube Channel
 - Binary, Hex, and 2's complement Review Sheet
- Simulators**
 - Cache Simulator
 - Pipeline Simulator
- Reference Material**
 - Green LEGv8 Cheat Sheet
 - C for C++ users by Ian Cooke
 - Symbol Table and Relocation Table for EECS 370
- GDB Content**
 - GDB Tutorial
 - GDB Reference Card

Instruction Set Architecture (ISA) Design Lectures

- **Lecture 2: ISA - storage types, binary and addressing modes**
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout

Basic Computer Model

- You know from 280 that computers have "memory"
 - Abstractly, a long array that holds values
- Every piece of data in a running program lives at a numerical **address** in memory
 - You can see the address in C by using the "&" operator



The screenshot displays a debugger interface with three main components:

- Console:** Shows the output "The address of x is 0x1000", where the address is circled in red.
- Memory:** A panel showing the stack. Under "The Stack", the entry for "main" shows the address "0x1000" (circled in red) and the value "3" (circled in green), followed by "x".
- Code:** A C program snippet on a dark background:

```
int main() {  
    int x = 3;  
    printf("The address of x is %p\n", &x);  
}
```

Annotations include a green circle around the value "3" and a red circle around the "&x" expression.

- Most programs work by loading values from memory to the processor, operating on those values, and writing values back into memory

Basic Memory Model

- 1st question in understanding how programs run on computers:
 - How are values actually represented in memory?
- Answer: binary

Aside: Decimal and Binary



- Humans represent numbers in base-10 (decimal) because we have 10 fingers (or "digits")
- The n^{th} digit corresponds to 10^n

$$\begin{aligned} & \text{1407} \\ &= 1 \cdot 10^3 + 4 \cdot 10^2 + 0 \cdot 10^1 + 7 \cdot 10^0 \\ &= 1000 + 400 + 00 + 7 \end{aligned}$$

Collection of 8
bits is called a
byte

- Computers are made of wires with either high or low voltages
- Internally represents values in base-2 (binary) since it has "binary digits"
 - (or bits for short)
- In binary, the n^{th} bit corresponds to 2^n

$$\begin{aligned} & \text{1101} \\ &= 1 \cdot 2^3 + 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 \\ &= 8 + 4 + 0 + 1 \\ &= 13 \end{aligned}$$



*Does Bart Simpson
count in octal?*

Aside: Hexadecimal

- A bunch of 0s and 1s is hard to read for humans
 - But translating to decimal and back is tricky
- Solution: Bases that are a power of 2 are easy to translate between, since a fixed group of bits corresponds to one digit
- In practice, base-16 or **hexadecimal** is used
 - Digits 0-9, plus letters A-F to represent 10-16

Aside: Hexadecimal

Represent binary using 0b. Hex
using 0x. If not specified, it's
decimal

- Every 4 bits corresponds to 1 hex digit (since $2^4=16$)

(binary)	0b	0010	0101	1010	1011
(hexadecimal)	0x	2	5	A	B

0x25AB

Operating on Binary Values

- All values are stored in binary, even when you specify the number in decimal
- It is often convenient to treat values as sequences of bits, rather than values
 - You will need to do this in P1a
- C provides "bitwise operators" to do this
 - Shift (" $<<$ " and " $>>$ ")
 - Bitwise boolean (" $\&$ ", " \mid ", " \wedge ", and " \sim ")

Shift Operators

- Shift a value x bits to the left via "<<"
- Inserts "x" zeros to the right (least significant)
- E.g.

```
int a = 60;
```

```
int s = a << 2;
```

Shift Operators

- Shift a value x bits to the left via " $<<$ "
- Inserts " x " zeros to the right (least significant)
- E.g.

```
int a = 60;    // 0b0011_1100
```

```
int s = a << 2; // 0b1111_0000
```

- "a" is still 60, "s" is 240
- Same idea for " $>>$ ", but to the right

shifting x to the left in
decimal \rightarrow multiplying
by 10^x

shifting x to the left in
binary \rightarrow multiplying by
 2^x

Bitwise operations

- Bitwise operations apply a Boolean operation on each bit of a value (or each pair of bits across two values)

```
int a = 60; // 0b0011_1100
```

```
int b = 13; // 0b0000_1101
```

```
int o = a | b;
```

Bitwise operations

- Bitwise operations apply a Boolean operation on each bit of a value (or each pair of bits across two values)

```
int a = 60; // 0b0011_1100
```

```
int b = 13; // 0b0000_1101
```

```
int o = a | b; // 0b0011_1101
```

- "a" and "b" are the same, "o" is 61
- **&** – and **|** – or **^** – xor **~** – not
- **Very different** from Boolean **&&**, **||**, etc
 - Why?

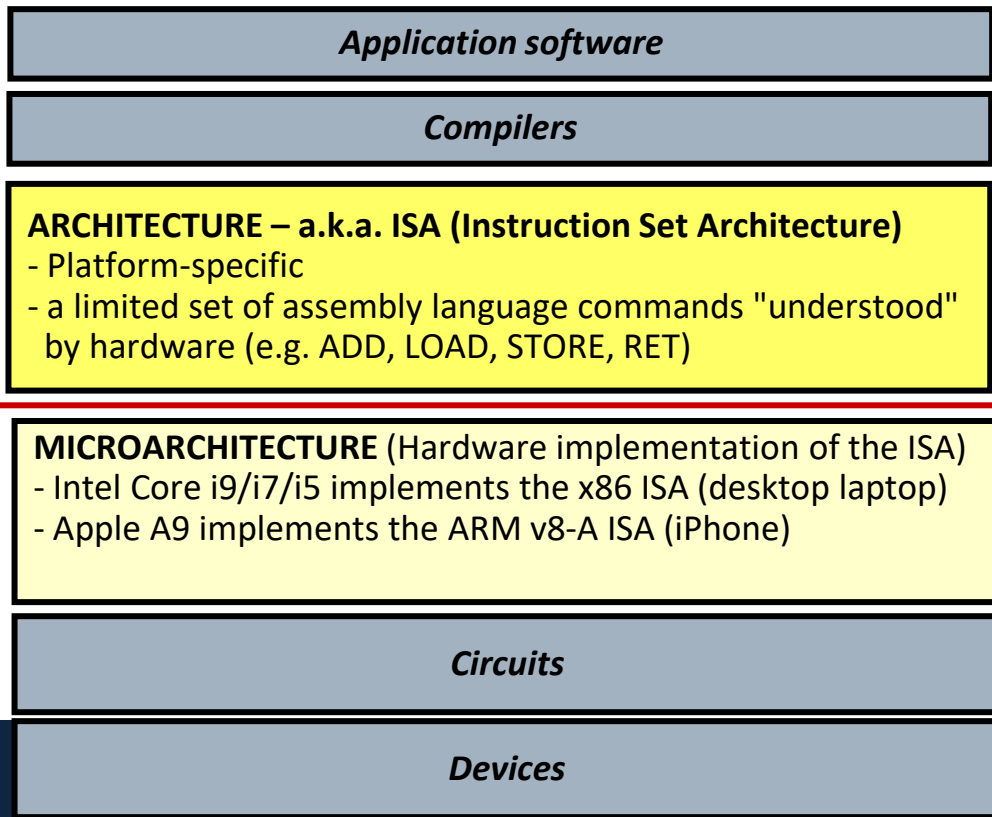
Different Data Types

- How does memory distinguish between different data types?
 - E.g. int, int *, char, float, double
- It doesn't! It's all just 0s and 1s!
- We'll see how to encode each of these later
- Exact length depends on architectures

Minimum Datatype Sizes

Type	Minimum size (bits)
char	8
int	16
long int	32
float	32
double	64

Where do ISAs come into the game ?



The
hw/sw
divide

How is Assembly Different from C/C++?

- C/C++: string multiple operations & operands in one statement:

$x = a + b + c - d$

- In assembly, you can

- only specify one of a fixed number of operations per statement, and
- have a fixed number of operands per statement

ADD r1, r2, r3 // equivalent to $r1 = r2 + r3$

- Complex Instruction Set Computing (CISC) ISAs focus on having many, complex instructions to make programming easier
 - E.g. x86, not discussed in this class
- Reduced Instruction Set Computing (RISC) ISAs focus on having fewer, simpler instructions to ease hardware design
 - E.g. LC2K, ARM (sorta), primary focus of this class

How is Assembly Different from C/C++?

- C/C++ instructions operate on **variables**

- e.g.

`x = i+j;`

- Practically unlimited

- Assembly instructions operate on **memory locations** or **registers**

- e.g.

`add r1, r2, r3`

`ld r1, 0x1000`

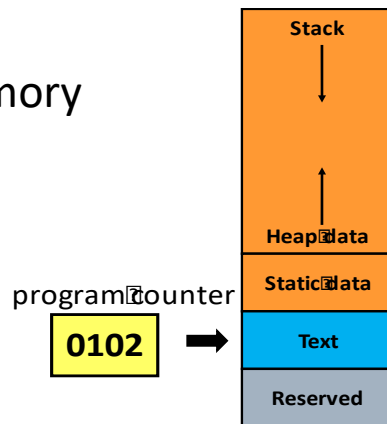
- Registers are basically a small number (~8-32) of fixed-length, hardware variables that have simple names like "r5"

Example Architectures

- ARMv8—LEGV8 subset from P+H text book
 - 32 registers (X0 – X31)
 - 64 bits in each register
 - Some have special uses e.g. X31 is always 0—XZR
- Intel x86 (not discussed much in this class)
 - 4 general purpose registers (eax, ebx, ecx, edx) 32 bits
 - Special registers: 3 pointer registers (si, di, ip), 4 segment (cs, ds, ss, es), 2 stack (sp, bp), status register (flags)
- LC2K (simple architecture made up for this class)
 - 8 registers, 32 bits each

How is Assembly Different from C/C++?

- C/C++: next line of code is executed until you get to:
 - function call
 - return statement
 - if statement or for/while loop
 - etc
- Assembly: a program counter (PC) keeps track of which memory address has the next instruction, gets incremented until
 - a "branch" or "jump" instruction
 - Used to change control flow (more later)




Traditional (von Neumann) Architecture

- We're used to thinking about program code and data being separate
- But in practice, both are stored in the same memory
 - That's why things like "function pointers" in C/C++ work
- This is the basis of the **von Neumann** Architecture

Traditional (von Neumann) Architecture


Here's the (endless) loop that hardware repeats forever:

- 
1. Fetch—get next instruction—use PC to find where it is in memory and place it in instruction register (IR)
 - PC is changed to “point” to the next instruction in the program
 2. Decode—control logic examines the contents of the IR to decide what instruction it should perform
 3. Execute—the outcome of the decoding process dictates
 - an arithmetic or logical operation on data
 - an access to data in the same memory as the instructions
 - OR a change to the contents of the PC

Let's execute this short program:

ADD X0, X1, X2
SUB X1, X2, X0

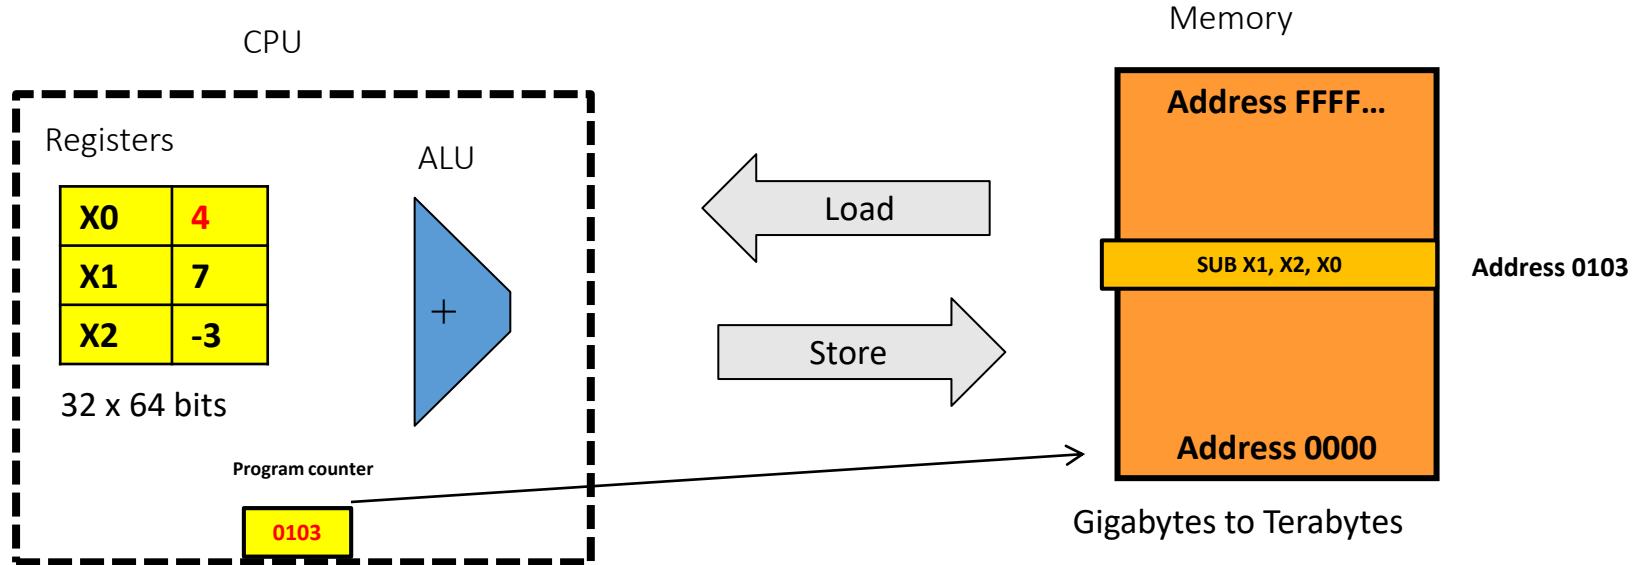
(Simplified) System Organization



Let's execute this short program:

ADD X0, X1, X2
SUB X1, X2, X0


(Simplified) System Organization



Let's execute this short program:

ADD X0, X1, X2
SUB X1, X2, X0

(Simplified) System Organization



Assembly Code – ARM Example

Poll: What are the final contents of X1, X2, and X3?

- What are the contents of the registers after executing the given assembly code (destination register is listed first in ARM)?

Program:

opcode	d	s1	s2imm
ADD	X3	X1	X2
ADDI	X3	X3	#3
SUB	X2	X3	X1

ADDI means "add immediate", the last field is a literal value, not a register index

Initial register file:

X1	25
X2	-4
X3	57

		(1) ADD X3, X1, X2	(2) ADDI X3, X3, #3	(3) SUB X2, X3, X1
X1	25	X1 25	X1 25	X1 25
X2	-4	X2 -4	X2 -4	X2 -1
X3	57	X3 21	X3 24	X3 24

How is Assembly Different from C/C++?

- No data types in assembly
- Everything is 0s and 1s: up to the programmer to interpret whether these bits should be interpreted as ints, bools, chars... or even instructions themselves!

```
char c = 'a';  
c++; // c is now 'b'
```

// results in the same assembly as

```
int x = 97;  
x++; // c is now 98
```

How is Assembly Different from C/C++?

- 1000011 could be interpreted as:

- Unsigned number: 67
- A character: 'C'
- Signed number: -61
- A float: 9.38e-44
- LC2K Instruction: 'add 0 0 3'

How did I get these?

Representing Values in Hardware

- Unsigned integers represented as we've seen
- Chars are represented as ASCII values
 - e.g. 'a' -> 97, 'b' -> 98, '#' -> 35
- What about negative numbers?
- Fractional numbers?

Negative Numbers

- There are many ways we could represent negative numbers
- Because it will eventually make our hardware simpler, the most common representation is 2's complement



No, not 2's *compliment*!

Two's Complement Representation

- Recall that 1101 in binary is 13 in decimal.

$$\begin{array}{cccc} 1 & 1 & 0 & 1 \\ 2^3 & 2^2 & 2^1 & 2^0 \end{array} = 8 + 4 + 1 = 13$$

- 2's complement numbers are very similar to unsigned binary numbers.
 - The only difference is that the first number is now negative.

$$\begin{array}{cccc} 1 & 1 & 0 & 1 \\ -2^3 & 2^2 & 2^1 & 2^0 \end{array} = -8 + 4 + 1 = -3$$

Fun with 2's Complement Numbers

- What is the range of representation of a 4-bit 2's complement number?
 - [-8, 7] (corresponding to 1000 and 0111)
- What is the range of representation of an n-bit 2's complement number?
 - $[-2^{(n-1)}, 2^{(n-1)} - 1]$
- Useful trick: You can negate a 2's complement number by inverting all the bits and adding 1.
 - 5 is represented as **0101**
 - Negate each bit: **1010**
 - Add 1: **1011** = $-8 + 2 + 1 = -5$

What about fractional numbers?


- One idea: fixed point notation
 - Have some bits represent numbers before decimal point, some bits represent numbers after decimal point
- Better idea: floating point notation
 - Inspired by scientific notation (e.g. 1.3×10^{-3})
 - Allows for larger range of numbers
 - We'll come back to this in a few weeks

Encoding Instructions

- So binary numbers can represent signed and unsigned numbers, chars, and fractional numbers
- But they must also represent instructions themselves!
 - After all, memory is just a collection of 1s and 0s
- We need a way of ***encoding*** instructions in order to store them in memory

Software program to machine code


*Example ISA
(simplified)*



Assembly Instruction Encoding

- Since the EDSAC (1949) almost all computers stored program instructions the same way they store data.
- Each instruction is encoded as a number

*Example ISA
(simplified)*



$$\begin{aligned} 011011010011001 &= 2^0 + 2^3 + 2^4 + 2^7 + 2^9 + 2^{10} + 2^{12} + 2^{13} \\ &= 13977 \end{aligned}$$

- This is the number stored in memory (in binary)!

Poll: How many different "operation codes" could be supported by this ISA? How many registers?

Next Time

- Finish Up ISAs
- LC2K details
- Lingering questions / feedback? I'll include an anonymous form at the end of every lecture: <https://bit.ly/3oXr4Ah>

