

Poll: Anything positive about
your weekend you'd like to
share?

EECS 370 - Lecture 4

ARM



Announcements

- Hope you had a meaningful MLK Day
 - Take some time to reflect on what it means for you
- HW 2
 - Posted on website, due next Monday
- P1
 - 3 parts, first part due next Thursday
- Always check course calendar for details on staff OH

Resources

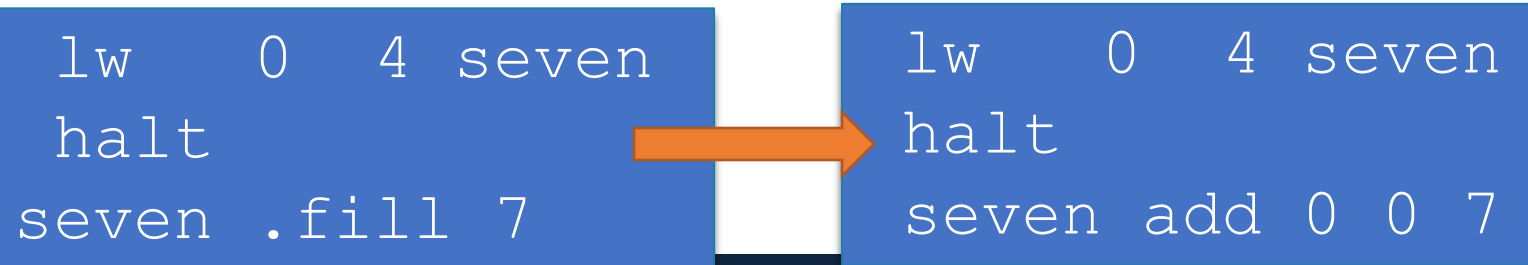
- Many resources on 370 website
 - <https://eecs370.github.io/#resources>
 - ARMv8 references
- Async discussion recordings

EECS 370: GREEN CARD FOR LEGv8						
Arithmetic Operations	Assembly code			Semantics	Comments	
add	ADD	Xd,	Xn,	Xm	$X5 = X2 + X7$	register-to-register
add & set flags	ADDSD	Xd,	Xn,	Xm	$X5 = X2 + X7$	flags NZVC
add immediate	ADDI	Xd,	Xn,	#uimm12	$X5 = X2 + \#19$	$0 \leq 12$ bit unsigned ≤ 4095
add immediate & set flags	ADDIS	Xd,	Xn,	#uimm12	$X5 = X2 + \#19$	flags NZVC
subtract	SUB	Xd,	Xn,	Xm	$X5 = X2 - X7$	register-to-register
subtract & set flags	SUBS	Xd,	Xn,	Xm	$X5 = X2 - X7$	flags NZVC
subtract immediate	SUBI	Xd,	Xn,	#uimm12	$X5 = X2 - \#20$	$0 \leq 12$ bit unsigned ≤ 4095
subtract immediate & set flags	SUBIS	Xd,	Xn,	Xm	$X5 = X2 - \#20$	flags NZVC
Data Transfer Operations	Assembly code			Semantics	Comments	
load register	LDUR	Xt,	[Xn, #simm9]	$X2 = M[X6, \#18]$	double word load into Xt from Xn + #simm9	
load signed word	LDURSW	Xt,	[Xn, #simm9]	$X2 = M[X6, \#18]$	word load to lower 32b Xt from Xn + #simm9; sign extend upper 32b	
load half	LDURH	Xt,	[Xn, #simm9]	$X2 = M[X6, \#18]$	$\frac{1}{2}$ word load to lower 16b Xt from Xn + #simm9; zero extend upper 48b	
load byte	LDURB	Xt,	[Xn, #simm9]	$X2 = M[X6, \#18]$	byte load to least 8b Xt from Xn + #simm9 zero extend upper 56b	
store register	STUR	Xt,	[Xn, #simm9]	$M[X5, \#12] = X4$	double word store from Xt to Xn + #simm9	
store word	STURW	Xt,	[Xn, #simm9]	$M[X5, \#12] = X4$	word store from lower 32b of Xt to Xn + #simm9	
store half word	STURH	Xt,	[Xn, #simm9]	$M[X5, \#12] = X4$	$\frac{1}{2}$ word load from lower 16b of Xt to Xn + #simm9	
store byte	STURB	Xt,	[Xn, #simm9]	$M[X5, \#12] = X4$	byte load from least 8b of Xt to Xn + #simm9	
	offset	#simm9 = -256 to +255			-256 \leq 9 bits signed immediate \leq +255	
move wide with zero	MOVZ	Xd,	#uimm16,	LSL N	$X9 = 0..0N0..0$	zeros out Xd then place a 16b (#uimm) into the first (N = 0)/second (N = 16)/third (N = 32)/fourth (N = 48) 16b slot of Xd
move wide with keep	MOVK	Xd,	#uimm16,	LSL N	$X9 = x..xNx..x$	place a 16b (#uimm) into the first (N = 0)/second (N = 16)/third (N = 32)/fourth (N = 48) 16b slot of Xd, without changing the other values (x's)
register aliases		$X28 = SP; X29 = FP; X30 = LR; X31 = XZR$				
Logical Operations	Assembly code			Semantics	Using C operations of & ^ < > >	
and	AND	Xd,	Xn,	Xm	$X5 = X2 \& X7$	bit-wise AND
and immediate	ANDI	Xd,	Xn,	#uimm12	$X5 = X2 \& \#19$	bit-wise AND with 0 ≤ 12 bit unsigned ≤ 4095
inclusive or	ORR	Xd,	Xn,	Xm	$X5 = X2 X7$	bit-wise OR
inclusive or immediate	ORRI	Xd,	Xn,	#uimm12	$X5 = X2 \#11$	bit-wise OR with 0 ≤ 12 bit unsigned ≤ 4095
exclusive or	EOR	Xd,	Xn,	Xm	$X5 = X2 \wedge X7$	bit-wise EOR
exclusive or immediate	EORI	Xd,	Xn,	#uimm12	$X5 = X2 \wedge \#57$	bit-wise EOR with 0 ≤ 12 bit unsigned ≤ 4095
logical shift left	LSL	Xd,	Xn,	#uimm6	$X1 = X2 << \#10$	shift left by a constant ≤ 63
logical shift right	LSR	Xd,	Xn,	#uimm6	$X5 = X3 >> \#20$	shift right by a constant ≤ 63
Unconditional branches	Assembly code			Semantics	Also known as Jumps	
branch	B	#simm26	goto PC + #1200		PC relative branch PC + 26b offset; $-2^{25} \leq \#simm26 \leq 2^{25}-1$; 4b instruction	
branch to register	BR	Xt	target in Xt		Xt contains a full 64b address	
branch with link	BL	#simm26	$X30 = PC + 4; PC = \#11000$		PC relative branch to PC + 26b offset; 16 million instructions; X30 = LR contains return from subroutine address	

Left-over questions:

"What happens when you execute a .fill instruction?"

- Hopefully, that never happens!
- Remember – (encoded) instructions and data coexist in memory
- .fill directive gives us a way to place a specific value somewhere in memory
 - It's intended to be the target of a load, the PC should never point there
- If PC **does** end up pointing there, it will be executed as if it was the corresponding encoded instruction
 - Which in this case is...?



Left-over questions:

- Remember, you can ask questions in Slido or end-of-lecture form:



Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- **Lecture 3 : LC2K**
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout



Labels in LC2K

- Labels are used in lw/sw instructions or beq instruction
- For lw or sw instructions, the assembler should compute offsetField to be equal to the address of the label
 - i.e. $\text{offsetField} = \text{address of the label}$
- For beq instructions, the assembler should translate the label into the numeric offsetField needed to branch to that label
 - i.e. $\text{PC} + 1 + \text{offsetField} = \text{address of the label}$

Labels in LC2K

- Labels are a way of referring to a line number in an assembly program.

```
loop  beq  3   4   end
      noop
      beq  0   0   loop
end    halt
```

- Here **loop** is 0 and **end** is 3.

```
// this is the
assembly for:
while(x != y) {
    x *= 2;
}
```

- What are the values of the labels here?

```
loop  beq  3   4   end
      add  3   3   3
tom   noop
      beq  0   0   loop
end    halt
```

Poll: What are the labels replaced with?

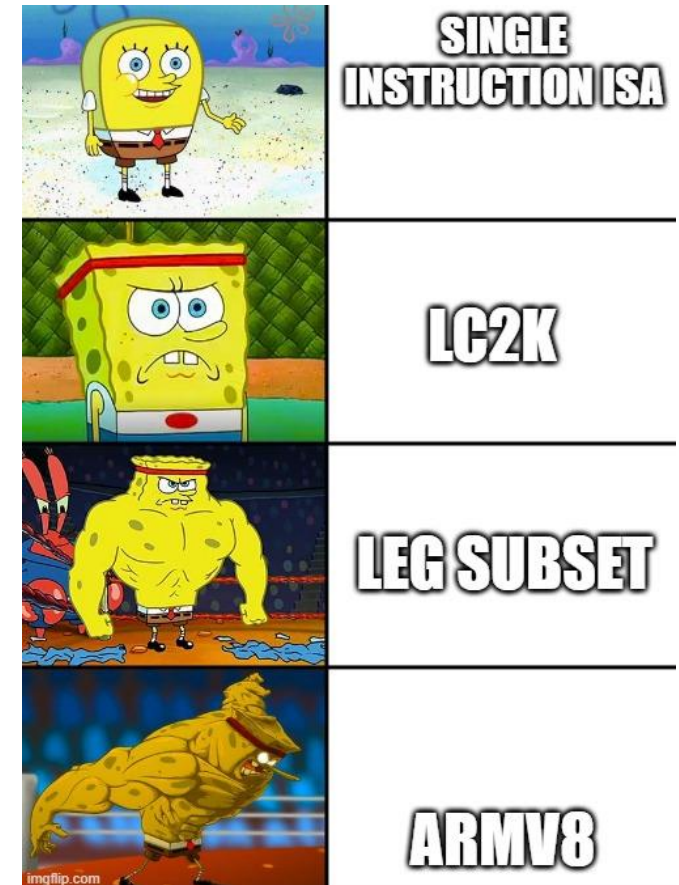
Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- **Lecture 4 : ARM**
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout



ARMv8 ISA

- LC2K is intended to be an extremely bare-bones ISA
 - "Bare minimum"
 - Easy to design hardware for, really annoying to program on (as you'll see in P1m)
 - Invented for our projects, not used anywhere in practice
- ARM (specifically v8) is a much more powerful ISA
 - Used heavily in practice (most smartphones, some laptops & supercomputers)
 - Subset (LEG) is focus of hw and lecture



ARM vs LC2K at a Glance

	LC2K	LEG
# registers	8	32
Register width	32 bits	64 bits
Memory size	2^{18} bytes	2^{64} bytes
# instructions	8	40-ish
Addressability	Word	Byte

We'll discuss what this means in a bit



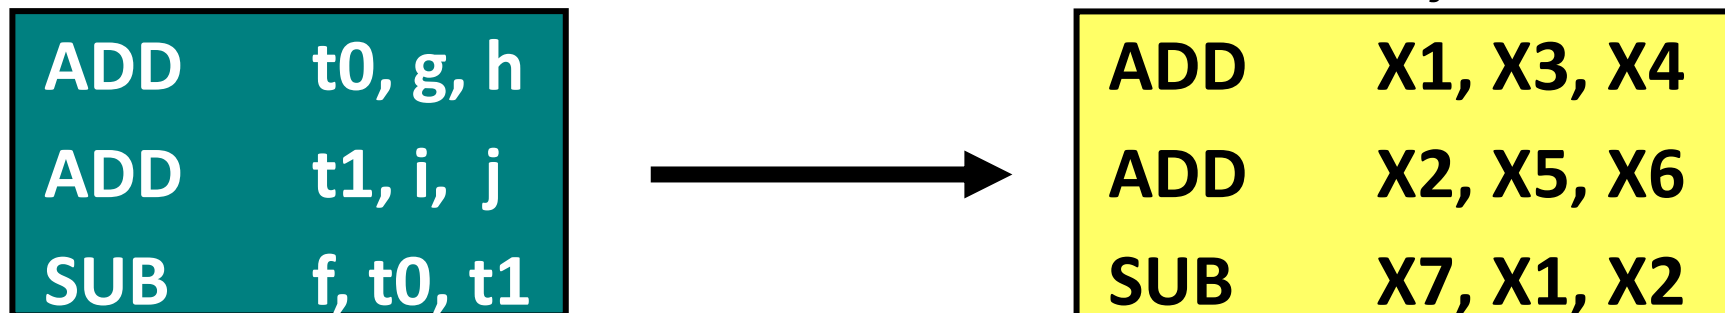
ARM Instruction Set—LEGv8 subset

- The main types of instructions fall into the familiar classes we saw with LC2K:
 1. Arithmetic
 - Add, subtract, (multiply not in LEGv8)
 2. Data transfer
 - Loads and stores—LDUR (load unscaled register), STUR, etc.
 3. Logical
 - AND, ORR, EOR, etc.
 - Logical shifts, LSL, LSR
 4. Conditional branch
 - CBZ, CBNZ, B.cond
 5. Unconditional branch (jumps)
 - B, BR, BL

LEGv8 Arithmetic Instructions

- Format: three operand fields
 - Dest. register usually the **first one** – check instruction format
 - ADD X3, X4, X7 // $X3 = X4 + X7$
 - **LC2K generally has the destination on the right!!!!**
 - E.g. add 1 2 3 // $r1 + r2 \rightarrow r3$
- C-code example: $f = (g + h) - (i + j)$

X1 \rightarrow t0
 X2 \rightarrow t1
 X3 \rightarrow g
 X4 \rightarrow h
 X5 \rightarrow i
 X6 \rightarrow j



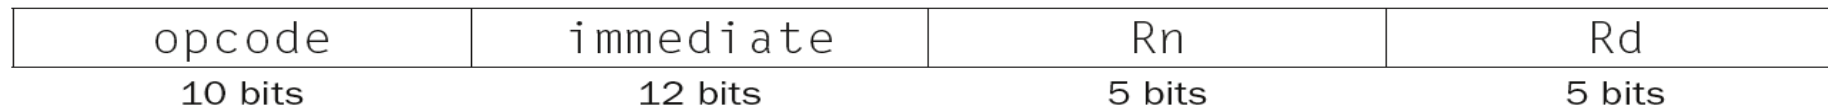
LEGv8 R-instruction Encoding

- Register-to-register operations
- Consider ADD X3, X4, X7
 - $R[Rd] = R[Rn] + R[Rm]$
 - $Rd = X3, Rn = X4, Rm = X7$
- Rm = second register operand
- ~~shamt = shift amount~~
 - not used in LEG for ADD/SUB and set to 0
- Rn = first register operand
- Rd = destination register
- ADD opcode is 10001011000, what are the other fields?

opcode	Rm	shamt	Rn	Rd
11 bits	5 bits	6 bits	5 bits	5 bits

I-instruction Encoding

- Format: second source operand can be a register or **i**mmediate—a constant in the instruction itself
- e.g., `ADD X3, X4, #10` //although we write "ADD", this is "ADDI"
- Format: 12 bits for immediate constants 0-4095



- Don't need negative constants because we have SUBI
- C-code example: $f = g + 10$
- C-code example: $f = g - 10$

```
ADDI    X7, X5, #10
```

```
SUBI    X7, X5, #10
```

LEGv8 Logical Instructions

- Logical operations are bit-wise
- For example assume
 - $X9 = 11111111\ 11111111\ 11111111\ 00000000\ 00000000\ 00000000\ 00001101\ 11000000$
 - $X13 = 00000000\ 00000000\ 11011010\ 00000000\ 00000000\ 00000000\ 00111100\ 00000000$
 - AND $X2, X13, X9$ would result in
 - $X2 = 00000000\ 00000000\ 11011010\ 00000000\ 00000000\ 00000000\ 00001100\ 00000000$
- AND and OR correspond to C operators `&` and `|`
- For immediate fields the 12 bit constant is padded with zeros to the left—zero extended

Category	InstructionExample		Meaning	Comments
Logical	and	AND $X1, X2, X3$	$X1 = X2 \& X3$	Three reg. operands; bit-by-bit AND
	inclusive or	ORR $X1, X2, X3$	$X1 = X2 X3$	Three reg. operands; bit-by-bit OR
	exclusive or	EOR $X1, X2, X3$	$X1 = X2 \wedge X3$	Three reg. operands; bit-by-bit XOR
	and immediate	ANDI $X1, X2, 20$	$X1 = X2 \& 20$	Bit-by-bit AND reg. with constant
	inclusive or immediate	ORRI $X1, X2, 20$	$X1 = X2 20$	Bit-by-bit OR reg. with constant
	exclusive or immediate	EORI $X1, X2, 20$	$X1 = X2 \wedge 20$	Bit-by-bit XOR reg. with constant
	logical shift left	LSL $X1, X2, 10$	$X1 = X2 \ll 10$	Shift left by constant
	logical shift right	LSR $X1, X2, 10$	$X1 = X2 \gg 10$	Shift right by constant

LEGv8 Shift Logical Instructions

- LSR X6, X23, #2

```

X23 = 11111111 11111111 11111111 00000000 00000000 00000000 11011010 00000010
X6  = 00111111 11111111 11111111 11000000 00000000 00000000 00110110 10000000
  
```

- C equivalent

- X6 = X23 >> 2;

- LSL X6, X23, #2

- What register gets modified?
- What does it contain after executing the LSL instruction?

Poll: Why is shifting so valuable?

- Makes multiplying easier
- Allows quicker 2s-complement conversions
- Allows for more complex branching behavior
- It's always a good time to get shifty

- In shift operations Rm is always 0—shamt is 6 bit unsigned

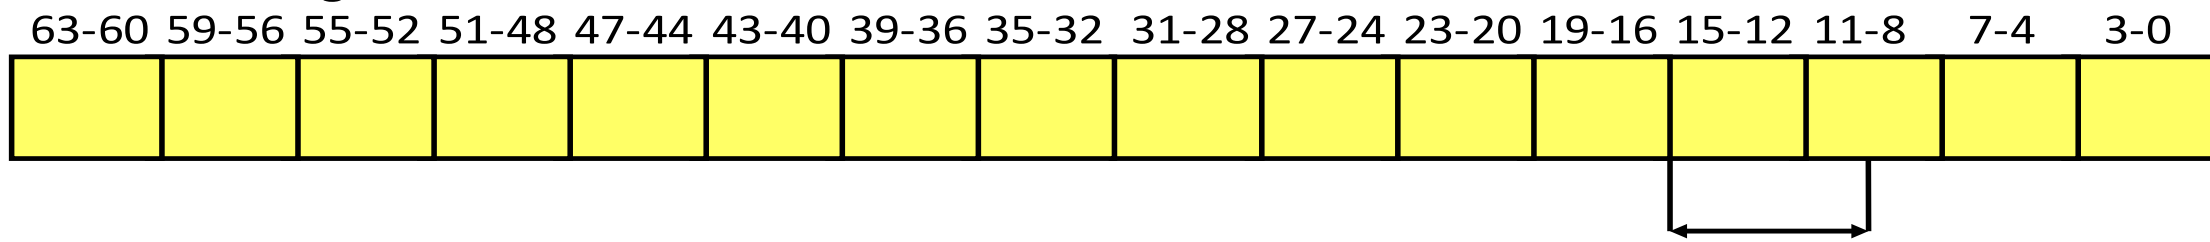
opcode	Rm	shamt	Rn	Rd
11 bits	5 bits	6 bits	5 bits	5 bits

Pseudo Instructions

- Instructions that use a shorthand “mnemonic” that expands to pre-existing assembly instruction
- Example:
 - `MOV X12, X2` // the contents of X2 copied to X12—X2 unchanged
- This gets expanded to:
 - `ORR X12, XZR, X2`
- What alternatives could we use instead of ORR?

Class Problem #1

- Show the C and LEGv8 assembly for extracting the value in bits 15:10 from a 64-bit integer variable



Assume the variable is in X1

```
x = x >> 10
```

```
x = x & 0x3F
```

Want these bits

Poll: Which operations did you use (select all)?

- a) and
- b) or
- c) add
- d) left shift
- e) right shift

Word vs Byte Addressing

- A **word** is a collection of bytes
 - Exact size depends on architecture
 - in LC2K and ARM, 4 bytes
 - **Double word** is 8 bytes
- LC2K is **word addressable**
 - Each **address** refers to a particular **word** in memory
 - Wanna move forward one int? Increment address by **one**
 - Wanna move forward one char? Uhhh...
- ARM (and most modern ISAs) is **byte addressable**
 - Each **address** refers to a particular **byte** in memory
 - Wanna move forward one int? Increment address by **four**
 - Wanna move forward one char? Increment address by **one**



LEGv8 Memory Instructions

- Like LC2K, employs base + displacement addressing mode
 - Base is a register
 - Displacement is 9-bit immediate ± 256 bytes—sign extended to 64 bits
- Unlike LC2K (which always transfers 4 bytes), we have several options in LEGv8

Category	Instruction	Example	Meaning	Comments
	load register	LDUR X1, [X2,40]	$X1 = \text{Memory}[X2 + 40]$	Doubleword from memory to register
	store register	STUR X1, [X2,40]	$\text{Memory}[X2 + 40] = X1$	Doubleword from register to memory
	load signed word	LDURSW X1, [X2,40]	$X1 = \text{Memory}[X2 + 40]$	Word from memory to register
	store word	STURW X1, [X2,40]	$\text{Memory}[X2 + 40] = X1$	Word from register to memory
	load half	LDURH X1, [X2,40]	$X1 = \text{Memory}[X2 + 40]$	Halfword memory to register
	store half	STURH X1, [X2,40]	$\text{Memory}[X2 + 40] = X1$	Halfword register to memory
	load byte	LDURB X1, [X2,40]	$X1 = \text{Memory}[X2 + 40]$	Byte from memory to register
	store byte	STURB X1, [X2,40]	$\text{Memory}[X2 + 40] = X1$	Byte from register to memory
	move wide with zero	MOVZ X1,20, LSL 0	$X1 = 20 \text{ or } 20 * 2^{16} \text{ or } 20 * 2^{32} \text{ or } 20 * 2^{48}$	Loads 16-bit constant, rest zeros
	move wide with keep	MOVK X1,20, LSL 0	$X1 = 20 \text{ or } 20 * 2^{16} \text{ or } 20 * 2^{32} \text{ or } 20 * 2^{48}$	Loads 16-bit constant, rest unchanged

D-Instruction fields

- **D**ata transfer
- opcode and op2 define data transfer operation
- address is the ± 256 bytes displacement
- Rn is the base register
- Rt is the destination (loads) or source (stores)
- More complicated modes are available in full ARMv8

Look over formatting
on your own

opcode	address	op2	Rn	Rt
11 bits	9 bits	2 bits	5 bits	5 bits

LEGv8 Memory Instructions

- Registers are 64 bits wide
- But sometimes we want to deal with non-64-bit entities
 - E.g. ints (32 bits), chars (8 bits)
- When we load smaller elements from memory, what do we set the other bits to?

- Option A: set to zero



- Option B: sign extend



- We'll need different instructions for different options

Load Instruction Sizes



How much data is retrieved from memory at the given address?

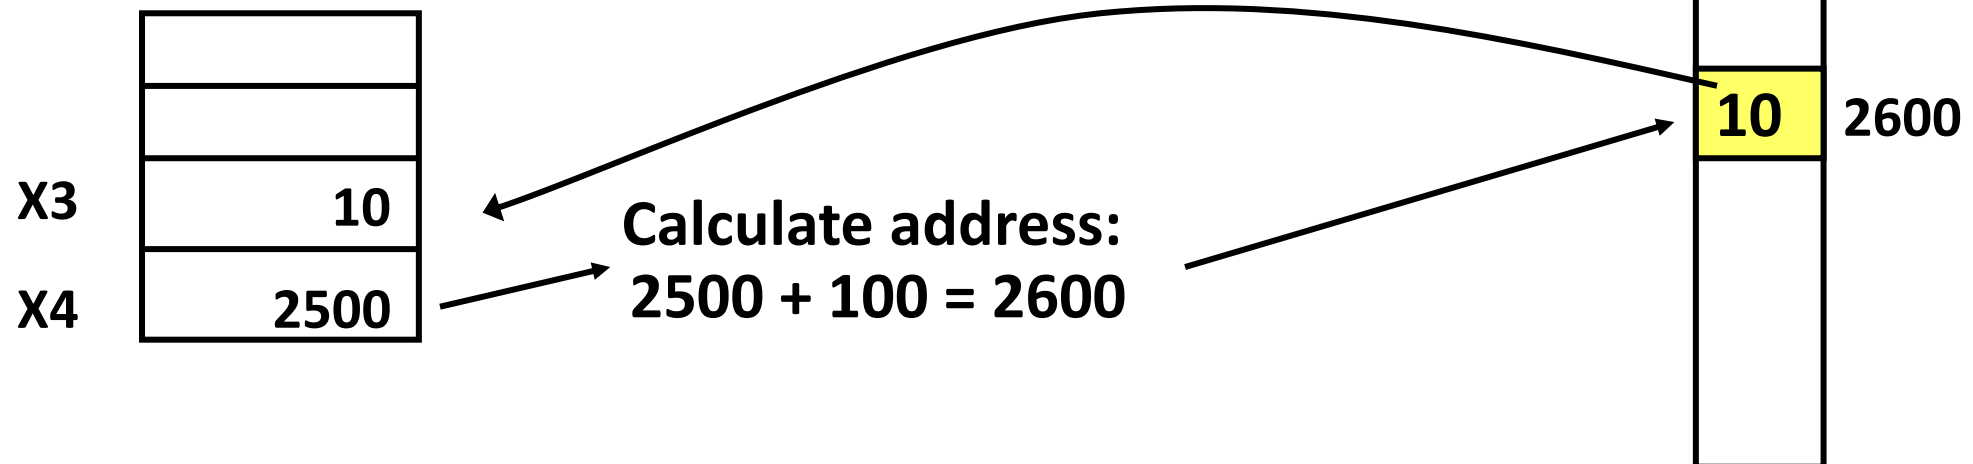
Desired amount of data to transfer?	Operation	Unused bits in register?	Example
64-bits (double word or whole register)	LDUR (Load unscaled to register)	N/A	0xFEDC_BA98_7654_3210
16-bits (half-word) into lower bits of reg	LDURH	Set to zero	0x0000_0000_0000_3210
8-bits (byte) into lower bits of reg	LDURB	Set to zero	0x0000_0000_0000_0010
32-bits (word) into lower bits of reg	LDURSW (load signed word)	Sign extend (0 or 1 based on most significant bit of transferred word)	0x0000_0000_7654_3210 or 0xFFFF_FFFF_F654_3210 (depends on bit 31)

Load Instruction in Action

```
struct {  
    int arr[25];  
    char c;  
} my_struct;
```

```
int func() {  
    my_struct.c++;  
    // load value from mem into reg  
    // then increment it  
}
```

LDURB X3, [X4, #100]



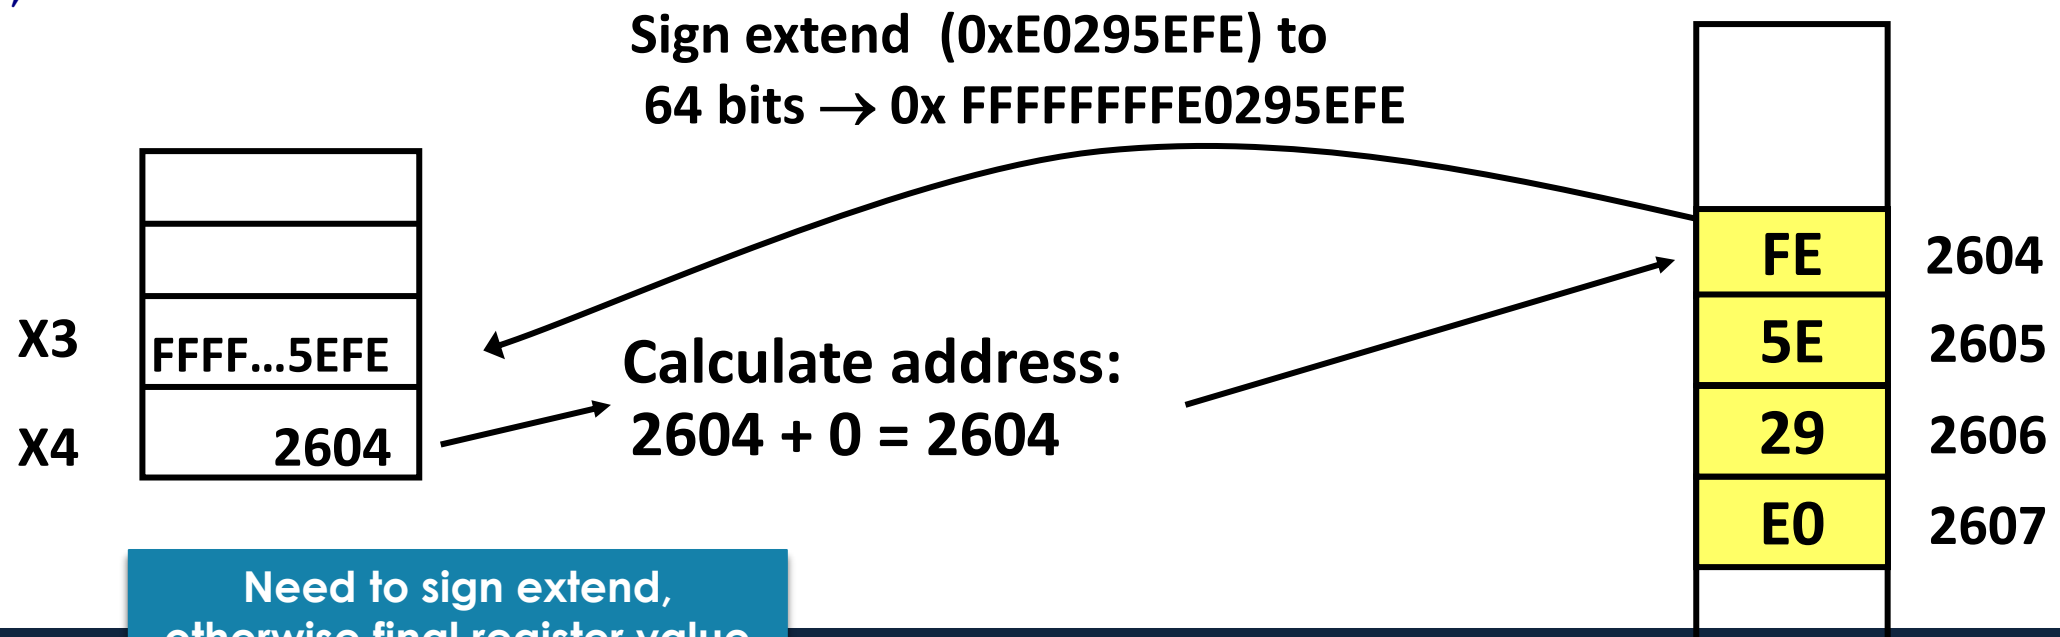
Load Instruction in Action – other example

```
int my_big_number = -534159618; // 0xE0295EFE in 2's complement
```

```
int inc_number() {
    my_big_number++;
    // load value from mem into reg
    // then increment it
};
```

TODO: add other instructions?

LDURSW X3, [X4, #0]



Need to sign extend,
otherwise final register value
will be positive!!!

But wait...

```
int my_big_number = -534159618; // 0xE0295EFE in 2's complement
```

- If I want to store this number in memory...

should it be stored like this?

FE	2604
5E	2605
29	2606
E0	2607

... or like this?

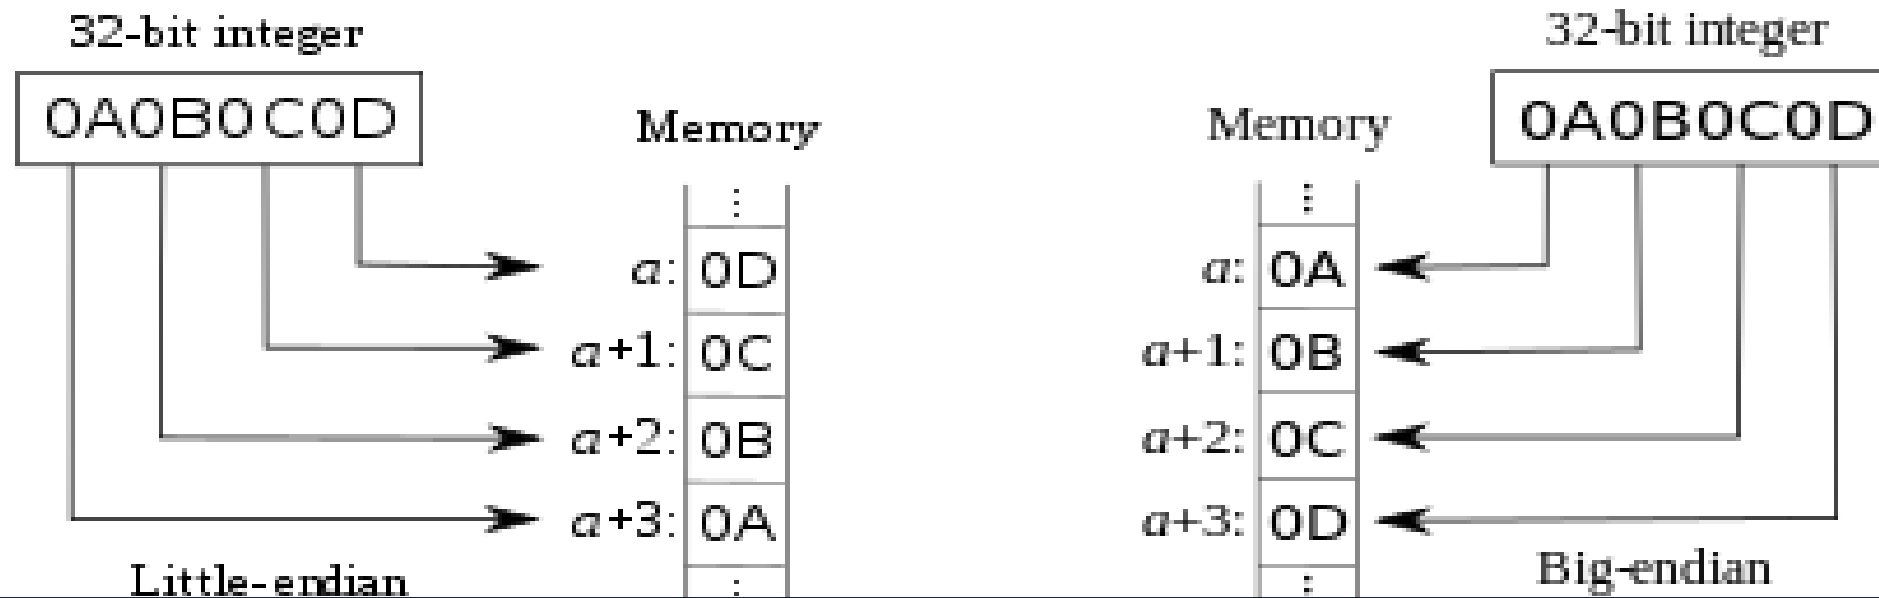
E0	2604
29	2605
58	2606
FE	2607

Poll: Which do you prefer?

- a) Big Endian
- b) Little Endian

Big Endian vs. Little Endian

- Endian-ness: ordering of bytes within a word
 - Little – Bigger address holds more significant bits
 - Big – Opposite, smaller address hold more significant bits
 - The Internet is big endian, x86 is little endian, LEG and ARMv8 can switch
 - But in general assume little endian. (Figures from Wikipedia)



Store Instructions

- Store instructions are simpler—there is no sign/zero extension to consider (do you see why?)

Desired amount of data to transfer?	Operation	Example
64-bits (double word or whole register)	STUR (Store unscaled register)	0xFEDC_BA98_7654_3210
16-bits (half-word) from lower bits of reg	STURH	0x0000_0000_0000_3210
8-bits (byte) from lower bits of reg	STURB	0x0000_0000_0000_0010
32-bits (word) from lower bits of reg	STURW	0x1111_1111_F654_3210

Next Time

- More examples on doing stuff in ARM assembly
 - Like if/else, while loops, etc
- Lingering questions / feedback? I'll include an anonymous form at the end of every lecture: <https://bit.ly/3oXr4Ah>

