

# **4. Instruction Set Architecture – ARM - from C to Assembly**

---

**EECS 370 – Introduction to Computer Organization – Winter 2023**

**EECS Department  
University of Michigan in Ann Arbor, USA**


# Announcements—Reminders

---

- ❑ Project 1.a due Thursday 1/26
- ❑ Project 1.s and 1.m due Thursday 2/2
- ❑ Homework 1 due Monday 1/23
- ❑ All due dates on calendar on web page

# Office hours

- ❑ Most of the office hours use an on-line queue.
  - We have three queues: One for BBB, one for the UGLI, one for on-line
- ❑ We also have group office hours. These don't have queues and are focused on questions about lecture material, homework, and project specifications and high-level project issues.
  - See the website.



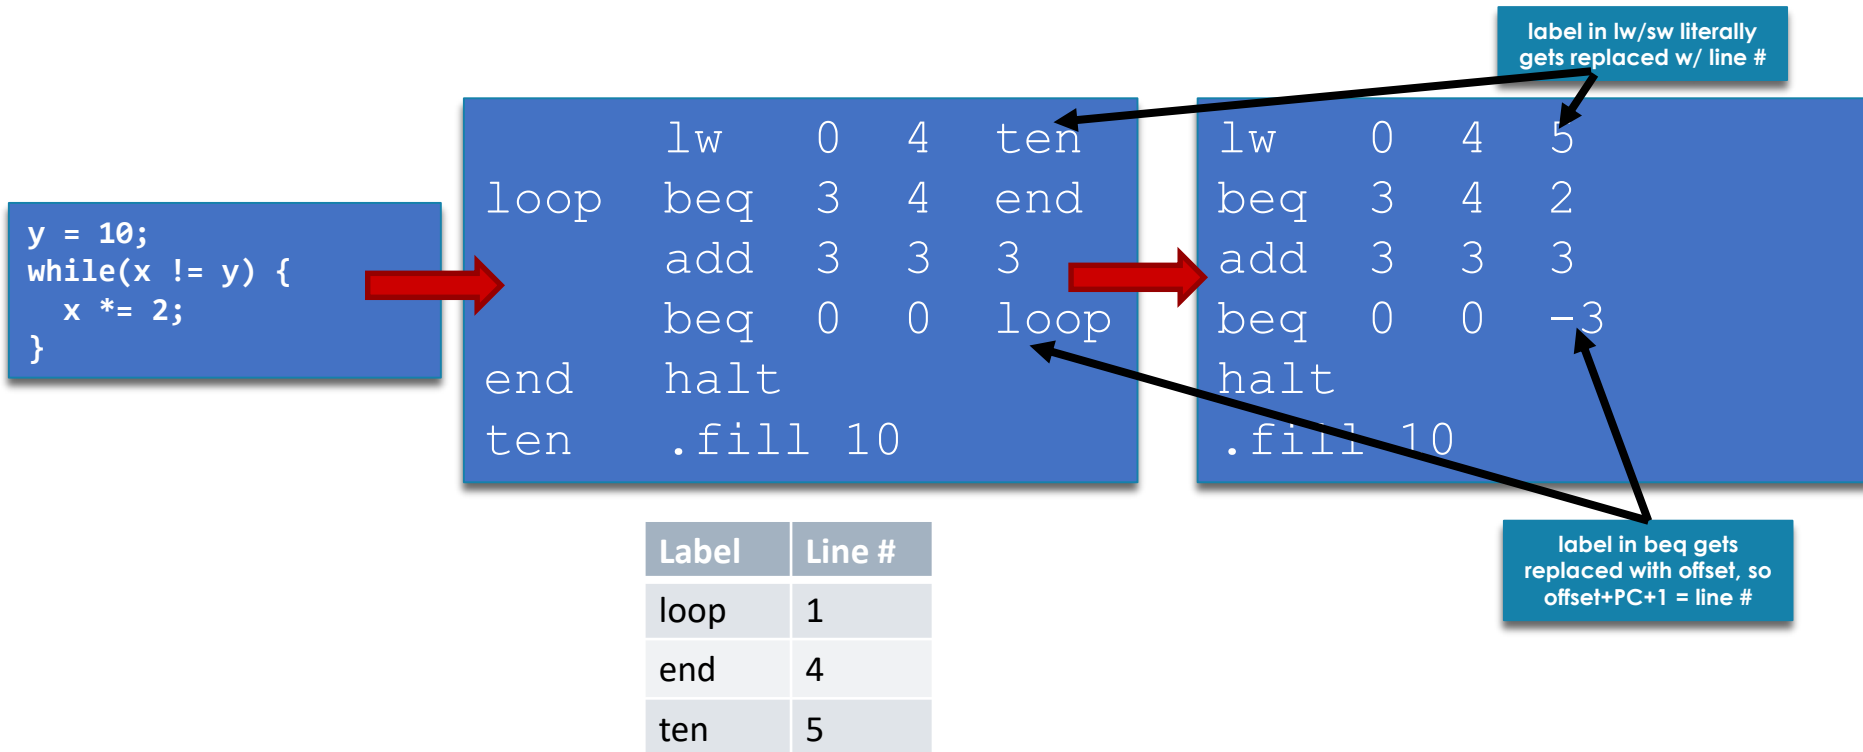
# Instruction Set Architecture (ISA) Design Lectures

---

- ❑ Lecture 2: Storage types
- ❑ Lecture 3 : Addressing modes and LC2K
- ❑ **Lecture 4 : ARM Assembly**
- ❑ Lecture 5 : More Assembly and Function Calls
- ❑ Lecture 6 : Caller/Callee and Linker

# Left-over questions:

*"I'm not getting labels & offsets for lw/sw/beq" (LC2K)*



## ARMv8 ISA 64-bit ISA



Chevrolet Corvette Z06

LS7  
2006 7.0L V-8 (LS7)

# ARMv8 ISA

- ❑ LC2K is intended to be an extremely bare-bones ISA
  - "Bare minimum"
  - Easy to design hardware for, really annoying to program on (as you'll see in P1m)
  - Invented for our projects, not used anywhere in practice
- ❑ ARM (specifically v8) is a much more powerful ISA
  - Used heavily in practice (most smartphones, some laptops & supercomputers)
  - Subset (LEG) is used in lecture and homework where LC2K cannot be.



**SINGLE  
INSTRUCTION ISA**


**LC2K**

**LEG SUBSET**

**ARMV8**

## ARMs and LEGs

- ❑ ARMv8 is the 64 bit version—all registers are 64 bits wide
- ❑ Addresses are calculated to be 64 bits too
- ❑ **BUT: Instructions are 32 bits**
- ❑ We use a (small) subset of the v8 ISA used in P+H
- ❑ It is referred to as the LEGv8 in keeping with the body part theme!





# ARM vs LC2K at a Glance

---

	LC2K	LEG
# registers	8	32
Register width	32 bits	64 bits
Memory size	$2^{18}$ bytes	$2^{64}$ bytes
# instructions	8	40-ish
Addressability	Word	Byte

We'll discuss what  
this means in a bit



# ARM Instruction Set—LEGv8 subset

- ❑ The main types of instructions fall into the familiar classes we saw with LC2K:
  1. Arithmetic
    - Add, subtract, multiply (not in LEGv8)
  2. Data transfer
    - Loads and stores—LDUR (load unscaled register), STUR, etc.
  3. Logical
    - AND, ORR, EOR, etc.
    - Logical shifts, LSL, LSR
  4. Conditional branch
    - CBZ, CBNZ, B.cond
  5. Unconditional branch (jumps)
    - B, BR, BL

# LEGv8 Arithmetic Instructions

- ❑ Format: three operand fields
  - Dest. register usually the first one – check instruction format
  - ADD X3, X4, X7 – Think ADD X3 = X4, X7
  - LC2K generally has the destination on the right!!!!
  - E.g. add 1 2 3 // r3 = r1 + r2

- ❑ C-code example:  $f = (g + h) - (i + j)$

X1 → t0


X2 → t1

X3 → g

X4 → h

X5 → i

X6 → j



## LEGv8 R-instruction Fields

- ❑ Register-to-register operations
- ❑ Consider ADD X3, X4, X7
  - $R[Rd] = R[Rn] + R[Rm]$
  - $Rd = X3, Rn = X4, Rm = X7$
- ❑  $Rm$  = second register operand
- ~~❑  $shamt$  = shift amount~~
  - not used in LEG for ADD/SUB and set to 0
- ❑  $Rn$  = first register operand
- ❑  $Rd$  = destination register
- ❑ ADD opcode is 10001011000, what are the other fields?

opcode	$Rm$	$shamt$	$Rn$	$Rd$
11 bits	5 bits	6 bits	5 bits	5 bits

# LEGv8 Arithmetic Operations

- ❑ Machine State—more on this concept as our understanding evolves
  1. Registers: 32 64 bit registers. X31 aliased to XZR which is always 0  
Can't use as a destination
  2. PC—Program counter
  3. FLAGS: NZVC record the results of (arithmetic) operations  
Negative, Zero, oVerflow, Carry—not present in LC2K

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	$X1 = X2 + X3$	Three register operands
	subtract	SUB X1, X2, X3	$X1 = X2 - X3$	Three register operands
	add immediate	ADDI X1, X2, 20	$X1 = X2 + 20$	Used to add constants
	subtract immediate	SUBI X1, X2, 20	$X1 = X2 - 20$	Used to subtract constants
	add and set flags	ADDs X1, X2, X3	$X1 = X2 + X3$	Add, set condition codes
	subtract and set flags	SUBs X1, X2, X3	$X1 = X2 - X3$	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	$X1 = X2 + 20$	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	$X1 = X2 - 20$	Subtract constant, set condition codes

## I-instruction fields

- ❑ Format: second source operand can be a register or **i**mmEDIATE—a constant in the instruction itself
- ❑ e.g., `ADDI X3, X4, #10`
- ❑ Format: 12 bits for immediate constants 0-4095

opcode	immediate	Rn	Rd
10 bits	12 bits	5 bits	5 bits

- ❑ Don't need negative constants because we have SUBI
- ❑ C-code example: `f = g + 10`

`ADDI X7, X5, #10`
- ❑ C-code example: `f = g - 10`

`SUBI X7, X5, #10`

# LEGv8 Logical Instructions

- ❑ Logical operations are bit-wise
- ❑ For example assume
  - $X9 = 11111111\ 11111111\ 11111111\ 00000000\ 00000000\ 00000000\ 00001101\ 11000000$
  - $X14 = 00000000\ 00000000\ 11011010\ 00000000\ 00000000\ 00000000\ 00111100\ 00000000$

AND X2, X13, X9 would result in

  - $X2 = 00000000\ 00000000\ 11011010\ 00000000\ 00000000\ 00000000\ 00001100\ 00000000$
- ❑ AND and OR correspond to C operators `&` and `|`
- ❑ For immediate fields the 12 bit constant is padded with zeros to the left—zero extended

Category	Instruction	Example	Meaning	Comments
Logical	and	AND X1, X2, X3	$X1 = X2 \& X3$	Three reg. operands; bit-by-bit AND
	inclusive or	ORR X1, X2, X3	$X1 = X2   X3$	Three reg. operands; bit-by-bit OR
	exclusive or	EOR X1, X2, X3	$X1 = X2 \wedge X3$	Three reg. operands; bit-by-bit XOR
	and immediate	ANDI X1, X2, 20	$X1 = X2 \& 20$	Bit-by-bit AND reg. with constant
	inclusive or immediate	ORRI X1, X2, 20	$X1 = X2   20$	Bit-by-bit OR reg. with constant
	exclusive or immediate	EORI X1, X2, 20	$X1 = X2 \wedge 20$	Bit-by-bit XOR reg. with constant
	logical shift left	LSL X1, X2, 10	$X1 = X2 \ll 10$	Shift left by constant
	logical shift right	LSR X1, X2, 10	$X1 = X2 \gg 10$	Shift right by constant

## LEGv8 Shift Logical Instructions

- ❑ LSR X6, X23, #2
 

X23 = 11111111 11111111 11111111 00000000 00000000 00000000 11011010 00000010

X6 = 00111111 11111111 11111111 11000000 00000000 00000000 00110110 10000000
- ❑ C equivalent
  - X6 = X23 >> 2;
- ❑ CLASS Question: LSL X6, X23, #2 ?
  - What register gets modified?
  - What does it contain after executing the LSL instruction?
- ❑ In shift operations Rm is always 0—shamt is 6 bit unsigned

opcode	Rm	shamt	Rn	Rd
11 bits	5 bits	6 bits	5 bits	5 bits




## Pseudo Instructions

---

- ❑ Instructions that use a shorthand “mnemonic” that expands to an assembly instruction
- ❑ Example:
  - `MOV X12, X2` // the contents of X2 copied to X12—X2 unchanged
- ❑ This gets expanded to:
  - `ORR X12, XZR, X2`
- ❑ What alternatives could we use instead of ORR?

## Class Problem

- Show the C and LEGv8 assembly for extracting the value in bits 15:10 from a 64-bit integer variable



Want these bits

Assume the variable is in X1

```
x = x >> 10
x = x & 0x3F
```

```
LSR    X1, X1, #10
ANDI   X1, X1, #0x3F
```

# Word vs Byte Addressing

---

## ❑ A **word** is a collection of bytes

- Exact size depends on architecture
- in LC2K and ARM, 4 bytes
  - **Double word** is 8 bytes

## ❑ LC2K is **word addressable**

- Each address refers to a particular **word** in memory
- Wanna move forward one int? Increment address by **one**
- Wanna move forward one char? Uhhh...



## ❑ ARM (and most modern ISAs) is **byte addressable**

- Each address refers to a particular **byte** in memory
- Wanna move forward one int? Increment address by **four**
- Wanna move forward one char? Increment address by **one**

# LEGv8 Memory Instructions

- ❑ Employs base + displacement addressing mode
  - Base is a register
  - Displacement is 9-bit immediate  $\pm 256$  bytes—sign extended to 64 bits

Category	Instruction	Example	Meaning	Comments
	load register	LDUR X1, [X2,40]	$X1 = \text{Memory}[X2 + 40]$	Doubleword from memory to register
	store register	STUR X1, [X2,40]	$\text{Memory}[X2 + 40] = X1$	Doubleword from register to memory
	load signed word	LDURSW X1,[X2,40]	$X1 = \text{Memory}[X2 + 40]$	Word from memory to register
	store word	STURW X1, [X2,40]	$\text{Memory}[X2 + 40] = X1$	Word from register to memory
	load half	LDURH X1, [X2,40]	$X1 = \text{Memory}[X2 + 40]$	Halfword memory to register
	store half	STURH X1, [X2,40]	$\text{Memory}[X2 + 40] = X1$	Halfword register to memory
	load byte	LDURB X1, [X2,40]	$X1 = \text{Memory}[X2 + 40]$	Byte from memory to register
	store byte	STURB X1, [X2,40]	$\text{Memory}[X2 + 40] = X1$	Byte from register to memory
	move wide with zero	MOVZ X1,20, LSL 0	$X1 = 20 \text{ or } 20 * 2^{16} \text{ or } 20 * 2^{32} \text{ or } 20 * 2^{48}$	Loads 16-bit constant, rest zeros
	move wide with keep	MOVK X1,20, LSL 0	$X1 = 20 \text{ or } 20 * 2^{16} \text{ or } 20 * 2^{32} \text{ or } 20 * 2^{48}$	Loads 16-bit constant, rest unchanged

## D-Instruction fields

- ❑ Data transfer
- ❑ opcode and op2 define data transfer operation
- ❑ address is the  $\pm 256$  bytes displacement
- ❑ Rn is the base register
- ❑ Rt is the destination (loads) or source (stores)
- ❑ More complicated modes are available in full ARMv8

opcode	address	op2	Rn	Rt
11 bits	9 bits	2 bits	5 bits	5 bits

## LEGv8 Memory Instructions

- ❑ Registers are 64 bits wide
- ❑ But sometimes we want to deal with non-64-bit entities
  - E.g. ints (32 bits), chars (8 bits)
- ❑ When we load smaller elements from memory, what do we set the other bits to?

- Option A: set to zero



- Option B: sign extend



- ❑ We'll need different instructions for different options

# Load Instruction Sizes

---

How much data is retrieved from memory at the given address?

- ❑ LDUR X3, [X4, #100]
  - Load (unscaled) to register—retrieve a double word (64 bits) from address (X4+100)
- ❑ LDURH X3, [X4, #100]
  - Load halfword (16 bits) from address (X4+100) to the low 16 bits of X3—top 48 bits of X3 are set zero
- ❑ LDURB X3, [X4, #100]
  - Load byte (8 bits) from address (X4+100) and put in the low 8 bits of X3—zero extend the destination register X3 (top 56 bits)
- ❑ What about loading words?
- ❑ LDURSW X3, [X4, #100]
  - retrieve a word (32 bits) from address (X4+100) and put in lower half of X3—top 32 bits of X3 are sign extended

## LEGv8 Data Transfer Instructions--Stores

- ❑ Store instructions are simpler—there is no sign/zero extension to consider
- ❑ STUR X3, [X4, #100]
  - Store (unscaled) register—write the double word (64 bits) in register X3 to the 8 bytes at address (X4+100)
- ❑ STURW X3, [X4, #100]
  - Store word—write the word (32 bits) in the low 4 bytes of register X3 to the 4 bytes at address (X4+100)
- ❑ STURH X3, [X4, #100]
  - Store half word—write the half word (16 bits) in the low 2 bytes of register X3 to the 2 bytes at address (X4+100)
- ❑ STURB X3, [X4, #100]
  - Store byte—write the least significant byte (8 bits) in register X3 to the byte at address (X4+100)



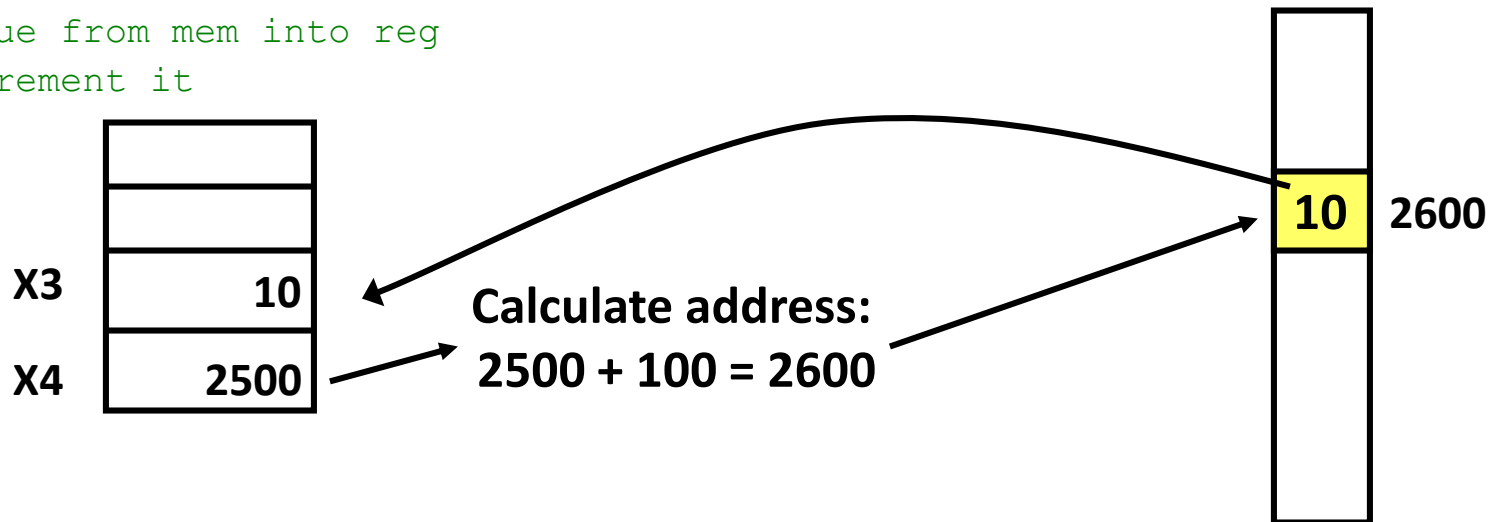
# Load Instruction in Action

❑ LDURB X3, [X4, #100] // load byte

Retrieves 8-bit value from memory location (X4+100) and puts the result into X3. The other 56 most significant bits are zero extended

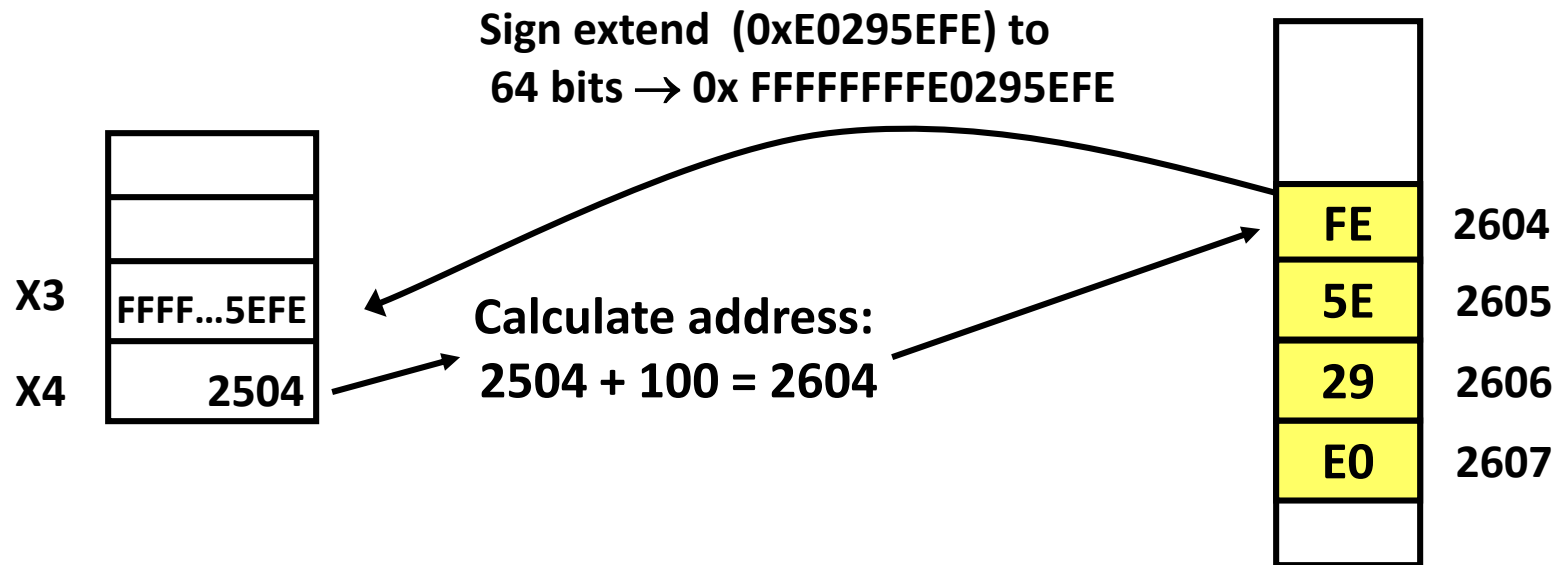
```
struct {  
    int arr[25];  
    unsigned char c;  
} my_struct;
```

```
int func() {  
    my_struct.c++;  
    // load value from mem into reg  
    // then increment it  
}
```



## Load Instruction in Action – other example

- LDURSW X3, [X4, #100] // load signed word  
Retrieves 4 byte value from memory location (X4+100) and puts the result into X3 (sign extended)



- We have mixed decimal and hex to keep the numbers easy to read
- Recall that E = 1110 and thus is treated as a negative 2's complement number

## But wait...

---

```
int my_big_number = -534159618; // 0xE0295EFE in 2's complement
```

- If I want to store this number in memory...

should it be stored like this?


FE	2604
5E	2605
29	2606
E0	2607

... or like this?

E0	2604
29	2605
58	2606
FE	2607

# Big Endian vs. Little Endian

- Endian-ness: ordering of bytes within a word
  - Little - increasing numeric significance with increasing memory addresses
  - Big – The opposite, most significant byte first
  - The Internet is big endian, x86 is little endian, LEG and ARMv8 can switch
    - But in general assume little endian. (Figures from Wikipedia)



# Store Instructions

- Store instructions are simpler—there is no sign/zero extension to consider (do you see why?)


Desired amount of data to transfer?	Operation	Example
64-bits (double word or whole register)	STUR (Store unscaled register)	0xFEDC_BA98_7654_3210
16-bits (half-word) from lower bits of reg	STURH	0x0000_0000_0000_3210
8-bits (byte) from lower bits of reg	STURB	0x0000_0000_0000_0010
32-bits (word) from lower bits of reg	STURW	0x1111_1111_F654_3210

## Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STUR     X3, [X5, #100]
STURB    X4, [X5, #102]
```

register file



Memory  
(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107


little endian

## Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STUR     X3, [X5, #100]
STURB    X4, [X5, #102]
```

register file



Memory

(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

little endian

## Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STUR     X3, [X5, #100]
STURB    X4, [X5, #102]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory

(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

little endian



## Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STUR    X3, [X5, #100]
STURB   X4, [X5, #102]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory

(each location is 1 byte)

0xFF	100
0x00	101
0x00	102
0x00	103
0x00	104
0x00	105
0x00	106
0x00	107

little endian

## Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STUR     X3, [X5, #100]
STURB   X4, [X5, #102]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory

(each location is 1 byte)

0xFF	100
0x00	101
<b>0x02</b>	102
0x00	103
0x00	104
0x00	105
0x00	106
0x00	107

little endian

## Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STURB   X3, [X5, #103]
LDURSW  X4, [X5, #100]
```

register file



*We shown the registers as blank. What do they actually contain before we run the snippet of code*

Memory  
(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107


little endian

## Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STURB   X3, [X5, #103]
LDURSW  X4, [X5, #100]
```

register file



Memory

(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

little endian

## Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STURB   X3, [X5, #103]
LDURSW  X4, [X5, #100]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory

(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

little endian

## Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STURB   X3, [X5, #103]
LDURSW  X4, [X5, #100]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory

(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0xFF	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

little endian

## Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STURB   X3, [X5, #103]
LDURSW  X4, [X5, #100]
```

register file

X3	0x0000000000000000FF
X4	0xFFFFFFFFFFFFFFFF0302

Memory

(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0xFF	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

little endian

# Converting C to assembly – example 1

Write ARM assembly code for the following C expression:

**C:** `a = b + names[ i ];`

Assume that **a** is in X1, **b** is in X2, **i** is in X3, and the address of the array **names** is in X4. The array holds 64-bit integers

```
LSL      X5, X3, #3          // calculate array offset i*8 (why #3 if we want 8?)
ADD      X6, X4, X5          // calculate address of names[i]
LDUR     X4, [X6, #0]        // load names[i]
ADD      X1, X2, X4          // calculate b + names[i]
```

CLASS QUESTION: What if names was an array of **int** (4 bytes) type?



## Converting C to assembly – example 2

Write ARM assembly code for the following C expression:

(assume an int is 4 bytes)

```
struct { int a; unsigned char b, c; } y;  
y.a = y.b + y.c;
```

Assume that a pointer to **y** is in X1.

```
LDURB  X2, [X1, #4]  // load y.b
```

```
LDURB  X3, [X1, #5]  // load y.c
```

```
ADD     X4, X2, X3    // calculate y.b+y.c
```

```
STURW  X4, [X1, #0]  // store y.a
```

How do you determine the offsets for the struct sub-fields?