

**Poll: Who is responsible for  
erasing a whiteboard in a public  
space?**

- a) The person who is done using it
- b) The person who is about to use it

# EECS 370 - Lecture 6

## Function Calls



# Announcements

- P1
  - Part a due Thu @ 11:59 via Autograder
  - Project 1 s + m due Thu 1/26



- HW 2
  - Posted on website, due Mon 2/6 at 8 pm
- Get exam conflicts and SSD accommodations sent to us **ASAP**
  - Forms listed on the website

# Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- **Lecture 6 : Converting C to assembly – functions**
- Lecture 7 : Translation software; libraries, memory layout



# ARM/LEGv8 Sequencing Instructions

- Sequencing instructions change the flow of instructions that are executed
  - This is achieved by modifying the program counter (PC)
- Unconditional branches are the most straightforward they ALWAYS change the PC and thus “jump” to another instruction out of the usual sequence
- Conditional branches

If (condition\_test) goto target\_address

*condition\_test* examines the four flags from the processor status word (SPSR)

*target\_address* is a 19 bit signed word displacement on current PC

# LEGv8 Conditional Instructions

- Two varieties of conditional branches
  1. One type compares a register to see if it is equal to zero.
  2. Another type checks the condition codes set in the status register.

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Talked about CBZ and CBNZ last time
- Let's look at B.cond now

# LEGv8 Conditional Instructions

- Motivation:
  - Some types of branches makes sense to check if a certain value is zero or not
    - while(a)
  - But not all:
    - if(a > b)
    - if(a == b)
  - Using an extra **program status register** to check for various conditions allows for a greater breadth of branching behavior

# LEGv8 Conditional Instructions Using FLAGS

- FLAGS: NZVC record the results of (arithmetic) operations  
Negative, Zero, oVerflow, Carry—not present in LC2K
- We explicitly set them using the “set” modification to ADD/SUB etc.
- Example: ADDS causes the 4 flag bits to be set according as the outcome is negative, zero, overflows, or generates a carry

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	$X1 = X2 + X3$	Three register operands
	subtract	SUB X1, X2, X3	$X1 = X2 - X3$	Three register operands
	add immediate	ADDI X1, X2, 20	$X1 = X2 + 20$	Used to add constants
	subtract immediate	SUBI X1, X2, 20	$X1 = X2 - 20$	Used to subtract constants
	add and set flags	ADDS X1, X2, X3	$X1 = X2 + X3$	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	$X1 = X2 - X3$	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	$X1 = X2 + 20$	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	$X1 = X2 - 20$	Subtract constant, set condition codes



# ARM Condition Codes Determine Direction of Branch

- In LEGv8 only ADDS / SUBS / ADDIS / SUBIS / CMP /CMPI set the condition codes FLAGS or condition codes in PSR—the program status register
- Four primary condition codes evaluated:
  - N – set if the result is **negative** (i.e., bit 63 is non-zero)
  - Z – set if the result is **zero** (i.e., all 64 bits are zero)
  - ~~• C – set if last addition/subtraction had a **carry**/borrow out of bit 63~~
  - ~~• V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)~~
- Don't worry about the C and V for this class



# ARM Condition Codes Determine Direction of Branch--continued

	Encoding	Name (& alias)	Meaning (integer)	Flags
→	0000	EQ	Equal	Z==1
→	0001	NE	Not equal	Z==0
	0010	HS (CS)	Unsigned higher or same (Carry set)	C==1
	0011	LO (CC)	Unsigned lower (Carry clear)	C==0
	0100	MI	Minus (negative)	N==1
	0101	PL	Plus (positive or zero)	N==0
	0110	VS	Overflow set	V==1
	0111	VC	Overflow clear	V==0
	1000	HI	Unsigned higher	C==1 && Z==0
	1001	LS	Unsigned lower or same	!(C==1 && Z==0)
→	1010	GE	Signed greater than or equal	N==V
→	1011	LT	Signed less than	N!=V
→	1100	GT	Signed greater than	Z==0 && N==V
→	1101	LE	Signed less than or equal	!(Z==0 && N==V)
→	1110	AL	Always	Any
	1111	NV <sup>†</sup>		

Need to know  
the 7 with the  
red arrows

# Conditional Branches: How to use

- CMP instruction lets you compare two registers.
  - Could also use SUBS etc.
    - That could save you an instruction.
- B.cond lets you branch based on that comparison.

- Example:

```
CMP    X1, X2  
B.GT   Label1
```

- Branches to Label1 if X1 is greater than X2.

# Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

# Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

## Using Labels

```
CMP X1, X2  
B.NE L1  
ADD X1, X1, #1  
B L2  
L1: ADD X2, X2, #1  
L2: ...
```

Note that conditions in assembly are often the inverse of the "if" condition. Why?

## Without Labels

```
CMP X1, X2  
B.NE 3  
ADD X1, X1, #1  
B 2  
ADD X2, X2, #1
```

Assemblers must deal with labels and assign displacements

# Loop—Example

// assume all variables are long long integers (64 bits or 8 bytes)  
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

# of branch instructions  
=  $3 \cdot 10 + 1 = 31$

a.k.a. while-do template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	CMPI	X1, #10
	B.EQ	endLoop
	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endif:	ADDI	X1, X1, #1
	B	Loop1
endLoop:		

# Extra Example: Do-while Loop

Look through  
this on your  
own

// assume all variables are long long integers (64 bits or 8 bytes)  
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

# of branch instructions  
=  $2 \times 10 = 20$

a.k.a. do-while template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endif:	ADDI	X1, X1, #1
	CMPI	X1, #10
	B.LT	Loop1
endLoop:		

# Extra Problem – For Your Reference

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in X1  
// struct {int val; struct node *next;} node;  
// struct node *ptr;
```

```
if ((ptr != NULL) && (ptr->val > 0))  
    ptr->val++;
```

# Extra Problem

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in X1
// struct {int val; struct node *next;} node;
// struct node *ptr;
```

```
if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```

```
cmp r1, #0
beq Endif
ldursw r2, [r1, #0]
cmp r2, #0
b.le Endif
add r2, r2, #1
str r2, [r1, #0]
Endif : ....
```



# Branching far away

- Underlying philosophy of ISA design: **make the common case fast**
- Most branches target nearby instructions
  - Displacement of 19 bits is usually enough
- BUT what if we need to branch really far away (more than  $2^{19}$  words)?

CBZ     X15, FarLabel

- The assembler is smart enough to replace that with

CBNZ   X15, L1

B       FarLabel

L1:     .....

- The simple branch instruction (B) has a 26 bit offset which spans about 64 million instructions!
- In LC2K, we can do a similar thing by using JALR instead of BEQ

# Unconditional Branching Instructions

Unconditional branch	branch	B	2500	go to PC + 10000	Branch to target address; PC-relative
	branch to register	BR	X30	go to X30	For switch, procedure return
	branch with link	BL	2500	X30 = PC + 4; PC + 10000	For procedure call PC-relative

- There are three types of unconditional branches in the LEGv8 ISA.
  - The first (**B**) is the PC relative branch with the 26 bit offset from the last slide.
  - The second (**BR**) jumps to the address contained in a register (X30 above)
  - The third (**BL**) is like our PC relative branch but it does something else.
    - It sets X30 (always) to be the current PC+4 before it branches.
- Why is BL storing PC+4 into a register?

# Branch with Link (BL)

- Branch with Link is the branch instruction used to call functions
  - Functions need to know where they were called from so they can return.
    - In particular they will need to return to right after the function call
    - Can use “BR X30”
- Say that we execute the instruction BL #200 when at PC 1000.
  - What address will be branched to?
  - What value is stored in X30?
  - How is that value in X30 useful?

**Poll**

# Converting function calls to assembly code

C: factorial(5);

- Need to pass parameters to the called function—factorial
- Need to save return address of caller so we can get back
- Need to save register values (why?)
- Need to jump to factorial

Execute instructions for factorial()  
Jump to return address

- Need to get return value (if used)
- Restore register values

# Task 1: Passing parameters

- Where should you put all of the parameters?
  - Registers?
    - Fast access but few in number and wrong size for some objects
  - Memory?
    - Good general solution but slow
- ARMv8 solution—and the usual answer:
  - Both
    - Put the first few parameters in registers (if they fit) (X0 – X7)
    - Put the rest in memory on the call stack— **important concept**

# Call stack

- ARM conventions (and most other processors) allocate a region of memory for the “call” stack
  - This memory is used to manage all the storage requirements to simulate function call semantics
    - Parameters (that were not passed through registers)
    - Local variables
    - Temporary storage (when you run out of registers and need somewhere to save a value)
    - Return address
    - Etc.
- Sections of memory on the call stack [**stack frames**] are allocated when you make a function call, and de-allocated when you return from a function

# The stack grows as functions are called

FUNCTION CALLS

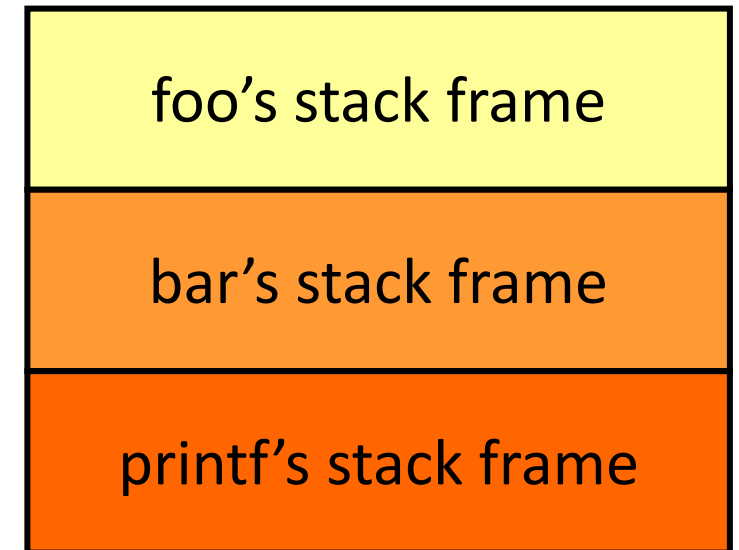
```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```


inside foo



bar calls printf



foo calls bar



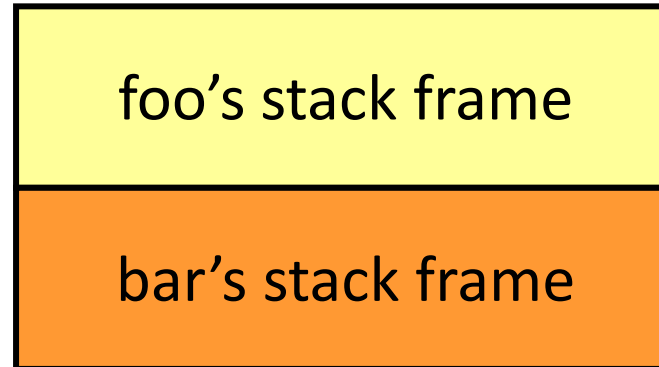
# The stack shrinks as functions return

FUNCTION CALLS


```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

printf returns



bar returns





# Stack frame contents

FUNCTION CALLS

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}  
  
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

foo's stack frame

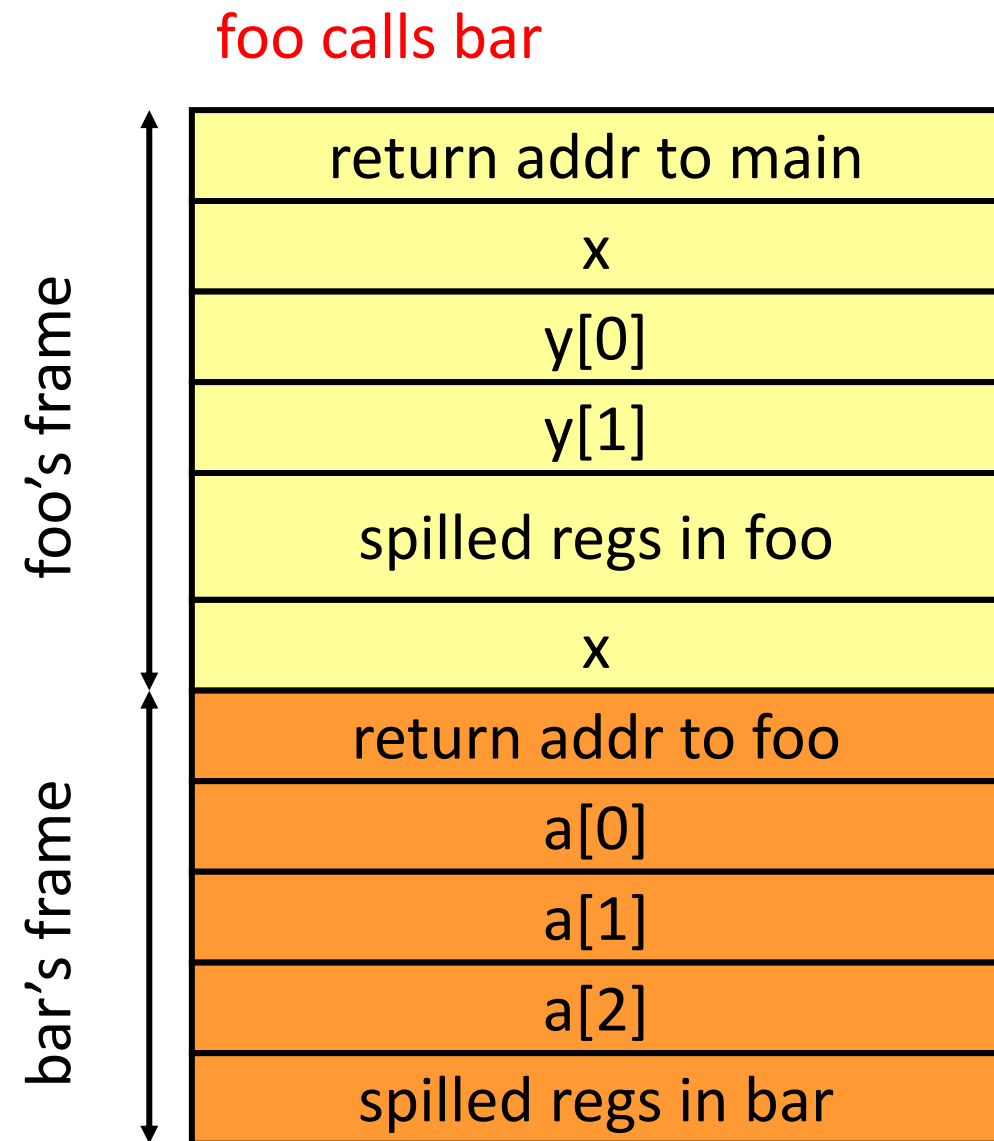
return addr to main
x
y[0]
y[1]
spilled registers in foo

# Stack frame contents (2)

```
void foo()
{
    int x, y[2];
    bar(x);
}

void bar(int x)
{
    int a[3];
    printf();
}
```

Spill data—not enough room in x0-x7 for  
params and also caller and callee saves




## Review: Where do the variables go?

# Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```



# Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```

w goes in static, as it's a global

x goes on the stack, as it's a parameter

y goes in static, 1 copy of this!!


p goes on the stack

allocate 10 bytes on heap, ptr  
set to the address

string goes in static, pointer  
to string on stack, p goes on  
stack

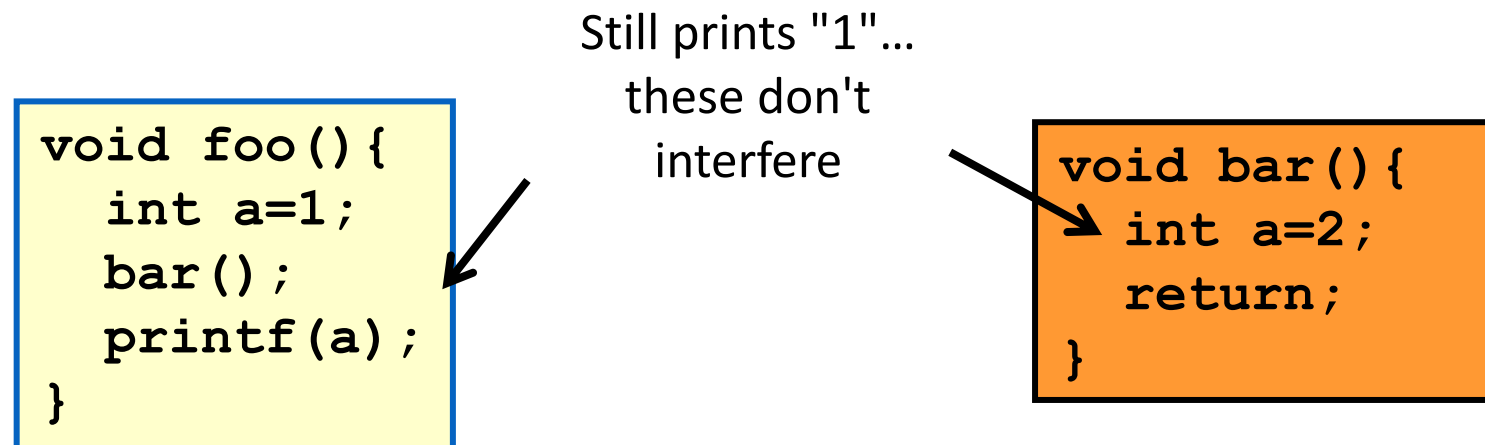
The addresses of local variables  
will be different depending on  
where we are in the call stack

```
ld r13, r1, STACK  
// Here, r13 is holding the offset of  
how far we are down the stack
```



# What about registers?

- Higher level languages (like C/C++) provide many abstractions that don't exist at the assembly level
- E.g. in C, each function has its own local variables
  - Even if different function have local variables with the same name, they are independent and guaranteed not to interfere with each other!



# What about registers?

- But in assembly, all functions share a small set (e.g. 32) of registers
  - Called functions will overwrite registers needed by calling functions

```
main: movz X0, #1  
      bl  foo  
      bl  printf
```

foo() overwrites  
X0 if we don't  
do something!!

```
foo: movz X0, #2  
     br  X30
```

- "Someone" needs to save/restore values when a function is called to ensure this doesn't happen

# Two Possible Solutions

- Either the **called** function **saves** register values before it overwrites them and **restores** them before the function returns (**callee** saved)...

```
main: movz X0, #1
      bl  foo
      bl  printf
```

```
foo:  stur X0, [stack]
      movz X0, #2
      ldur X0, [stack]
      br  X30
```

- Or the **calling** function **saves** register values before the function call and **restores** them after the function call (**caller** saved)...

```
main: movz X0, #1
      stur X0, [stack]
      bl  foo
      ldur X0, [stack]
      bl  printf
```

```
foo:  movz X0, #2
      br  X30
```

# Another example

## Original C Code

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

No need to  
save r2/r3.  
Why?

## Additions for Caller-save

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    save r1 to stack  
    save r4 to stack  
    bar();  
    restore r1  
    restore r4  
    d = a+d;  
    return();  
}
```

Assume bar() will  
overwrite registers  
holding a,d

## Additions for Callee-save

```
void foo() {  
    int a,b,c,d;  
    save r1  
    save r2  
    save r3  
    save r4  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    bar();  
    d = a+d;  
    restore r1  
    restore r2  
    restore r3  
    restore r4  
    return();  
}
```

bar() will save a,b, but  
now foo() must save  
main's variables



# “caller-save” vs. “callee-save”

- Caller-save

- What if bar() doesn't use r1/r4?
- No harm done, but wasted work

```
void foo(){  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    save r1 to stack  
    save r4 to stack  
    bar();  
    restore r1  
    restore r4  
    d = a+d;  
    return();  
}
```


- Caller-save

- What if main() doesn't use r1-r4?
- No harm done, but wasted work

```
void foo(){  
    int a,b,c,d;  
    save r1  
    save r2  
    save r3  
    save r4  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    bar();  
    d = a+d;  
    restore r1  
    restore r2  
    restore r3  
    restore r4  
    return();  
}
```

# Another helpful visual

CALLER-CALLEE



# Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Caller-saved
  - Only needs saving if value is “live” across function call
  - **Live** = contains a useful value: Assign value before function call, use that value after the function call
  - In a leaf function (a function that calls no other function), caller saves can be used without saving/restoring

a, d are live

b, c are NOT  
live

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

# Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Callee-saved
  - Only needs saving at beginning of function and restoring at end of function
  - Only save/restore it if function overwrites the register

Only use r1-  
r4

No need to  
save other  
registers

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

# Caller versus Callee

- Which is better??
- Neither is obviously better – depends on specific code
- Common solution: designate some registers as caller-saved, some as callee-saved
- More next time

# Next Time

- Finish Up Function Calls
- Talks about linking – the final puzzle piece of software
- Lingering questions / feedback? I'll include an anonymous form at the end of every lecture: <https://bit.ly/3oXr4Ah>

