

21. Cache Organization: Wrapup & Practice Problems

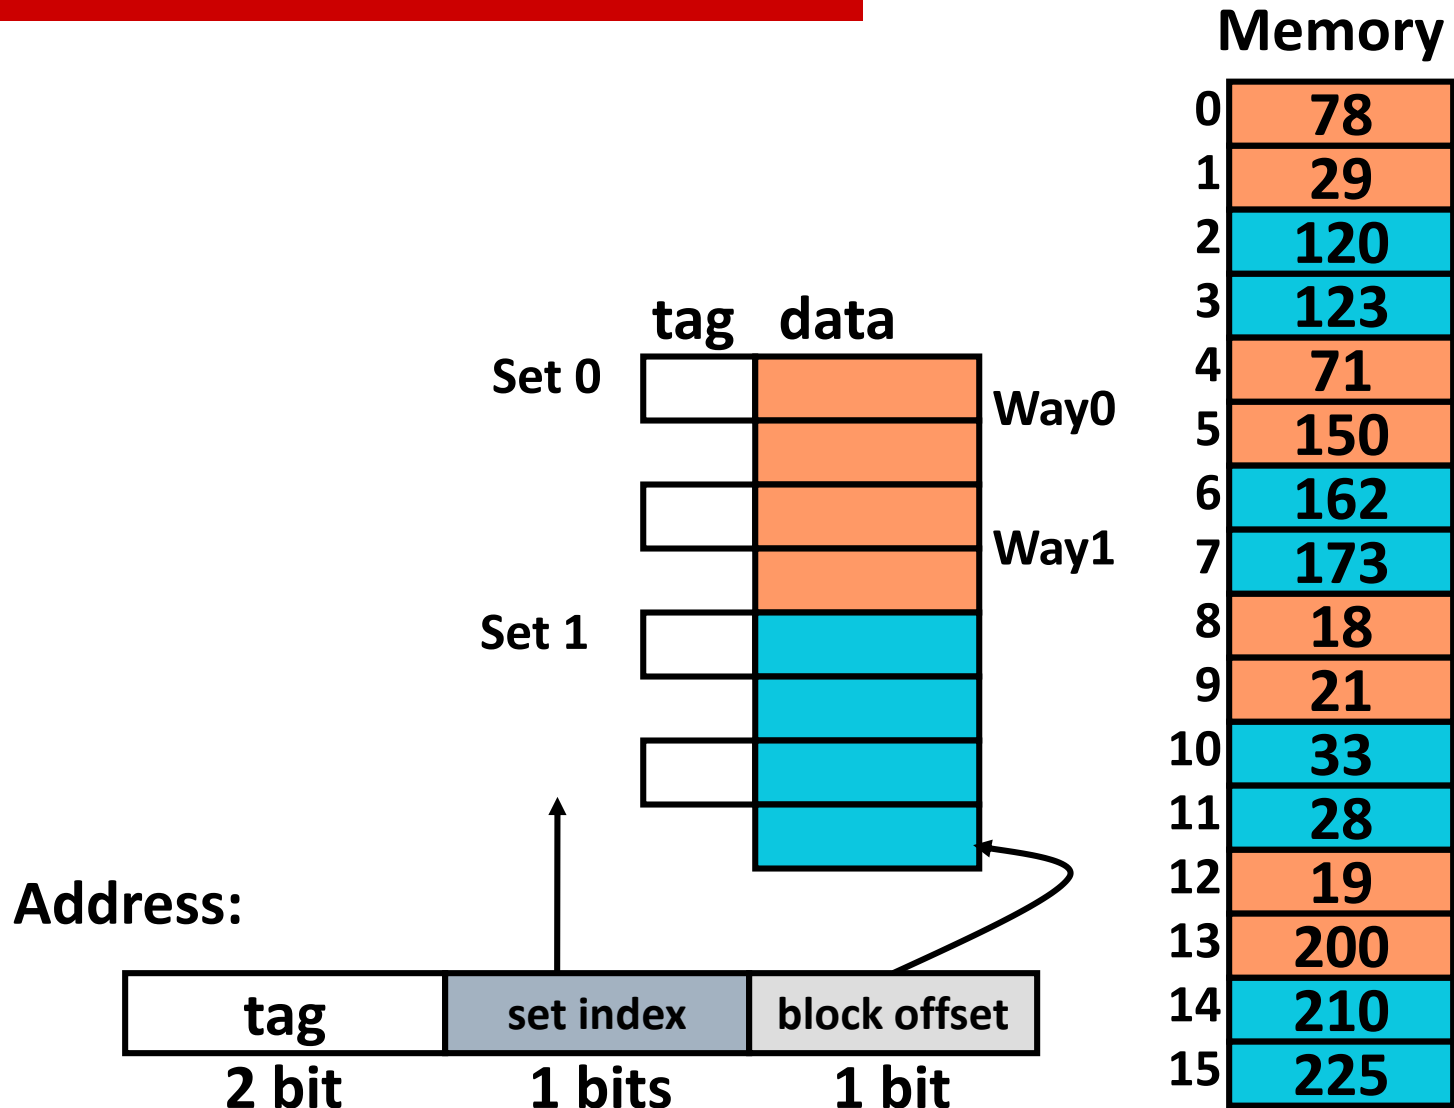
EECS 370 – Introduction to Computer Organization – Winter 2023

**EECS Department
University of Michigan in Ann Arbor, USA**

Assignments

- ❑ HW5 due Monday April 3rd
- ❑ Project 4 due Thursday April 13th

Review: Set-associative cache



Review: The **3C's** of Cache Misses

- ❑ First reference to an address
 - **Compulsory** miss
 - Also sometimes called a “**cold start**” miss
 - First reference to any block will always miss
- ❑ Cache is too small to hold all the data
 - **Capacity** miss
 - Would have had a hit with a large enough cache
- ❑ Replaced it from a busy set
 - **Conflict** miss
 - Would have had a hit with a fully associative cache

Review: Classifying Cache Misses

- ❑ Can you classify all cache misses?
 - **Compulsory miss?**
 - **Capacity miss?**
 - **Conflict miss?**
- ❑ Yes! (with a cache simulator)
 - Simulate with a cache of unlimited size (cache size = memory size) – Any misses must be **compulsory misses**
 - Simulate again with a fully associative cache of the intended size - Any new misses must be **capacity misses**
 - Simulate a third time, with the actual intended cache - Any new misses must be **conflict misses**

Review: Fixing cache misses

❑ Compulsory misses

- First reference to a address
- No way to completely avoid these
- Reduce by **increasing block size**
- This reduces the total number of blocks

❑ Capacity misses

- Would have a hit with a large enough cache
- Reduce by **building a bigger cache**

❑ Conflict misses


- Would have had a hit with a fully associative cache
- Cache does not have enough associativity
- Reduce by **increasing associativity**

Cache Parameters vs. Miss Rate

- ❑ Cache Size
- ❑ Block Size
- ❑ Associativity
- ❑ Replacement policy


Cache Size

- ❑ Cache size in the total data (not including tag) capacity
 - bigger can exploit temporal locality better
 - not ALWAYS better
- ❑ Too large a cache adversely affects hit & miss latency
 - smaller is faster => bigger is slower
 - access time may degrade critical path
- ❑ Too small a cache
 - doesn't exploit temporal locality well
 - useful data replaced often
- ❑ **Working set**: the whole set of data executing application references
 - **Within a time interval**




Block size (also called Line size)

- ❑ Block size is the data that is associated with an address tag
 - Sub-blocking: A block divided into multiple pieces (each with V bit)
 - Can improve “write” performance
- ❑ Too small blocks
 - don't exploit spatial locality well
 - have larger tag overhead
- ❑ Too large blocks
 - too few total # of blocks
 - likely-useless data transferred
 - Extra bandwidth/energy consumed



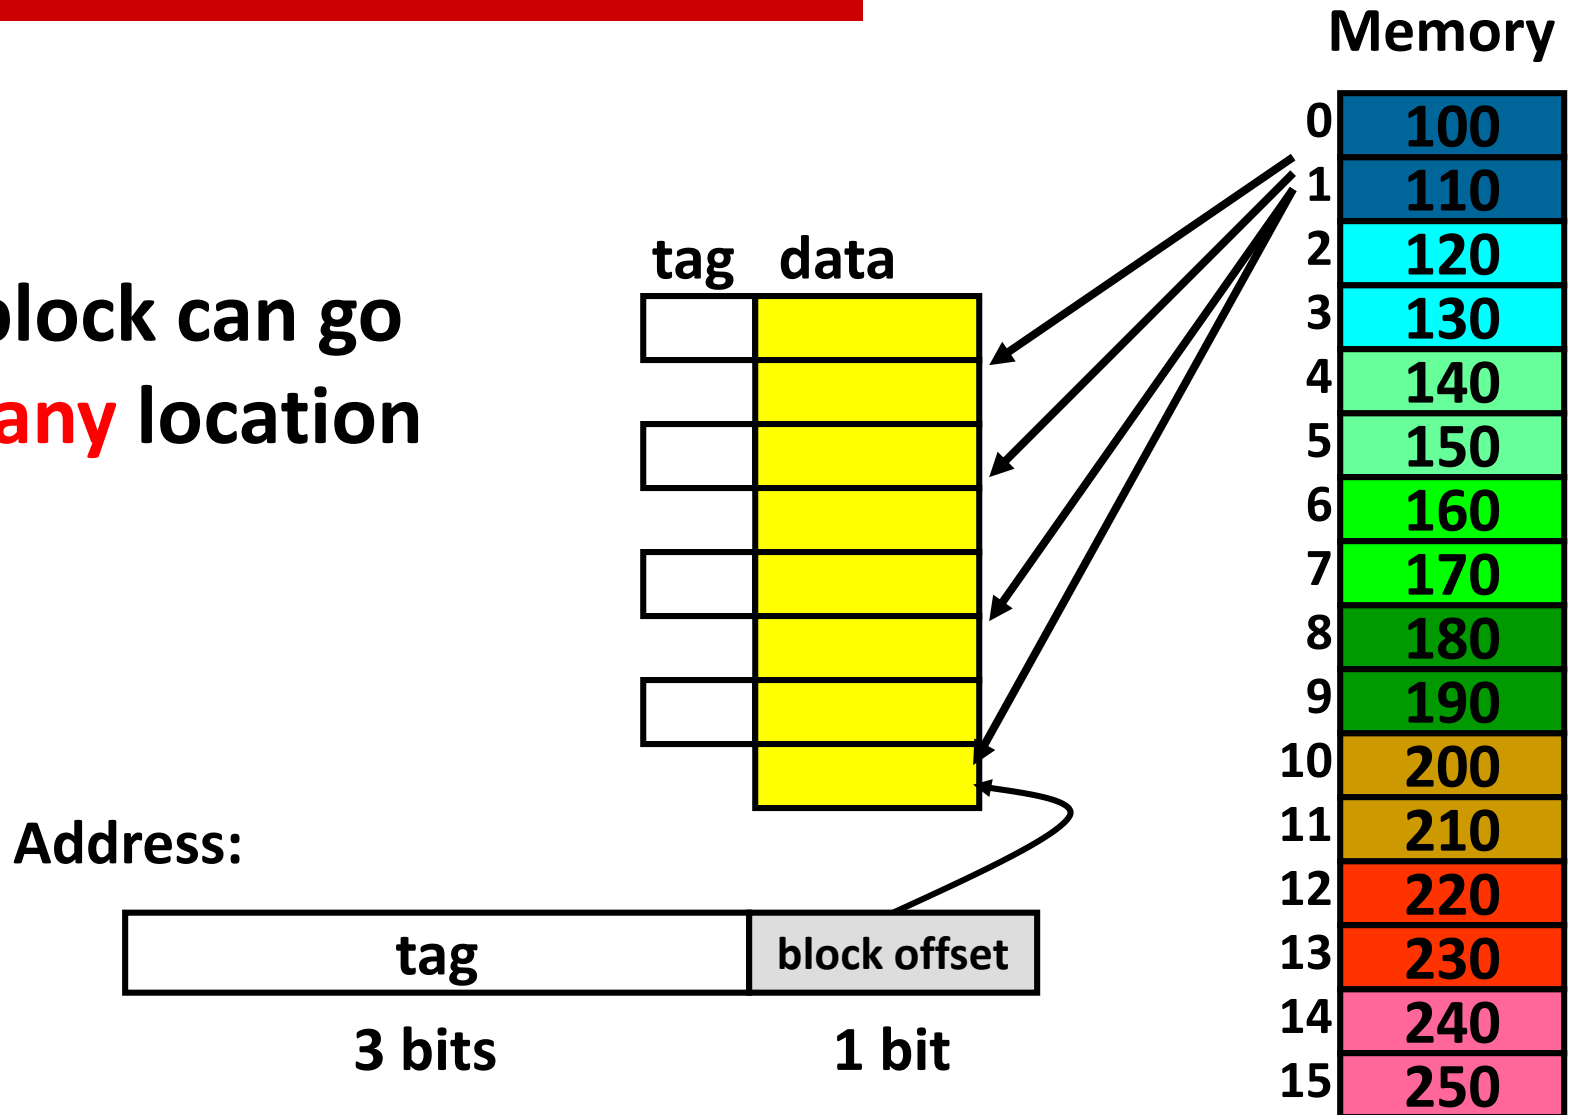
Associativity

- ❑ How many blocks can map to the same index (or set)?
- ❑ Larger associativity
 - lower miss rate, less variation among programs
 - diminishing returns
- ❑ Smaller associativity
 - lower cost
 - faster hit time
 - Especially important for L1 caches
- ❑ Power of 2 associativity?




Review: Fully-associative caches

A block can go
to **any** location




Review: Direct-mapped caches

A block can go to **one** location



Review: Set-associative cache



Practice Problem 1: CPI with caches

The *grinder* application run on the LC2k with full data forwarding and all branches predicted not-taken has the following instruction frequencies:

45% R-type 20% Branches 15% Loads 20% Stores

In *grinder*, 40% of branches are taken and 50% of LWs are followed by an immediate use.

The I-cache has a miss rate of 3% and the D-cache has a miss rate of 6% (no overlapping of misses). On a miss, the main memory is accessed and has a latency of 100 ns. The clock frequency is 500 MHz.

What is the CPI of *grinder* on the LC2k?

Practice Problem 1: Solution

The *grinder* application run on the LC2k with full data forwarding and all branches predicted not-taken has the following instruction frequencies:

45% R-type 20% Branches 15% Loads 20% Stores

In *grinder*, 40% of branches are taken and 50% of LWs are followed by an immediate use.

The I-cache has a miss rate of 3% and the D-cache has a miss rate of 6% (no overlapping of misses). On a miss, the main memory is accessed and has a latency of 100 ns. The clock frequency is 500 MHz.

What is the CPI of *grinder* on the LC2k?

Stalls per cache miss = 100 ns / 2ns = 50 cycles (500 Mhz → 2ns cycle time)

CPI = 1 + data hazard stalls + control hazard stalls + icache stalls + dcache stalls

CPI = 1 + 0.15*0.50*1 + 0.20*0.40*3 + 1*0.03*50 + 0.35*0.06*50

Practice Problem 2: Memory Usage

□ Say you have the following:

- A program that generates 2 Billion loads and 1 Billion stores, each 4 bytes in size.
- A cache with a 32-byte block which gets a 95% hit rate on that program.

□ How many bytes of memory would be read and written if:

- We had a write-through cache with a no-write allocate policy?
- We had a write-back cache with a write-allocate policy? (Assume 25% of all misses result in a dirty eviction)
- We had no cache?

Practice Problem 2: Solution (1/3)

□ Let's start with the no-cache case.

- All stores go to memory and are 4 bytes each
 - Writes: $1 \text{ billion stores} * 4 \text{ bytes} = 4 \text{ billion bytes}$
- All loads go to memory and are 4 bytes each.
 - Reads: $2 \text{ billion loads} * 4 \text{ bytes} = 8 \text{ billion bytes}$

□ Write-through, no allocate.

- All stores still go to memory and are still 4 bytes each.
 - Writes: $1 \text{ billion stores} * 4 \text{ bytes} = 4 \text{ billion bytes}$
- Only loads that miss in the cache go to memory. But they read the full cache block.
 - Reads: $2 \text{ billion loads} * 0.05 * 32 \text{ bytes} = 3.2 \text{ billion bytes}$

Practice Problem 2: Solution (2/3)

- ❑ Write-back, write-allocate (data reads)
 - *Store* misses result in a cache block being read.
 - Reads: $1 \text{ billion stores} * 0.05 * 32 \text{ bytes} = 1.6 \text{ billion bytes}$
 - *Load* misses result in a cache block being read.
 - Reads: $2 \text{ billion loads} * 0.05 * 32 \text{ bytes} = 3.2 \text{ billion bytes}$
 - So that is 4.8 billion bytes of data read.

Practice Problem 2: Solution (3/3)

- ❑ When do we write to memory?
 - Only on dirty evictions. Can be done by both loads and stores.
 - Recall we are assuming 25% of all evictions are of dirty data.

- ❑ Write-back, write-allocate (data writes)
 - *Store* misses result in dirty eviction 1/4 of the time.
 - Writes: 1 billion stores * 0.05 * 32 bytes * (.25) = 0.4 billion bytes
 - *Load* misses result in a cache block being read.
 - Writes: 2 billion loads * 0.05 * 32 bytes * (.25) = 0.8 billion bytes

- ❑ Note: could have just taken the read numbers and multiplied by 25%.

Practice Problem 3: CPI w/ Caches 2

- ❑ Given a 200 MHz processor with 8KB instruction and data caches and a with memory access latency of 20 cycles. Both caches are 2-way associative. A program running on this processor has a 95% icache hit rate and a 90% dcache hit rate. On average, 30% of the instructions are loads or stores. The CPI of this system, if caches were ideal would be 1.
- ❑ Suppose you have 2 options for the next generation processor, which do you pick?
 - **Option 1:** Double the clock frequency—assume this will increase your memory latency to 40 cycles. Also assume a base CPI of 1 can still be achieved after this change.
 - **Option 2:** Double the size of your caches, this will increase the instruction cache hit rate to 98% and the data cache hit rate to 95%. Assume the hit latency is still 1 cycle.

Practice Problem 3: Solution

Option 1: (double clock freq, base cycle time is 5 ns, so new cycle time is 2.5 ns)

$$\text{CPI} = \text{baseCPI} + \text{IcacheStallCPI} + \text{DcacheStallCPI}$$

$$\text{CPI} = 1.0 + 0.05 * 40 + 0.3 * 0.1 * 40 = 4.2$$

$$\text{Execution time} = 4.2 * \text{Ninstrs} * 2.5\text{ns} = 10.5\text{ns} * \text{Ninstrs}$$

Option 2 (icache/dcachel miss rates lowered to 2% and 5%)

$$\text{CPI} = \text{baseCPI} + \text{IcacheStallCPI} + \text{DcacheStallCPI}$$

$$\text{CPI} = 1.0 + 0.02 * 20 + 0.3 * 0.05 * 20 = 1.7$$

$$\text{Execution time} = 1.7 * \text{Ninstrs} * 5\text{ns} = 8.5\text{ns} * \text{Ninstrs}$$

Therefore, Option 2 is the better choice

Practice Problem 4: Guess that cache!

Here is the series of address references (in hex) to a cache of size 512 bytes. You are asked to **determine the configuration of the cache**. Assume 12-bit addresses

0x310 – Miss

0x30f – Miss

0x510 – Miss

0x31f – Hit

0x72d – Miss

0x72f – Hit

0x320 – Miss

0x520 – Miss

0x720 - Miss


Block size: ?


Associativity: ?

Number of sets: ?


Practice Problem 4: Solution

Here is the series of address references (in hex) to a cache of size 512 bytes. You are asked to **determine the configuration of the cache**. Assume 12-bit addresses

0x310 – Miss 

0x30f – Miss 

0x510 – Miss

0x31f – Hit 

0x72d – Miss

0x72f – Hit

0x320 – Miss

0x520 – Miss

0x720 – Miss

Determine block size

First hit must be brought in by another miss

Take closest address: 0x310, so know block size must be at least 16 bytes so 0x31f brought in when 0x310 miss occurs

Now, is the block size larger? Know that 0x30f was a miss, thus 0x310 and 0x30f not in the same block.


Thus, block size must be ≤ 16 bytes

Thus Block Size = 16 bytes


Practice Problem 4 : Solution (2)


Here is the series of address references (in hex) to a cache of size 512 bytes. You are asked to **determine the configuration of the cache**. Assume 12-bit addresses

Determine associativity

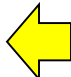
0x310 – Miss 


0x30f – Miss


0x510 – Miss 


0x31f – Hit 

0x72d – Miss

0x72f – Hit 

0x320 – Miss 

0x520 – Miss 

0x720 – Miss 

Assume direct mapped: 3-bit tag, 5-bit index, 4-bit offset.

If direct mapped 0x310 and 0x510 would both map to index 17, Thus 0x31f could not be a hit. So, not direct mapped.

Assume 2-way associative: 4-bit tag, 4-bit index, 4-bit offset

This fixes the green accesses, and allows 0x31f to be a hit.

What about > 2-way associative?

Now we also know that 0x720 is a miss even though 3 accesses earlier 0x72f was a hit, and thus it is in the cache. The intervening 2 accesses must kick it out, 0x320 and 0x520. Both go to set 2. If the associativity was > 2, then 0x720 would be a hit. So, must conclude that cache is 2-way associative.

Lastly, number of sets = $512 / (2 * 16) = 16$