

22. Virtual Memory: Basics

EECS 370 – Introduction to Computer Organization – Winter 2023

**EECS Department
University of Michigan in Ann Arbor, USA**

ANNOUNCEMENTS

- ❑ Project 4 due Thursday, April 13th
- ❑ Homework 6 due Monday, April 17th
- ❑ Just 3 more lectures of material and one review session!
 - ❑ Notice, we have no class on April 13th

Go check cache organization in your machine!

- ❑ For Linux you can use following command:

```
$:sudo dmidecode -t cache
```

DMI—desktop management interface
(you may need to install **dmidecode** on your machine)

Cache Organization Comparison/Review

Block size = 2 bytes, total cache size = 8 bytes for all caches


1. Fully associative (4-way associative)

V d tag data




2. Direct mapped

V d tag data



3. 2-way associative

V d tag data



Quick review questions


- ❑ If you have a 16KB cache with 32-byte cache lines that is four-way set-associative on a computer with 32-bit addresses:
 - How many sets do you have?
 - How many bits do you need for the Offset? Index? Tag?

- ❑ Describe why we have an Icache and a Dcache on nearly all computers rather than just one unified cache.

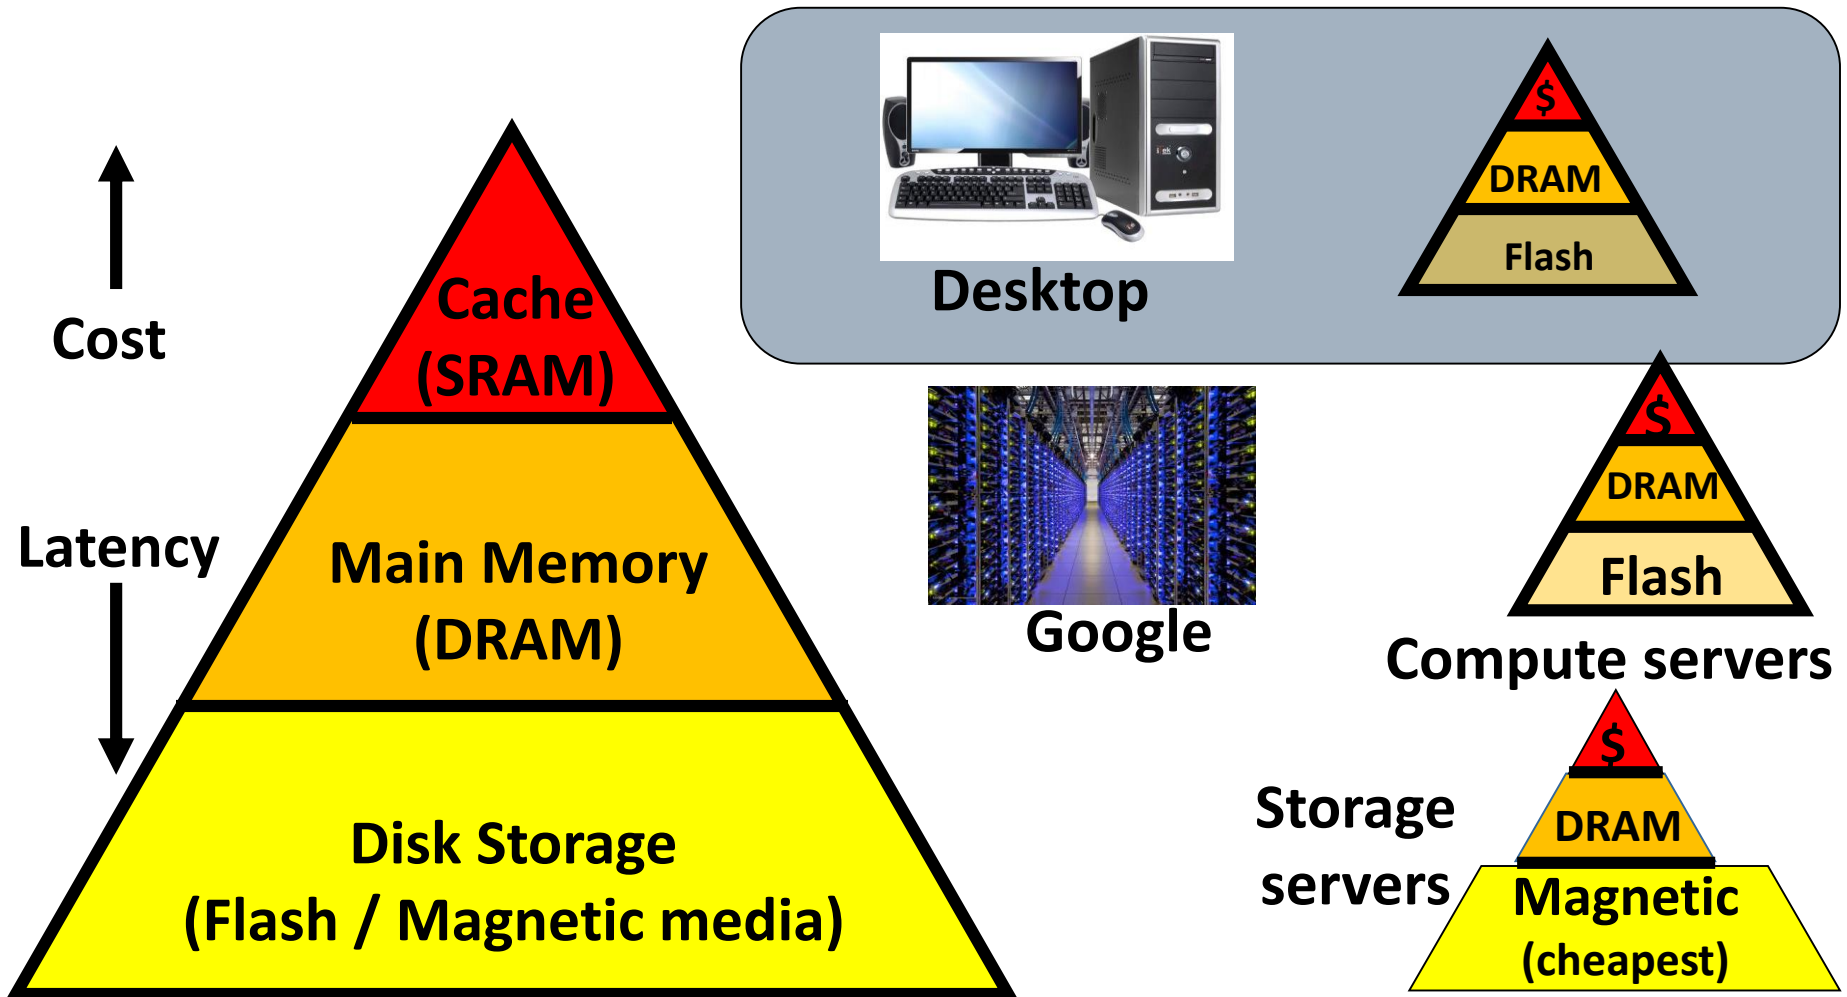
- ❑ Describe the primary advantage of write-back caches over write-through caches.
 - Give an example when that isn't the case (that is, where write-through caches do better than write-back caches on the advantage you identified).

VIRTUAL MEMORY

Storage Hierarchy



Memory and Storage Hierarchy



Memory: the issues(s)

- ❑ We run many programs on a same machine
 - Each of them may require GBs of storage
 - Unrelated programs must not have access to each other's storage

- ❑ DRAM is too expensive to buy 100s GB, but disk space is not...
 - We want our system to work even if it requires more DRAM than we bought.
 - We also don't want a program that works on a machine with 2048 MB of DRAM to stop working if we try to run it on a machine with only 512 MB of main memory.

- ❑ And, it would be nice to be able to enforce different policies on different portions of the memory (e.g.: read-only, etc)

Solution 1: User control


- ❑ Leave the problem to the programmer
 - Assume the programmer knows the exact configuration of the machine.
 - Programmer must either make sure the program fits in memory, or break the program up into pieces that do fit and load each other off the disk when necessary

- ❑ Not a bad solution in some domains
 - The hardware design is simple
 - (Original) PlayStation, engine control unit in a car, etc.
 - Systems with severe design constraints (e.g. RTOS)

Solution 2: Overlays

- ❑ A little automation to help the programmer
 - build the application in **overlays**
 - Two pieces of code/data may be overlaid iff
 - They are not active at the same time
 - They are placed in the same memory region
- ❑ Managing overlays is performed by the compiler
 - Good compilers may determine overlay regions
 - Compiler adds code to read the required overlay memory off the disk when necessary
- ❑ The hardware design is still simple (most of the time)

Overlay example



Solution 3: Virtual memory

- ❑ Build new hardware and software that automatically translates each memory reference from a

virtual address

(which the programmer sees as an array of bytes)


to a

physical address

(which the hardware uses to either index DRAM or identify where the storage resides on disk)

Basics of Virtual Memory

- ❑ Any time you see the word virtual in computer science and architecture it means “using a level of indirection”.
- ❑ Virtual memory hardware changes the virtual address the programmer sees into the physical one the memory chips see.

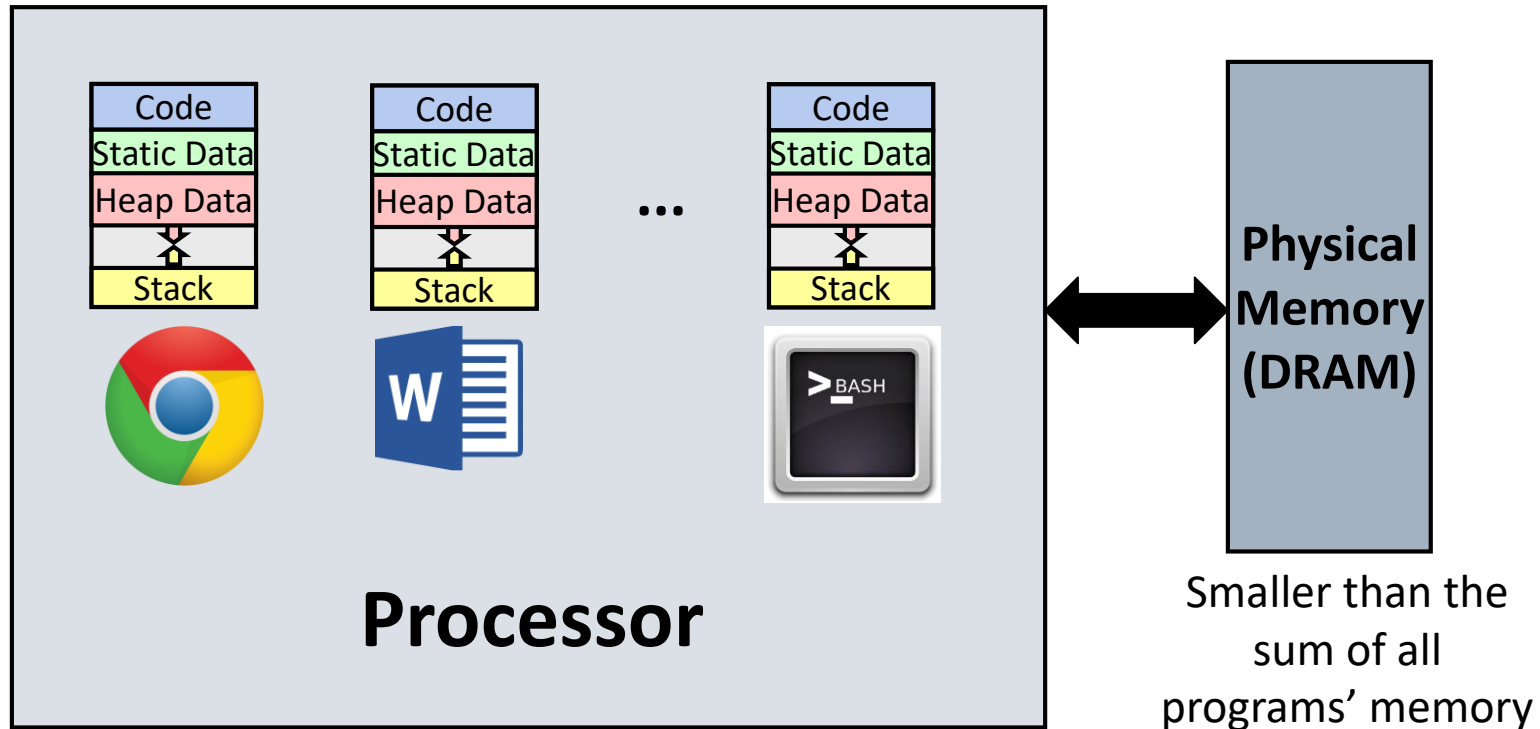


Why Virtual Memory?

- ❑ Virtual memory enables multiple programs to share the physical memory.

- ❑ Provides following 3 capabilities to the programs:
 1. **Transparency**
 - Don't need to know how other programs are using memory
 2. **Protection**
 - No program can modify the data of any other program
 3. **Programs not limited by DRAM capacity**
 - Each program can have more data than DRAM size

Revisit real system view—multitasking





Managing Virtual Memory—VM

- ❑ Managed by hardware logic *and* operating system software.
 - Hardware for speed
 - Software for flexibility and because disk storage is controlled by the operating system

- ❑ The hardware must be designed to support VM*

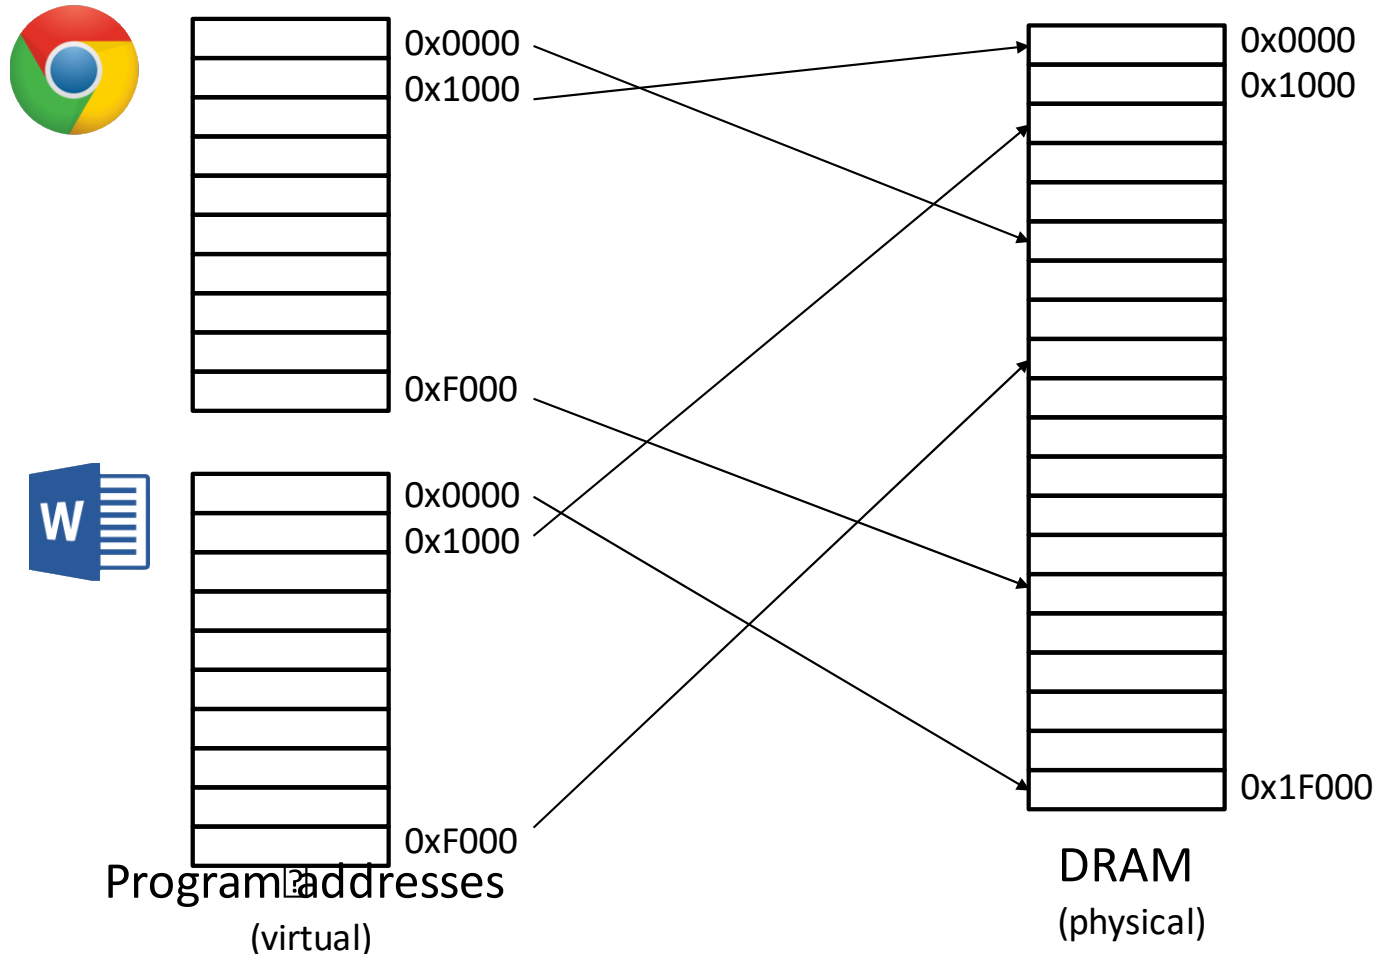
*(do not confuse VM for “Virtual Machine”, which is another concept entirely)

1. How to achieve transparency & protection?




Program addresses
(virtual)

1. How to achieve **transparency & protection**?



1. How to achieve transparency & protection?



Page Table

- ❑ Page tables are maintained by the operating system
- ❑ Each process has its own page table
- ❑ Contains address translation information i.e., virtual address → physical address
- ❑ Page tables themselves are kept in memory by OS, and OS knows the physical address of the page tables
 - No address translation is required by the OS for accessing the page tables

2. How to be **not limited by DRAM capacity?**


- ❑ Use disk as temporary space in case memory capacity is exhausted
 - This temporary space in disk is called **swap partition** in Linux-based systems
 - For fun check swap space in a linux system by:

\$: **top**

```
Tasks: 662 total,  1 running, 661 sleeping,  0 stopped,  0 zombie
%Cpu(s):  0.1 us,  0.0 sy,  0.0 ni, 99.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem: 32704372 total, 10813444 used, 21890928 free, 1018840 buffers
KiB Swap: 35162108 total,  89248 used, 35072860 free. 7053764 cached Mem
```

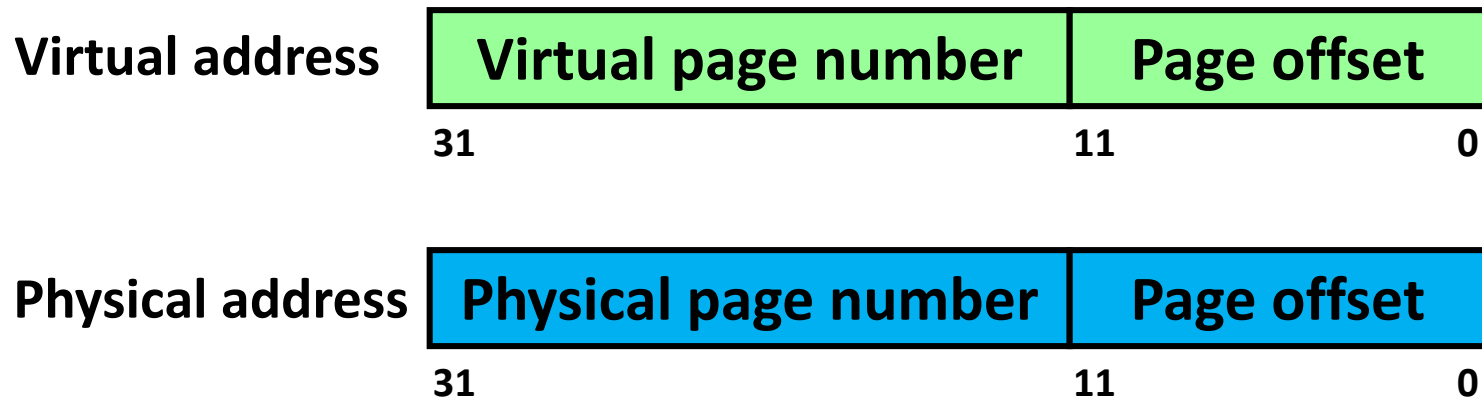
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
60256	nehaag	20	0	25356	3356	2444	R	6.0	0.0	0:00.02	top
1	root	20	0	38424	9040	2780	S	0.0	0.0	1:56.96	init
2	root	20	0	0	0	0	S	0.0	0.0	0:02.21	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	6:20.75	ksoftirqd/0

2. How to be **not limited by DRAM capacity?**



Virtual memory terminology

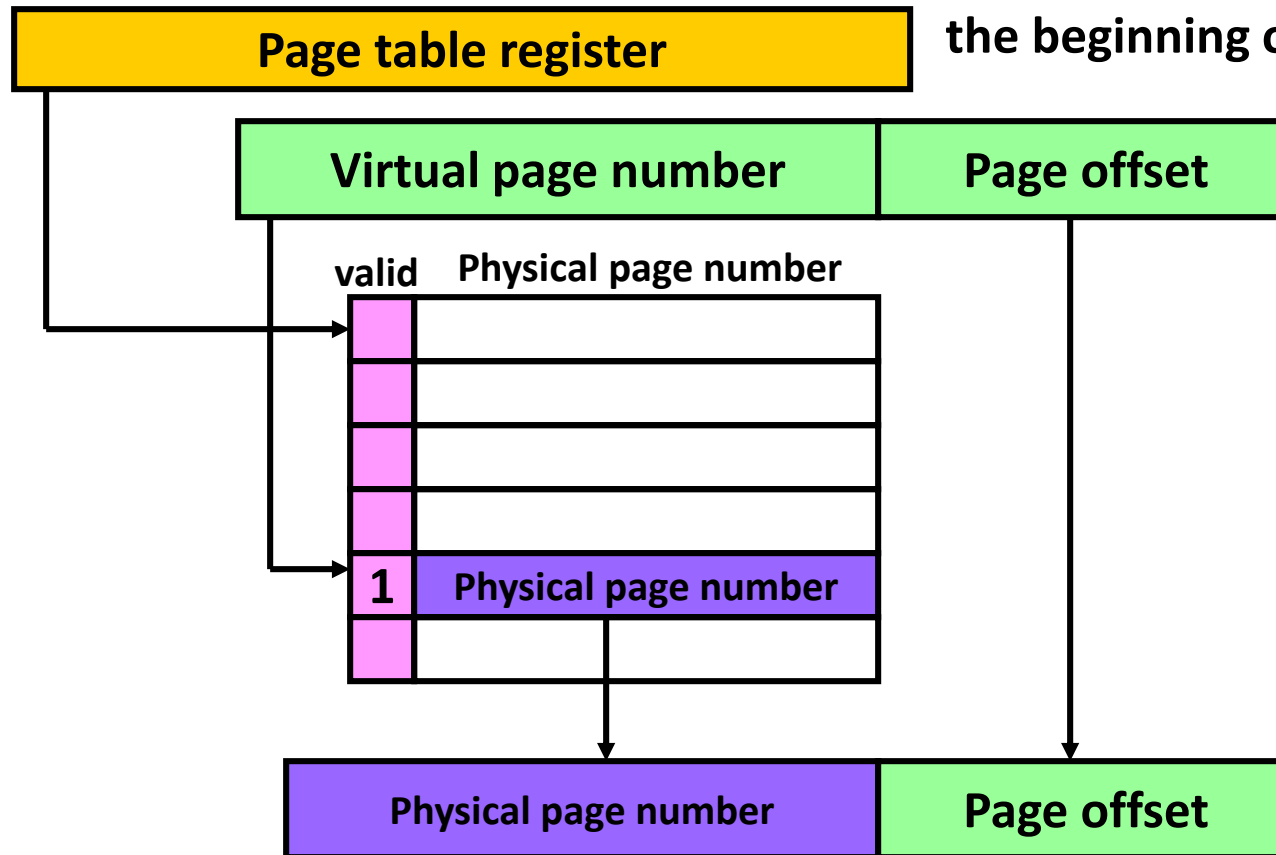
- ❑ Divide memory in chunks of **Pages** (e.g., 4KB for x86)
 - Size of physical page = size of virtual page
 - A virtual address consists of
 - A virtual page number
 - A page offset field (low order bits of the address)



- ❑ **Virtual Page** accesses that are not found in physical memory (DRAM) are called **Page Faults**

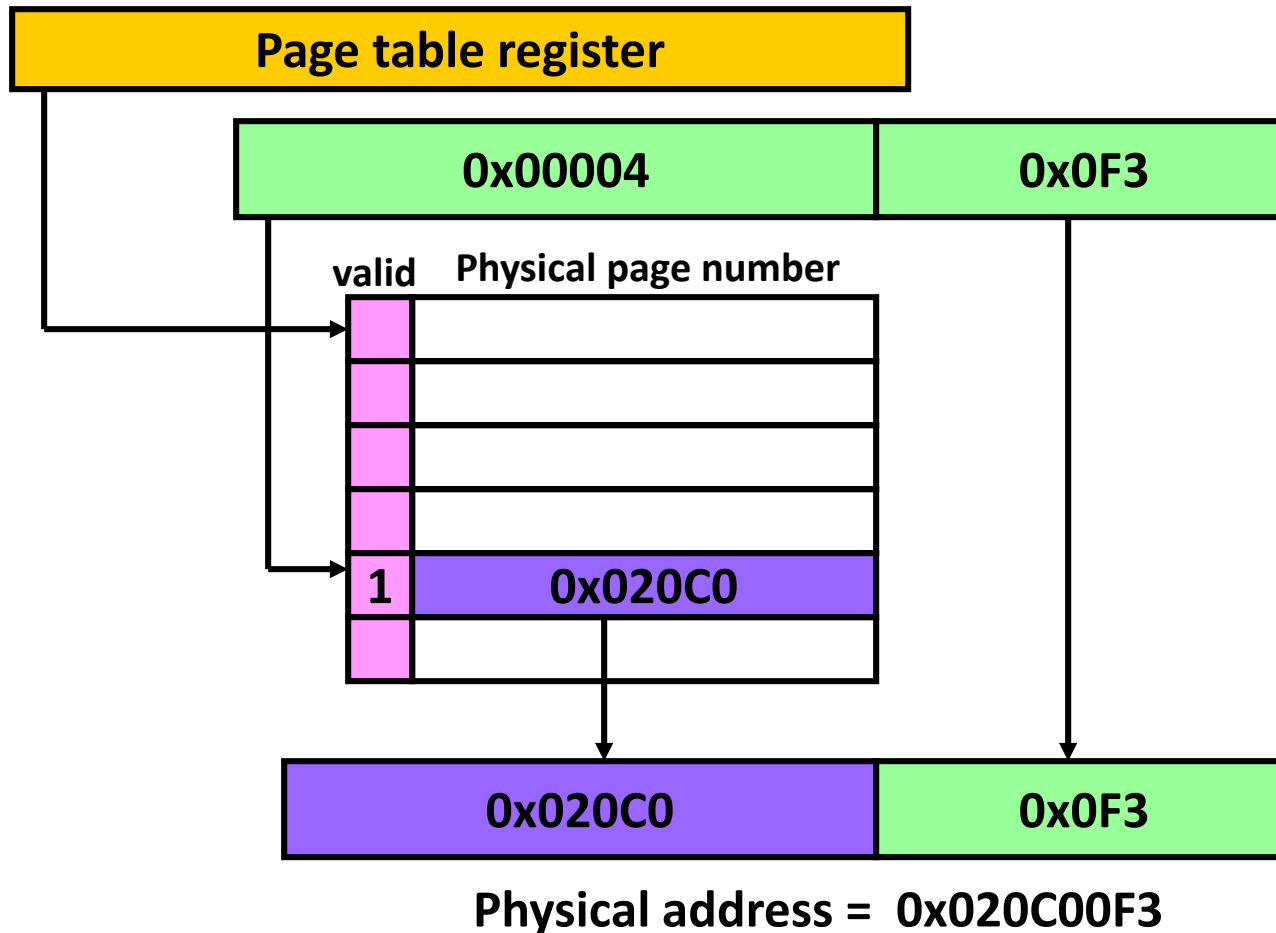
Page table components

The *Page table register* points to the beginning of the page table

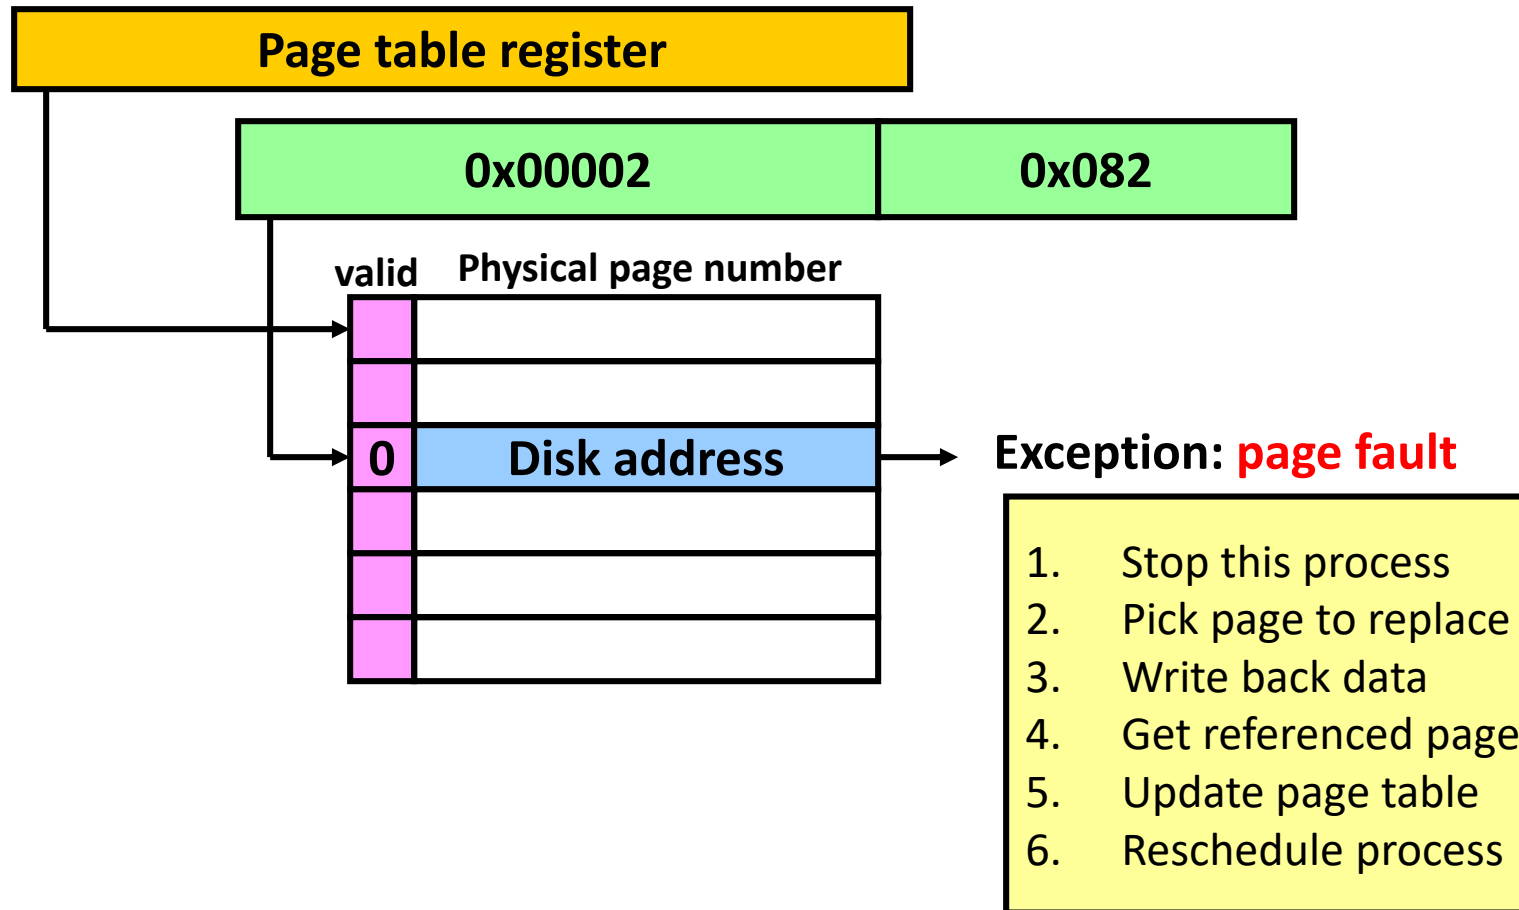


Page table components - Example

Virtual address = 0x000040F3



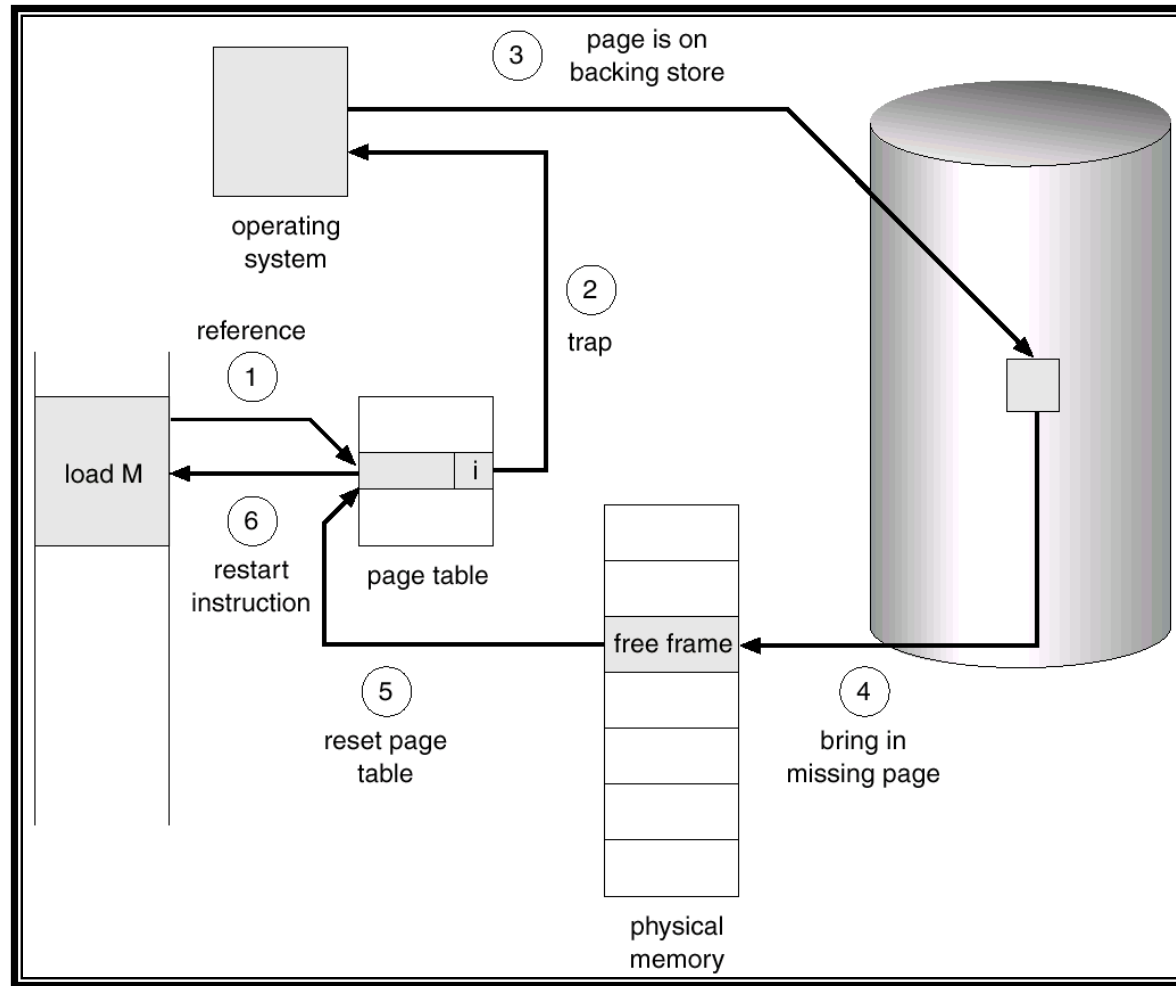
Page faults



How do we find it on disk?

- ❑ That is not a hardware problem! Go take EECS 482! 😊
- ❑ This is the operating system's job. Most operating systems partition the disk into logical devices (C: , D: , /home, etc.)
- ❑ They also have a hidden partition to support the disk portion of virtual memory
 - **Swap partition** on UNIX machines
 - You then index into the correct page in the swap partition.

Page faults



Class Problem

❑ Given the following:

- 4KB page size, physical memory of 16KB, page table stored in physical page 0 and can never be evicted, 20 bit, byte-addressable virtual address space.
- The page table initially has virtual page 0 in physical page 1, virtual page 1 in physical page 2 and no valid data in other physical pages.

❑ Fill in the table on the next slide for each reference

- Note: like caches we'll use LRU when we need to replace a page.

Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C			
0x01F0C			
0x20F0C			
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			

The page table initially has
virtual page 0 in physical
page 1, virtual page 1 in
physical page 2 and no valid
data in other physical pages.

4KB page size,
physical memory of 16KB,
page table stored in physical
page 0 and can never be
evicted, 20 bit, byte-
addressable virtual address
space.

Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C			
0x20F0C			
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			

The page table initially has
virtual page 0 in physical
page 1, virtual page 1 in
physical page 2 and no valid
data in other physical pages.

Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C			
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			

The page table initially has
virtual page 0 in physical
page 1, virtual page 1 in
physical page 2 and no valid
data in other physical pages.

Class Problem (continued)

The page table initially has
virtual page 0 in physical
page 1, virtual page 1 in
physical page 2 and no valid
data in other physical pages.

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C	0x20	Y (into 3)	0x3F0C
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			

Class Problem (continued)

The page table initially has
virtual page 0 in physical
page 1, virtual page 1 in
physical page 2 and no valid
data in other physical pages.

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C	0x20	Y (into 3)	0x3F0C
0x00100	0x0	N	0x1100
0x00200	0x0	N	0x1200
0x30000	0x30	Y (into 2)	0x2000
0x01FFF	0x1	Y (into 3)	0x3FFF
0x00200	0x0	N	0x1200