

EECS 370 - Lecture 23

Multi-Level Page Tables

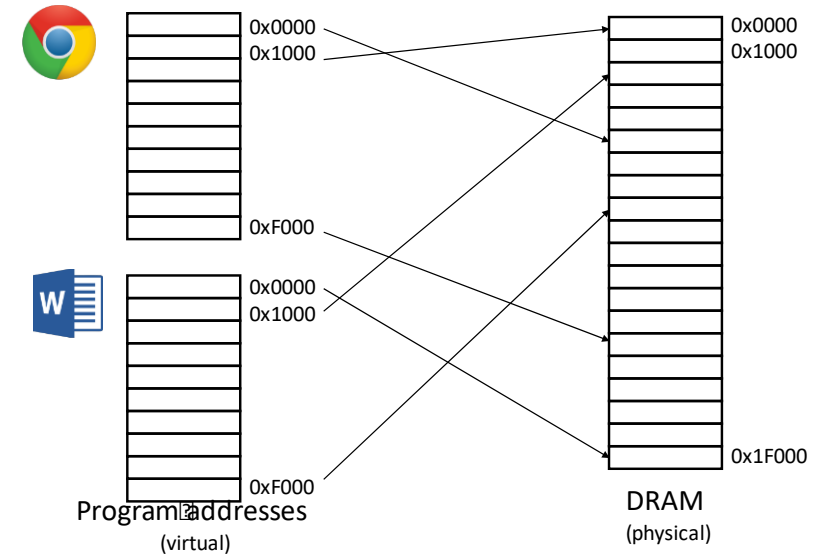
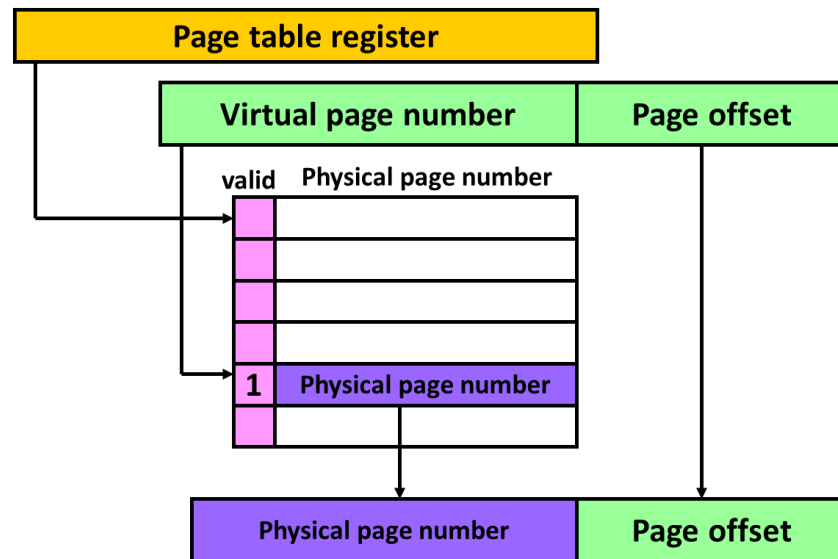


Announcements

- P4
 - Last project!
 - Due Thu (4/13)
- HW 6
 - Last homework!
 - Due Monday (4/17)
- Final exam
 - ...Last exam!
 - Thu (4/20) @ 10:30 am
 - Sample exams will be posted by early next week
 - (One sample exam given in homework)

Reminder

- We use a “page table” to translate “virtual” addresses (specified by the programmer / assembler) to “physical” addresses (sent to memory)
 - Allows multiple programs to run without interfering with each other
 - Allows programs to use more storage than available in DRAM
 - By mapping virtual addresses to disk instead of DRAM



Class Problem (continued)

4KB page size,
physical memory of 16KB,
page table stored in physical
page 0 and can never be
evicted, 20 bit, byte-
addressable virtual address
space.

The page table initially has
virtual page 0 in physical
page 1, virtual page 1 in
physical page 2 and no valid
data in other physical pages.

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C	0x20	Y (into 3)	0x3F0C
0x00100	0x0	N	0x1100
0x00200	0x0	N	0x1200
0x30000	0x30	Y (into 2)	0x2000
0x01FFF	0x1	Y (into 3)	0x3FFF
0x00200	0x0	N	0x1200

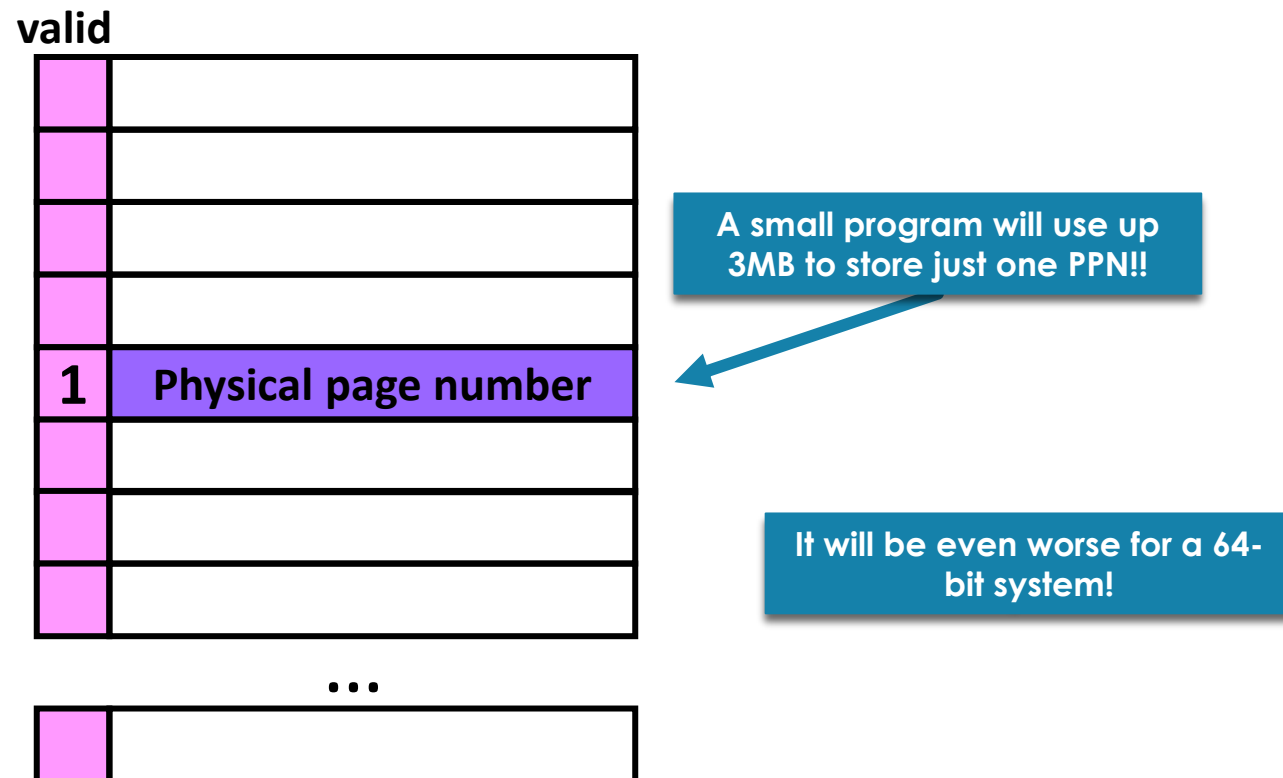


Size of the page table

- How big is a page table entry?
 - For 32-bit virtual address:
 - If the machine can support $1\text{GB} = 2^{30}$ bytes of physical memory and we use pages of size $4\text{KB} = 2^{12}$,
 - then the physical page number is $30-12 = 18$ bits.
Plus another **valid bit** + other useful stuff (**read only, dirty, etc.**)
 - Let say about 3 bytes.
- How many entries in the page table?
 - 1 entry per virtual page
 - ARM virtual address is 32 bits – 12 bit page offset = 20
 - Total number of virtual pages = 2^{20}
- Total size of page table = Number of virtual pages
 - * Size of each page table entry
 - = $2^{20} \times 3$ bytes ~ 3 MB

How can you organize the page table?

1. Single-level page table occupies continuous region in physical memory
 - Previous example always takes 3MB regardless of how much virtual memory is used



How can you organize the page table?

2. Option 2: Use a multi-level page table

- 1st level page table (much smaller!) holds addresses 2nd level page tables
 - 2nd level page tables hold translation info, or 3rd level page tables if we wanna go deeper
 - Only allocate space for 2nd level page tables that are used

valid	PPN
1	0x1
1	0x2
1	0x3

Single-level: Tons of wasted space!

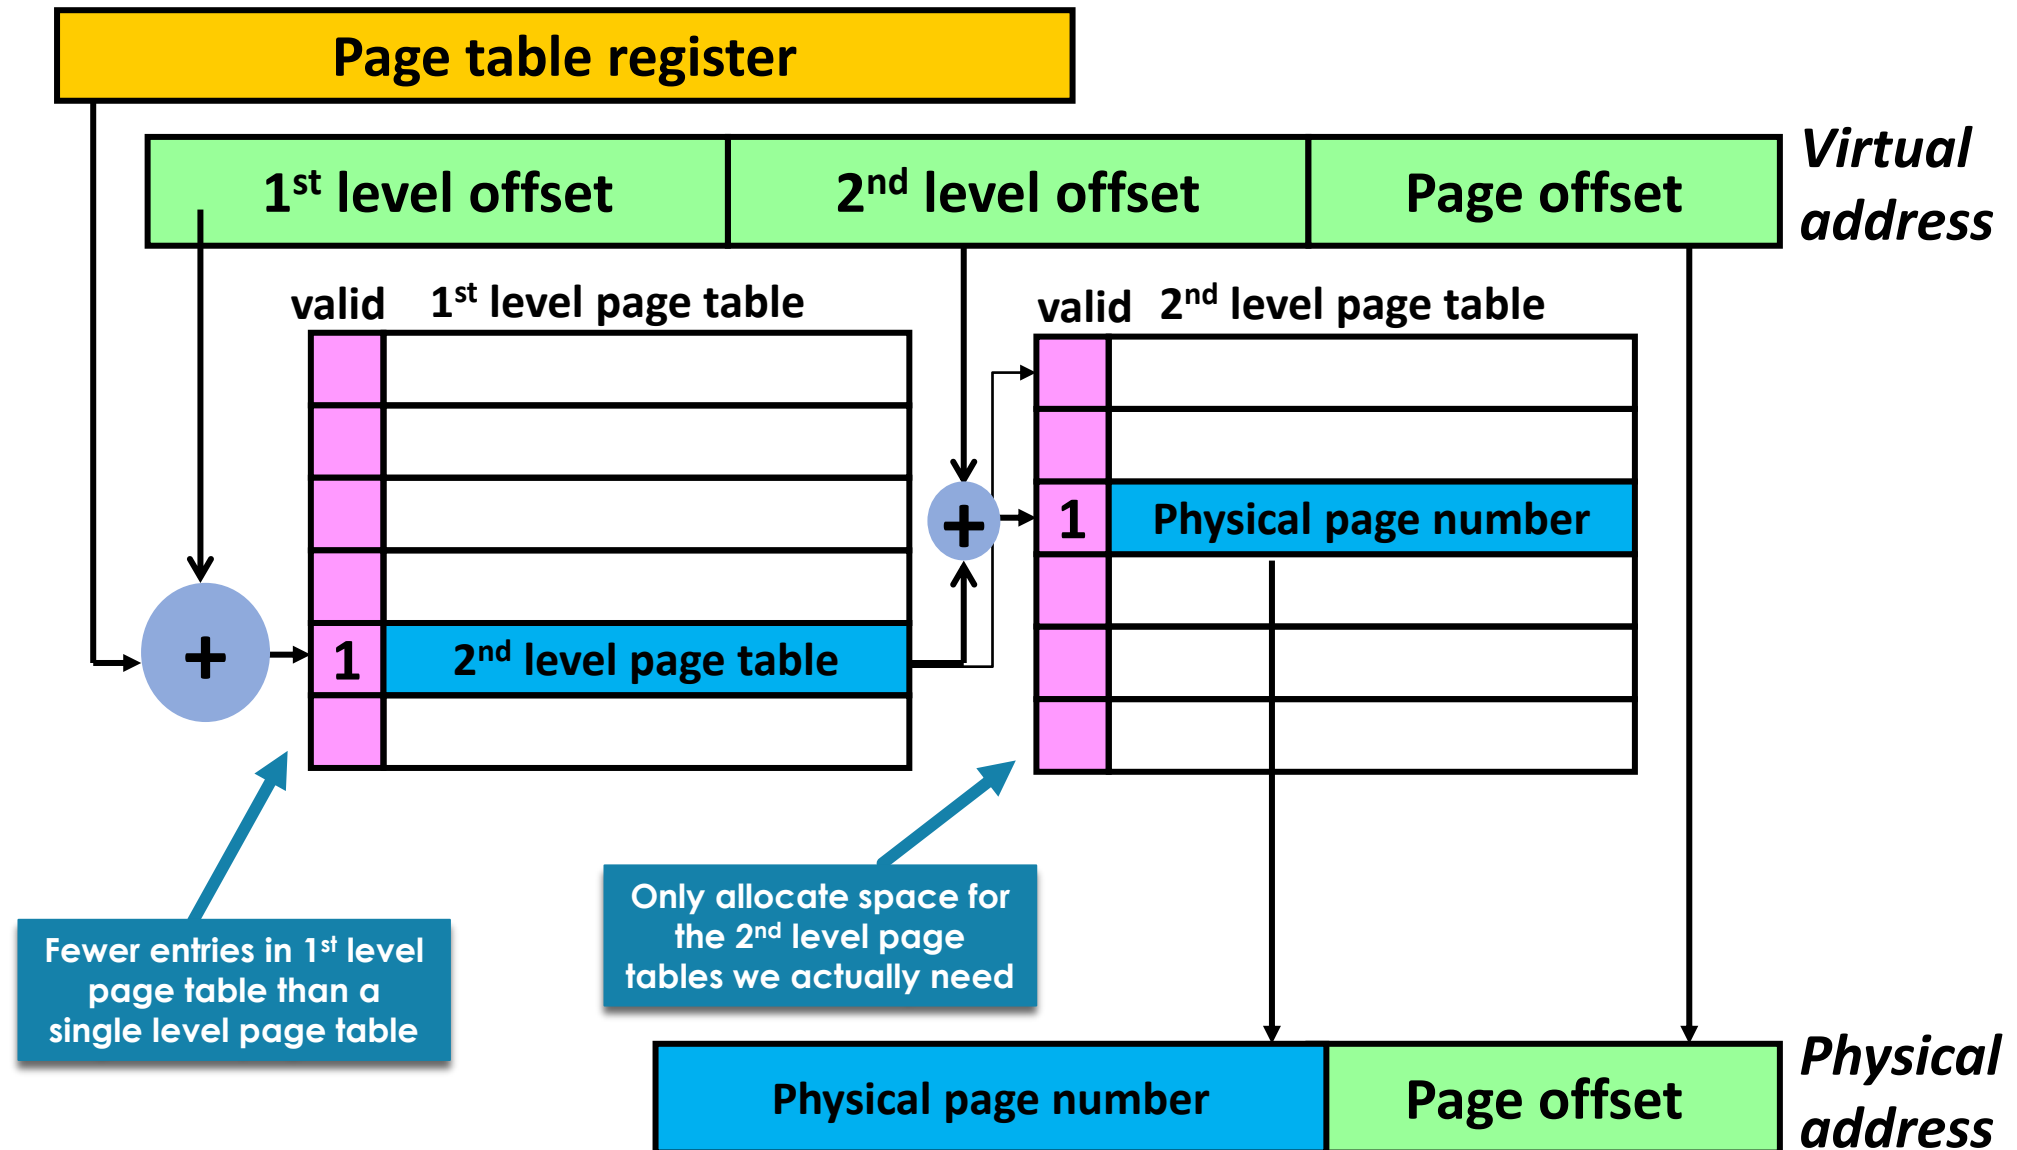
valid	2 nd level page table
1	0x1000

valid	PPN
1	0x1
1	0x2
1	0x3


Multi-level: Size is proportional to amount of memory used

Common case: most programs use small portion of virtual memory space

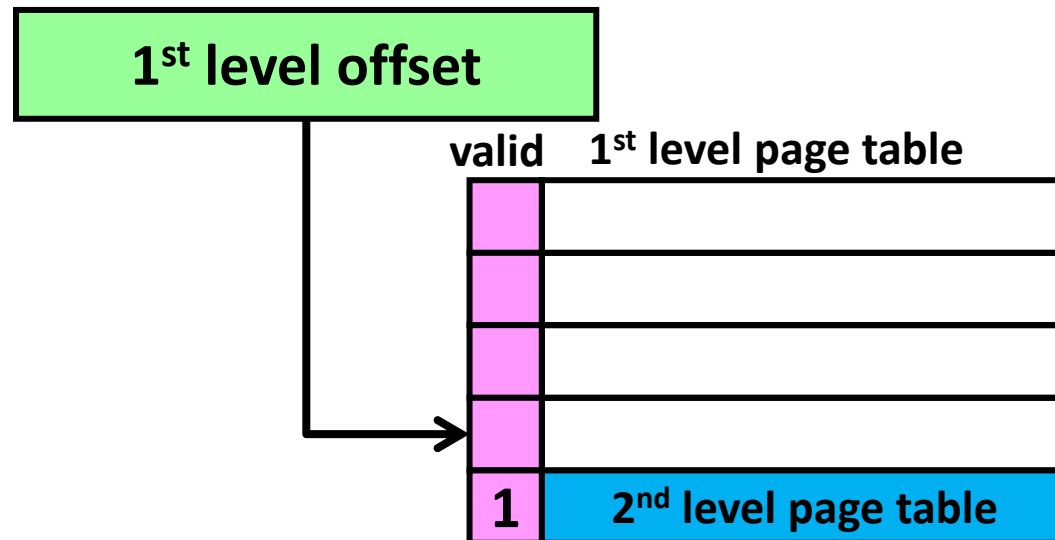
Hierarchical page table




Hierarchical page table – 32bit Intel x86



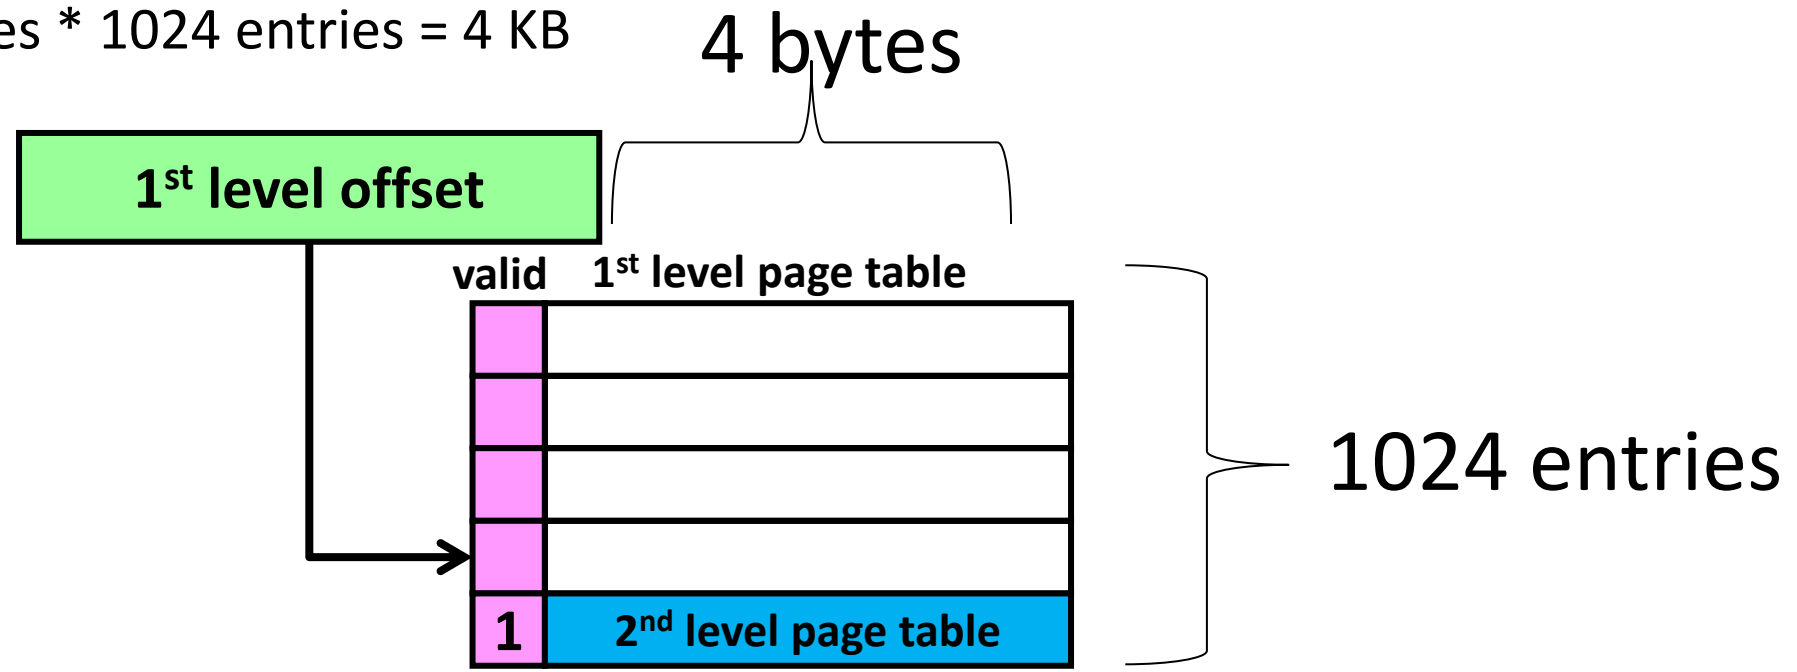
- How many bits in the virtual 1st level offset field? 10
- How many bits in the virtual 2nd level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1st level page table? $2^{10}=1024$



Hierarchical page table – 32bit Intel x86

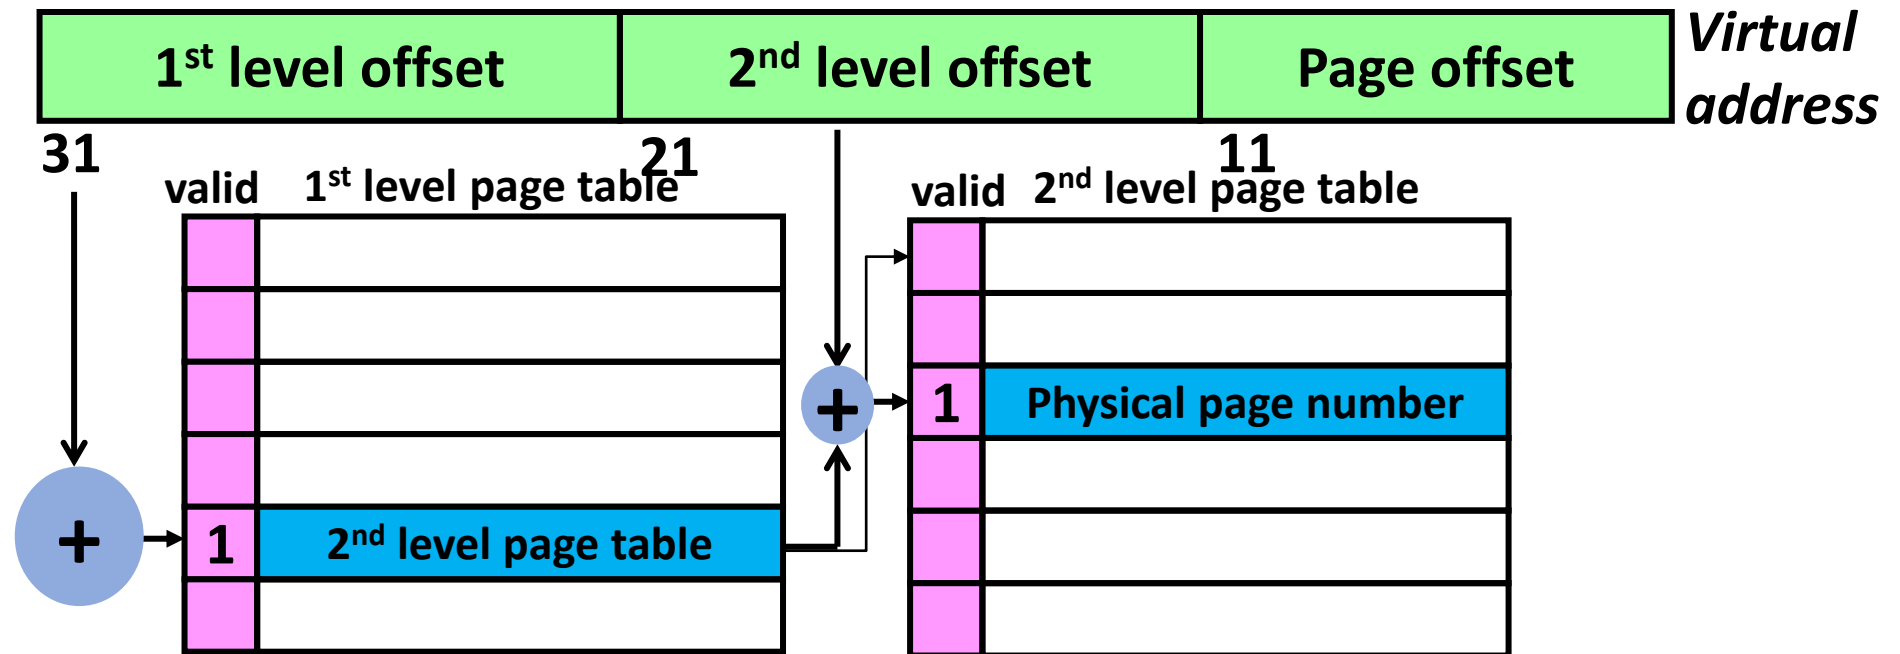


- How many bytes for each entry in the 1st level page table?
 - 4 bytes (address of 2nd level page)
- Total size of 1st level page table
 - 4 bytes * 1024 entries = 4 KB




Hierarchical page table

- How many entries in the 2nd level of the page table?
 - $2^{10} = 1024$
- How many bytes for each VPN in a 2nd level table?
 - Let's round up to 4 bytes




Hierarchical page table – 32bit Intel x86



- How many bits in the virtual 1st level offset field?
- How many bits in the virtual 2nd level offset field?
- How many bits in the page offset?
- How many entries in the 1st level page table?
- How many bytes for each entry in the 1st level page table?
- How many entries in the 2nd level of the page table?
- How many bytes for each entry in a 2nd level table?
- **What is the total size of the page table?**

(here n is number of valid entries in the 1st level page table)

Hierarchical page table – 32bit Intel x86



- How many bits in the virtual 1st level offset field? 10
- How many bits in the virtual 2nd level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1st level page table? $2^{10}=1024$
- How many bytes for each entry in the 1st level page table? 4
- How many entries in the 2nd level of the page table? $2^{10}=1024$
- How many bytes for each entry in a 2nd level table? ~ 4
- **What is the total size of the page table?**
(here n is number of valid entries in the 1st level page table) **$4K+n*4K$**

Class Problem (32 bit x86)

- What is the least amount of memory that could be used? When would this happen?
- What is the most memory that could be used? When would this happen?
- How much memory is used for this memory access pattern:
0x00000ABC
0x00000ABD
0x10000ABC
0x20000ABC
- How much memory if we used a single-level page table with 4KB pages? Assume entries are rounded to the nearest word (4B)

Class Problem (32 bit x86)

- What is the least amount of memory that could be used? When would this happen?
 - 4KB for 1st level page table. Occurs when no memory has been accessed (before program runs)
- What is the most memory that could be used? When would this happen?
 - 4KB for 1st level page table
+ 1024*4KB for all possible 2nd level page tables
= 4100KB (which slightly greater than 4096KB)
 - Occurs when program uses all virtual pages ($= 2^{20}$ pages)

Class Problem (32 bit x86)

- How much memory is used for this memory access pattern:

0x00000ABC // Page fault

0x00000ABD

0x10000ABC // Page fault

0x20000ABC // Page fault

- 4KB for 1st level page table + 3*4KB for each 2L page table
= 16 KB

Class Problem (32 bit x86)

- How much memory if we used a single-level page table with 4KB pages? Assume entries are rounded to the nearest word (4B)

Class Problem (32 bit x86)

- How much memory if we used a single-level page table with 4KB pages? Assume entries are rounded to the nearest word (4B)
 - 32 bits – 12 bit page offset = 20 bits of address
 - 2^{20} (num entries) * 4B
=> 4MB

Class Problem – Multi-level VM

- Design a two-level virtual memory system of a byte addressable processor with **24-bit long addresses**. No cache in the system. **256Kbytes of memory** installed, and no additional memory can be added.
 - **Virtual memory page: 512 Bytes**. Each page table entry must be an integer number of bytes, and must be the smallest size required to fit the physical page number + 1 bit to mark valid-entry
 - We want each second-level page table to fit exactly in one memory page, and 1st level page table entries are 3 bytes each (a memory address pointer to a 2L page table).
- Compute:
 - Number of entries in each 2nd level page table;
 - Number of virtual address bits used to index the 2nd level page table;
 - Number of virtual address bits used to index the 1st level page table;
 - Size of the 1st level page table.

Class Problem – Multi-level VM

Class Problem – Multi-level VM

Page Offset: 9 bits (512B page size)

Physical address = 18b (256KB Mem size)

Physical page number = 18b (256KB mem size) – 9b (offset)

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

2nd level page table entry size: 9b (physical page number) + 1b = ~ 2 bytes

2nd level page table **fits exactly in 1 page**

#entries in 2nd level page table is 512 bytes / 2 bytes = 256

#entries in 2nd level page table = 256 → Virtual page bits = 8b

Virtual 1st level page bits = 24 – 8 – 9 = 7b

1st level page table size = $2^7 * 3 \text{ bytes} = 384\text{B}$

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b

Page Replacement Strategies

- Page table indirection enables a fully associative mapping between virtual and physical pages.
- How do we implement LRU in OS?
 - True LRU is expensive, but LRU is a heuristic anyway, so approximating LRU is fine
 - Keep a “***accessed***” ***bit per page***, cleared occasionally by the operating system. Then pick any “unaccessed” page to evict

Other VM Translation Functions

- Page data location
 - Physical memory, disk, uninitialized data
- Access permissions
 - Read only pages for instructions
 - **This is how your system detects segmentation faults**
- Gathering access information
 - Identifying dirty pages by tracking stores
 - Identifying accesses to help determine LRU candidate

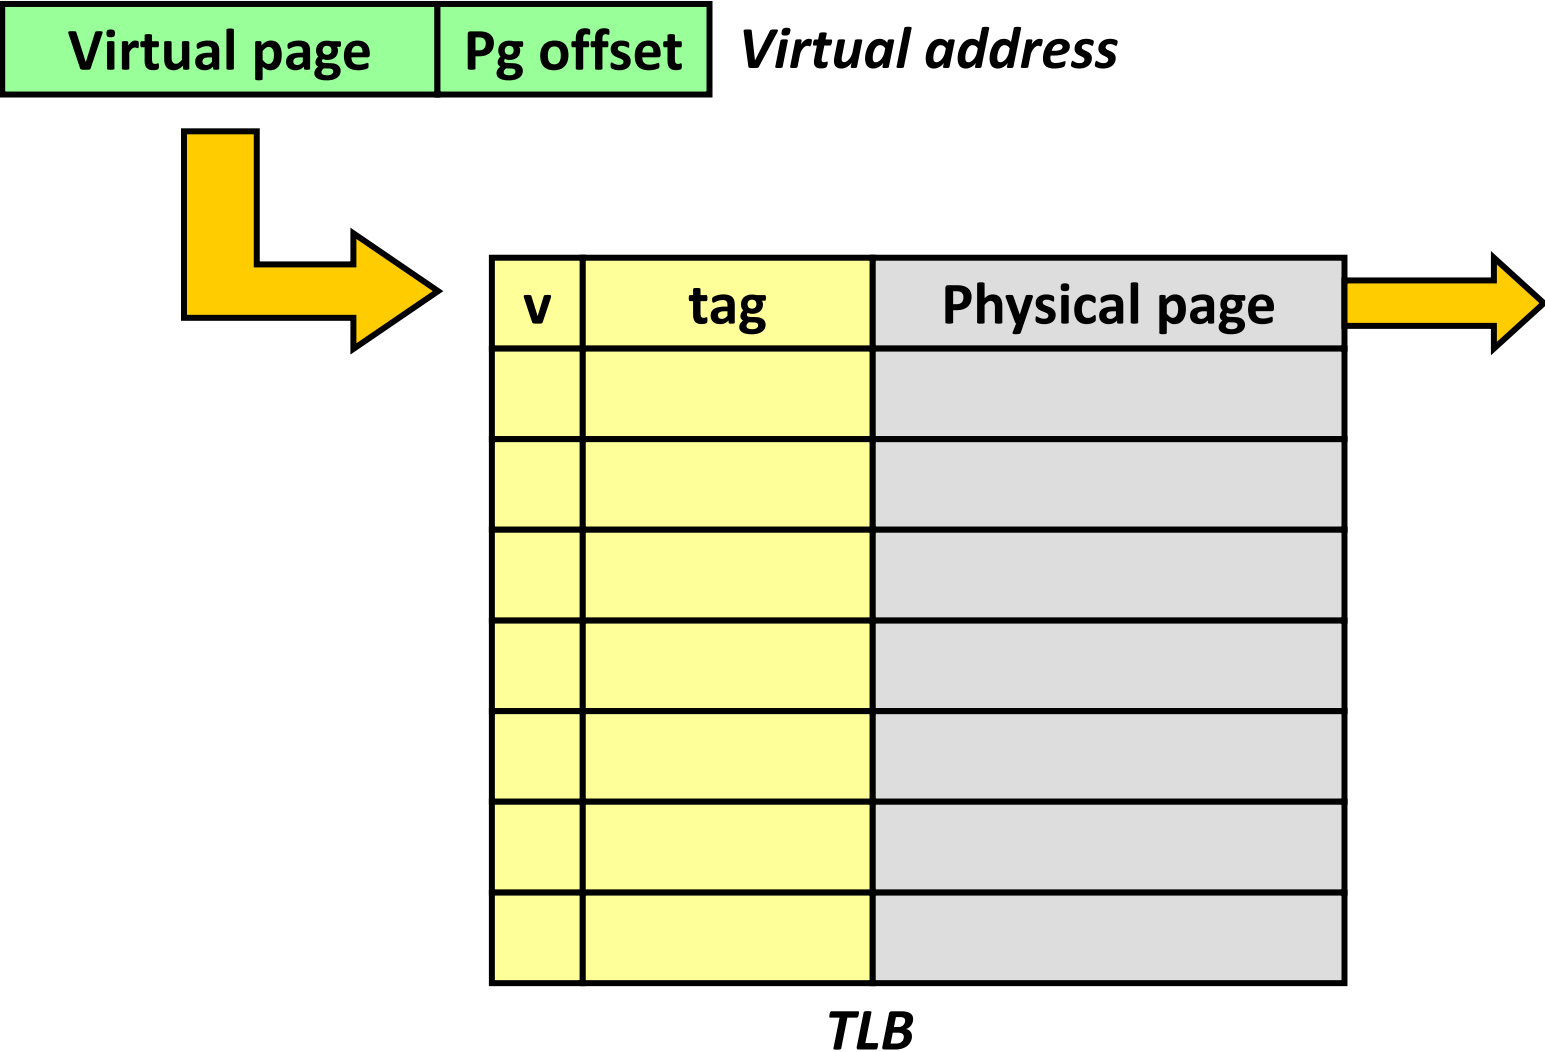
Performance of Virtual Memory

- To translate a virtual address into a physical address, we must first access the page table in physical memory
- Then we access physical memory again to get the data
 - A load instruction performs at least 2 memory reads
 - A store instruction performs at least 1 read and then a write
- Above lookup steps are **SLOW**

Translation look-aside buffer (TLB)

- We fix this performance problem by avoiding main memory in the translation from virtual to physical pages.
- Buffer common translations in a **Translation Look-aside Buffer (TLB)**, a fast cache memory dedicated to storing a small subset of valid V-to-P translations.
- 16-512 entries common.
- Generally has low miss rate ($< 1\%$).

Translation look-aside buffer (TLB)



Where is the TLB lookup?

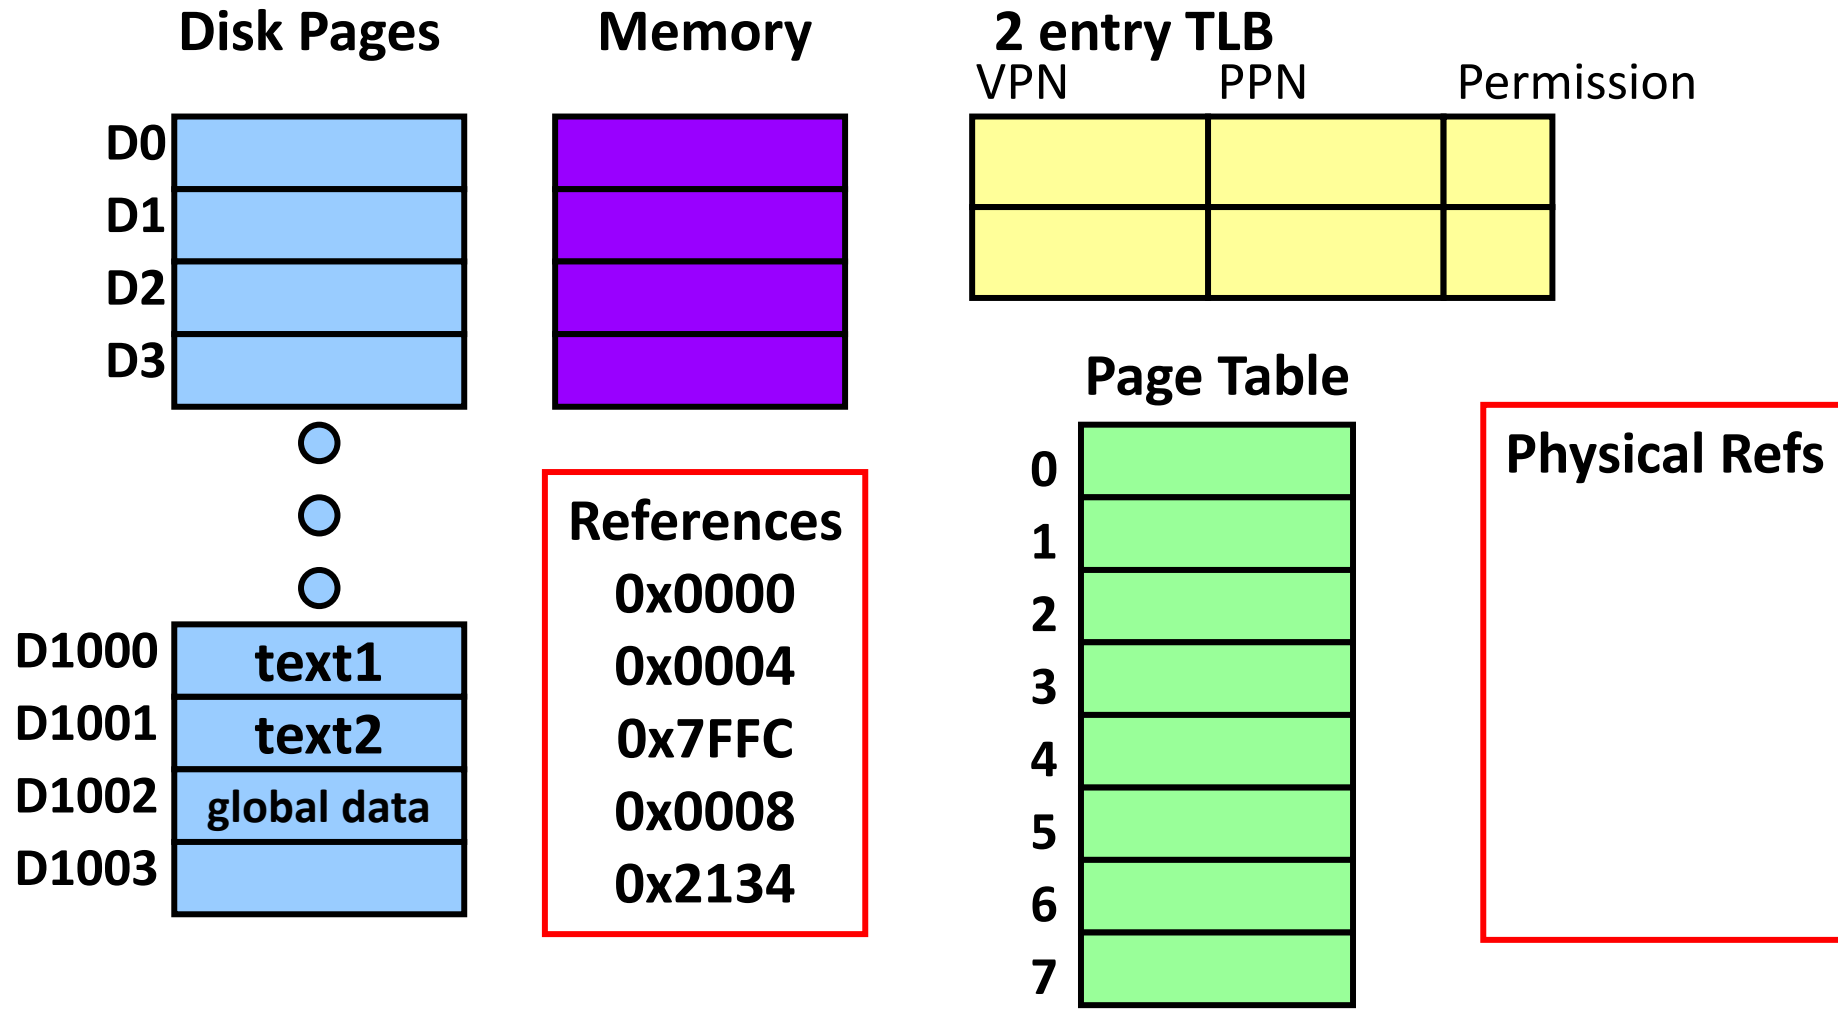
- We put the TLB lookup in the pipeline after the virtual address is calculated and before the memory reference is performed
 - This may be before or during the data cache access.
 - To be discussed next time
 - In case of a TLB miss, we need to perform the virtual to physical address translation during the memory stage of the pipeline.

Putting it all together

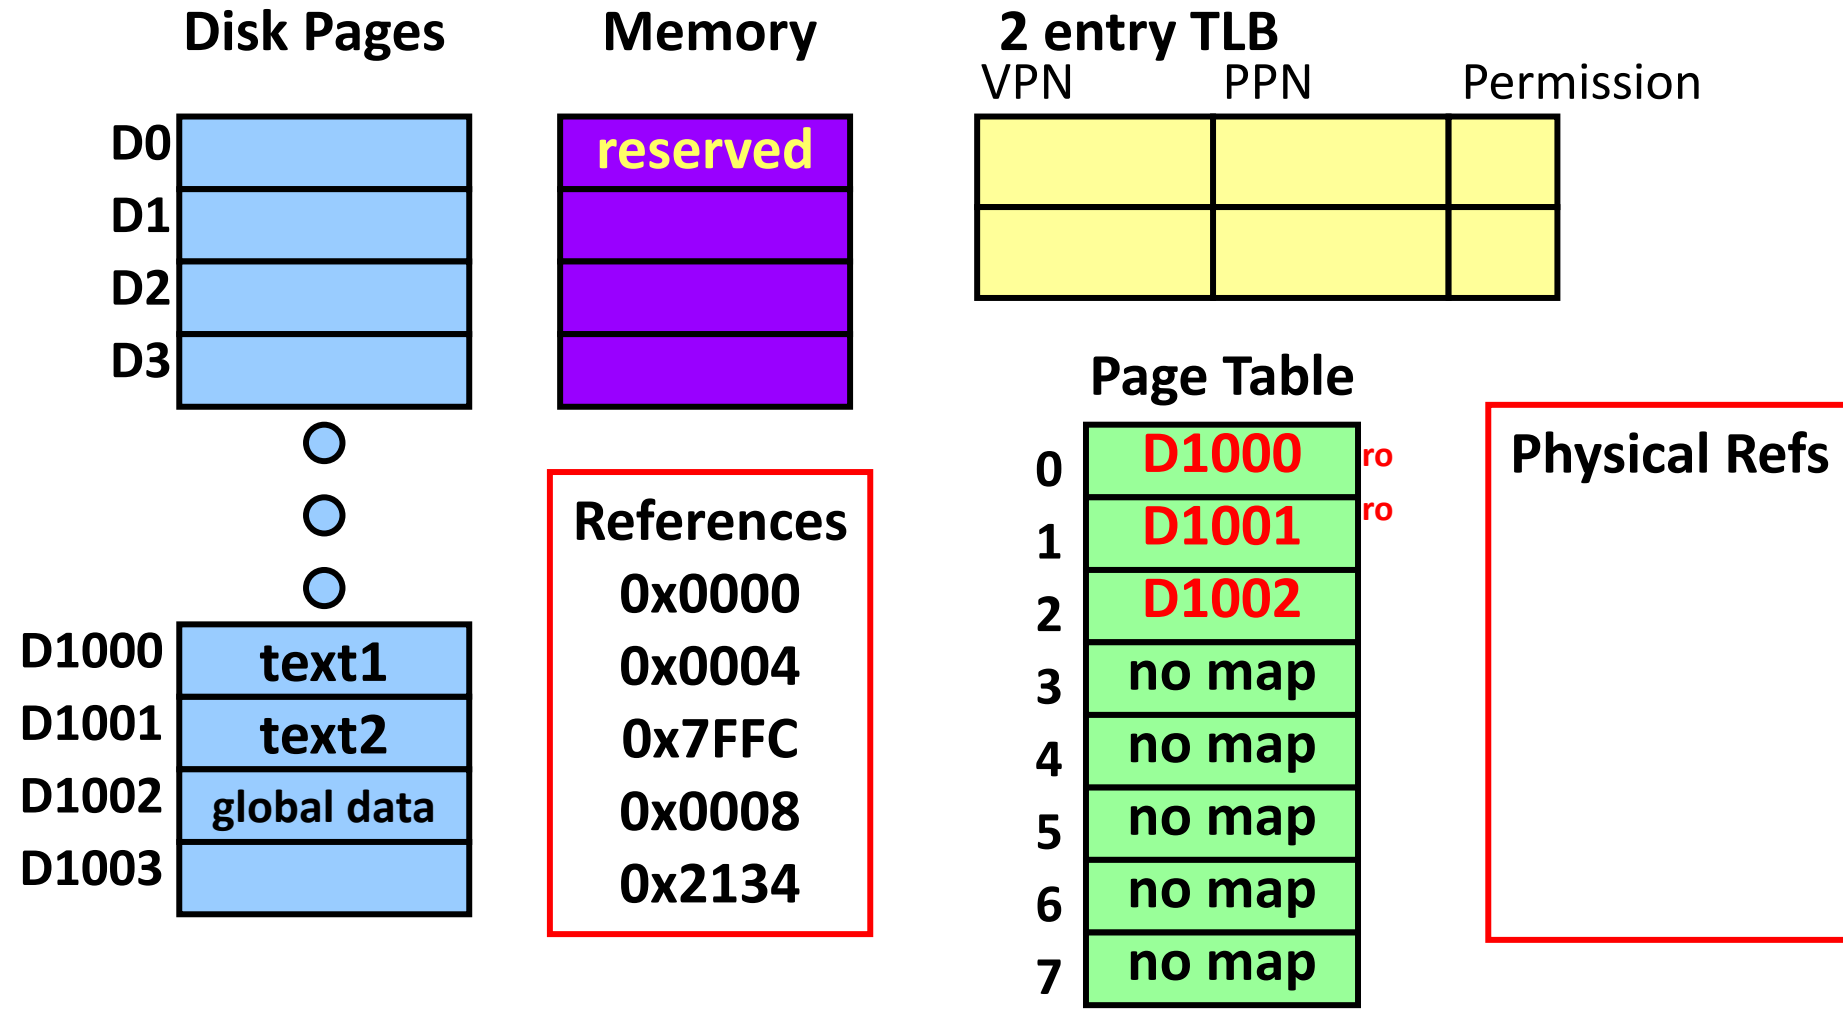
- Loading your program in memory
 - Ask operating system to create a new process
 - Construct a page table for this process
 - Mark all page table entries as invalid with a pointer to the disk image of the program
 - That is, point to the executable file containing the binary.
 - Run the program and get an immediate page fault on the first instruction.

Loading a program into memory

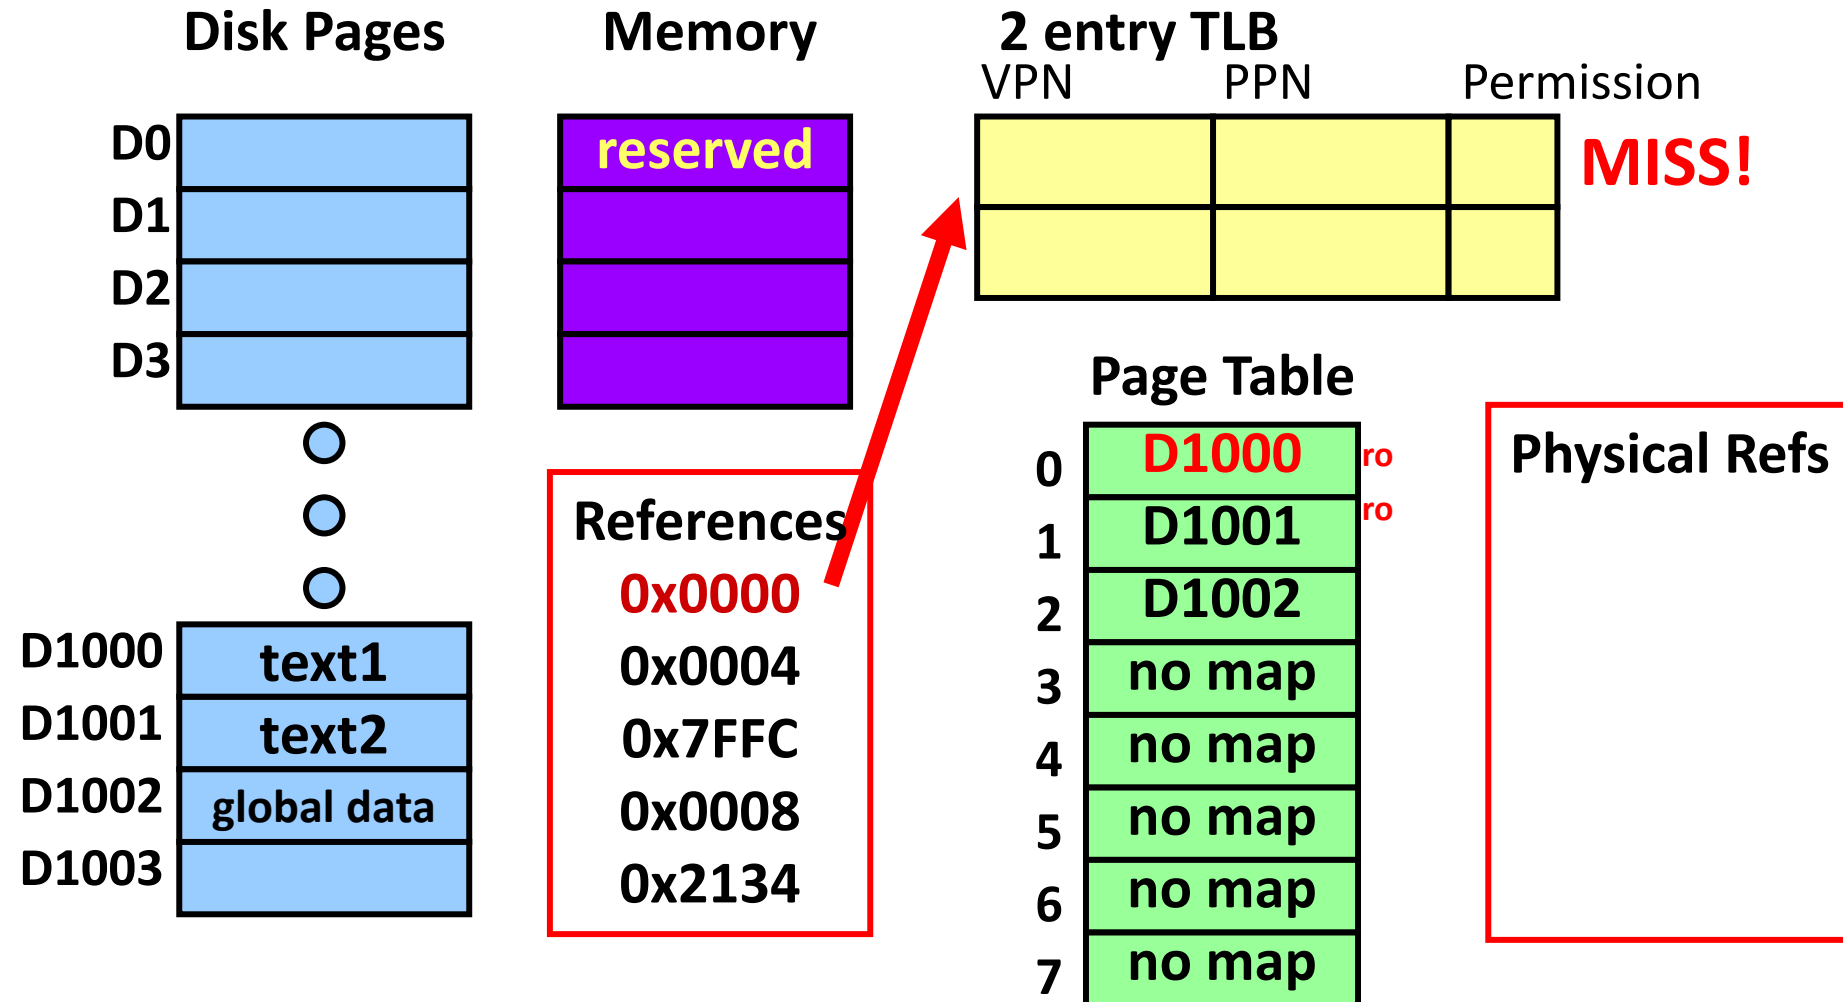
- ❑ Page size = 4 KB, Page table entry size = 4 B
- ❑ Page table register points to physical address 0x0000



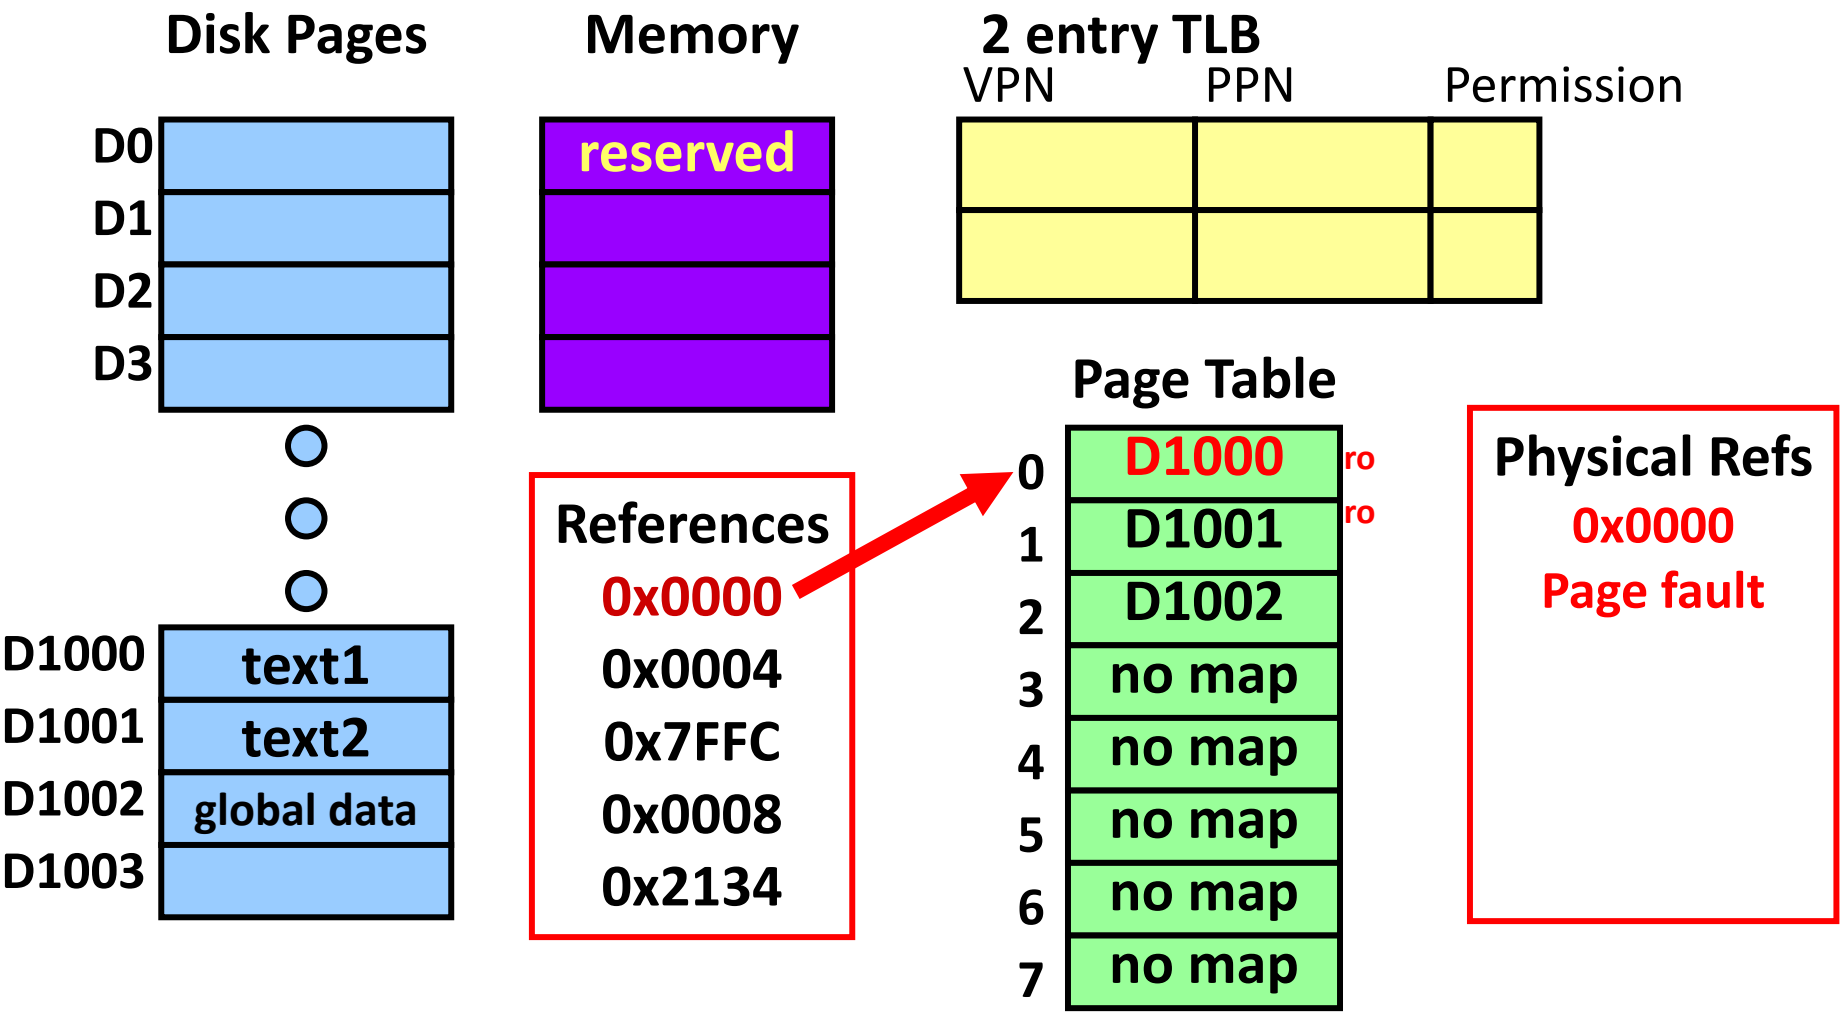
Step 1: Read executable header & initialize page table



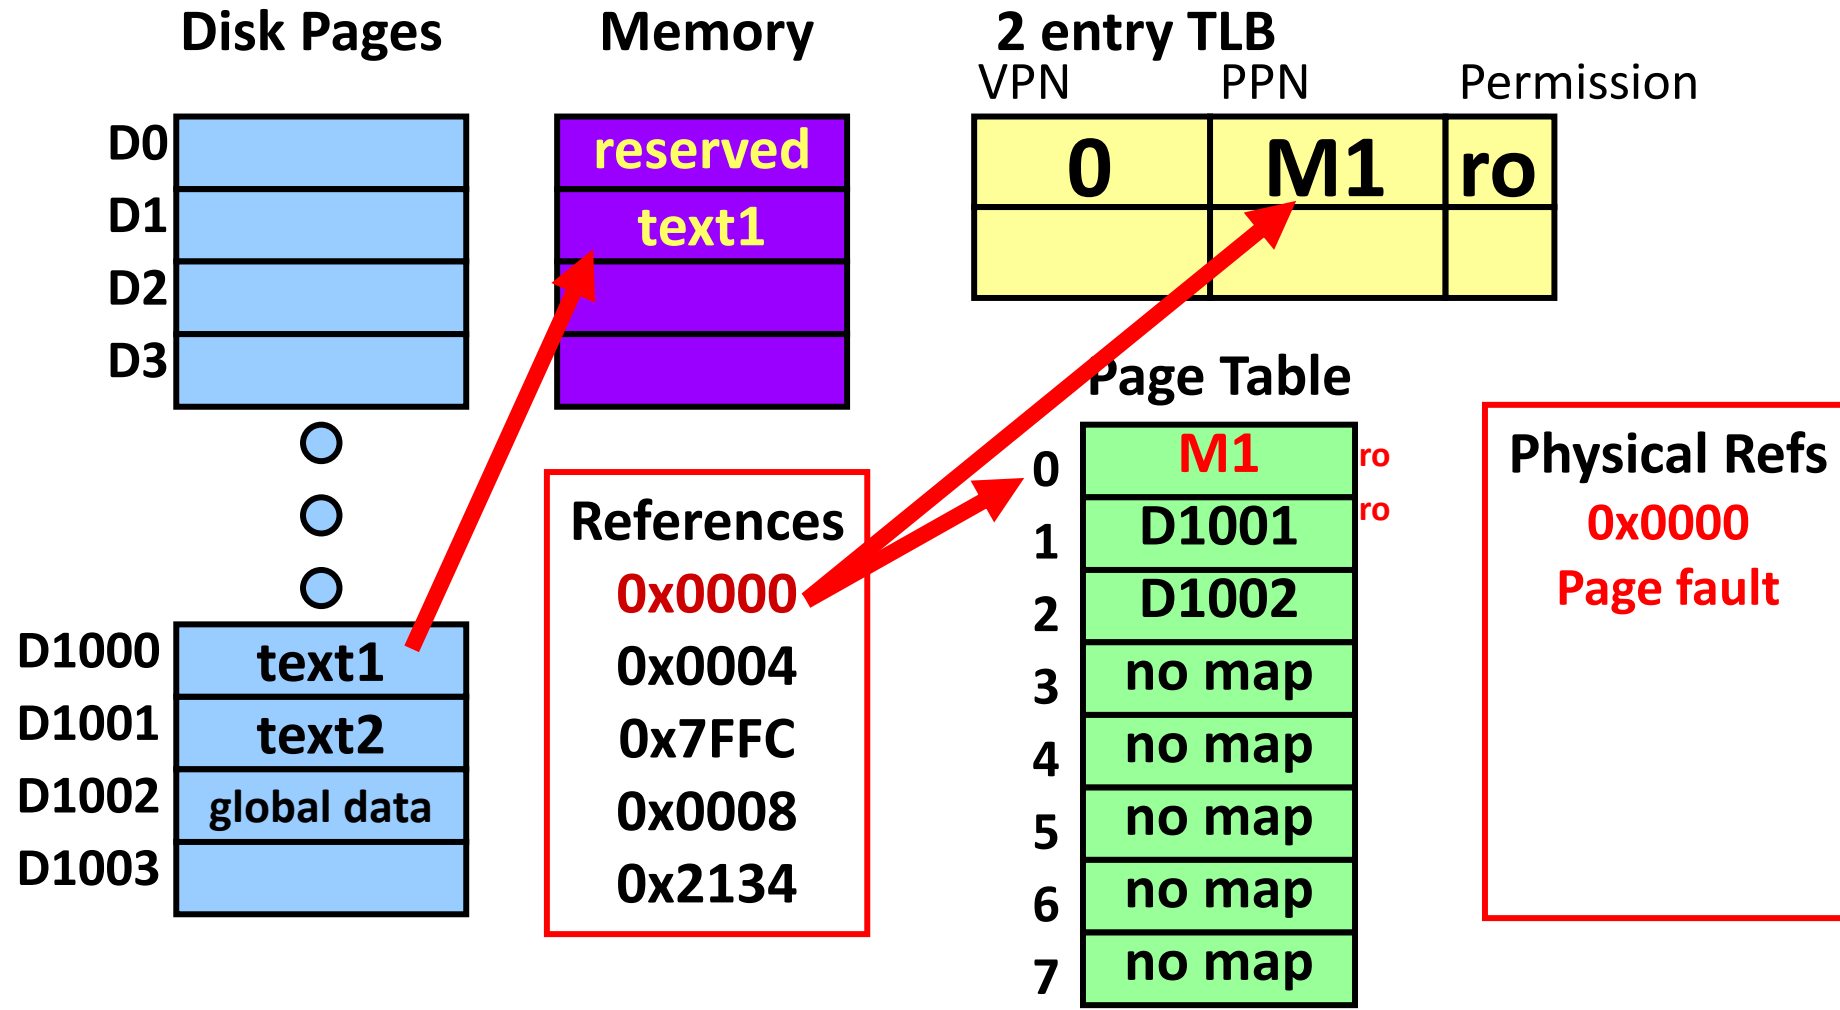
Step 2: Load PC from header & start execution



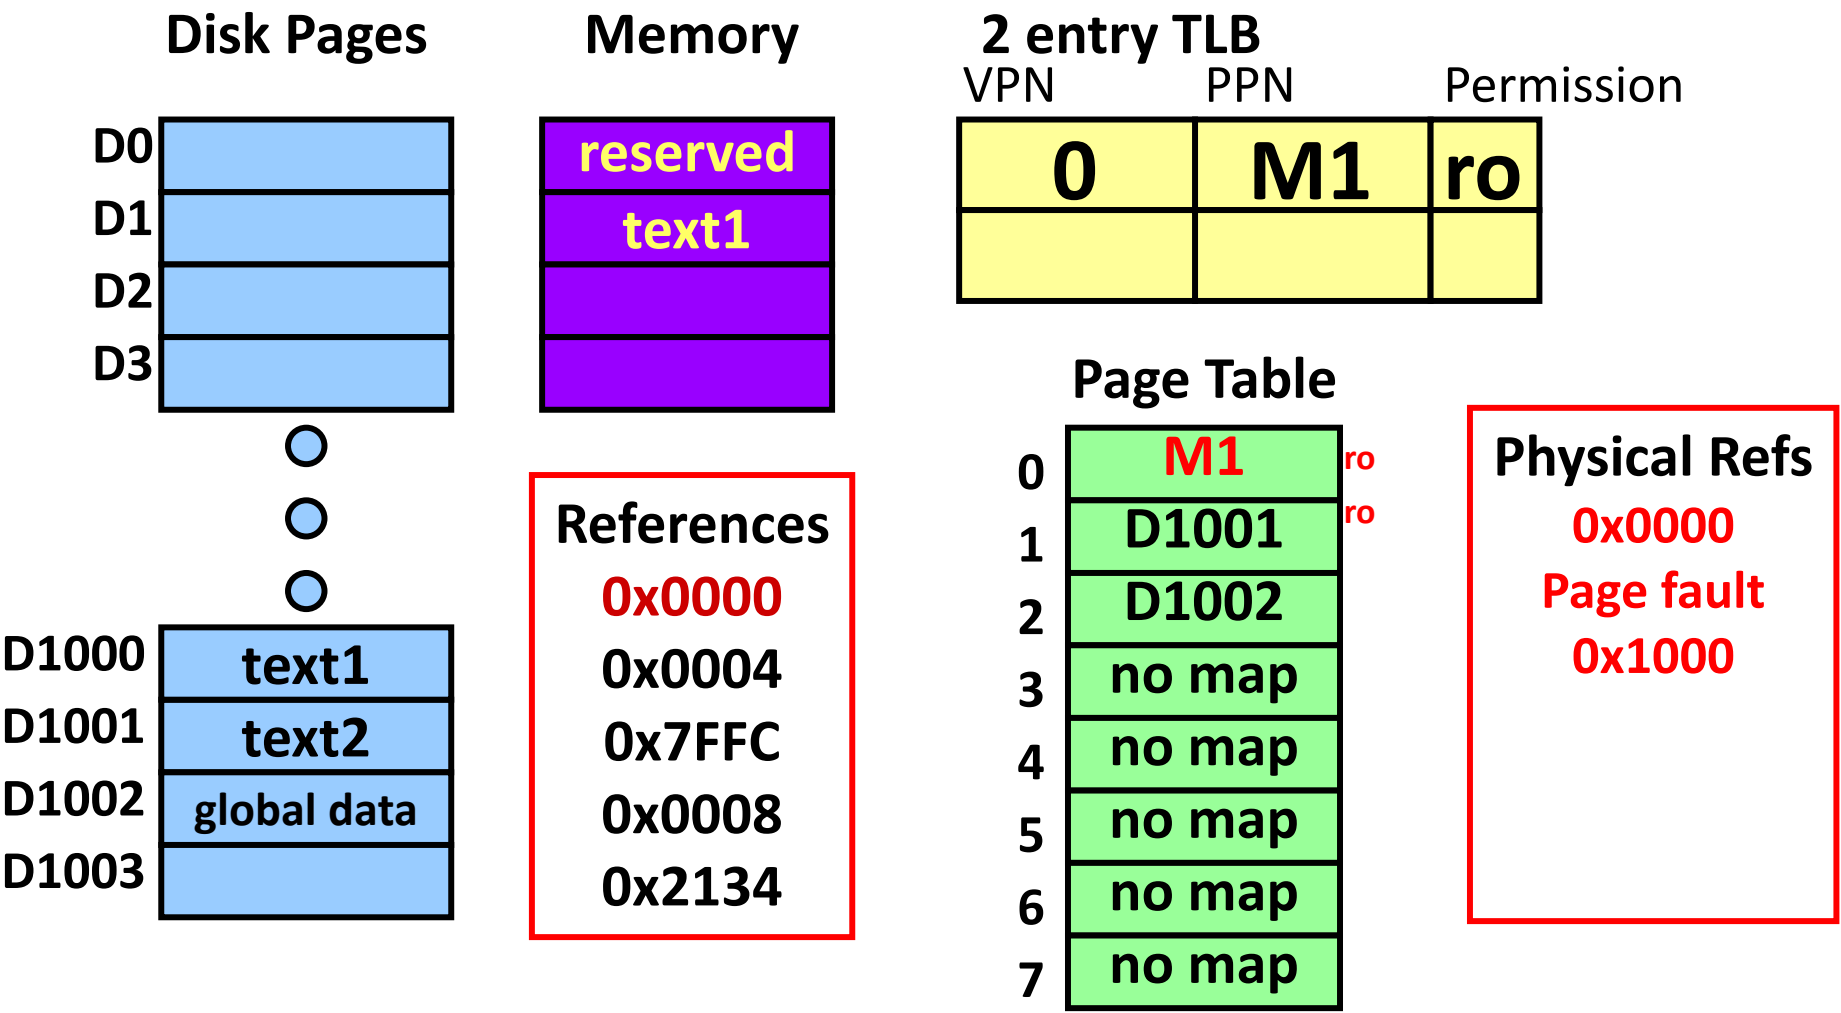
Fetching instruction 0000



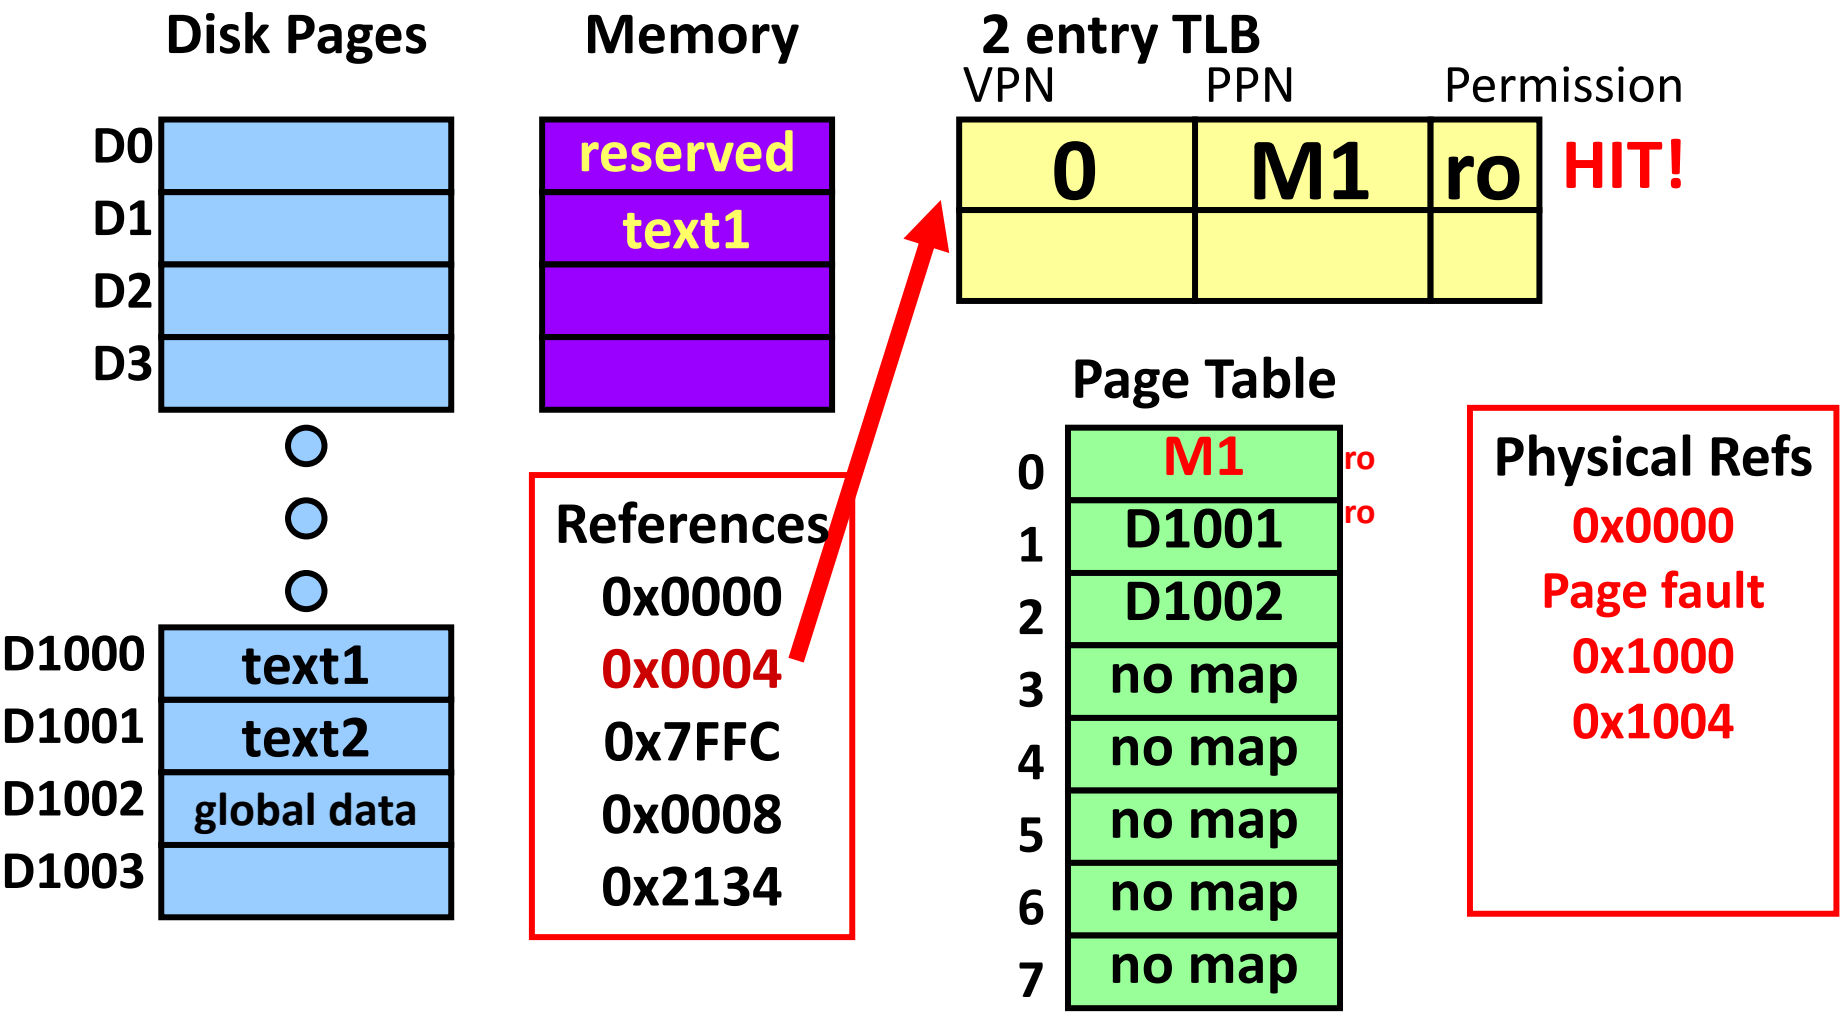
Fetching instruction 0000



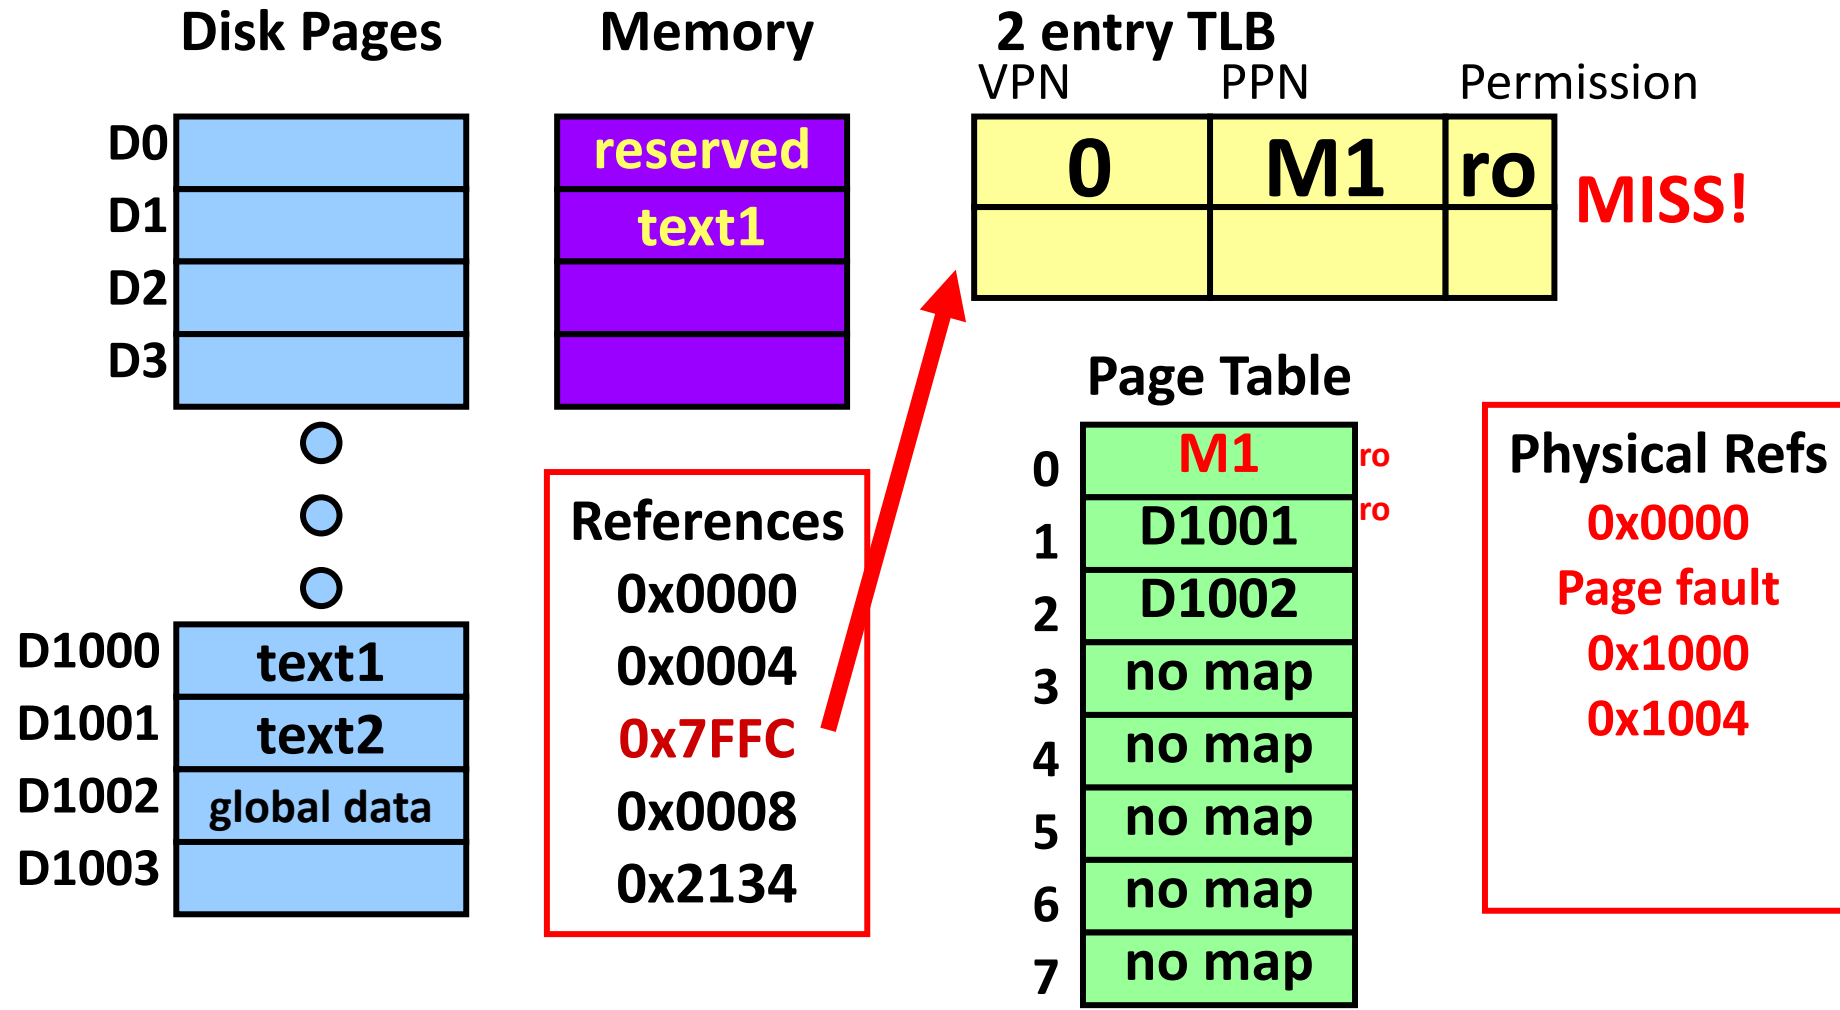
Fetching instruction 0000



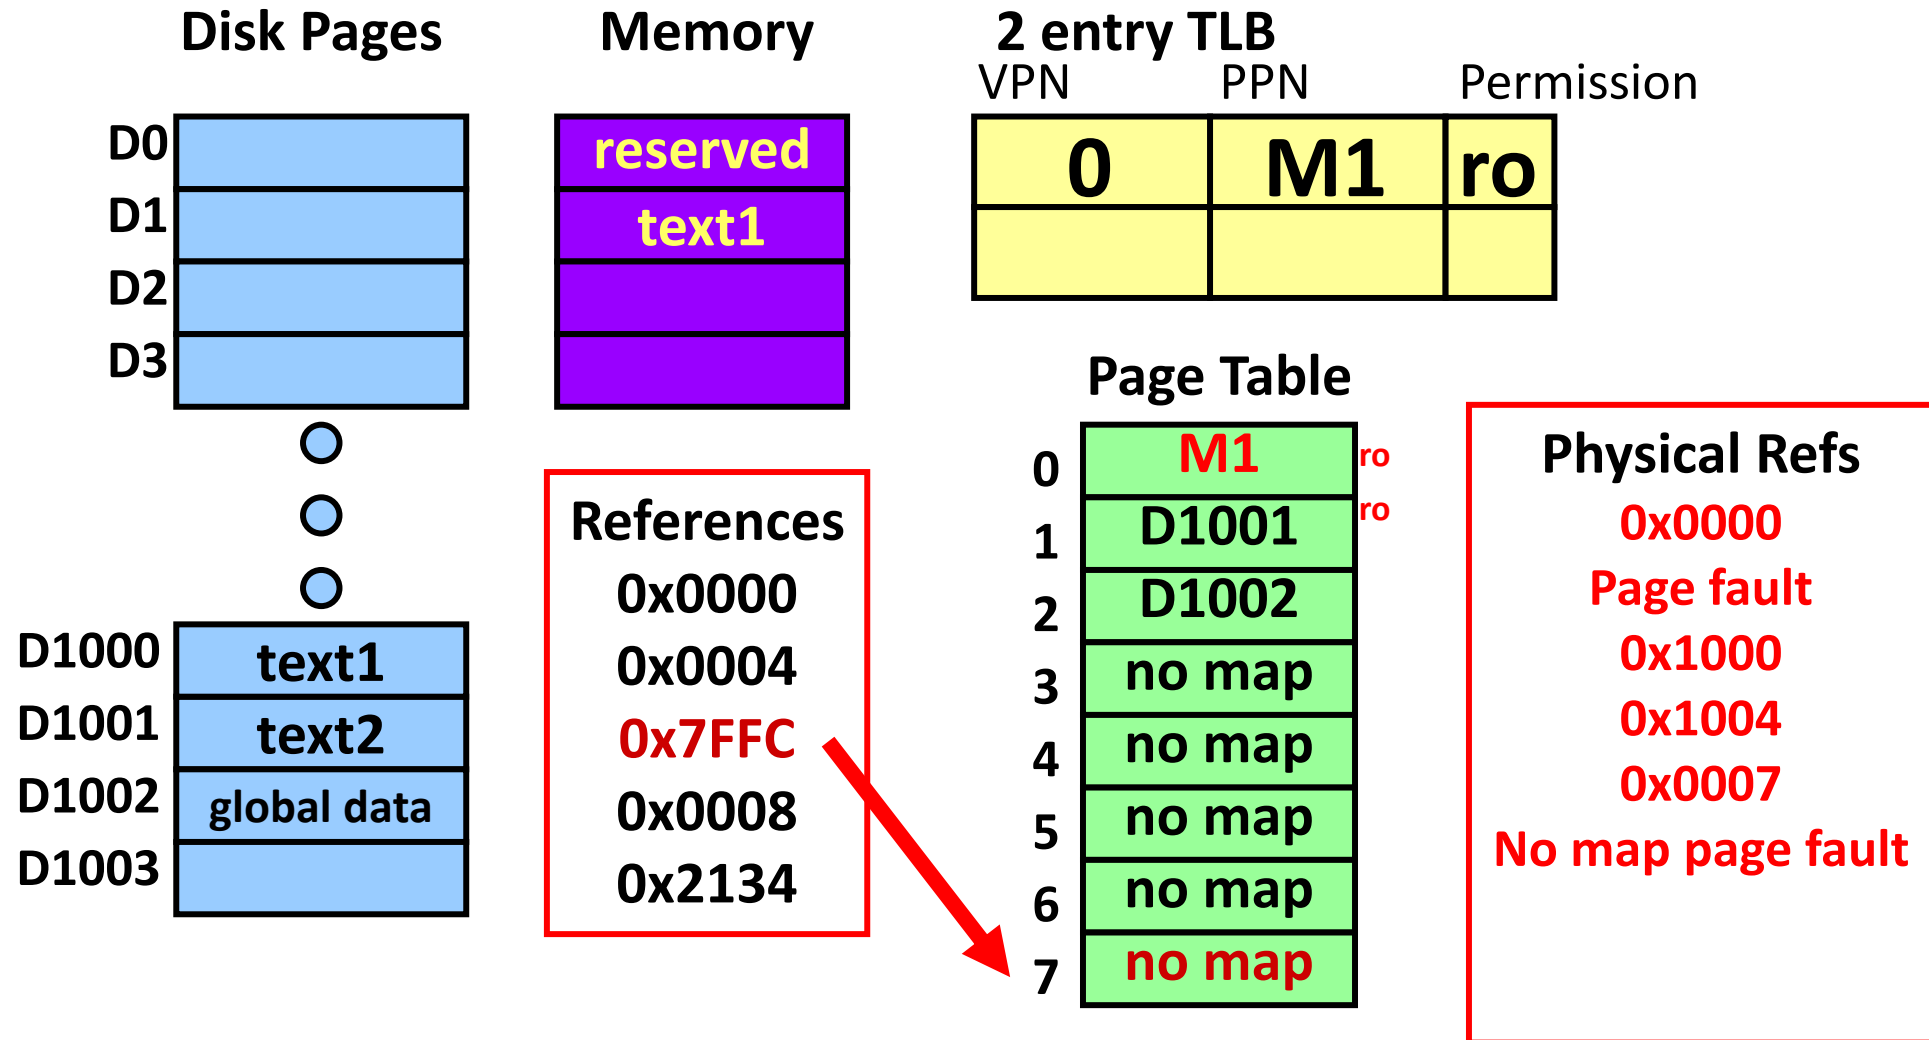
Fetching instruction 0004



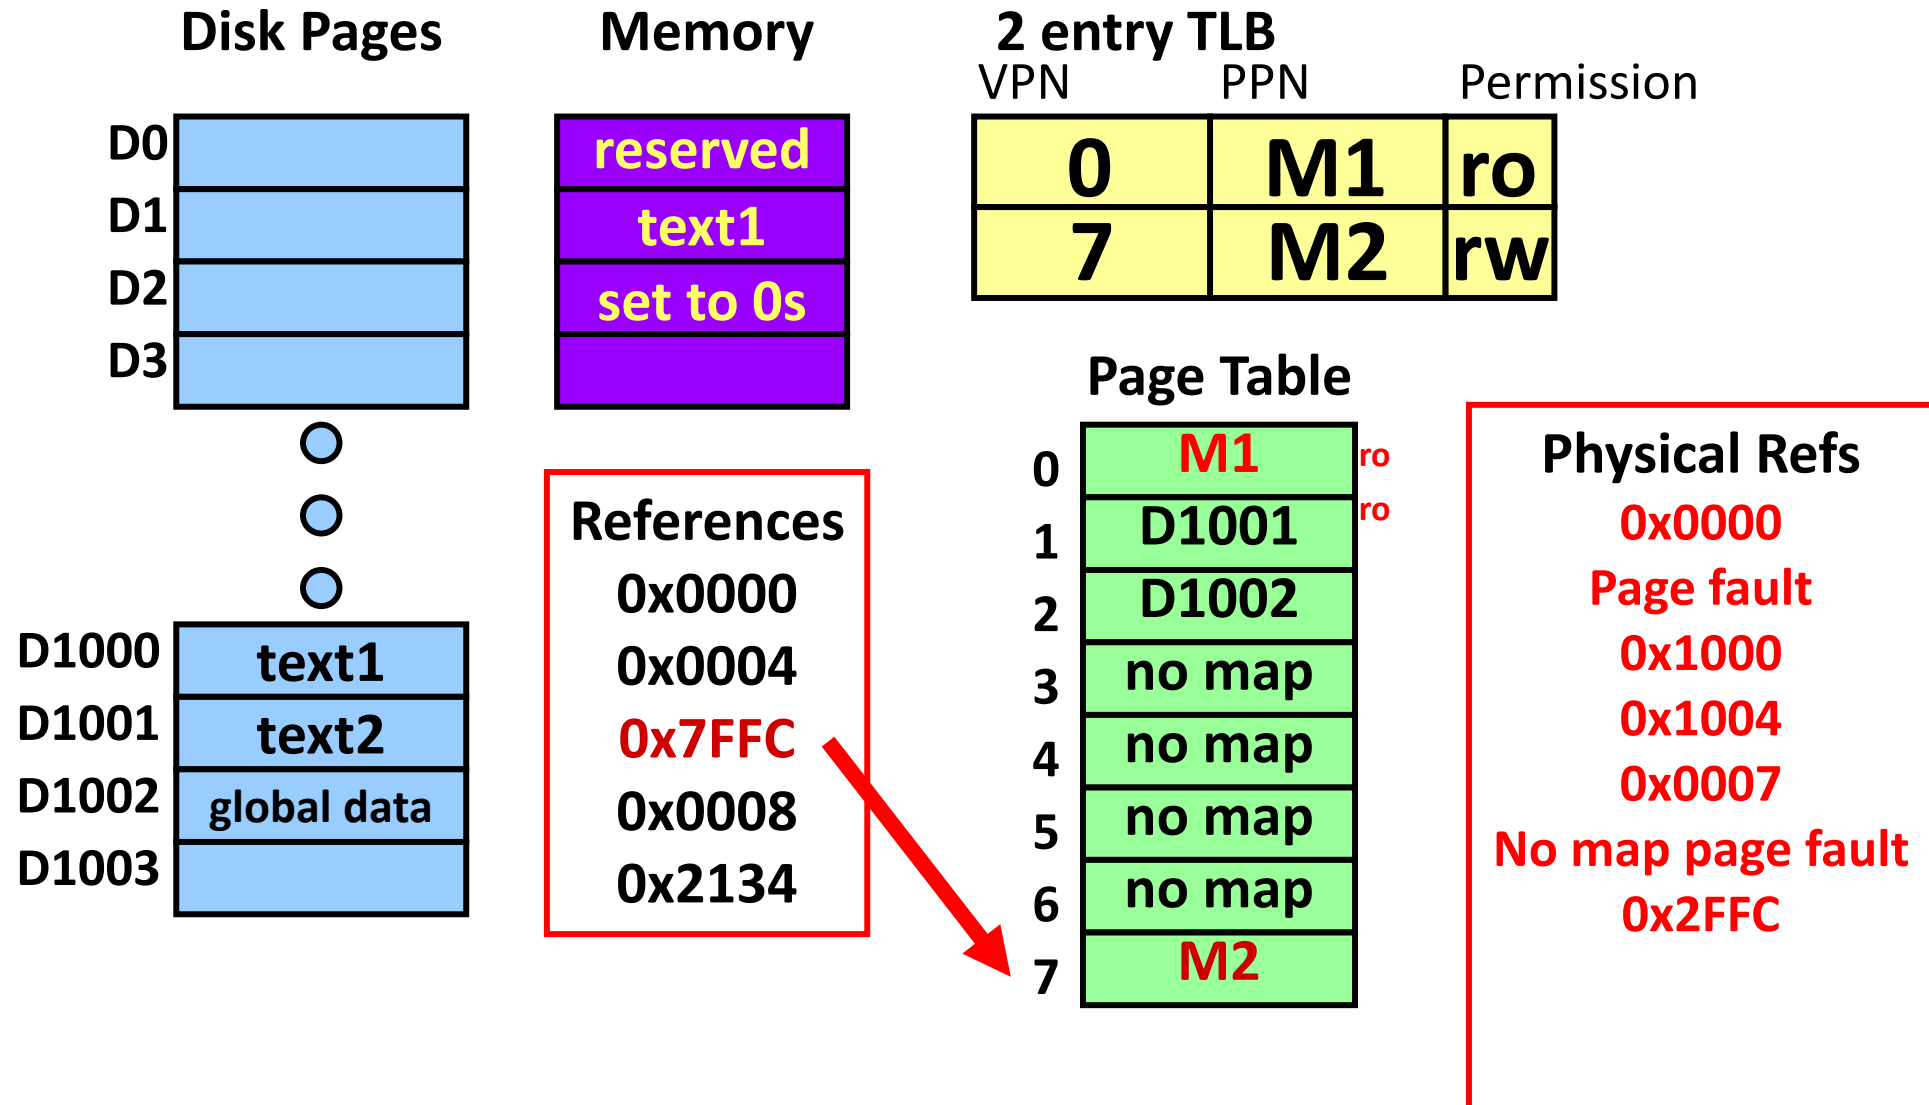
Reference 7FFC



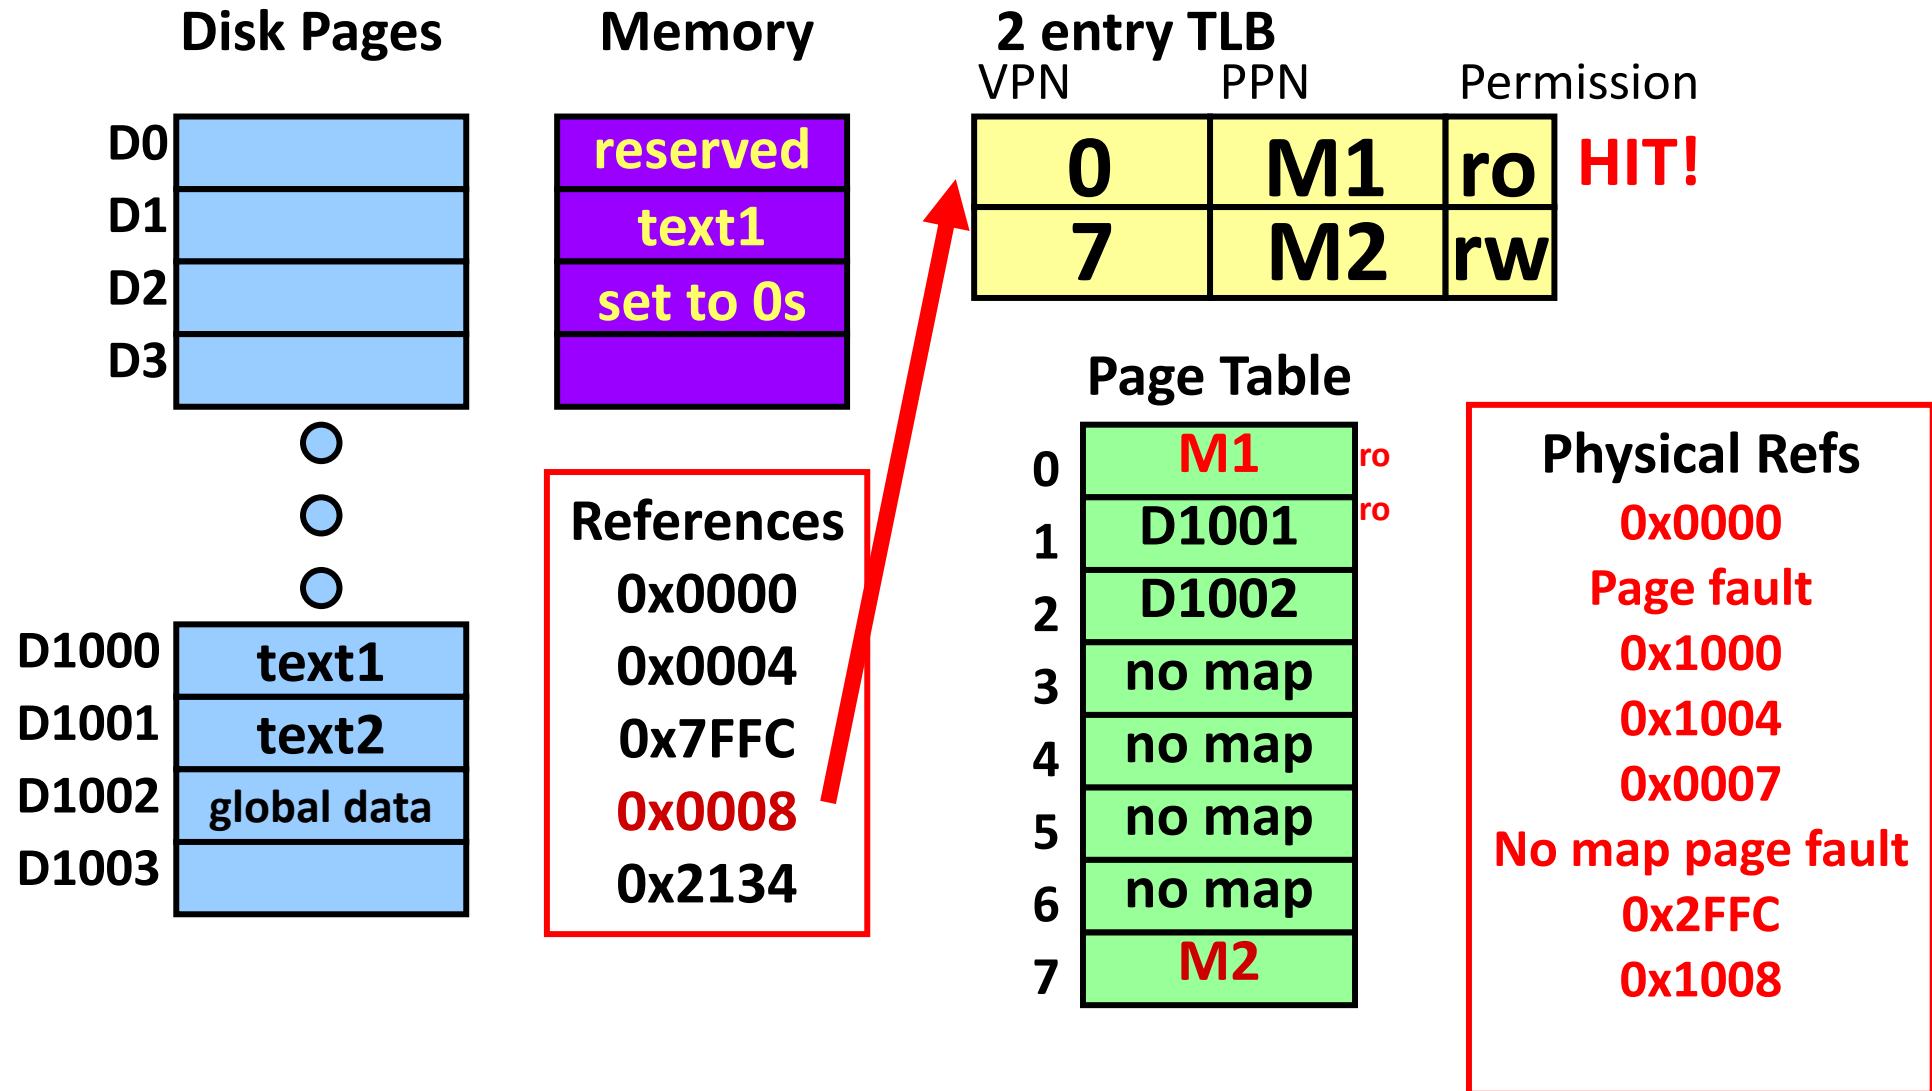
Reference 7FFC



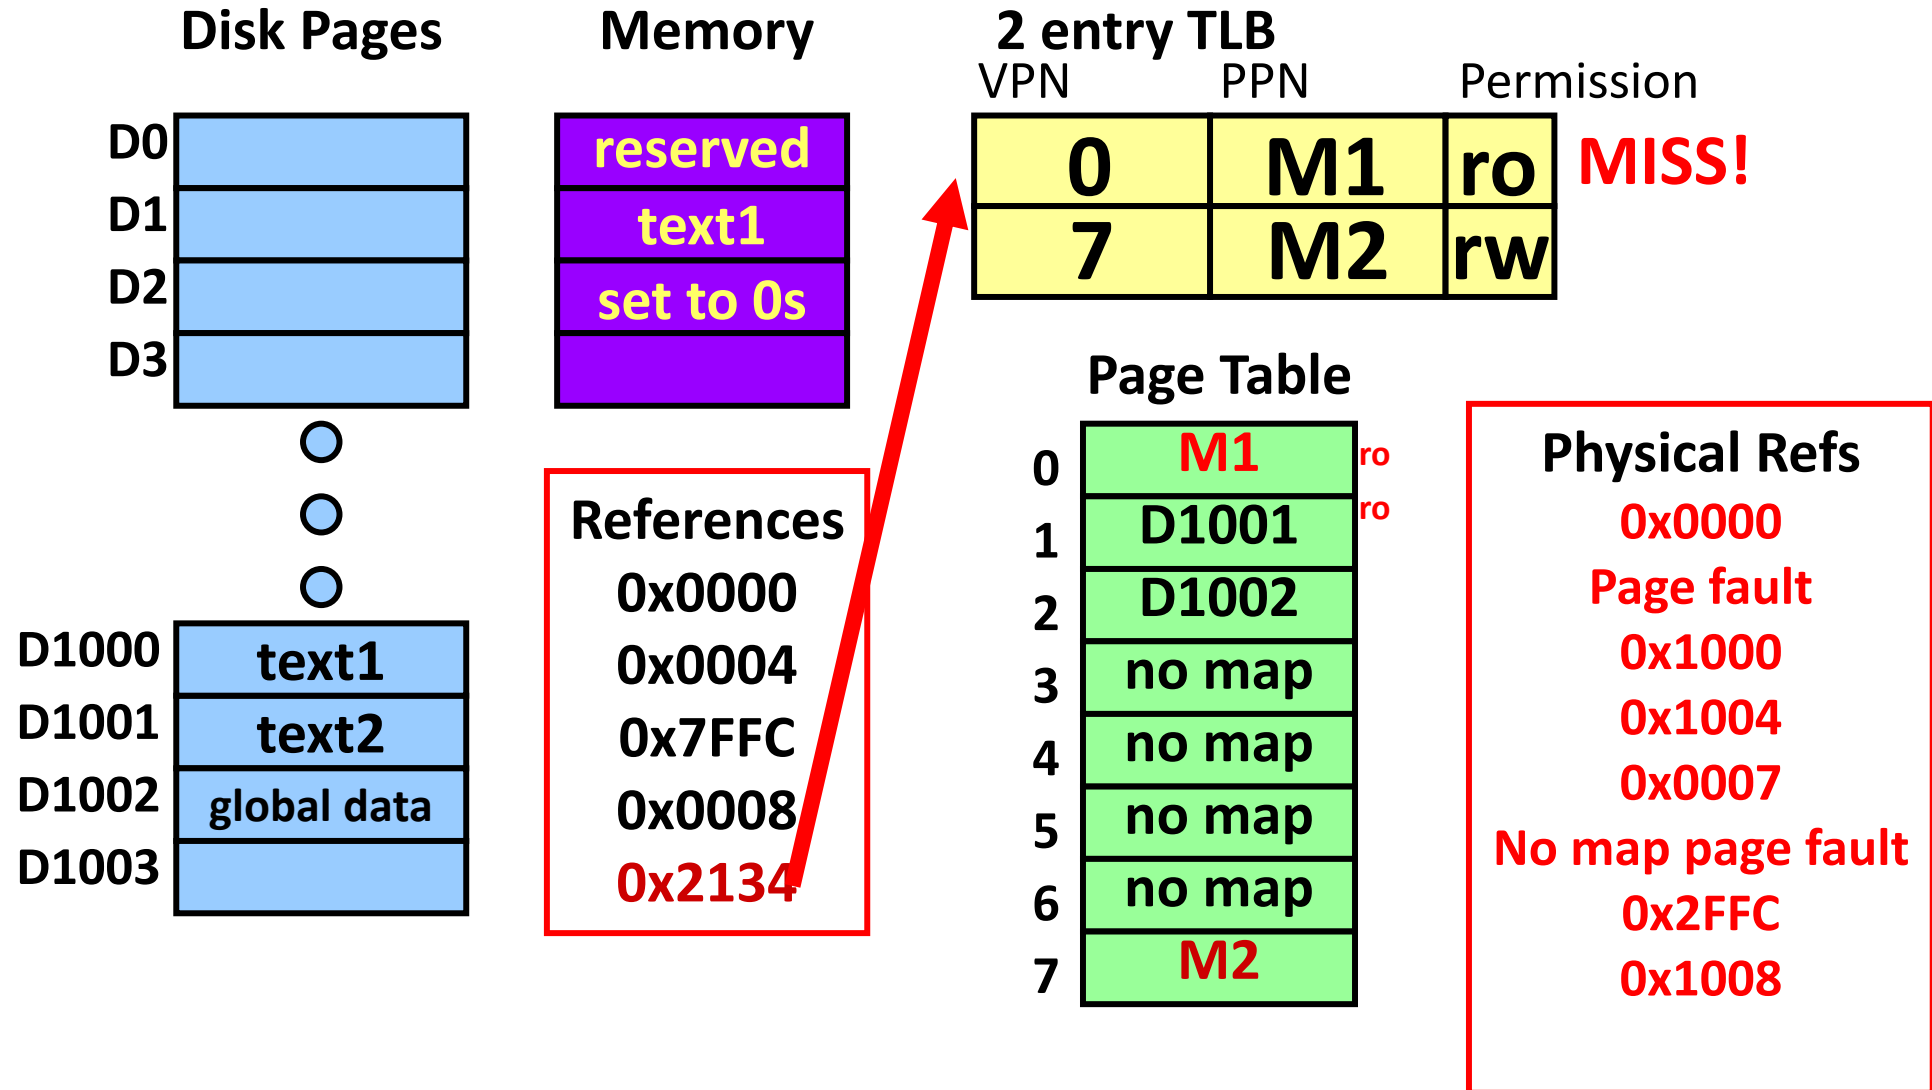
Reference 7FFC



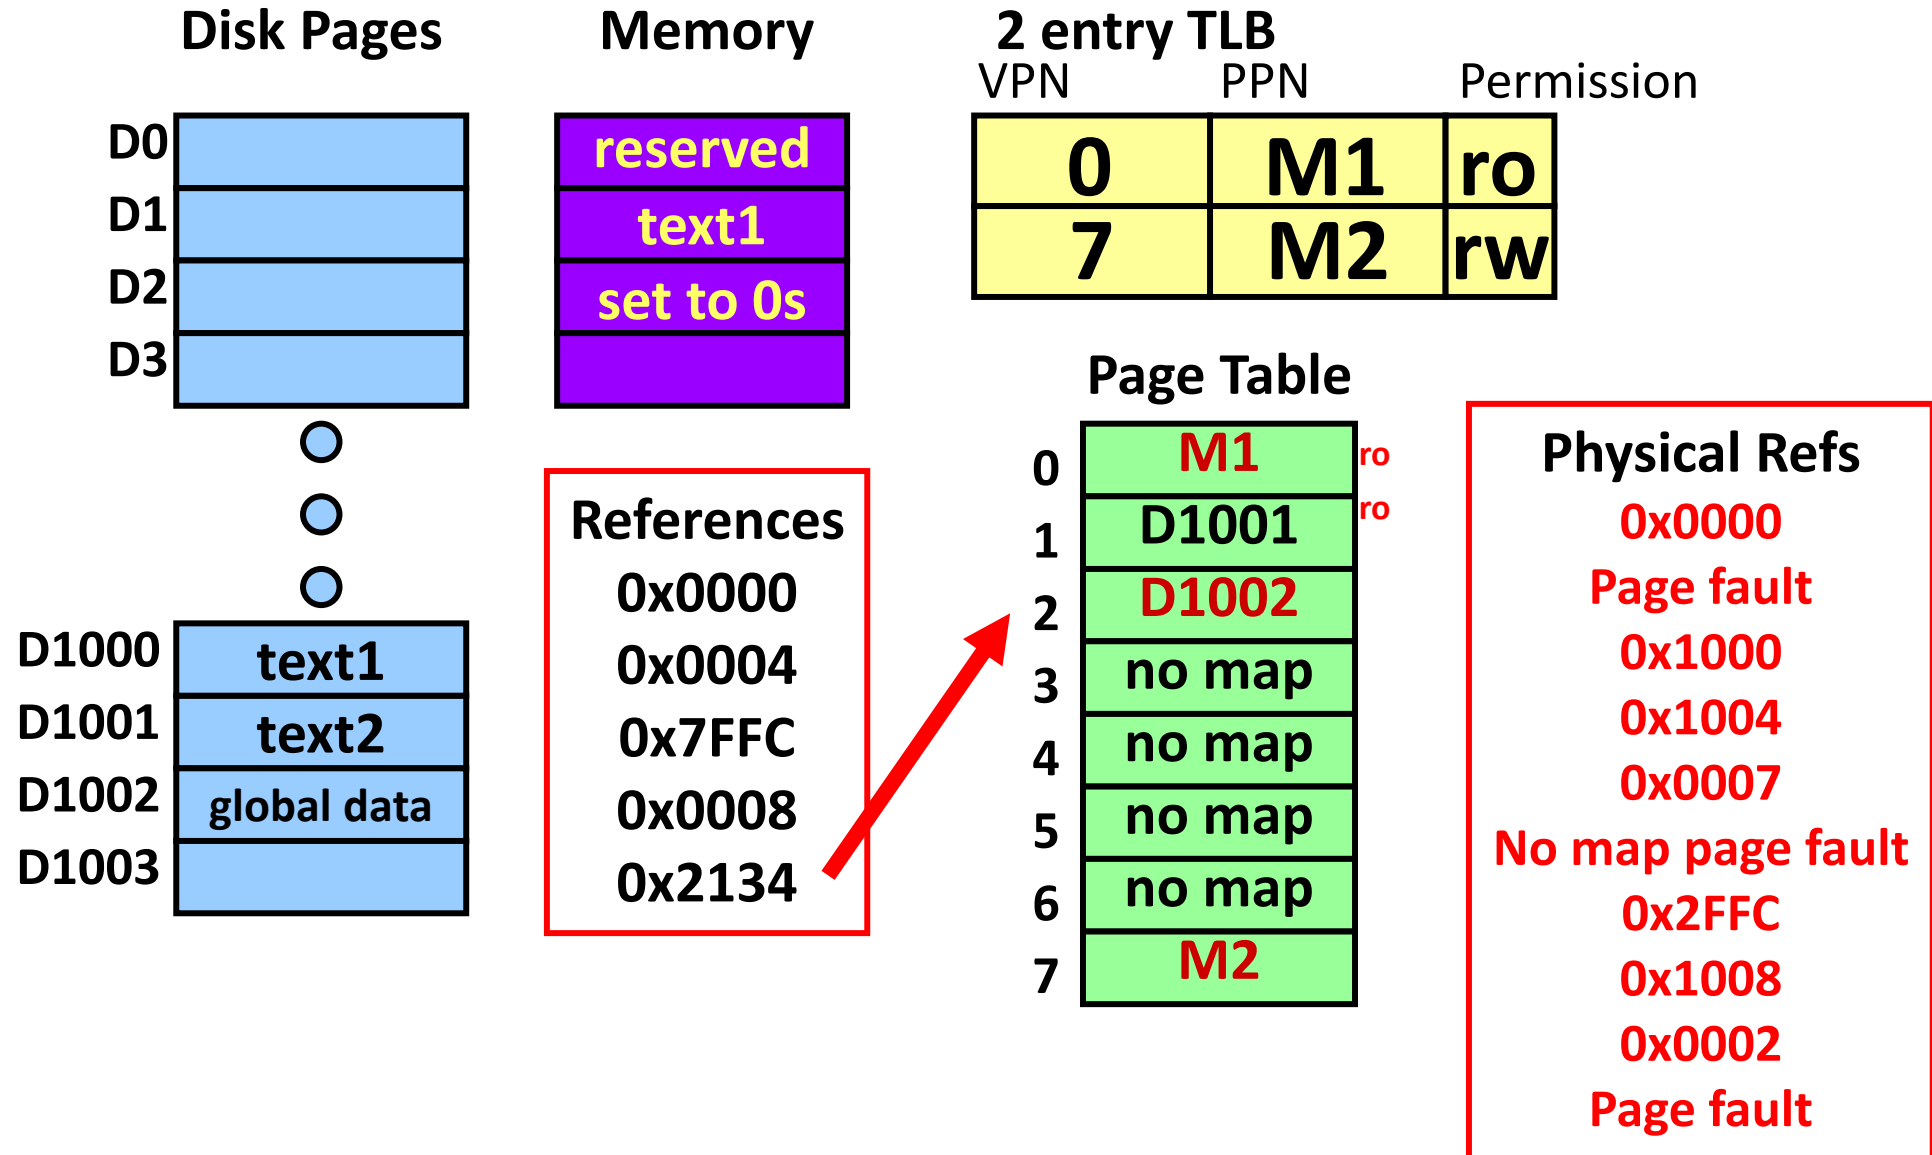
Fetching instruction 0008



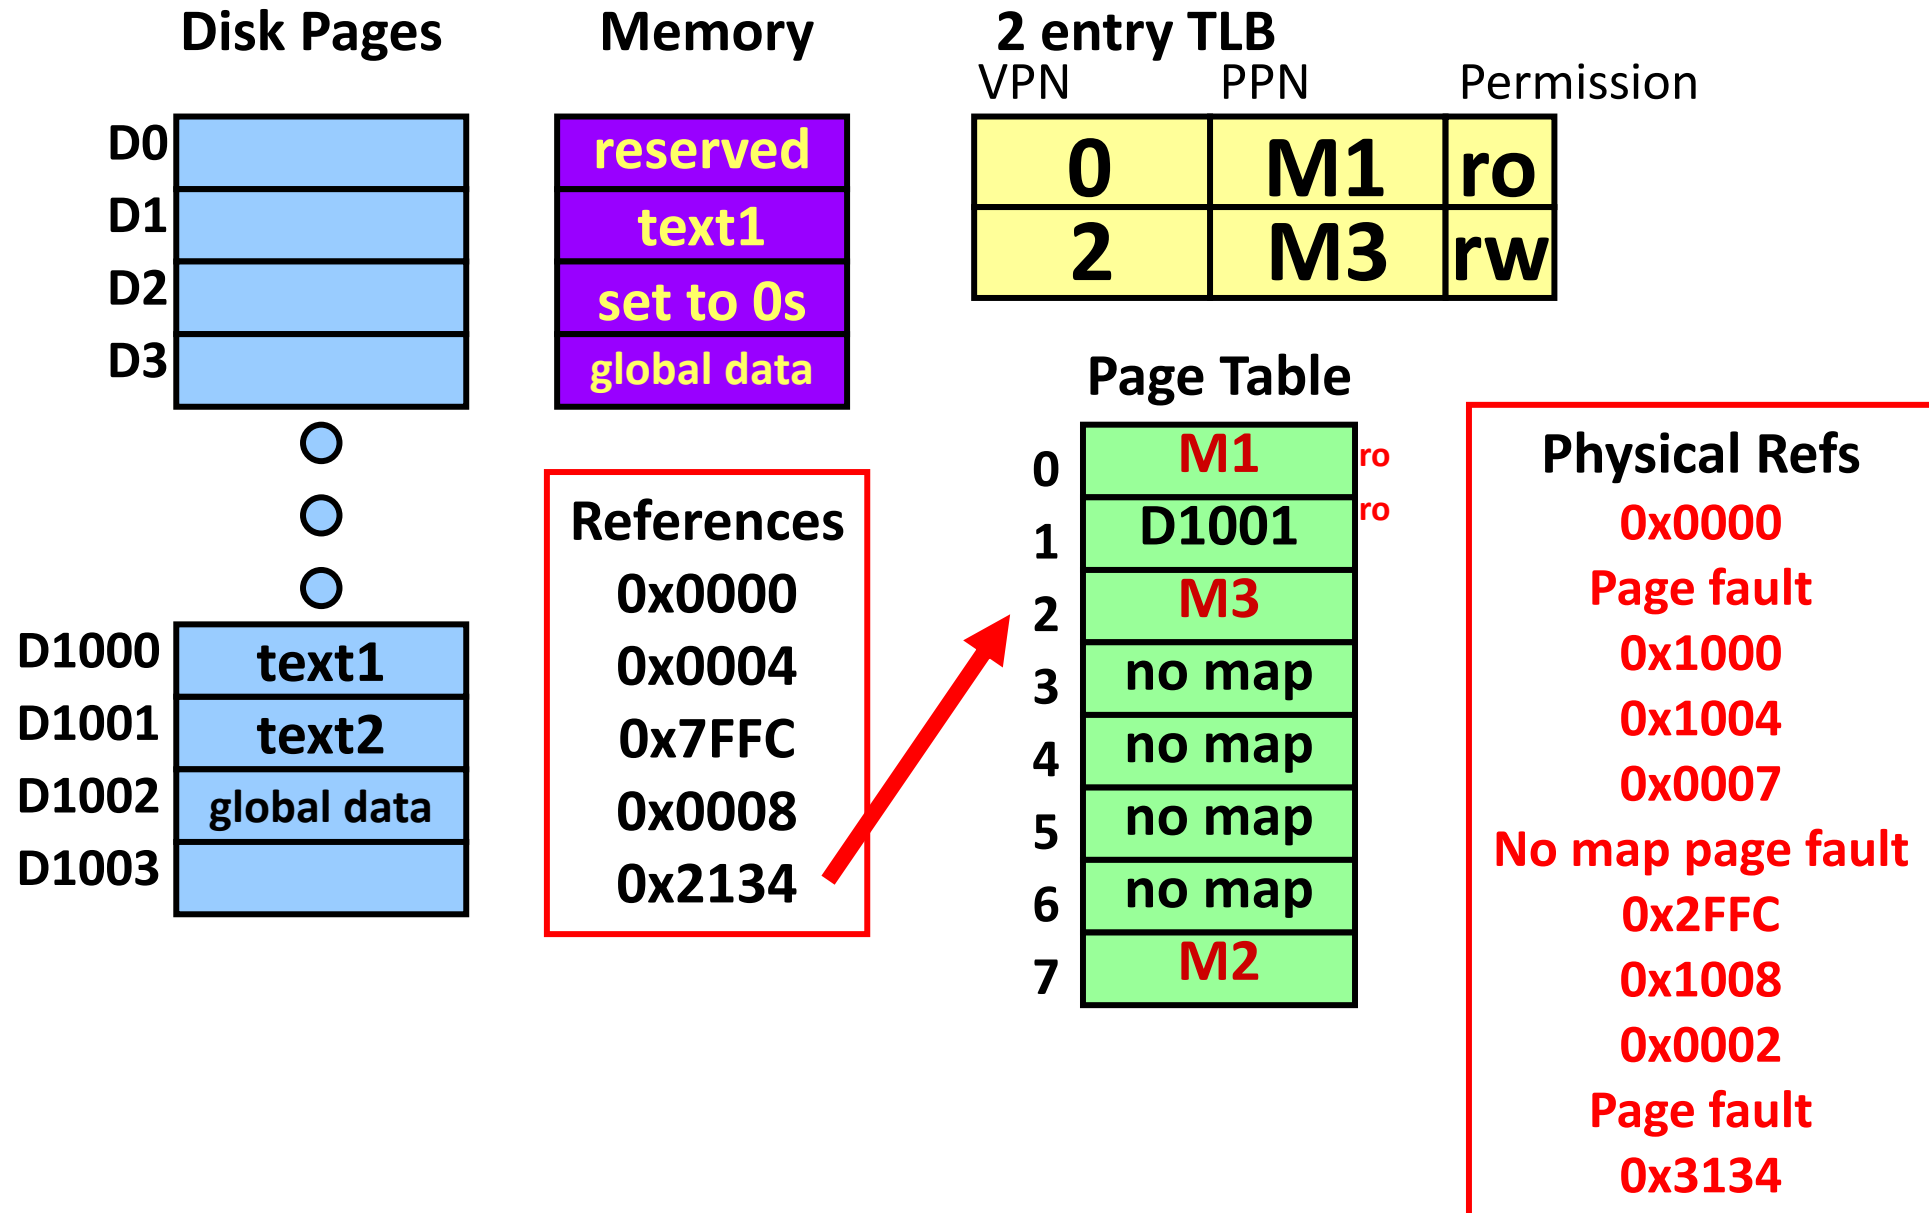
Reference 2134



Reference 2134



Reference 2134



Next time

- Speeding up virtual memory
 - Adding caches
- Wrap up class!
- Lingering questions / feedback? I'll include an anonymous form at the end of every lecture: <https://bit.ly/3oXr4Ah>

