

# **23. Virtual Memory: Design**

---

**EECS 370 – Introduction to Computer Organization – Winter 2023**

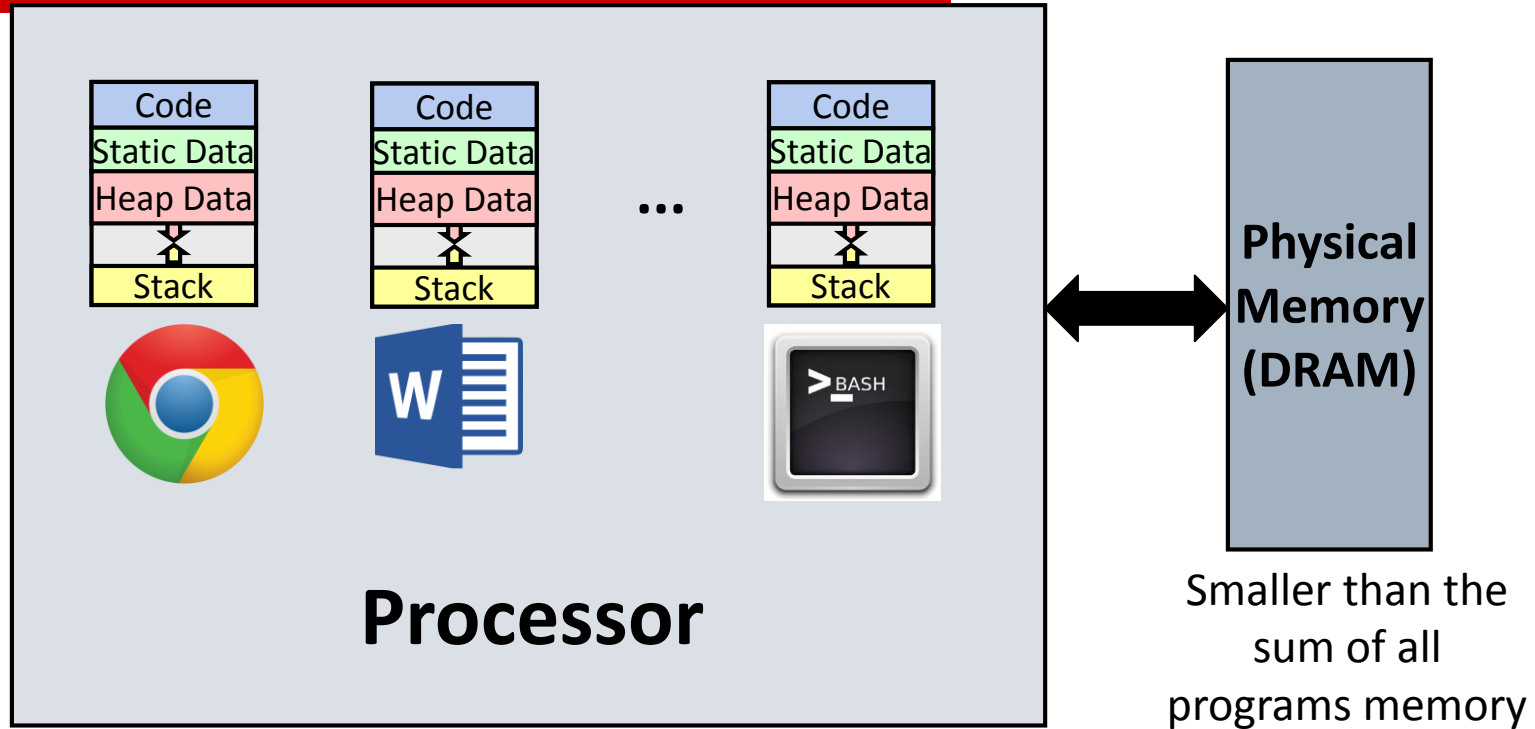
**EECS Department  
University of Michigan in Ann Arbor, USA**

# ANNOUNCEMENTS

---


- ❑ Project 4 due Thursday, April 13<sup>th</sup>
- ❑ Homework 6 due Monday, April 17<sup>th</sup>
- ❑ Just **2** more lectures of material and one review session!
  - ❑ Notice, we have no class on April 13<sup>th</sup>

# Review: Why Virtualize Memory?




1. Transparency
2. Protection
3. Programs not limited by DRAM Capacity

# How to achieve **transparency & protection**?



0x0000  
0x1000

0xF000



0x0000  
0x1000

0xF000

Program addresses  
(virtual)




0x0000  
0x1000



0x1F000

DRAM  
(physical)

# How to achieve transparency & protection?



Page table for Chrome



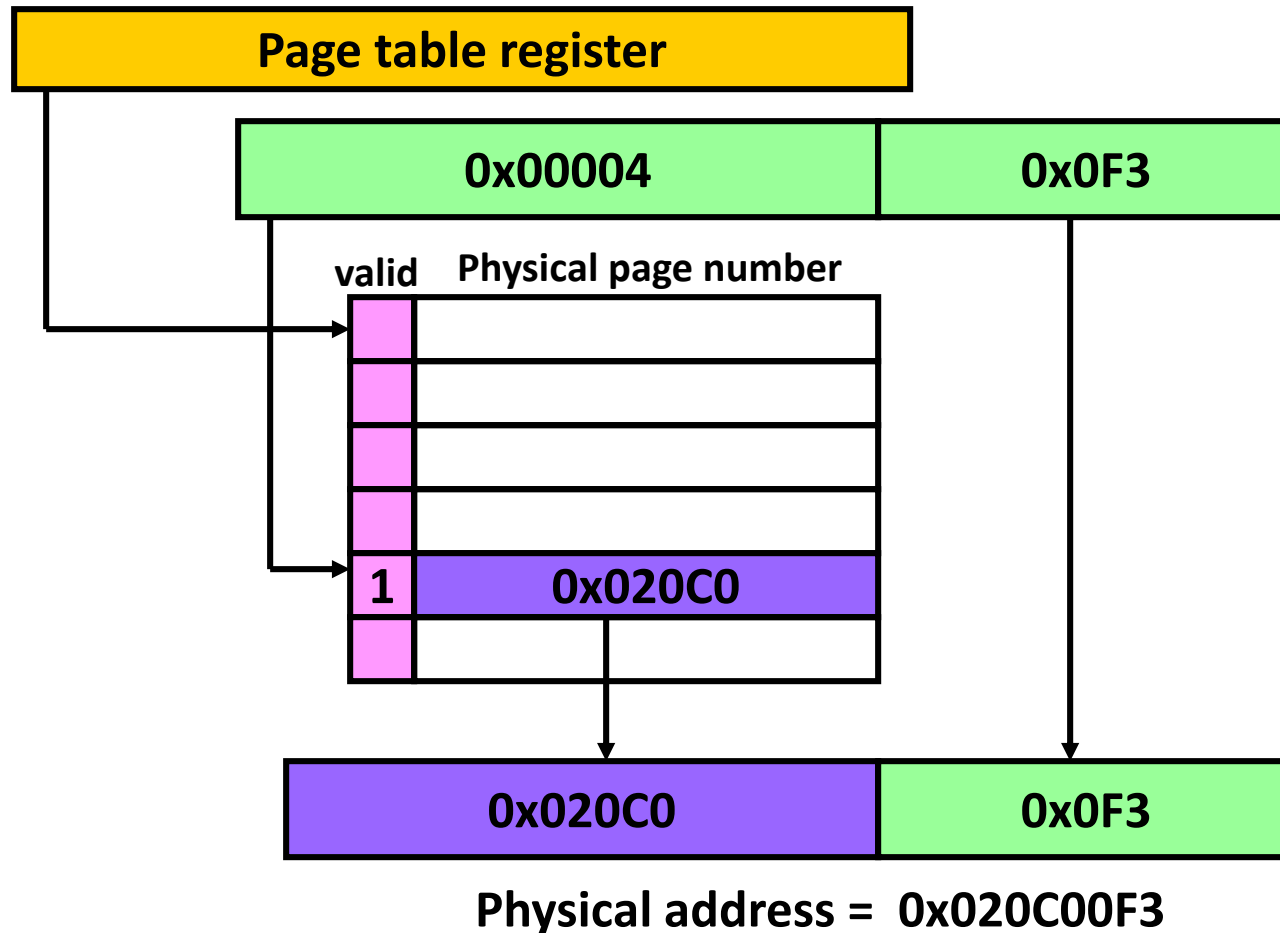
Page table for Word

Program addresses  
(virtual)

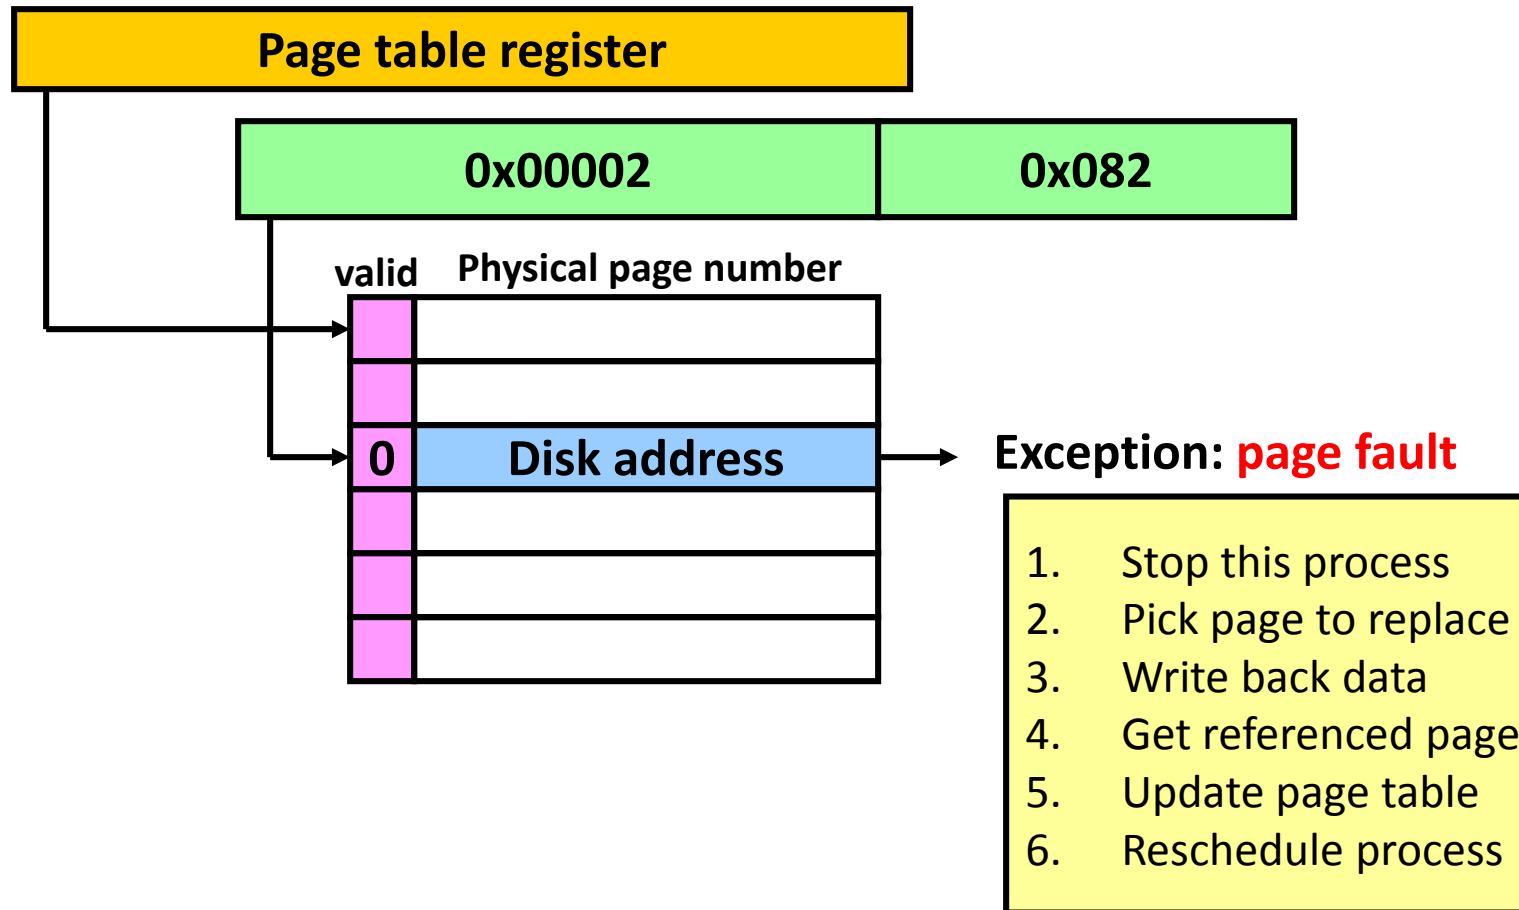
DRAM  
(physical)

# Review: Page table lookup

Virtual address = 0x000040F3



# Review: Page faults



# Class Problem

---

## ❑ Given the following:

- 4KB page size, physical memory of 16KB, page table stored in physical page 0 and can never be evicted, 20 bit, byte-addressable virtual address space.
- The page table initially has virtual page 0 in physical page 1, virtual page 1 in physical page 2 and no valid data in other physical pages.

## ❑ Fill in the table on the next slide for each reference

- Note: like caches we'll use LRU when we need to replace a page.



4KB page size,  
physical memory of 16KB,  
page table stored in physical  
page 0 and can never be  
evicted, 20 bit, byte-  
addressable virtual address  
space.

# Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C			
0x01F0C			
0x20F0C			
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			

The page table initially has  
virtual page 0 in physical  
page 1, virtual page 1 in  
physical page 2 and no valid  
data in other physical pages.

# Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C			
0x20F0C			
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			

The page table initially has  
virtual page 0 in physical  
page 1, virtual page 1 in  
physical page 2 and no valid  
data in other physical pages.

# Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C			
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			

The page table initially has  
virtual page 0 in physical  
page 1, virtual page 1 in  
physical page 2 and no valid  
data in other physical pages.

# Class Problem (continued)

---

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C	0x20	Y (into 3)	0x3F0C
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			

# Class Problem (continued)

The page table initially has  
virtual page 0 in physical  
page 1, virtual page 1 in  
physical page 2 and no valid  
data in other physical pages.

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C	0x20	Y (into 3)	0x3F0C
0x00100	0x0	N	0x1100
0x00200	0x0	N	0x1200
0x30000	0x30	Y (into 2)	0x2000
0x01FFF	0x1	Y (into 3)	0x3FFF
0x00200	0x0	N	0x1200

# Size of the page table

---

## ❑ How big is a page table entry?

- For 32-bit virtual address:
  - If the machine can support  $1\text{GB} = 2^{30}$  bytes of physical memory and we use pages of size  $4\text{KB} = 2^{12}$ ,
  - then the physical page number is  $30-12 = 18$  bits.  
Plus another valid bit + other useful stuff (read only, dirty, etc.)
  - Let say about 3 bytes.

## ❑ How many entries in the page table?

- ARM virtual address is 32 bits – 12 bit page offset = 20
- Total number of virtual pages =  $2^{20}$

## ❑ Total size of page table = Number of virtual pages

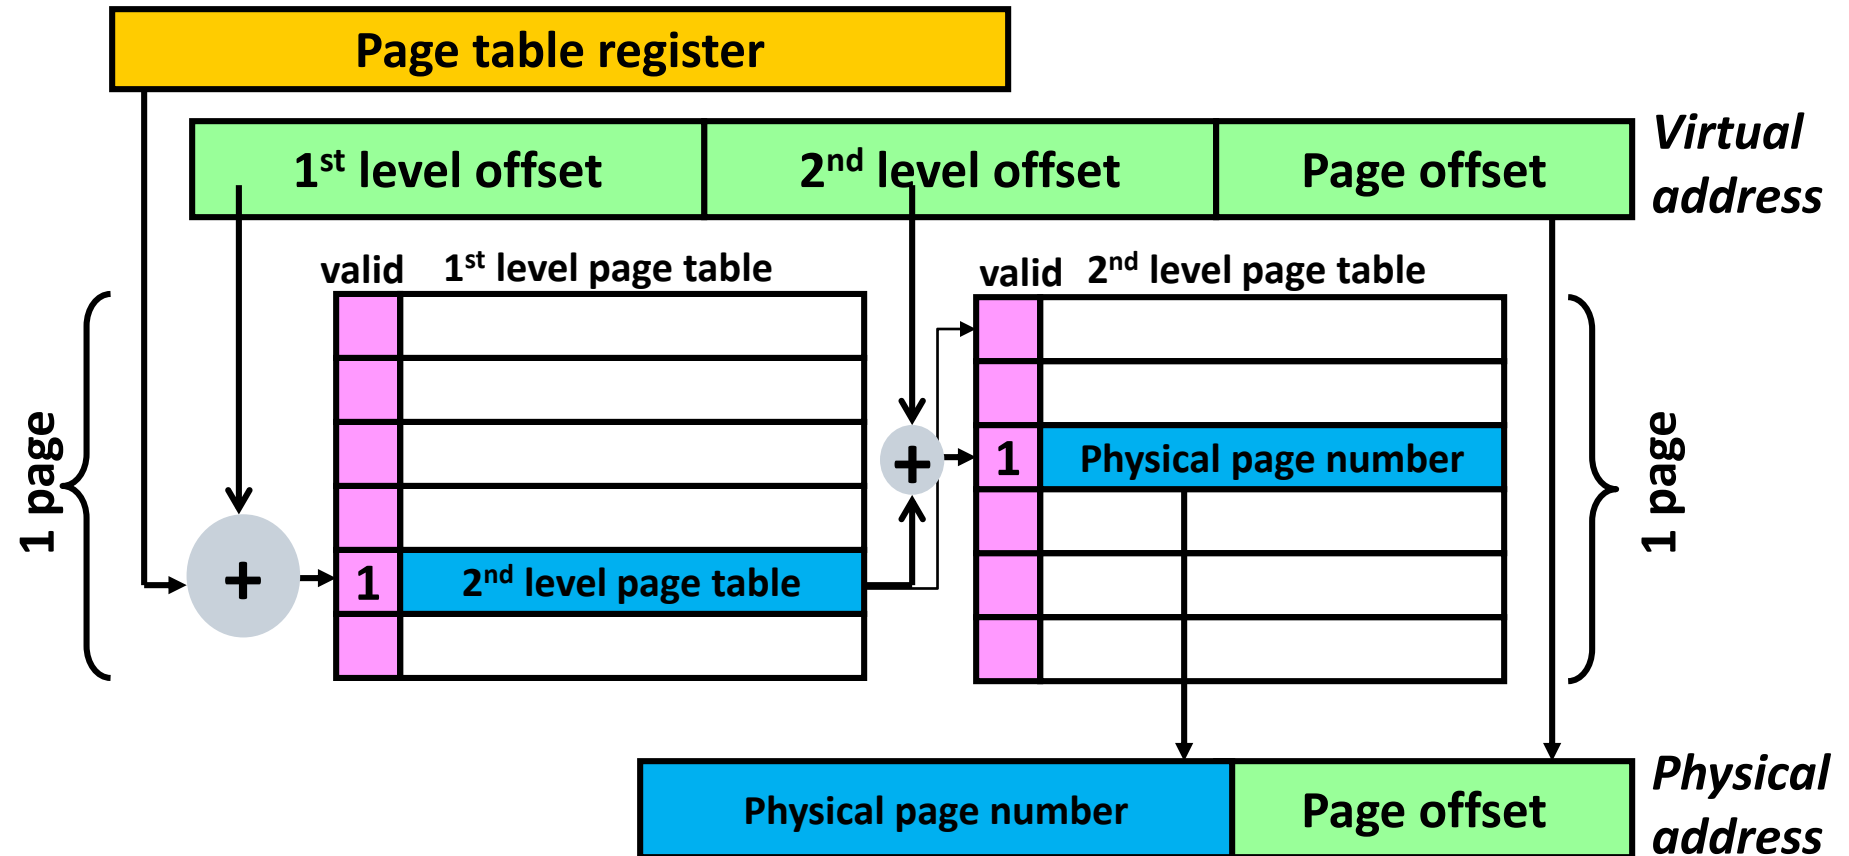
\* Size of each page table entry  
=  $2^{20} \times 3$  bytes  $\sim 3$  MB

# How can you organize the page table?

---


1. Single-level page table occupies continuous region in physical memory
  - Previous example always takes 3MB regardless of how much virtual memory is used
2. Use a multi-level page table! (Build a hierarchical page table and keep translations only if used in the program)
  - 1<sup>st</sup> level page table in physical memory
  - 2<sup>nd</sup> (or more) level page tables created only if needed
  - **Size is proportional to the amount of memory used**
  - Optimization for common case: *almost no programs use all the virtual pages at the same time*
  - **Common case: size of multi-level page table  $\ll$  single-level**
  - **Worst case: size of multi-level page table  $>$  single-level**

# Hierarchical page table





# Hierarchical page table – 32bit Intel x86



- ❑ How many bits in the virtual 1<sup>st</sup> level offset field? 10
- ❑ How many bits in the virtual 2<sup>nd</sup> level offset field? 10
- ❑ How many bits in the page offset? 12
- ❑ How many pages in the 1<sup>st</sup> level page table?  $2^{10}=1024$
- ❑ How many bytes for each entry in the 1<sup>st</sup> level page table? 4
- ❑ How many entries in the 2<sup>nd</sup> level of the page table?  $2^{10}=1024$
- ❑ How many bytes for each VPN in a 2<sup>nd</sup> level table? 4
- ❑ What is the total size of the page table?  $4K+n*4K$   
(here  $n$  is number of valid entries in the 1<sup>st</sup> level page table)

# Class Problem (32 bit x86)

---

- ❑ What is the least amount of memory that could be used? When would this happen?
- ❑ What is the most memory that could be used? When would this happen?
- ❑ How much memory is used for this memory access pattern:  
0x00000ABC  
0x00000ABD  
0x10000ABC  
0x20000ABC
- ❑ How much memory if we used a single-level page table with 4KB pages? Assume entries are rounded to the nearest word (4B)

# Class Problem (32 bit x86)

---

- ❑ What is the least amount of memory that could be used? When would this happen?
  - 4KB for 1<sup>st</sup> level page table. Occurs when no memory has been accessed (before program runs)
  
- ❑ What is the most memory that could be used? When would this happen?
  - 4KB for 1<sup>st</sup> level page table  
+ 1024\*4KB for all possible 2<sup>nd</sup> level page tables  
= 4100KB (which slightly greater than 4096KB)
  - Occurs when program uses all virtual pages ( $= 2^{20}$  pages)

# Class Problem (32 bit x86)

---

- How much memory is used for this memory access pattern:

0x00000ABC // Page fault

0x00000ABD

0x10000ABC // Page fault

0x20000ABC // Page fault

- 4KB for 1<sup>st</sup> level page table + 3\*4KB for each 2L page table  
= 16 KB

# Class Problem (32 bit x86)

---

- How much memory if we used a single-level page table with 4KB pages? Assume entries are rounded to the nearest word (4B)
- 32 bits – 12 bit page offset = 20 bits of address
  - $2^{20}$  (num entries) \* 4B  
=> 4MB

# Class Problem – Multi-level VM

---

- ❑ Design a two-level virtual memory system of a byte addressable processor with **24-bits long addresses**. No cache in the system. **256Kbytes of memory** installed, and no additional memory can be added.
  - **Virtual memory page: 512 Bytes**. Each page table entry must be an integer number of bytes, and must be the smallest size required to fit the physical page number + 1 bit to mark valid-entry
  - We want each second-level page table to fit exactly in one memory page, and 1<sup>st</sup> level page table entries are 3 bytes each (a memory address pointer to a 2L page table).
- ❑ **Compute:**
  - Number of entries in each 2<sup>nd</sup> level page table;
  - Number of virtual address bits used to index the 2<sup>nd</sup> level page table;
  - Number of virtual address bits used to index the 1<sup>st</sup> level page table;
  - Size of the 1<sup>st</sup> level page table.

# Class Problem – Multi-level VM

---

Page Offset: 9 bits (512B page size)

**Physical address = 18b (256KB Mem size)**

**Physical page number = 18b (256KB mem size) – 9b (offset)**

<b>Physical page number = 9b</b>	<b>Page offset = 9b</b>
----------------------------------	-------------------------

2<sup>nd</sup> level page table entry size: 9b (physical page number) + 1b ~ 2 bytes

2<sup>nd</sup> level page table **fits exactly in 1 page**

#entries in 2<sup>nd</sup> level page table is 512 bytes / 2 bytes = 256

#entries in 2<sup>nd</sup> level page table = 256 → Virtual page bits = 8b

Virtual 1<sup>st</sup> level page bits = 24 – 8 – 9 = 7b

1<sup>st</sup> level page table size =  $2^7 * 3 \text{ bytes} = 384\text{B}$

<b>1<sup>st</sup> level = 7b</b>	<b>2<sup>nd</sup> level = 8b</b>	<b>Page offset = 9b</b>
----------------------------------	----------------------------------	-------------------------

**Virtual address = 24b**

# Exercise using previous multi-level VM

Virtual Address	1 <sup>st</sup> level	2 <sup>nd</sup> level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C						
0x001F0C						
0x020F0C						

**1<sup>st</sup> level = 7b**

**2<sup>nd</sup> level = 8b**

**Page offset = 9b**

***Virtual address = 24b***

**Physical page number = 9b**

**Page offset = 9b**

***Physical address = 18b***

Assume memory for page tables are “somewhere else” in memory



# Exercise using previous multi-level VM

Virtual Address	1 <sup>st</sup> level	2 <sup>nd</sup> level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C						
0x020F0C						

**1<sup>st</sup> level = 7b**

**2<sup>nd</sup> level = 8b**

**Page offset = 9b**

***Virtual address = 24b***

**Physical page number = 9b**

**Page offset = 9b**

***Physical address = 18b***

Assume memory for page tables are “somewhere else” in memory

# Exercise using previous multi-level VM

Virtual Address	1 <sup>st</sup> level	2 <sup>nd</sup> level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C	0x00	0x0F	0x10C	Y	0x001	0x0030C
0x020F0C						

**1<sup>st</sup> level = 7b**

**2<sup>nd</sup> level = 8b**

**Page offset = 9b**

***Virtual address = 24b***

**Physical page number = 9b**

**Page offset = 9b**

***Physical address = 18b***

Assume memory for page tables are “somewhere else” in memory

# Exercise using previous multi-level VM

Virtual Address	1 <sup>st</sup> level	2 <sup>nd</sup> level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C	0x00	0x0F	0x10C	Y	0x001	0x0030C
0x020F0C	0x01	0x07	0x10C	Y	0x002	0x0050C

**1<sup>st</sup> level = 7b**

**2<sup>nd</sup> level = 8b**

**Page offset = 9b**

***Virtual address = 24b***

**Physical page number = 9b**

**Page offset = 9b**

***Physical address = 18b***

Assume memory for page tables are “somewhere else” in memory

# Page Replacement Strategies

---

- ❑ Page table indirection enables a fully associative mapping between virtual and physical pages.
  
- ❑ How do we implement LRU in OS?
  - True LRU is expensive, but LRU is a heuristic anyway, so approximating LRU is fine
  - Keep a “***accessed***” ***bit per page***, cleared occasionally by the operating system. Then pick any “unaccessed” page to evict

# Other VM Translation Functions

---

- ❑ Page data location
  - Physical memory, disk, uninitialized data
- ❑ Access permissions
  - Read only pages for instructions
- ❑ Gathering access information
  - Identifying dirty pages by tracking stores
  - Identifying accesses to help determine LRU candidate

# Performance of Virtual Memory

---

- ❑ To translate a virtual address into a physical address, we must first access the page table in physical memory
- ❑ Then we access physical memory again to get the data
  - A load instruction performs at least 2 memory reads
  - A store instruction performs at least 1 read and then a write
- ❑ Above lookup steps are Slow

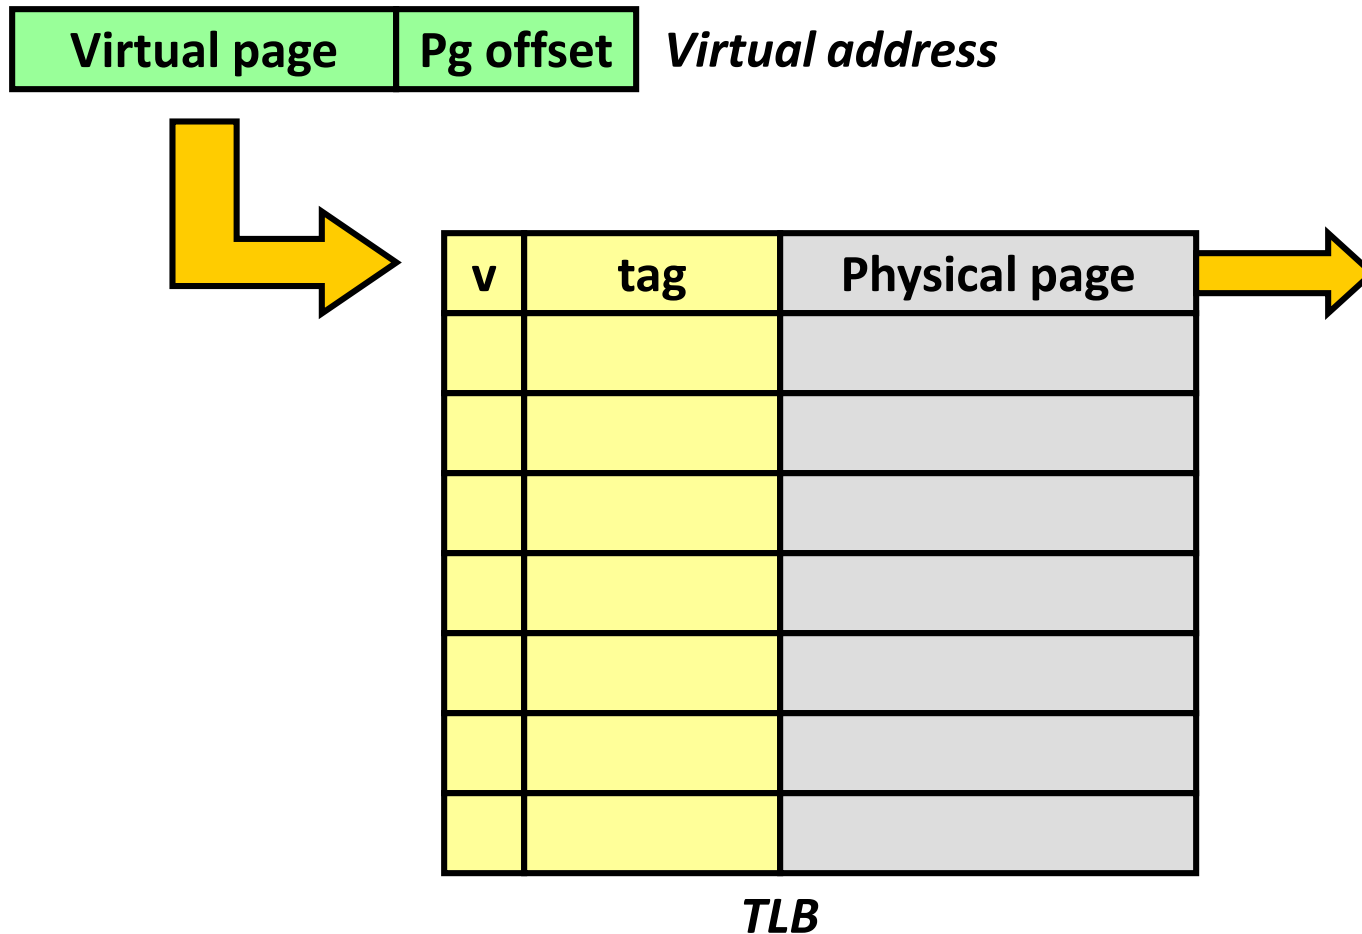
# Translation look-aside buffer (TLB)

---

- ❑ We fix this performance problem by avoiding main memory in the translation from virtual to physical pages.
- ❑ Buffer common translations in a **Translation Look-aside Buffer (TLB)**, a fast cache memory dedicated to storing a small subset of valid V-to-P translations.
- ❑ 16-512 entries common.
- ❑ Generally has low miss rate ( $< 1\%$ ).

# Translation look-aside buffer (TLB)

---





# Where is the TLB lookup?

---

- ❑ We put the TLB lookup in the pipeline after the virtual address is calculated and before the memory reference is performed.
  - This may be before or during the data cache access.
  - In case of a TLB miss, we need to perform the virtual to physical address translation during the memory stage of the pipeline.

# Putting it all together

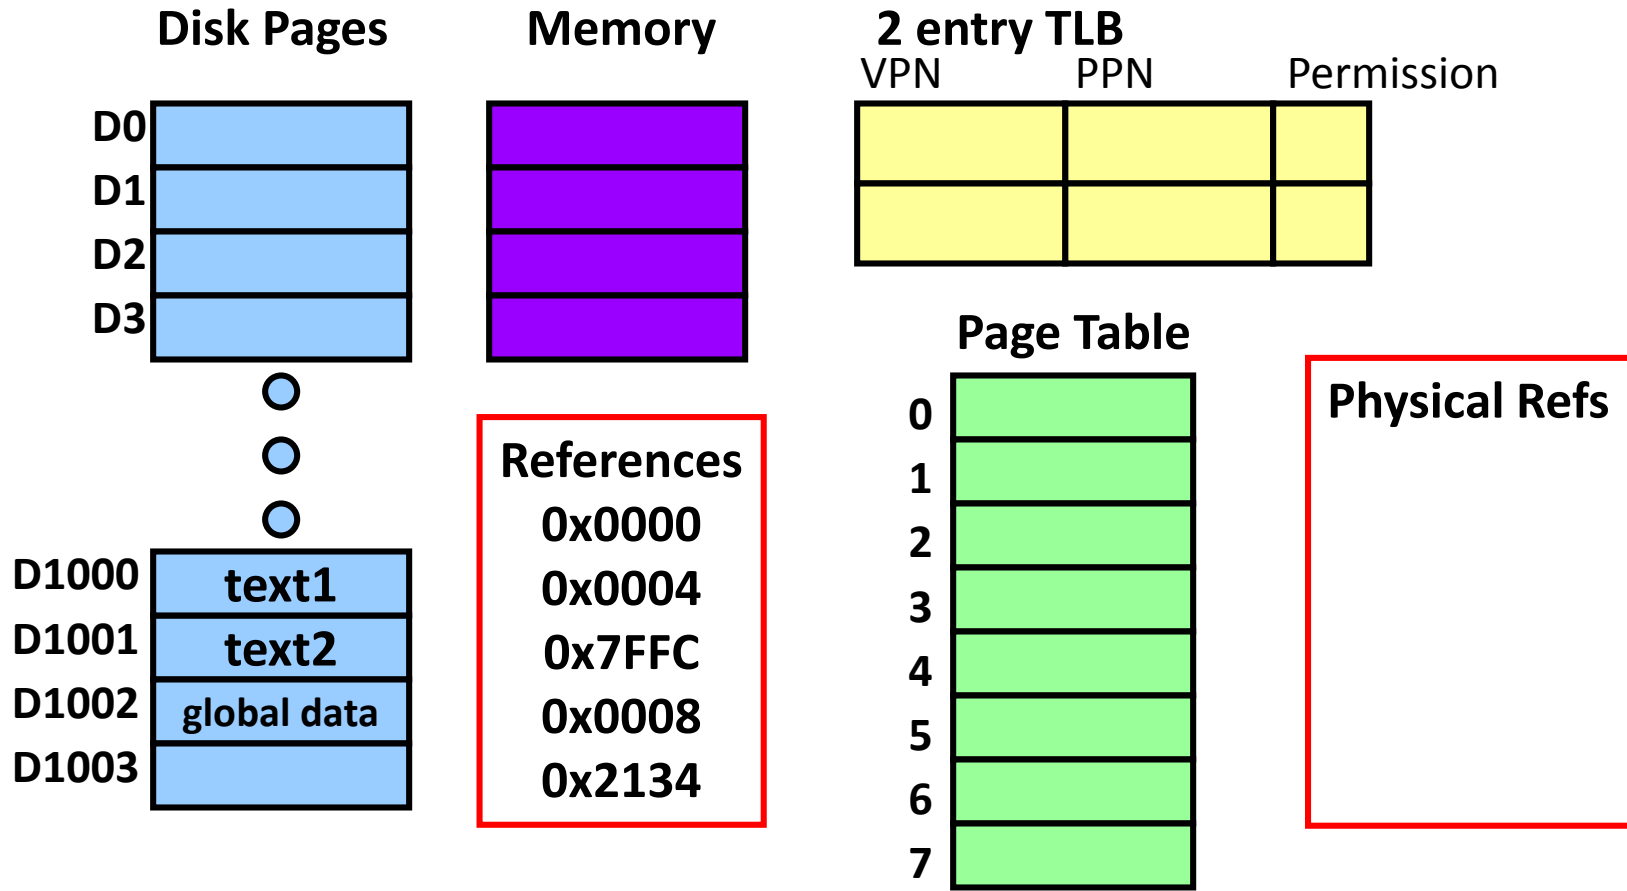
---

## ❑ Loading your program in memory

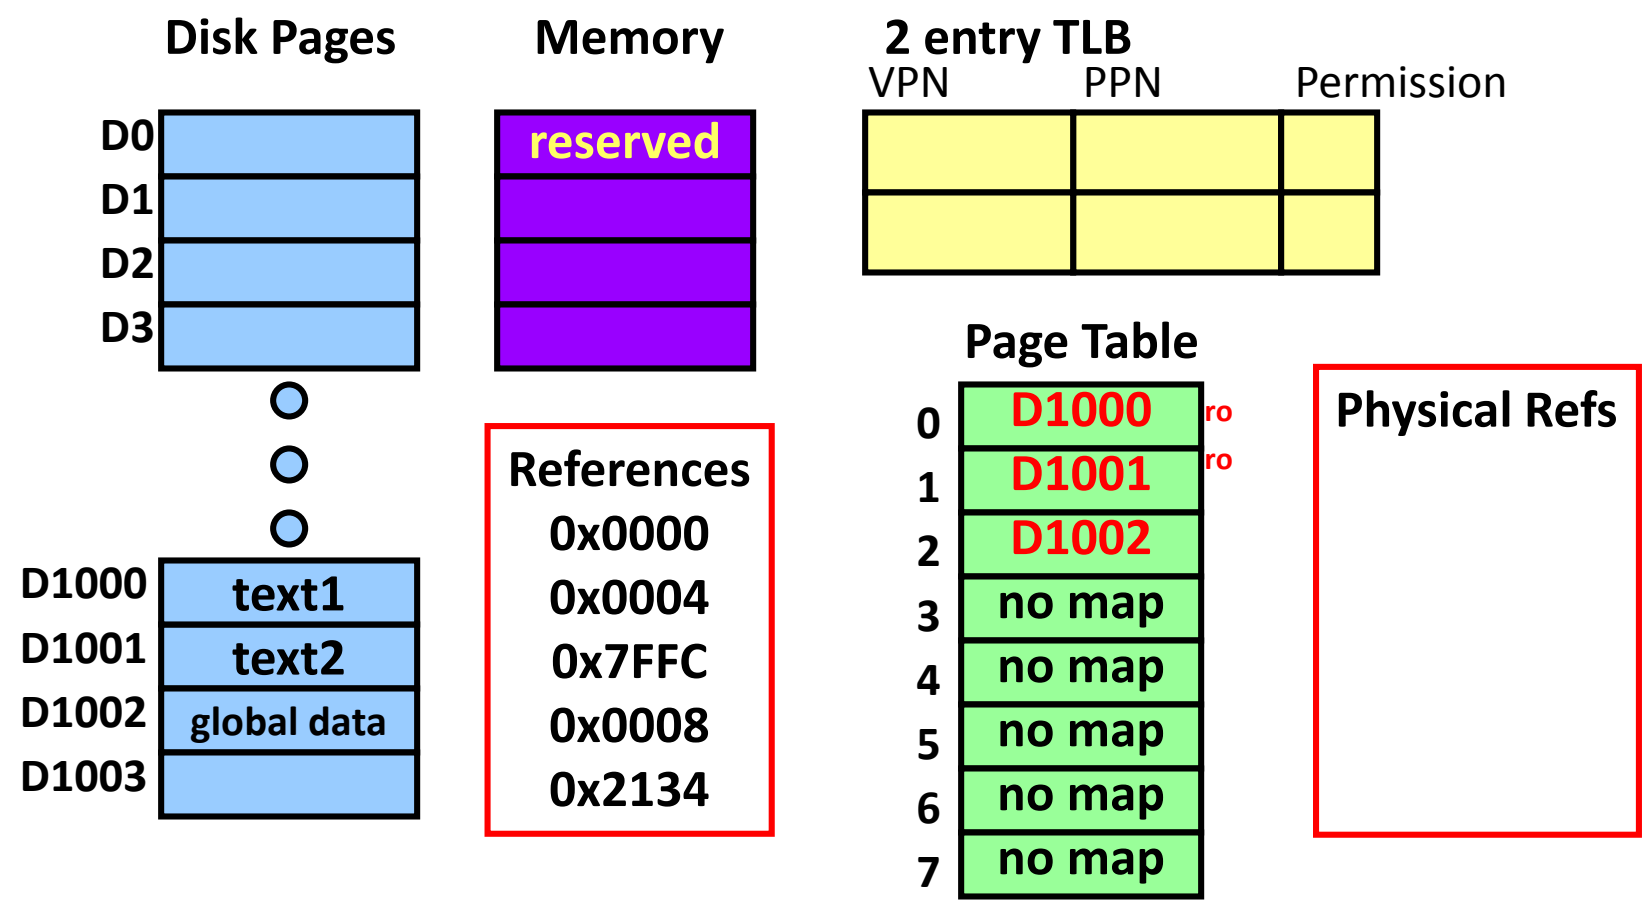
- Ask operating system to create a new process
- Construct a page table for this process
- Mark all page table entries as invalid with a pointer to the disk image of the program
  - That is, point to the executable file containing the binary.
- Run the program and get an immediate page fault on the first instruction.

# Loading a program into memory

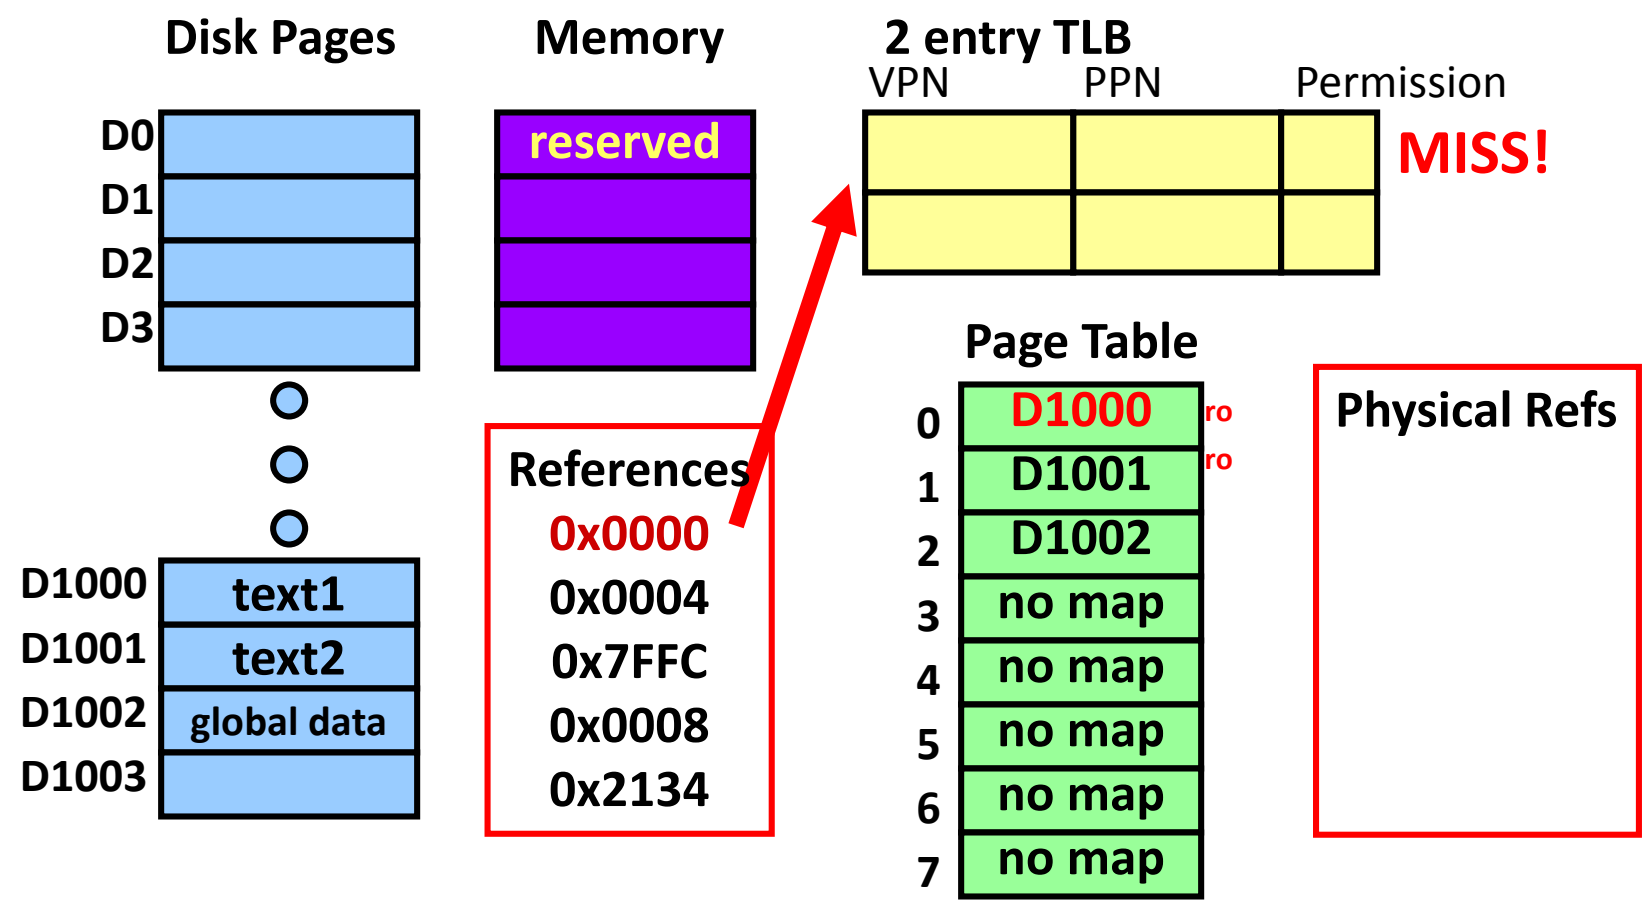
- ❑ Page size = 4 KB, Page table entry size = 4 B
- ❑ Page table register points to physical address 0x0000



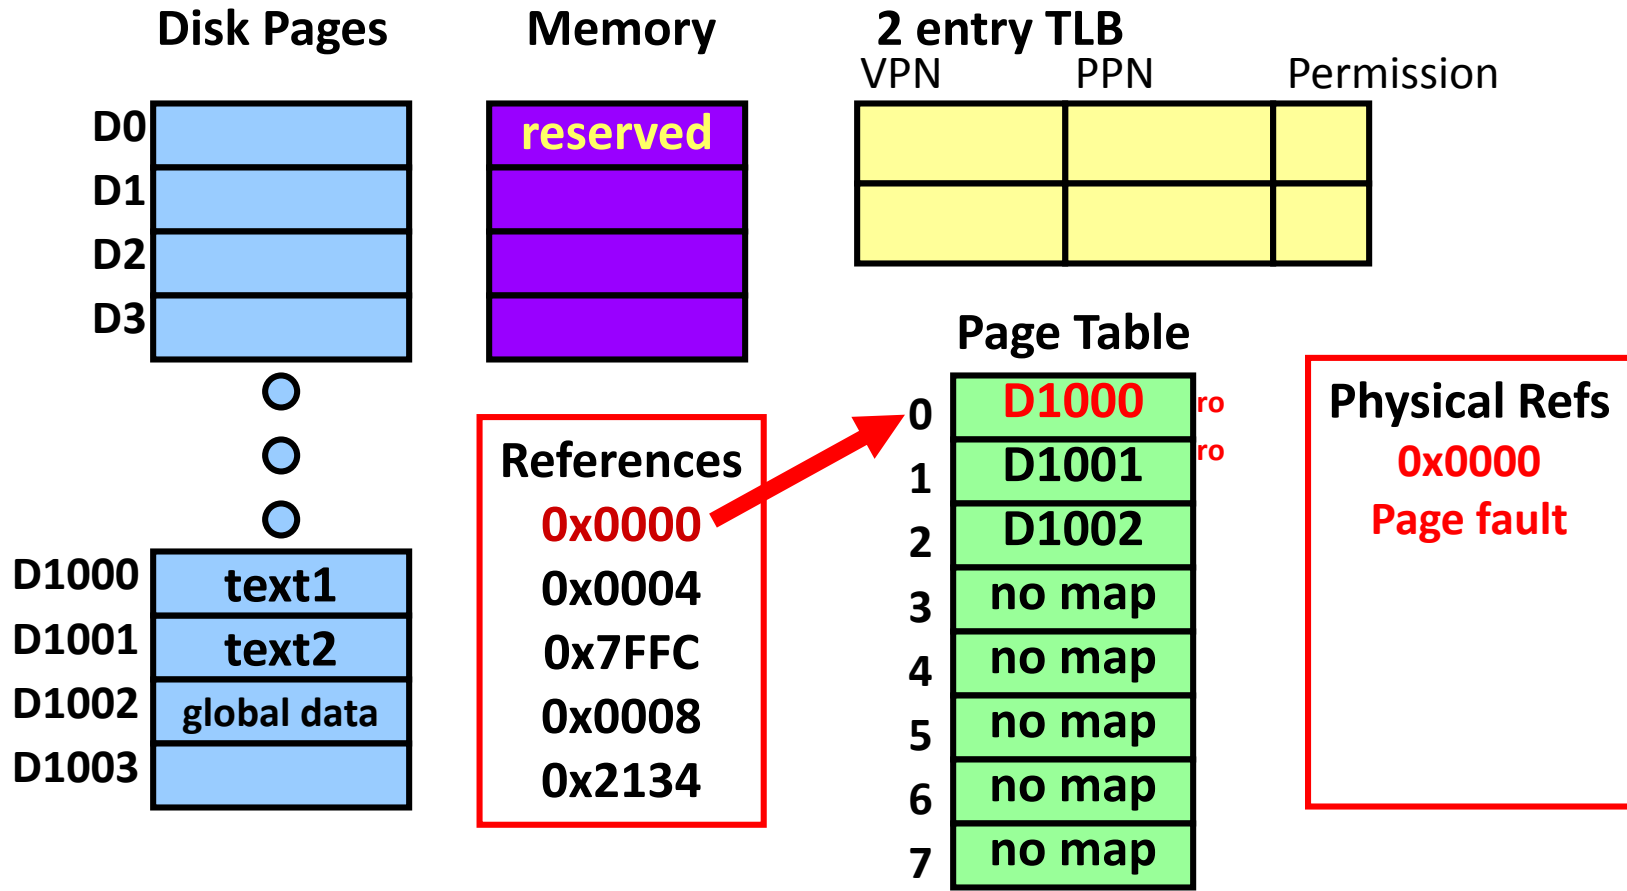
# Step 1: Read executable header & initialize page table



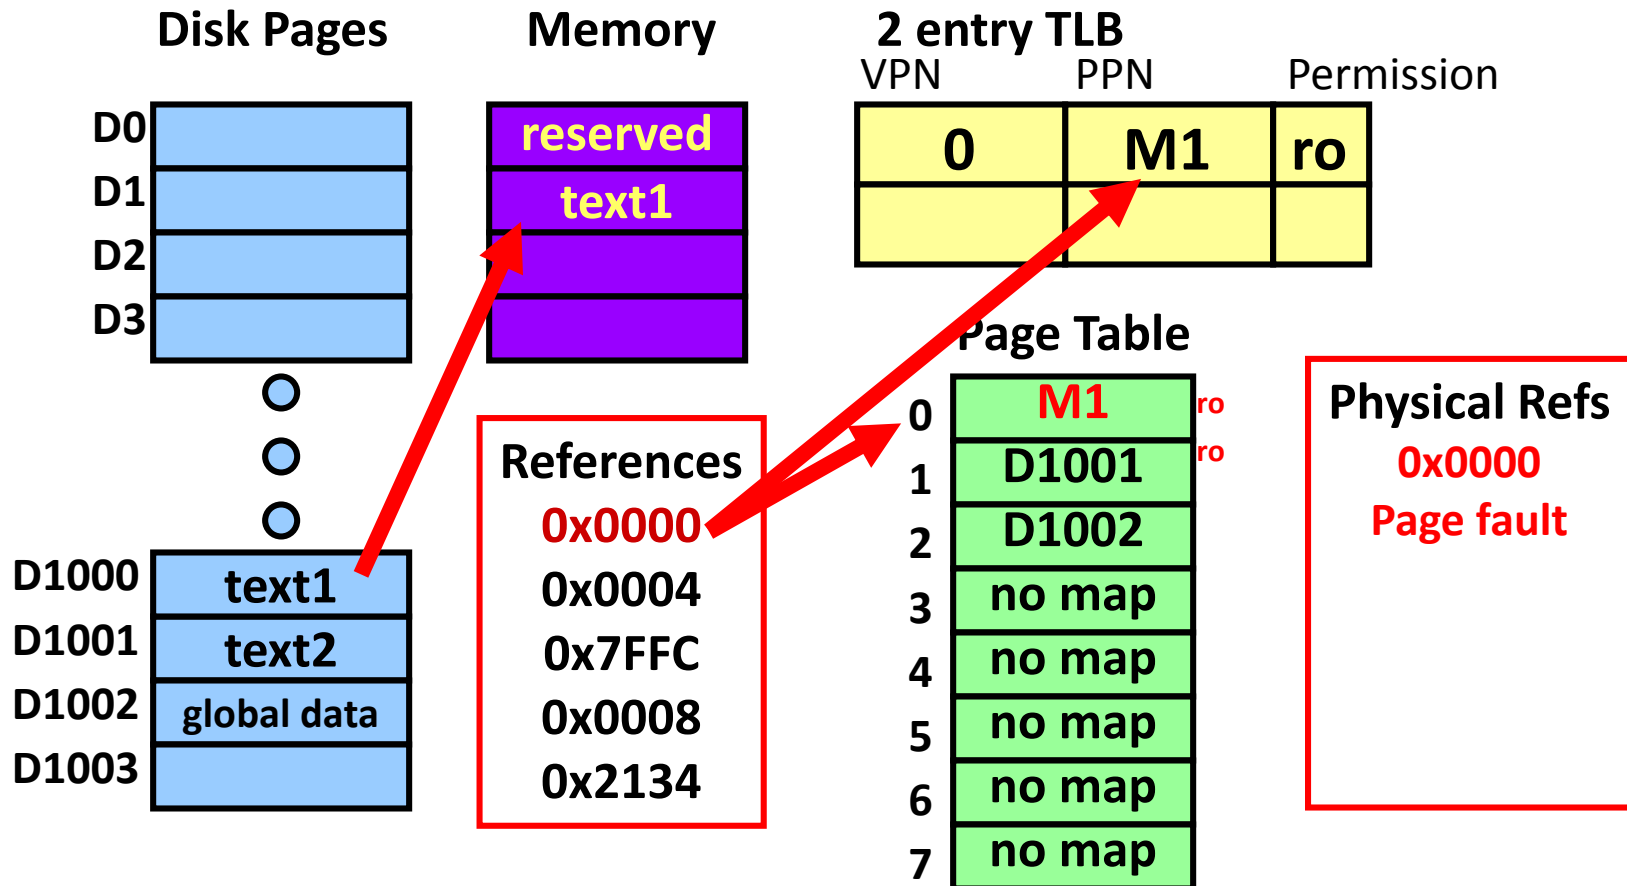
# Step 2: Load PC from header & start execution



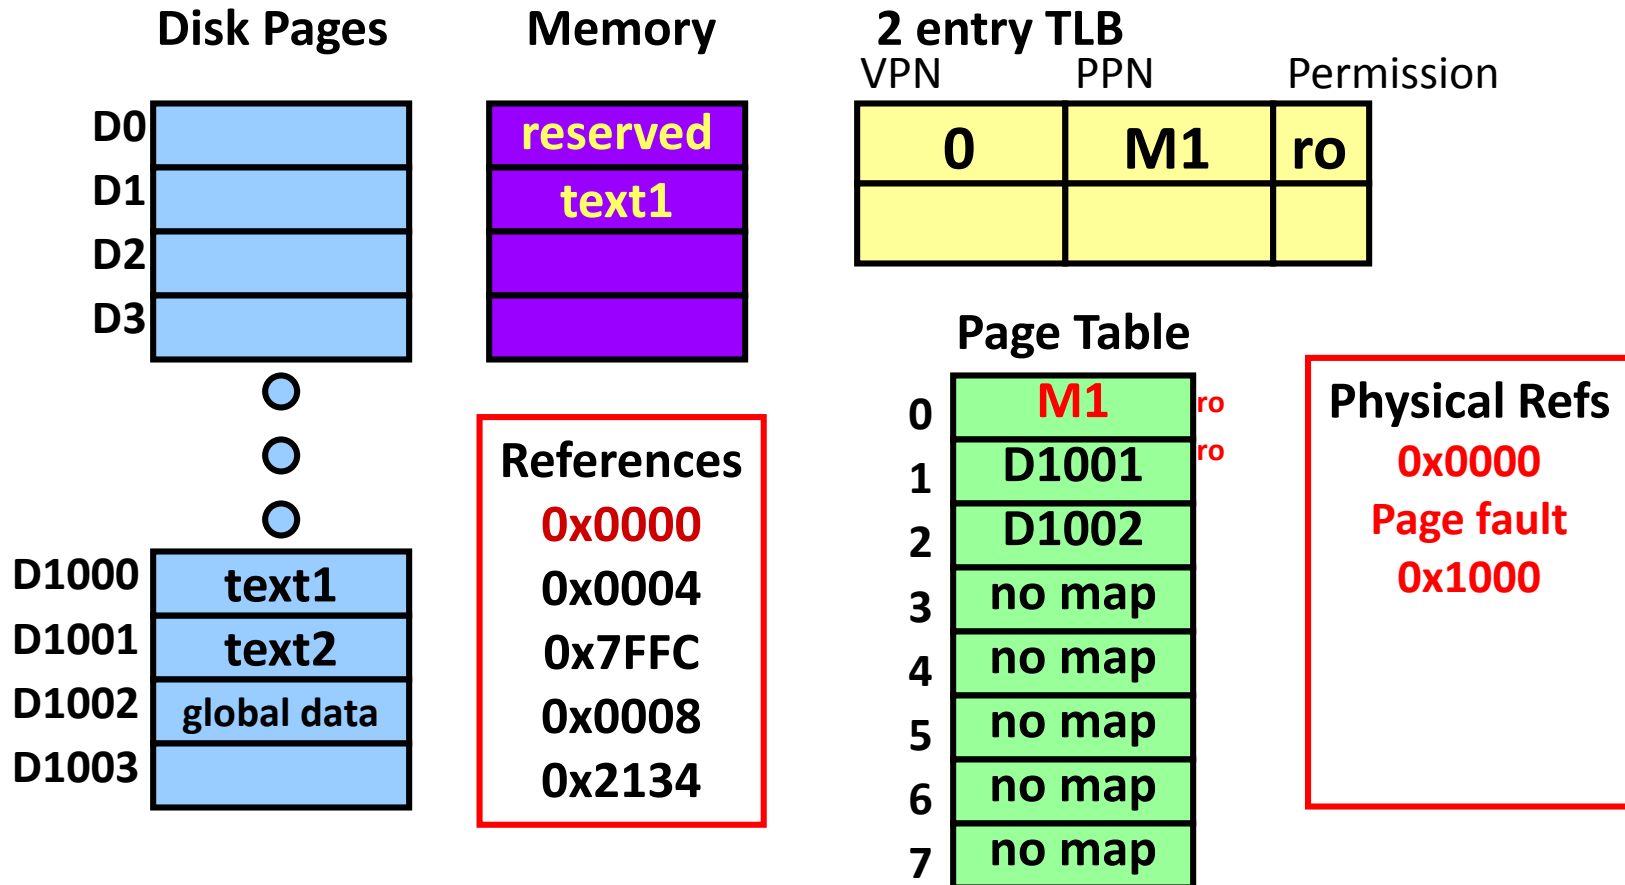
# Fetching instruction 0000



# Fetching instruction 0000

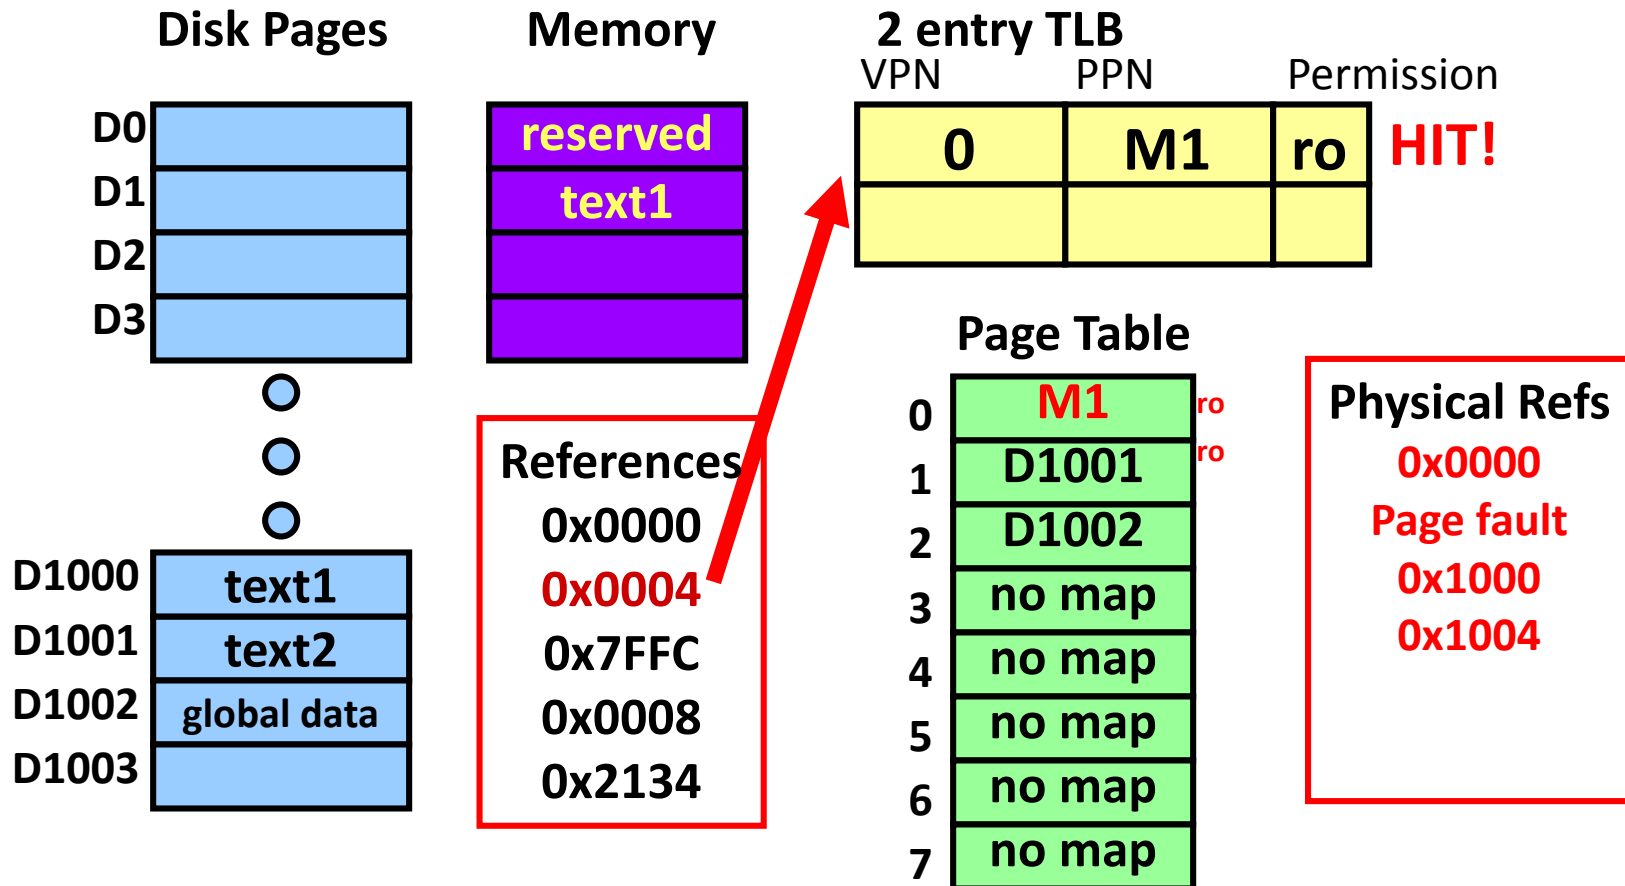


# Fetching instruction 0000

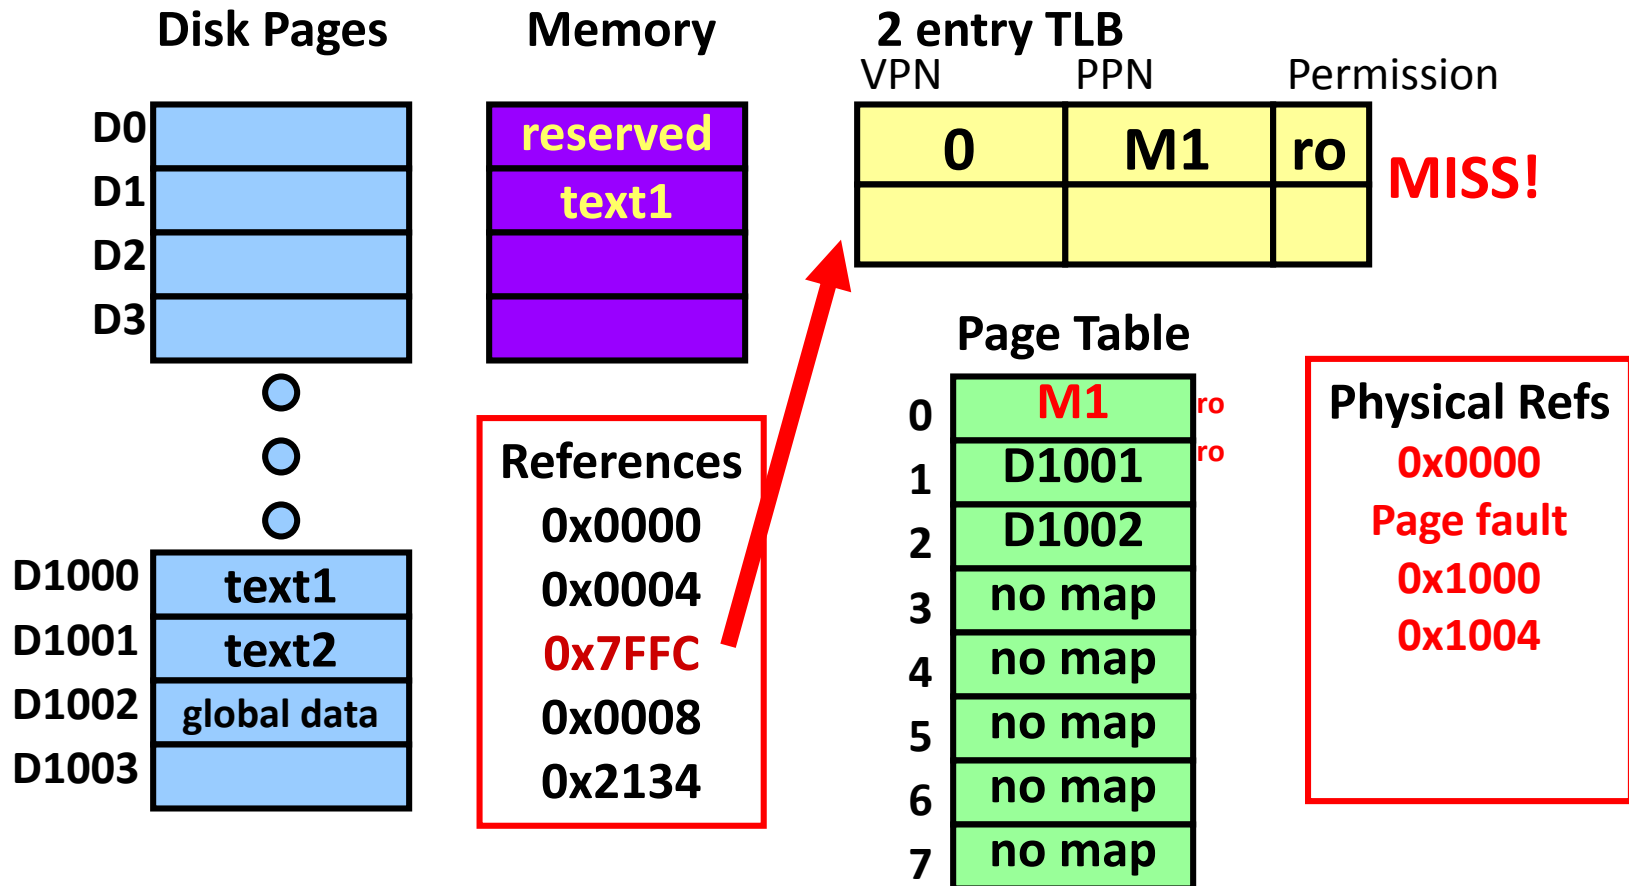




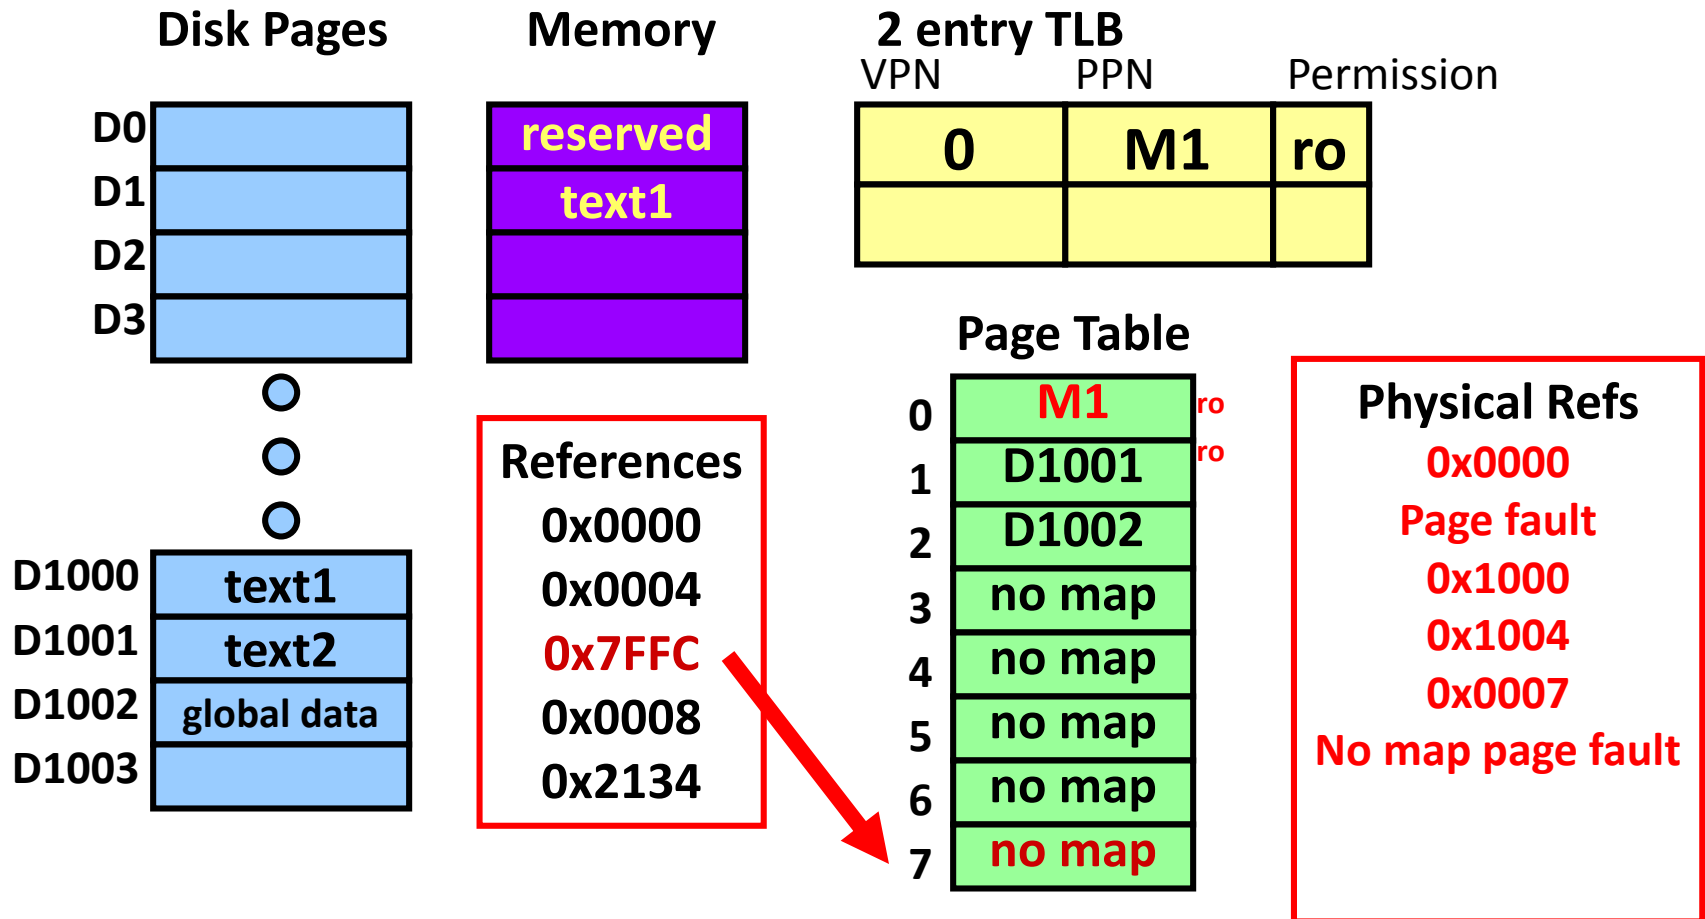
# Fetching instruction 0004



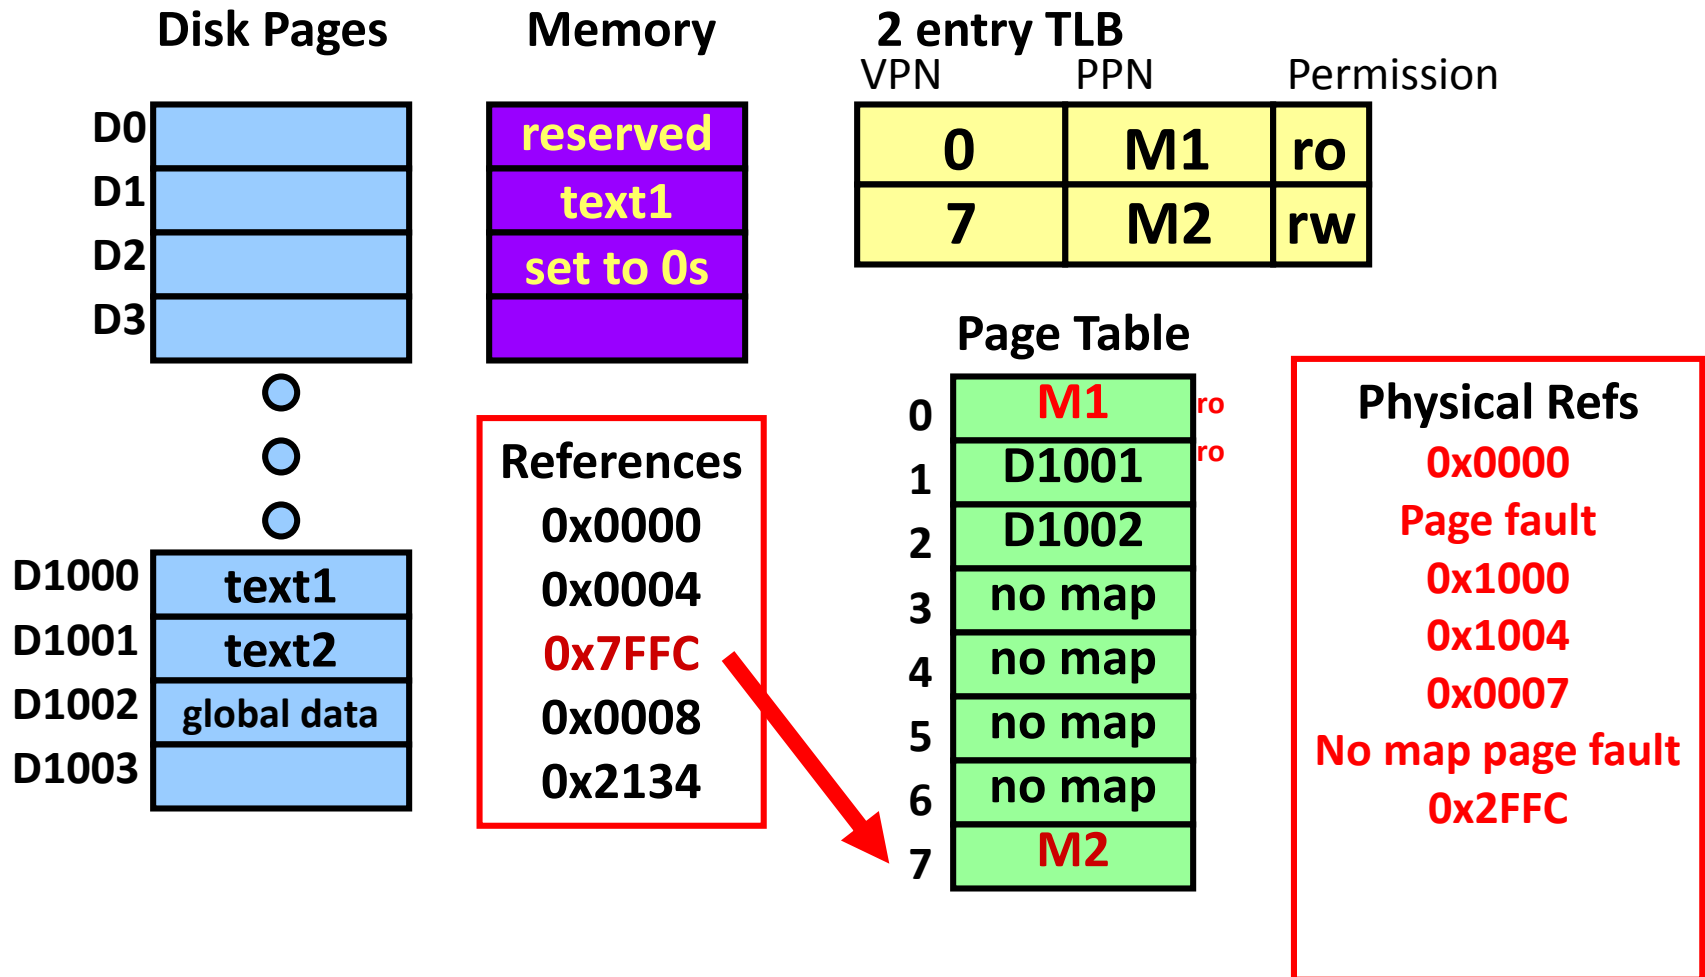
# Reference 7FFC



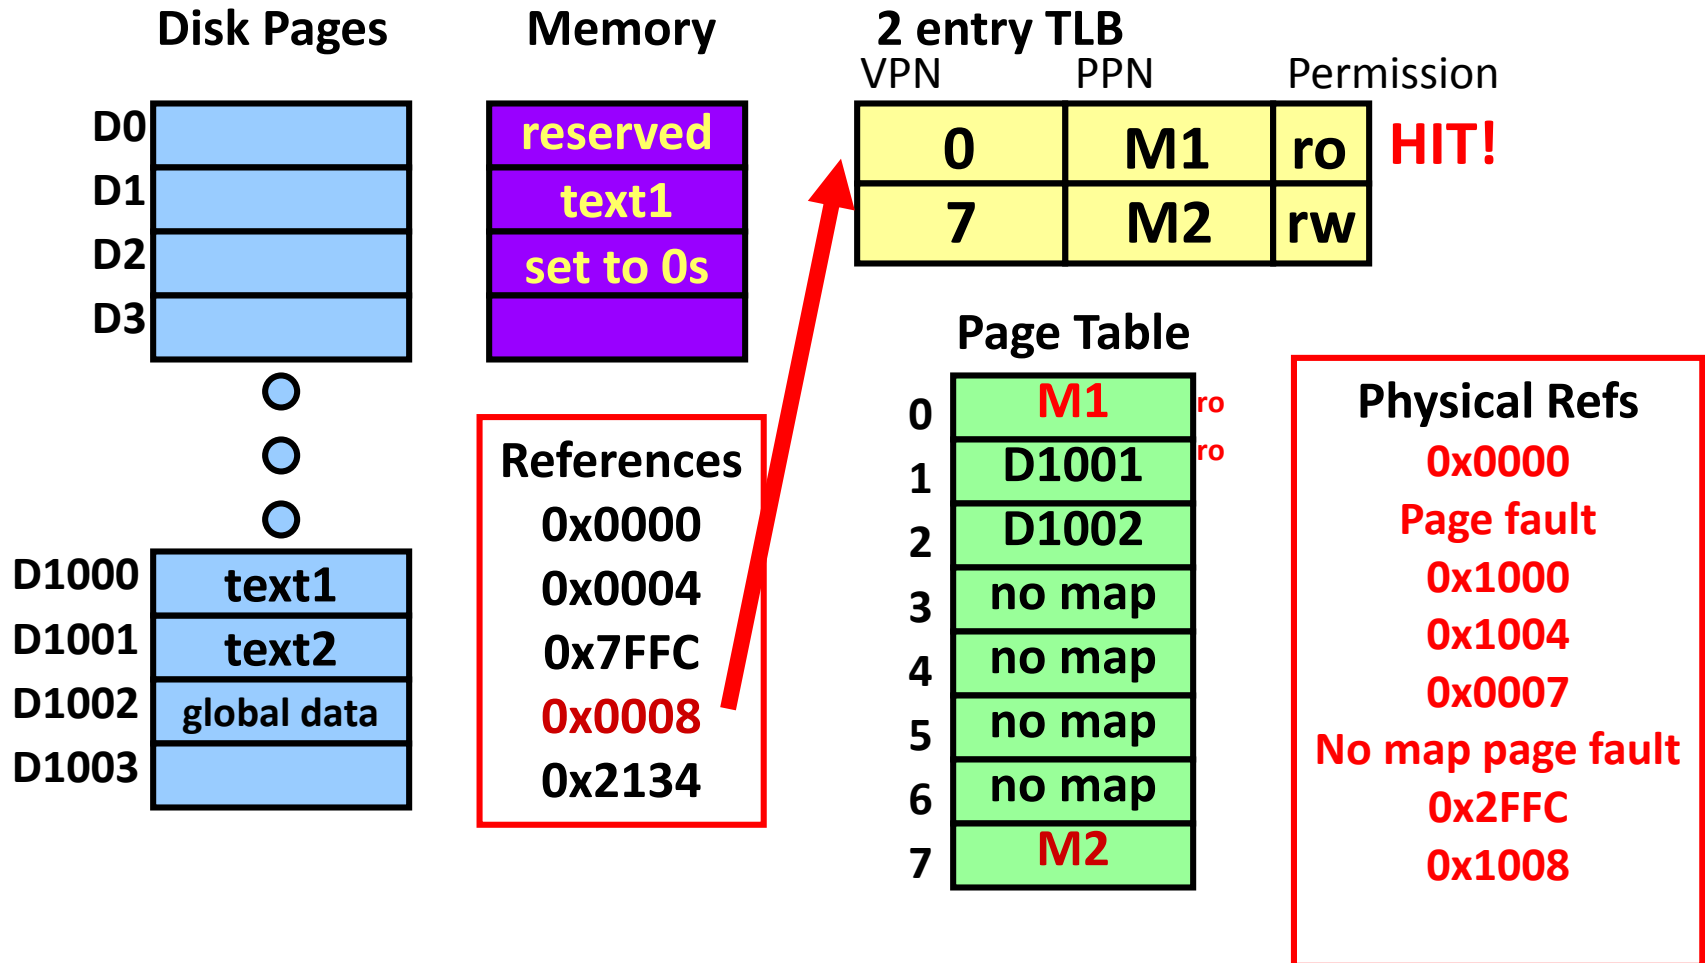
# Reference 7FFC



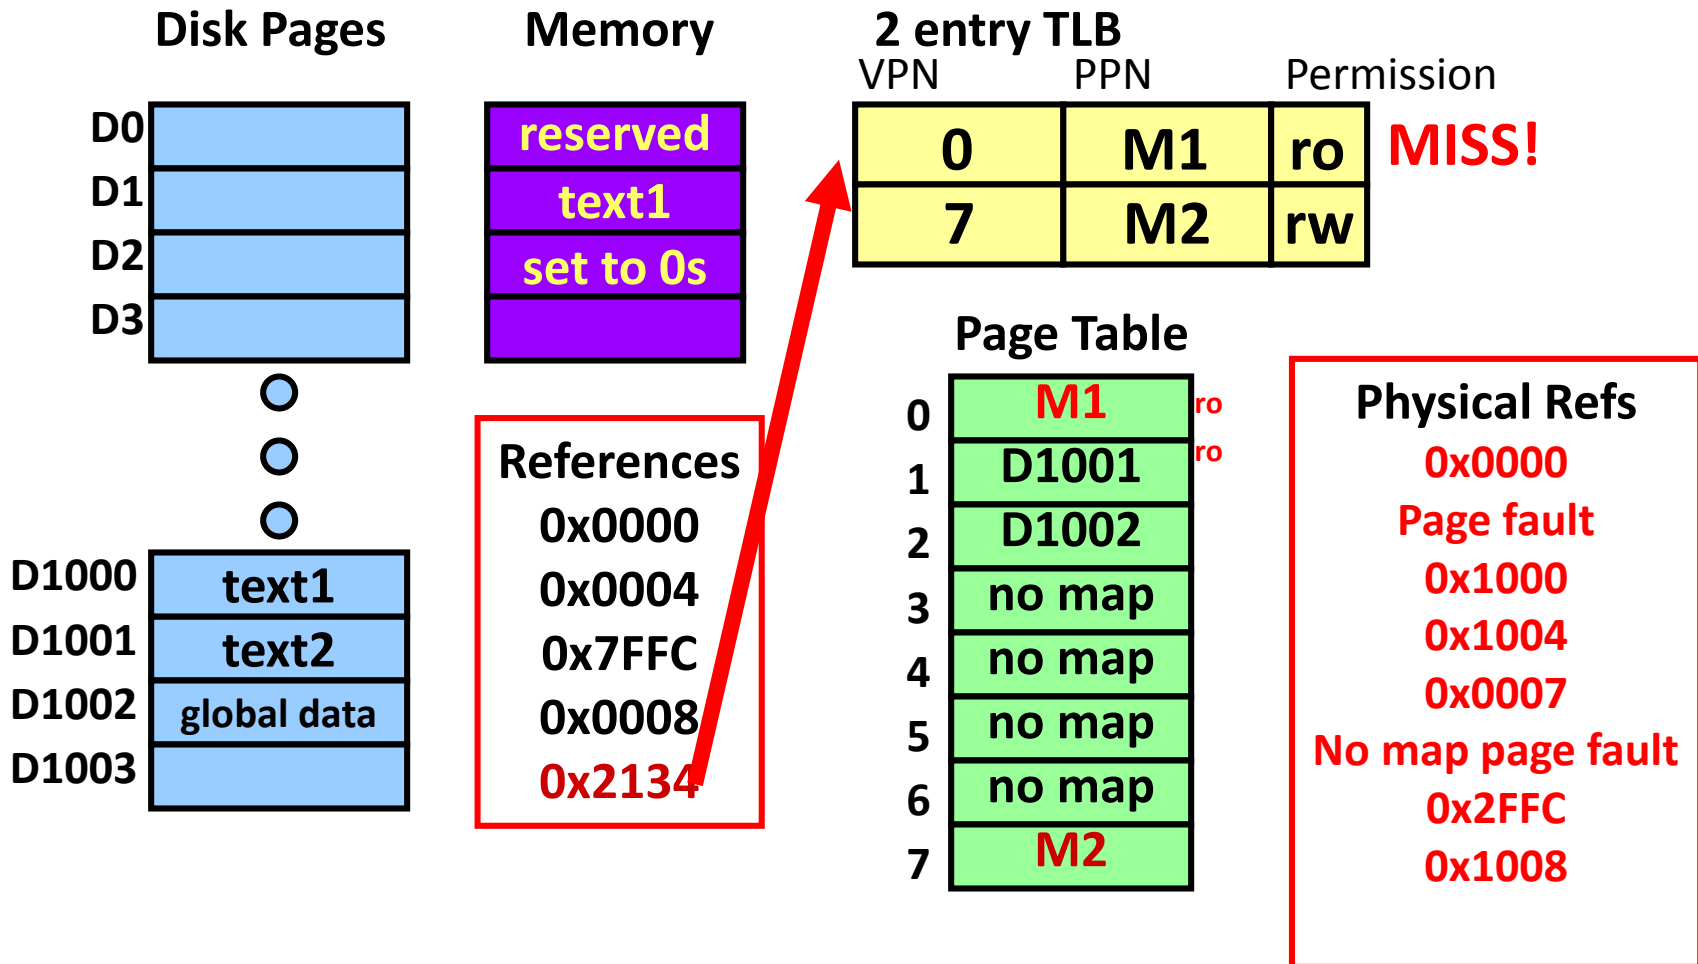
# Reference 7FFC



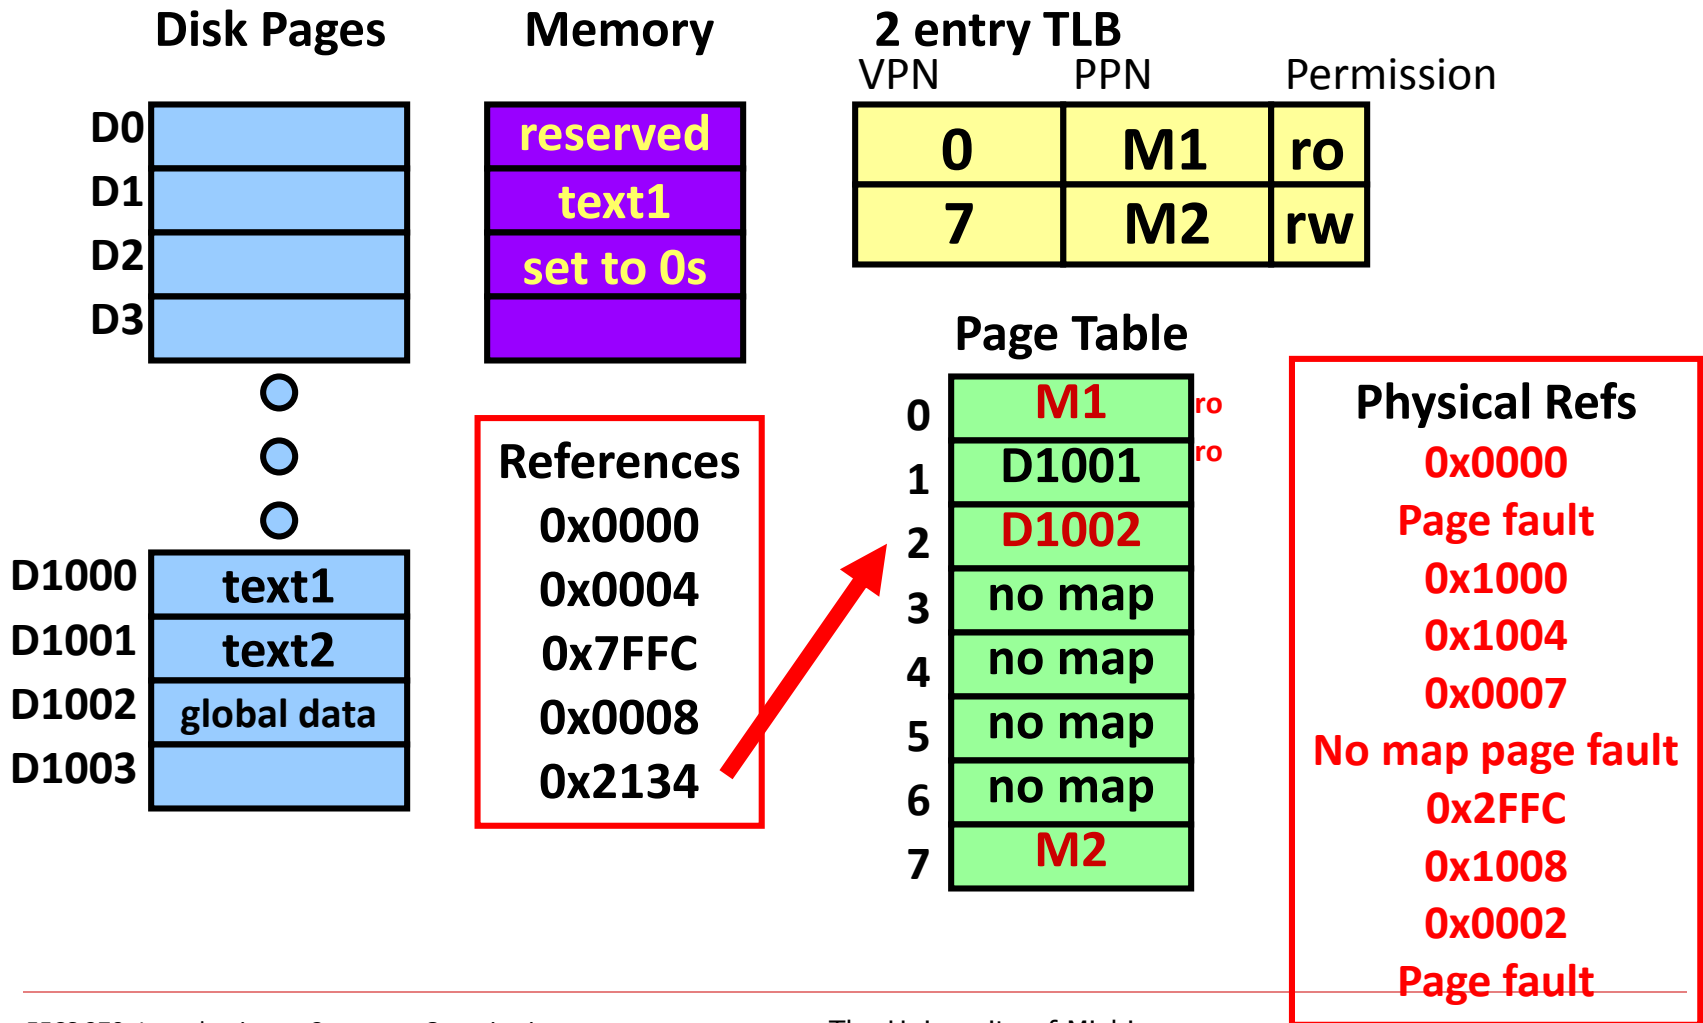
# Fetching instruction 0008



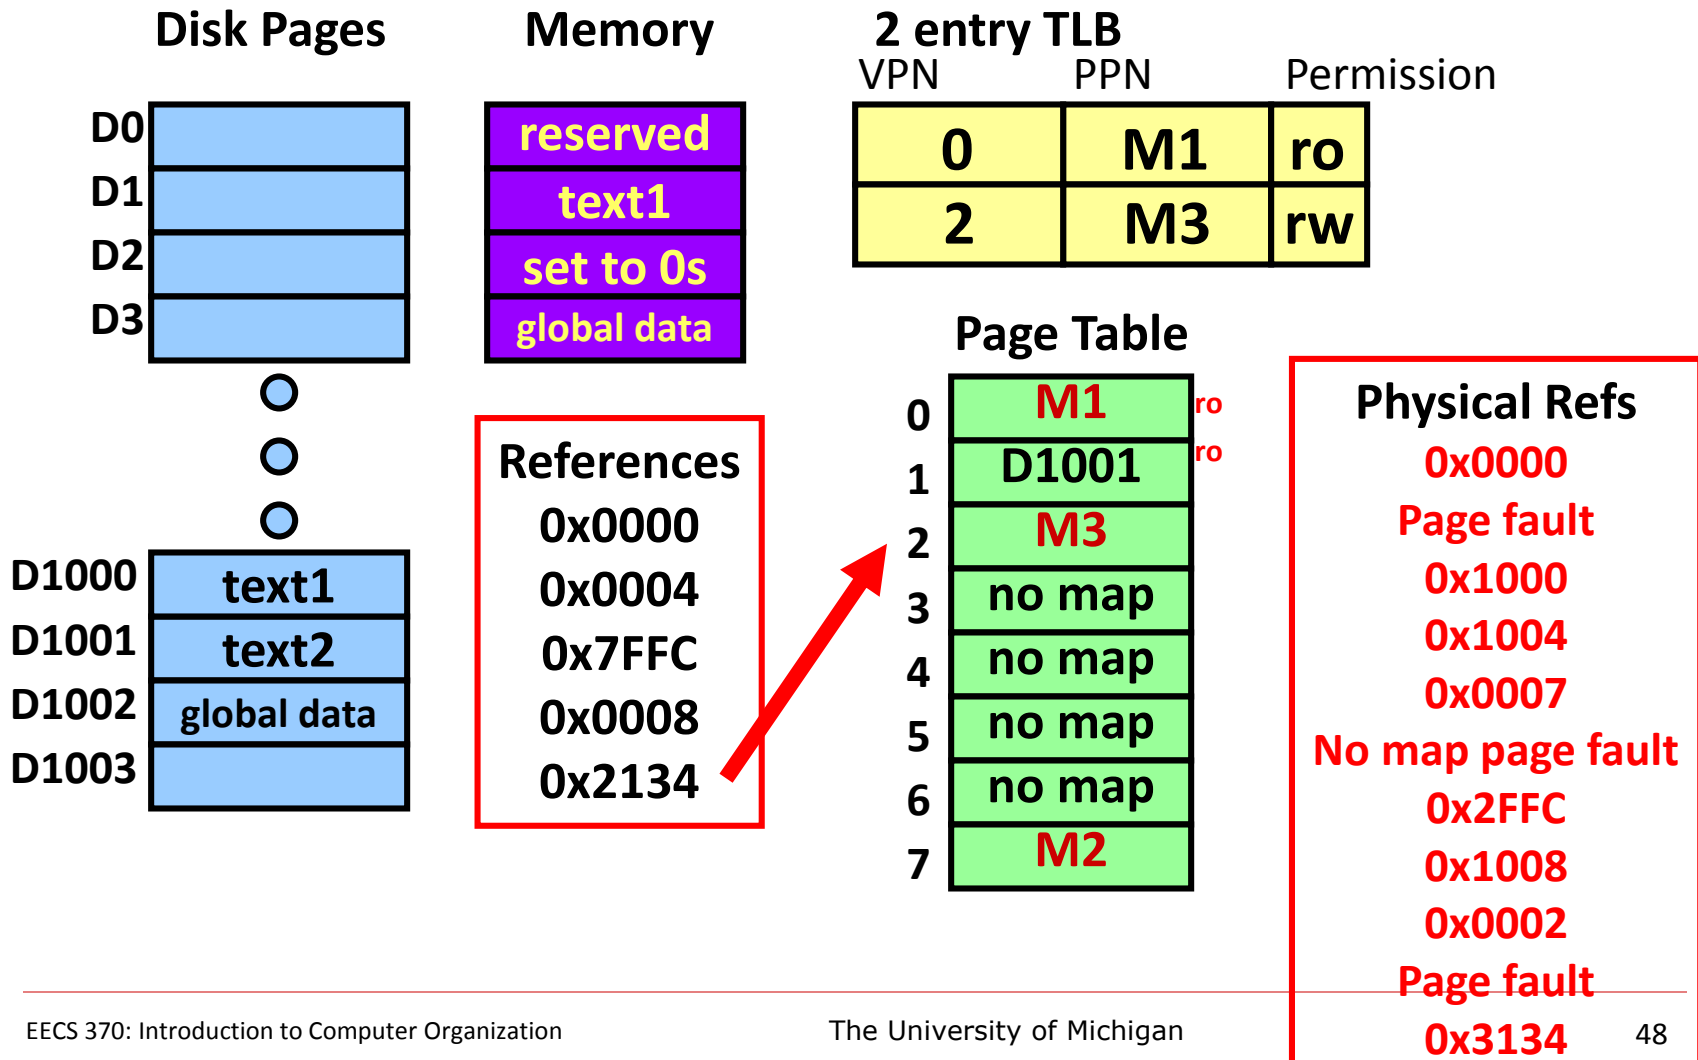
# Reference 2134



# Reference 2134



# Reference 2134





# Where is the TLB lookup?

---

- ❑ We put the TLB lookup in the pipeline after the virtual address is calculated and before the memory reference is performed
  - This may be *before* or *during* the data cache access
  - In case of a TLB miss, we need to perform the virtual to physical address translation during the memory stage of the pipeline.

# Next topic: Placing Caches in a VM System

---

- ❑ VM systems give us two different addresses:  
virtual and physical
- ❑ Which address should we use to access the data cache?
  - Physical address (after VM translations).
    - Delayed access.
  - Virtual address (before VM translation).
    - Faster access.
    - More complex.

# References (not part of Course Syllabus)

---

- ❑ See how Intel's memory management hardware works, Intel x86 Software Manual:  
<http://www.intel.com/content/www/us/en/architecture-and-technology/64-ia-32-architectures-software-developer-vol-3a-part-1-manual.html>
  - Chapter 4 is on Paging
- ❑ Linux page table management:  
<https://www.kernel.org/doc/gorman/html/understand/understand006.html>