

EECS 370 - Lecture 23

Making Virtual Memory
Fast + 370 Wrap-Up

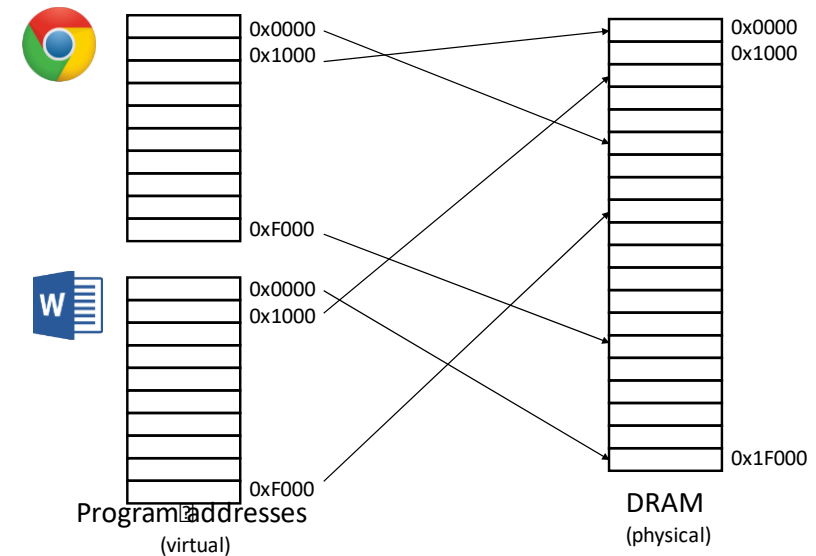
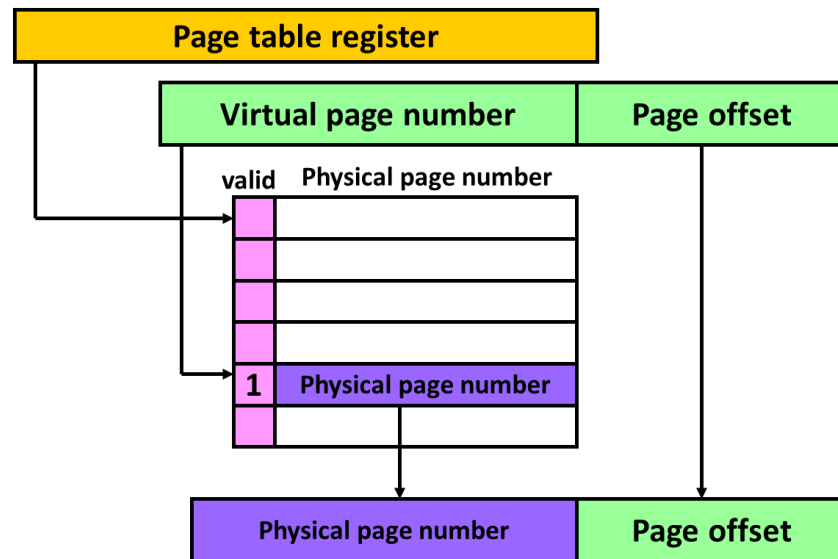


Announcements

- P4
 - Last project!
 - Due Thu (4/13)
- HW 6
 - Last homework!
 - Due Monday (4/17)
- Final exam
 - ...Last exam!
 - Thu (4/20) @ 10:30 am
 - Sample exams will be posted
 - (One sample exam given in homework)

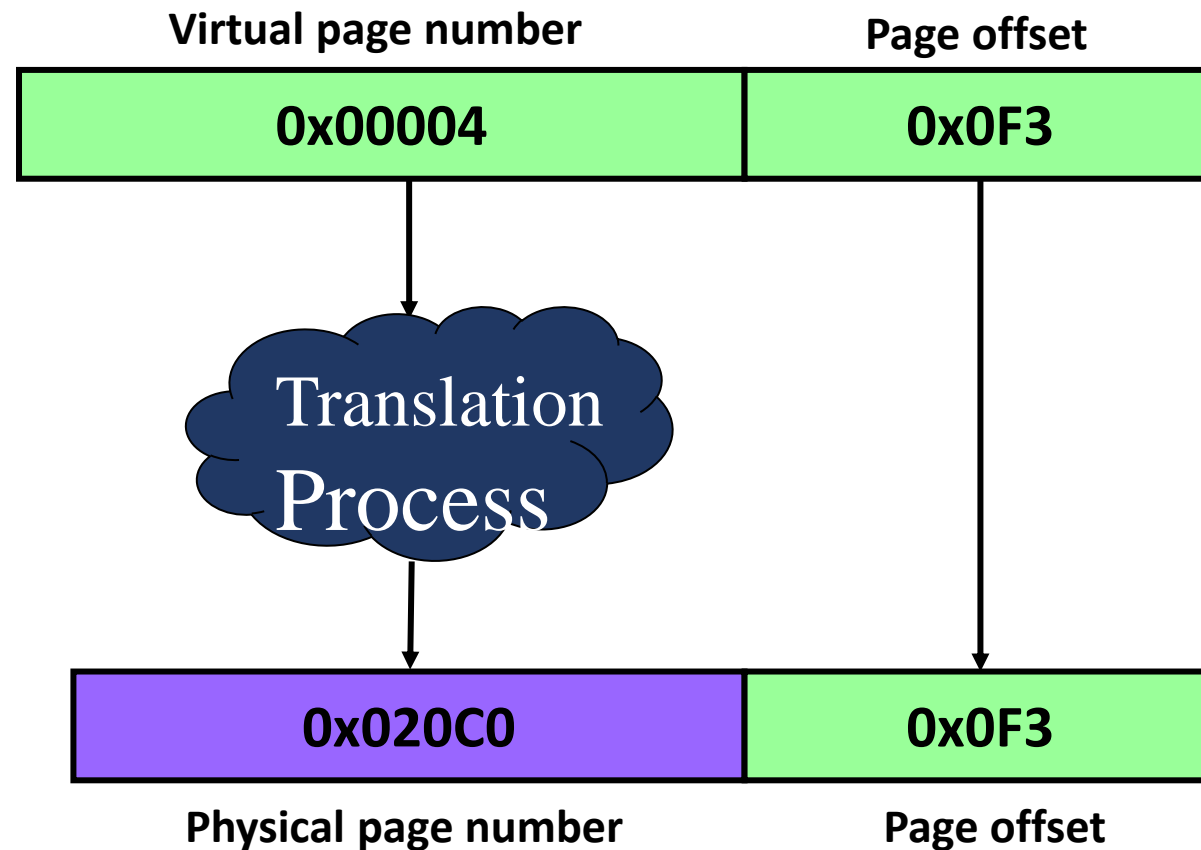
Reminder

- We use a “page table” to translate “virtual” addresses (specified by the programmer / assembler) to “physical” addresses (sent to memory)
 - Allows multiple programs to run without interfering with each other
 - Allows programs to use more storage than available in DRAM
 - By mapping virtual addresses to disk instead of DRAM



Review: Virtual to Physical Translation

Virtual address = 0x000040F3



Physical address = 0x020C00F3

Review

- To translate a virtual address into a physical address, we must first access the page table in physical memory
 - If it's an N-level page table, we must do N total loads before getting the physical page number
- Then we access physical memory again to get the data
 - A load instruction performs at least 2 memory reads
 - A store instruction performs at least 1 read and then a write
- Above lookups are **SLOW**

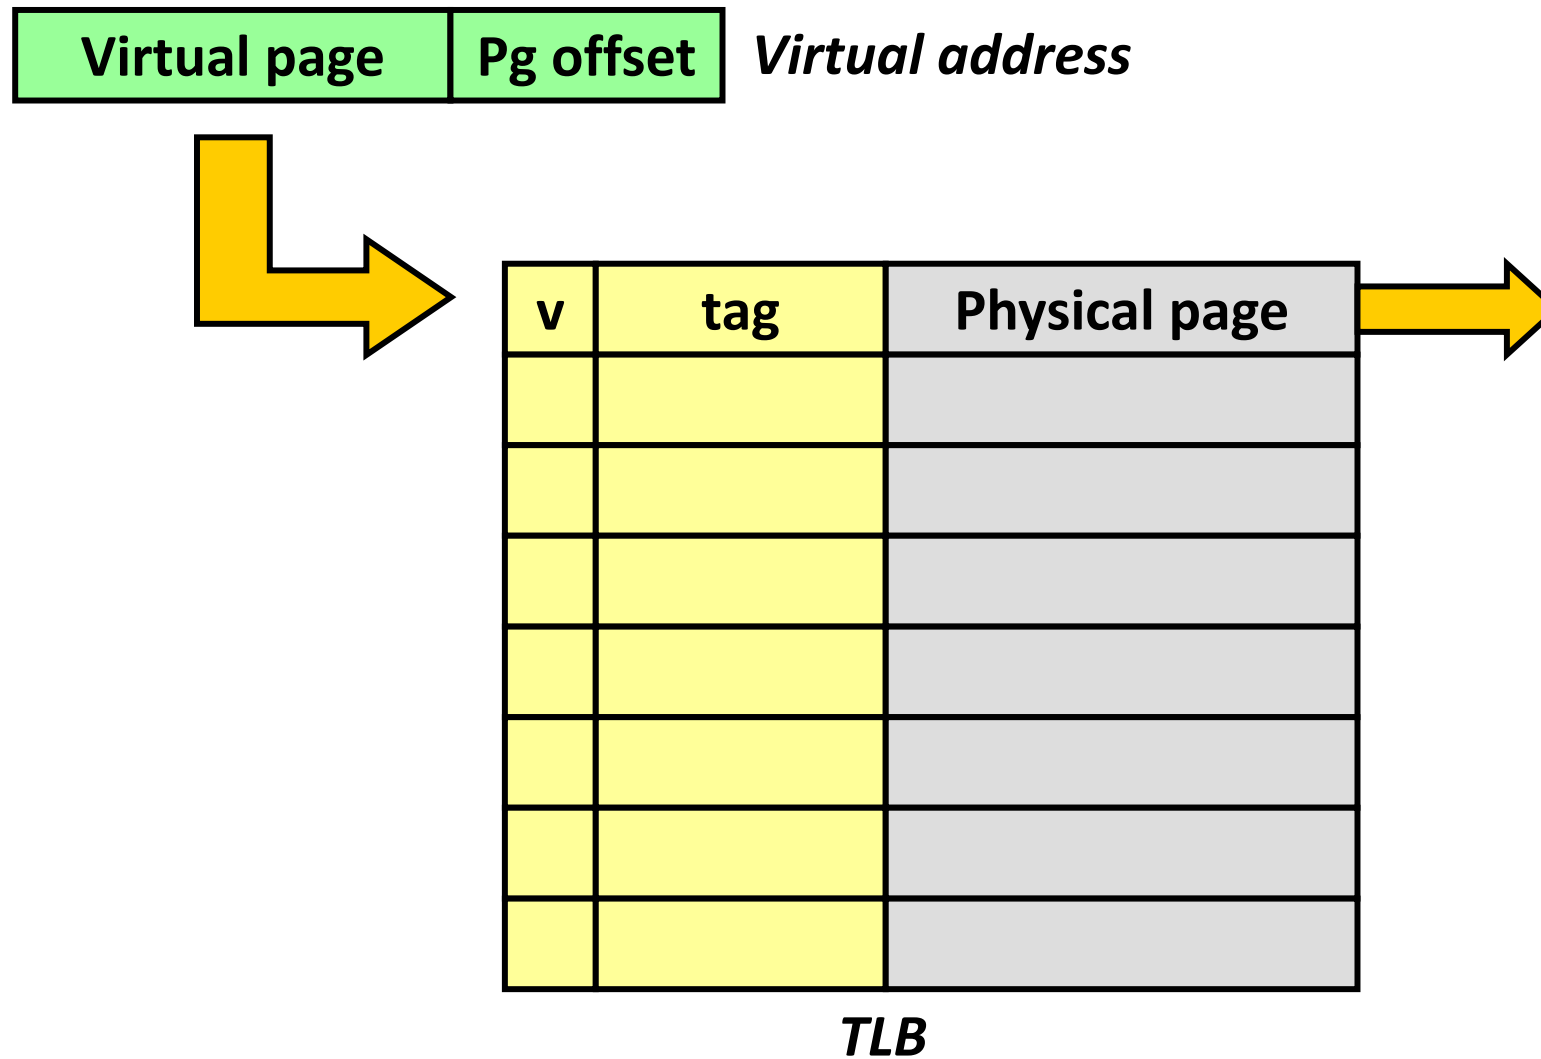
Translation look-aside buffer (TLB)

- We fix this performance problem by avoiding main memory in the translation from virtual to physical pages.
- Buffer common translations in a **Translation Look-aside Buffer (TLB)**, a fast cache memory dedicated to storing a small subset of valid V-to-P translations.
- 16-512 entries common.
- Generally has low miss rate ($< 1\%$).

Poll: Why can we make TLBs smaller than the cache hierarchy?

- a) Addresses are smaller than data
- b) TLB is accessed less frequently than caches
- c) Only need to store info about individual pages

Translation look-aside buffer (TLB)



Where is the TLB lookup?

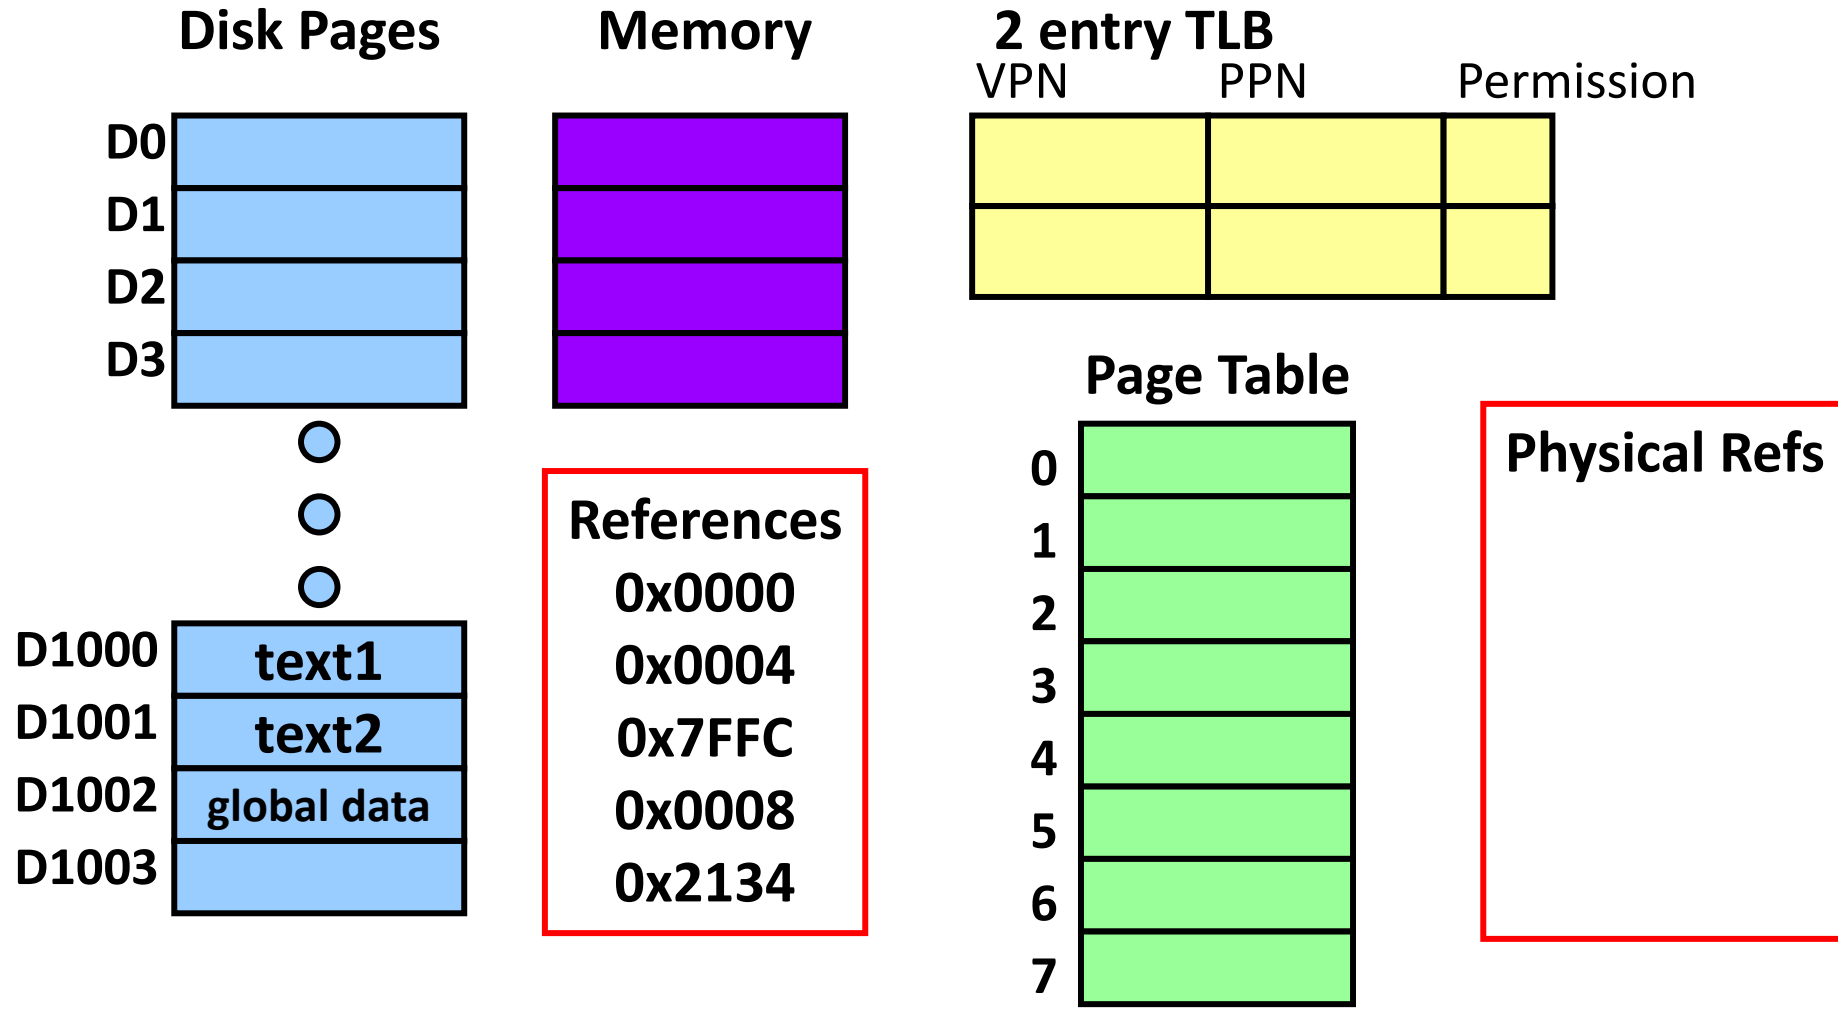
- We put the TLB lookup in the pipeline after the virtual address is calculated and before the memory reference is performed
 - This may be before or during the data cache access.
 - To be discussed next time
 - In case of a TLB miss, we need to perform the virtual to physical address translation during the memory stage of the pipeline.

Putting it all together

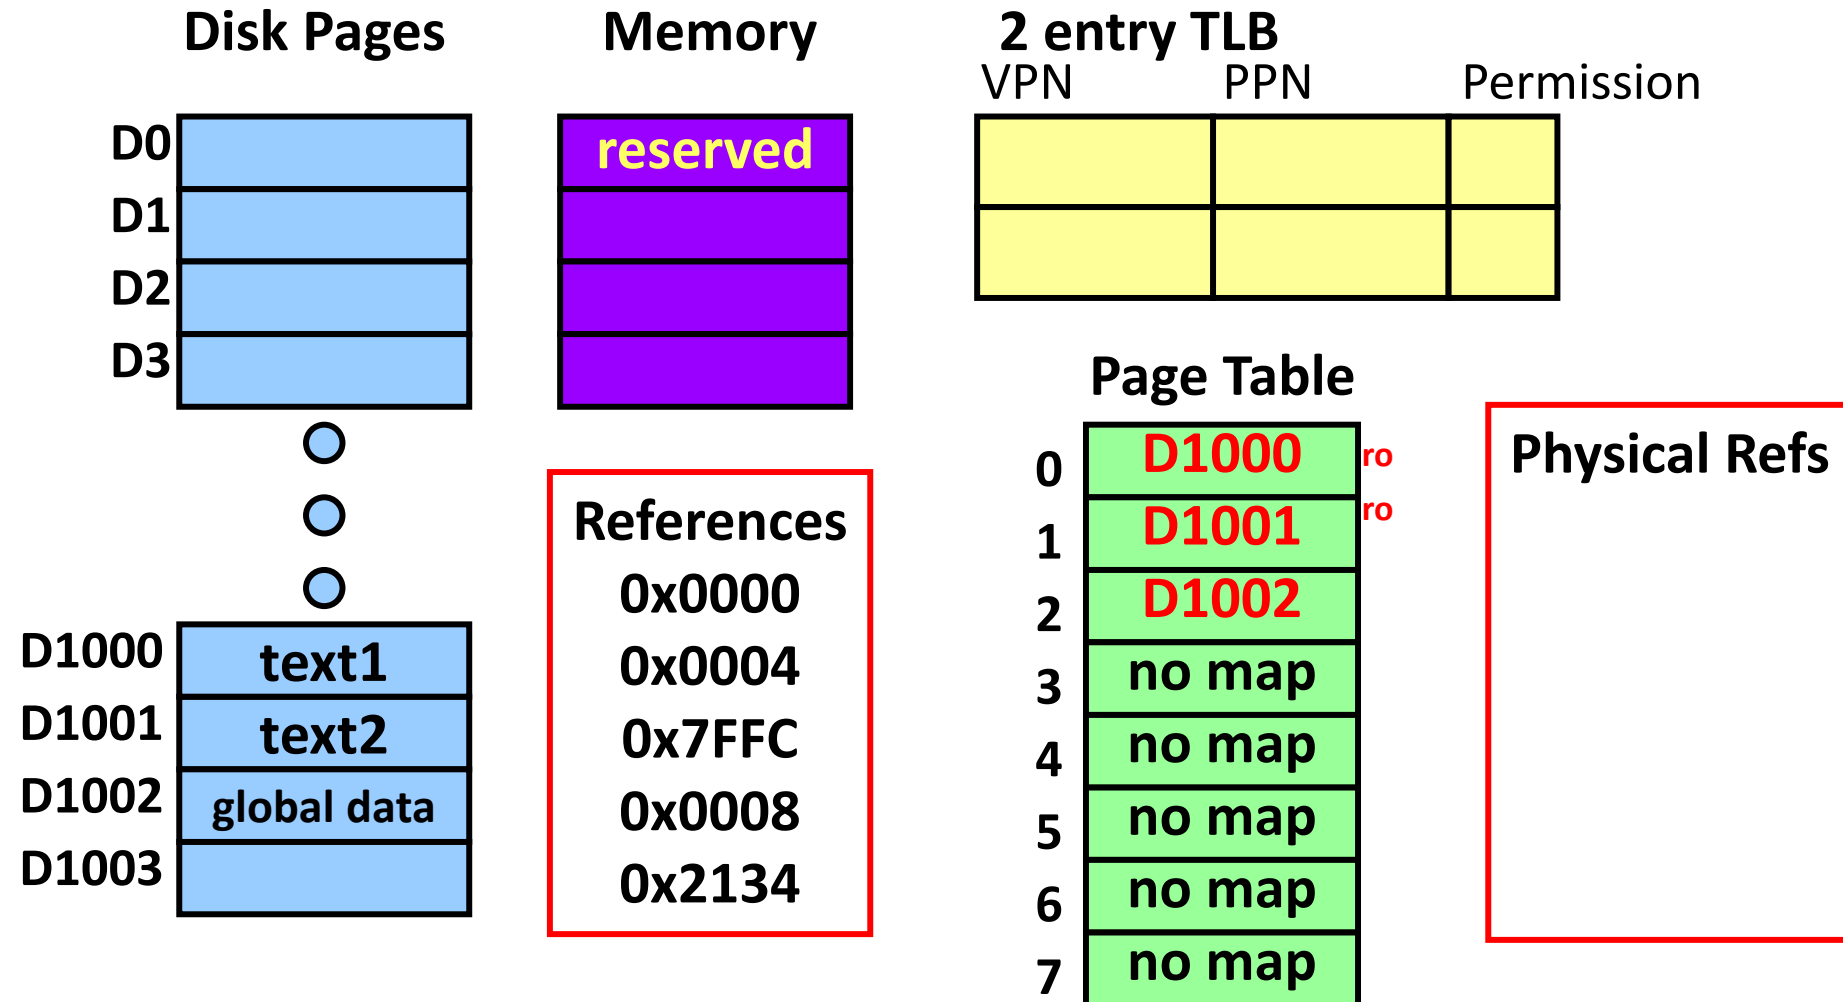
- Loading your program in memory
 - Ask operating system to create a new process
 - Construct a page table for this process
 - Mark all page table entries as invalid with a pointer to the disk image of the program
 - That is, point to the executable file containing the binary.
 - Run the program and get an immediate page fault on the first instruction.

Loading a program into memory

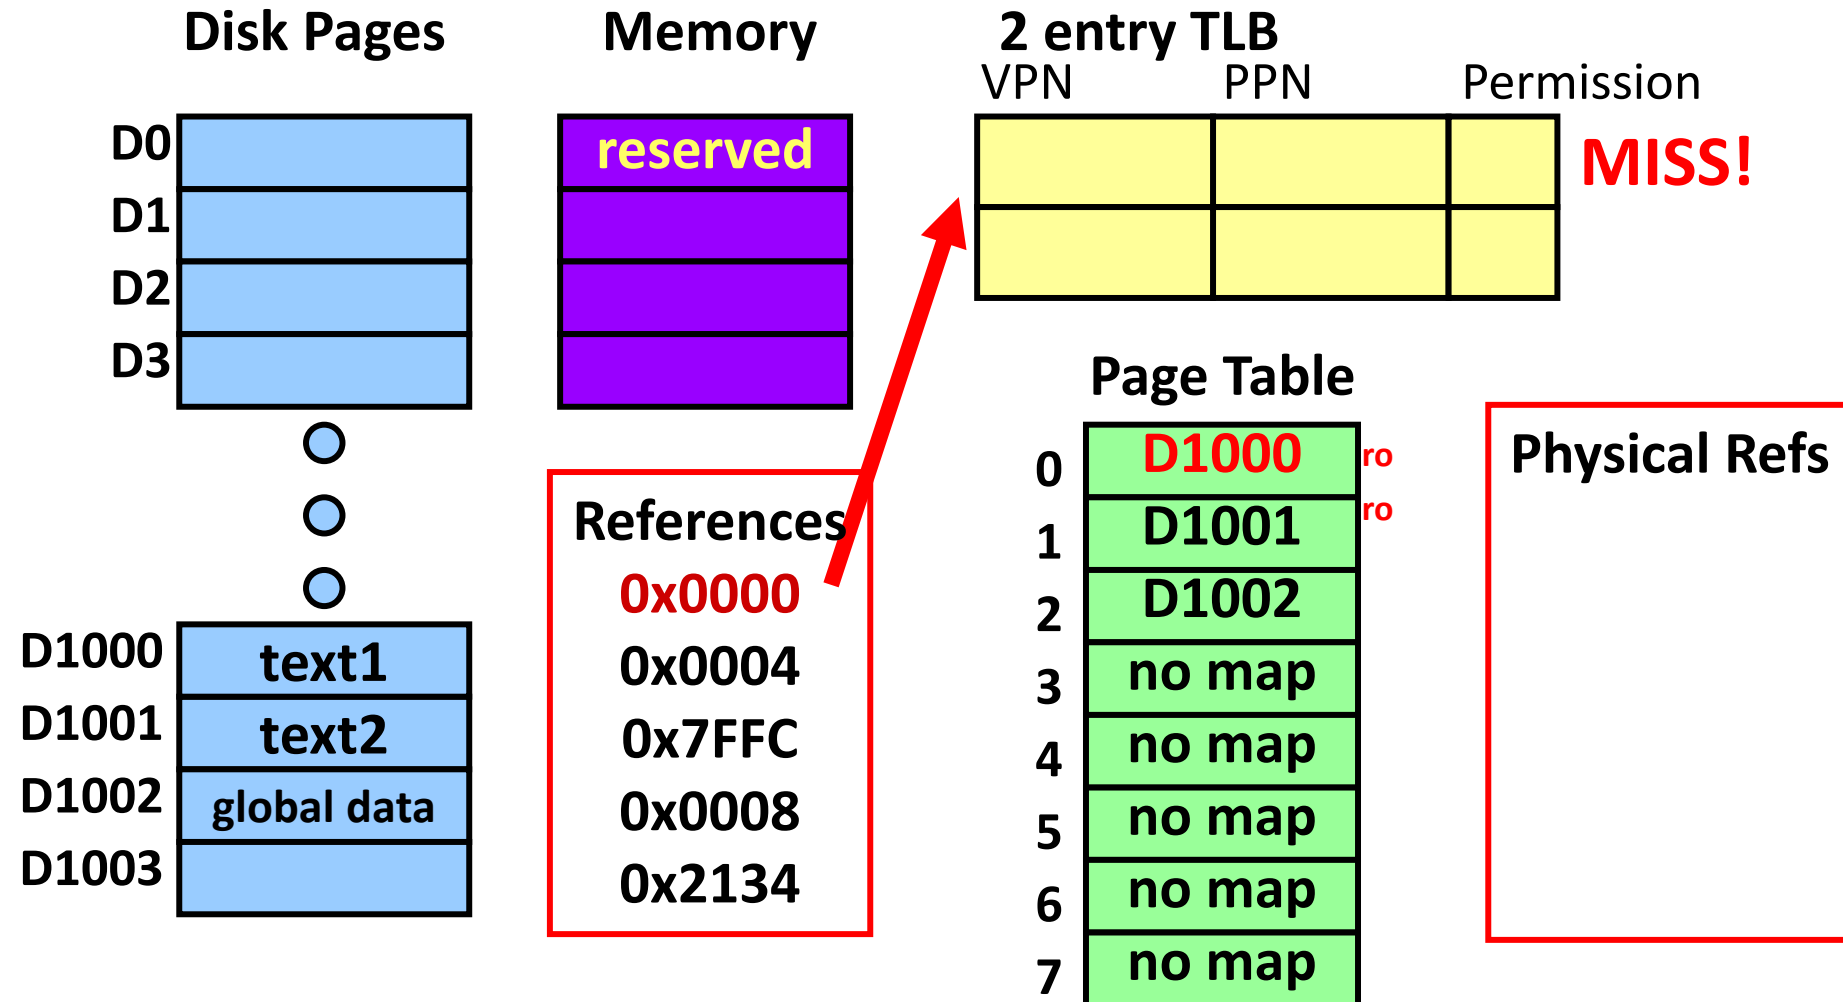
- ❑ Page size = 4 KB, Page table entry size = 4 B
- ❑ Page table register points to physical address 0x0000



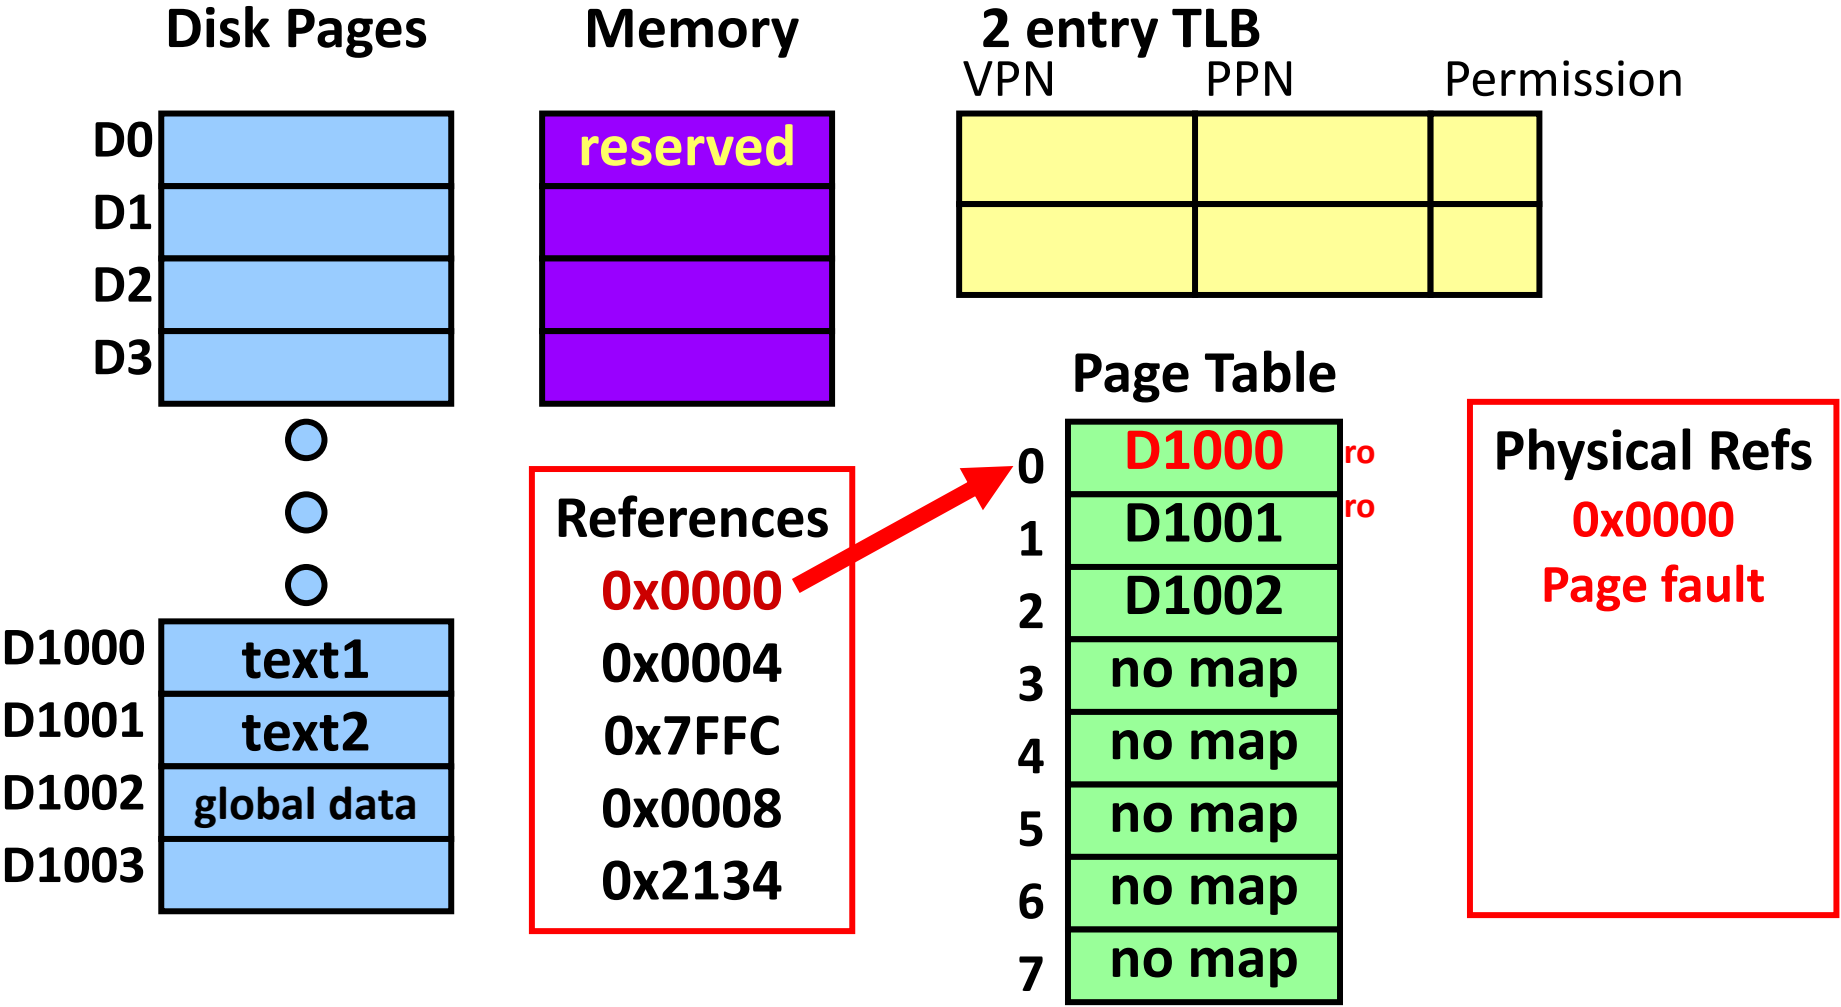
Step 1: Read executable header & initialize page table



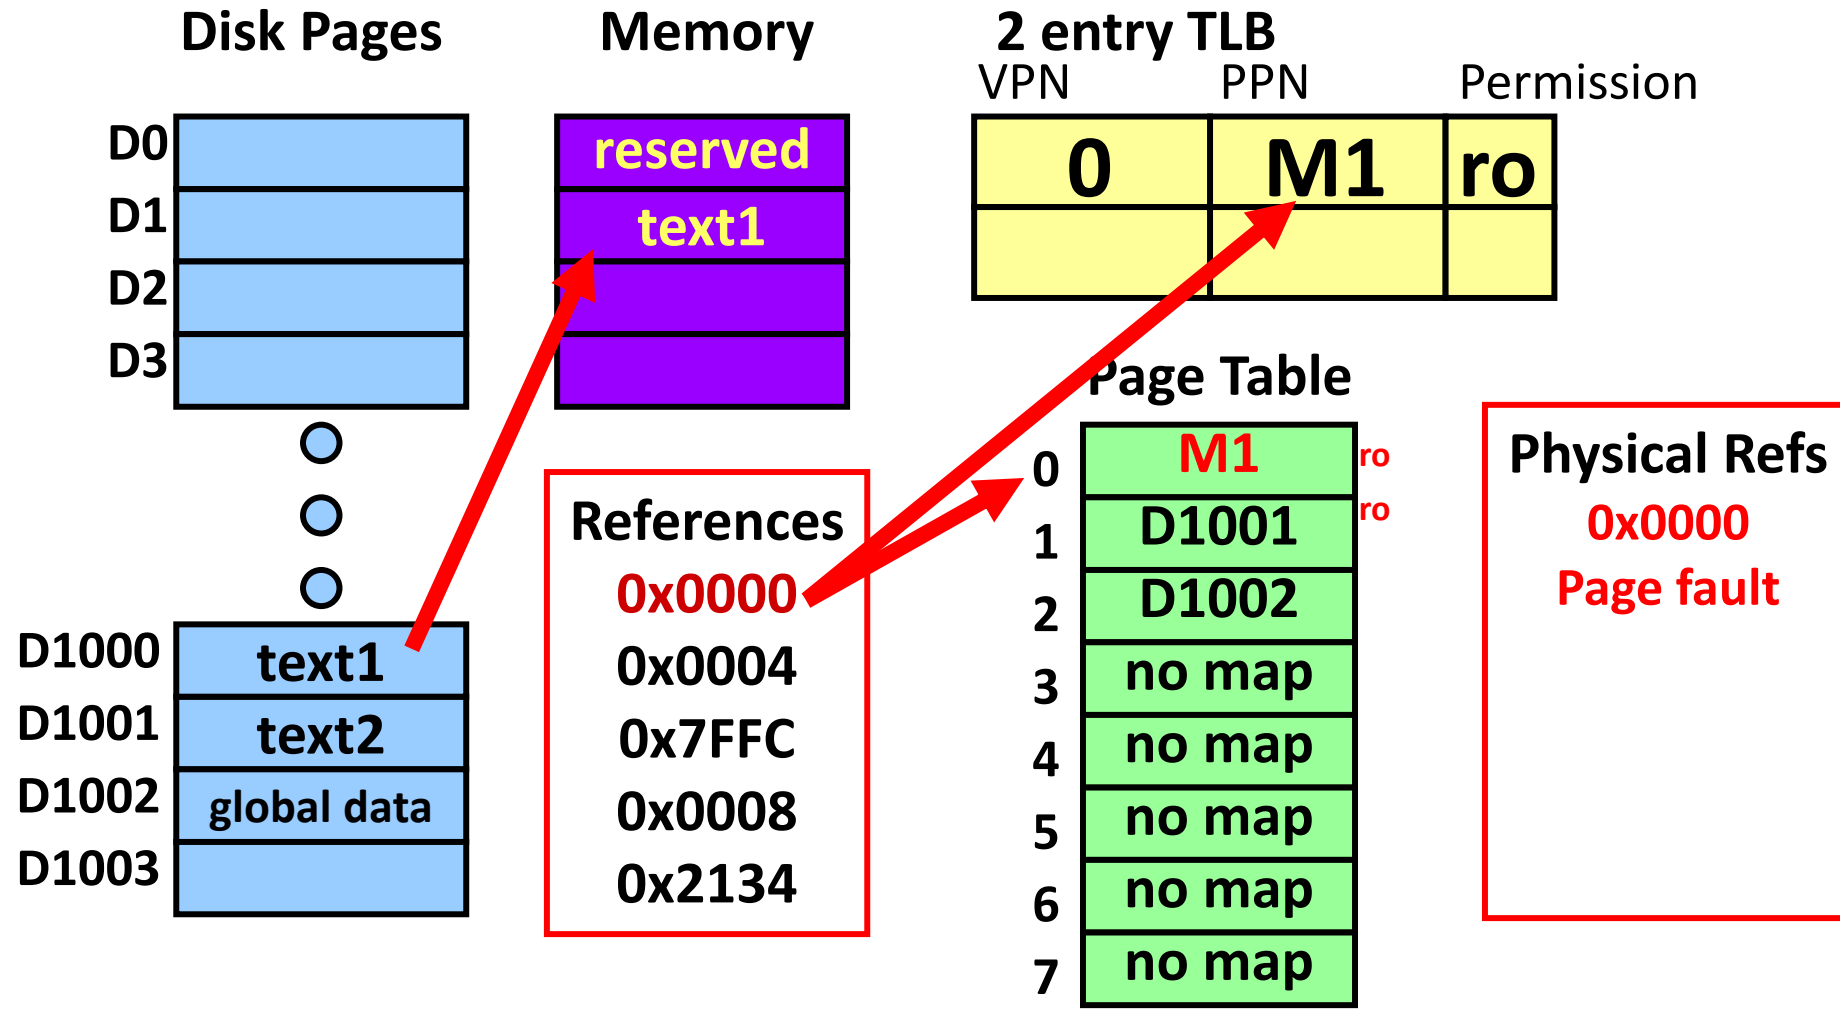
Step 2: Load PC from header & start execution



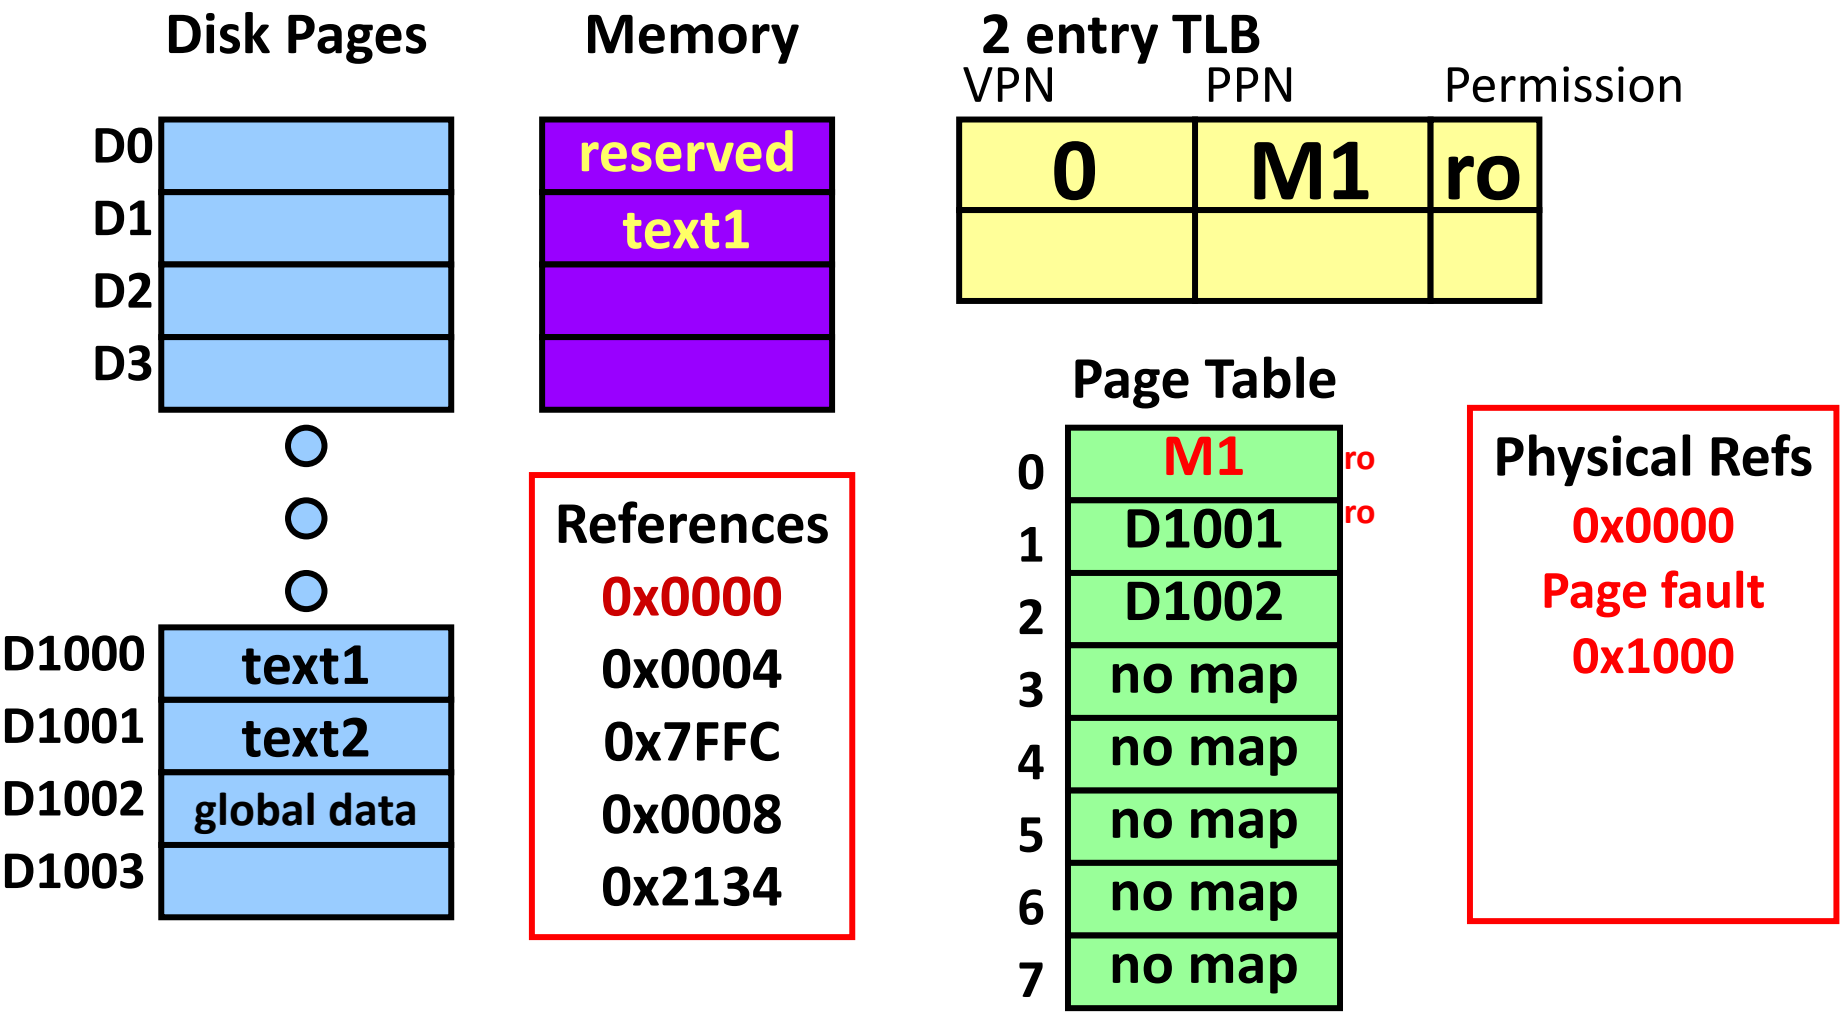
Fetching instruction 0000



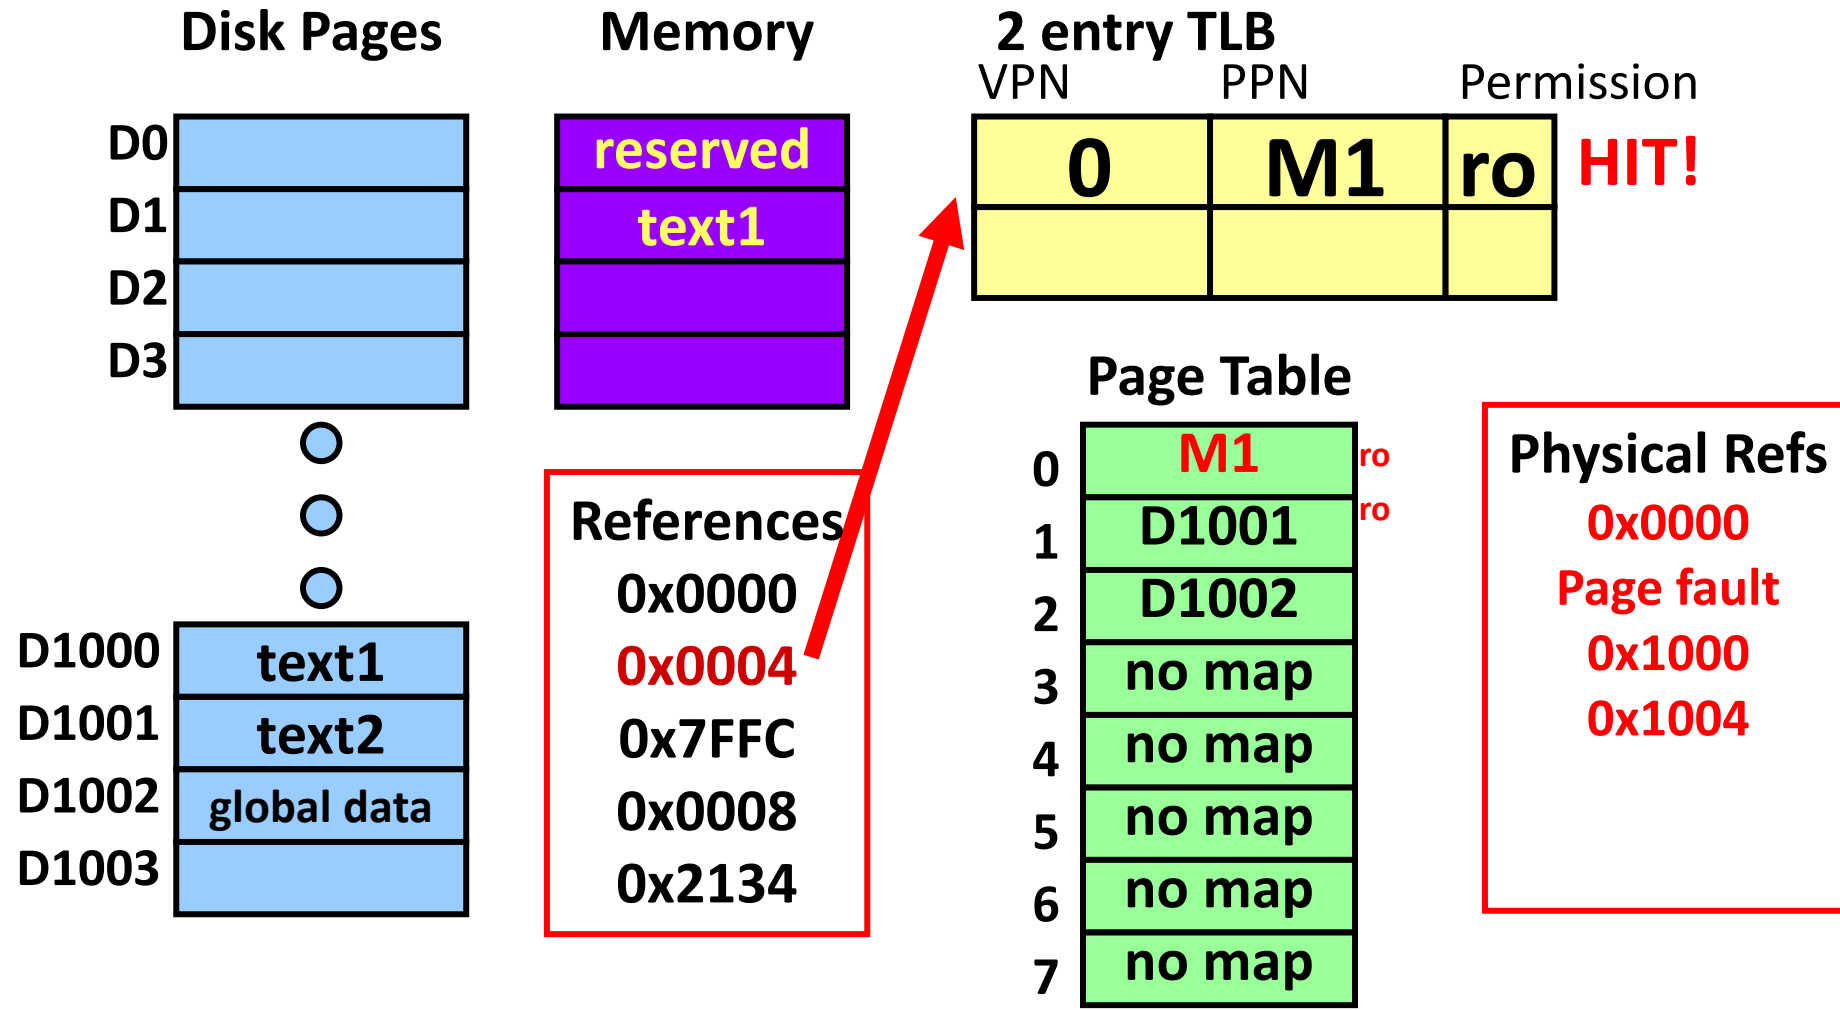
Fetching instruction 0000



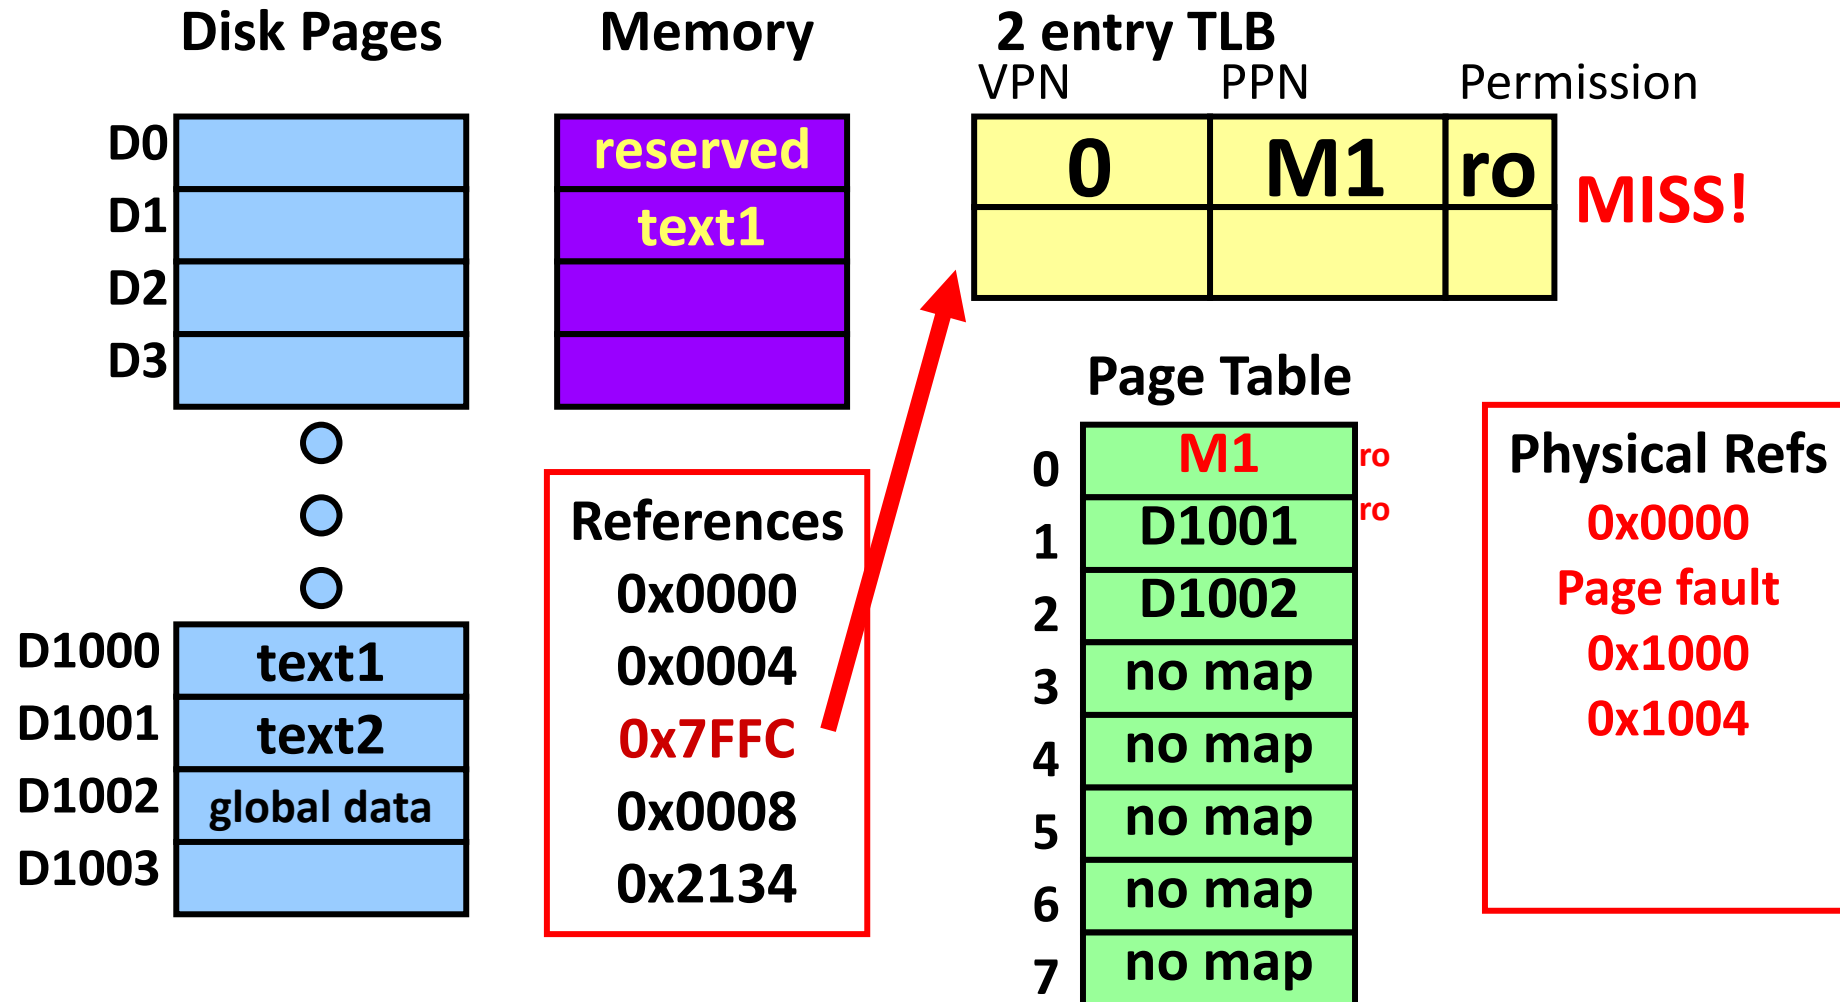
Fetching instruction 0000



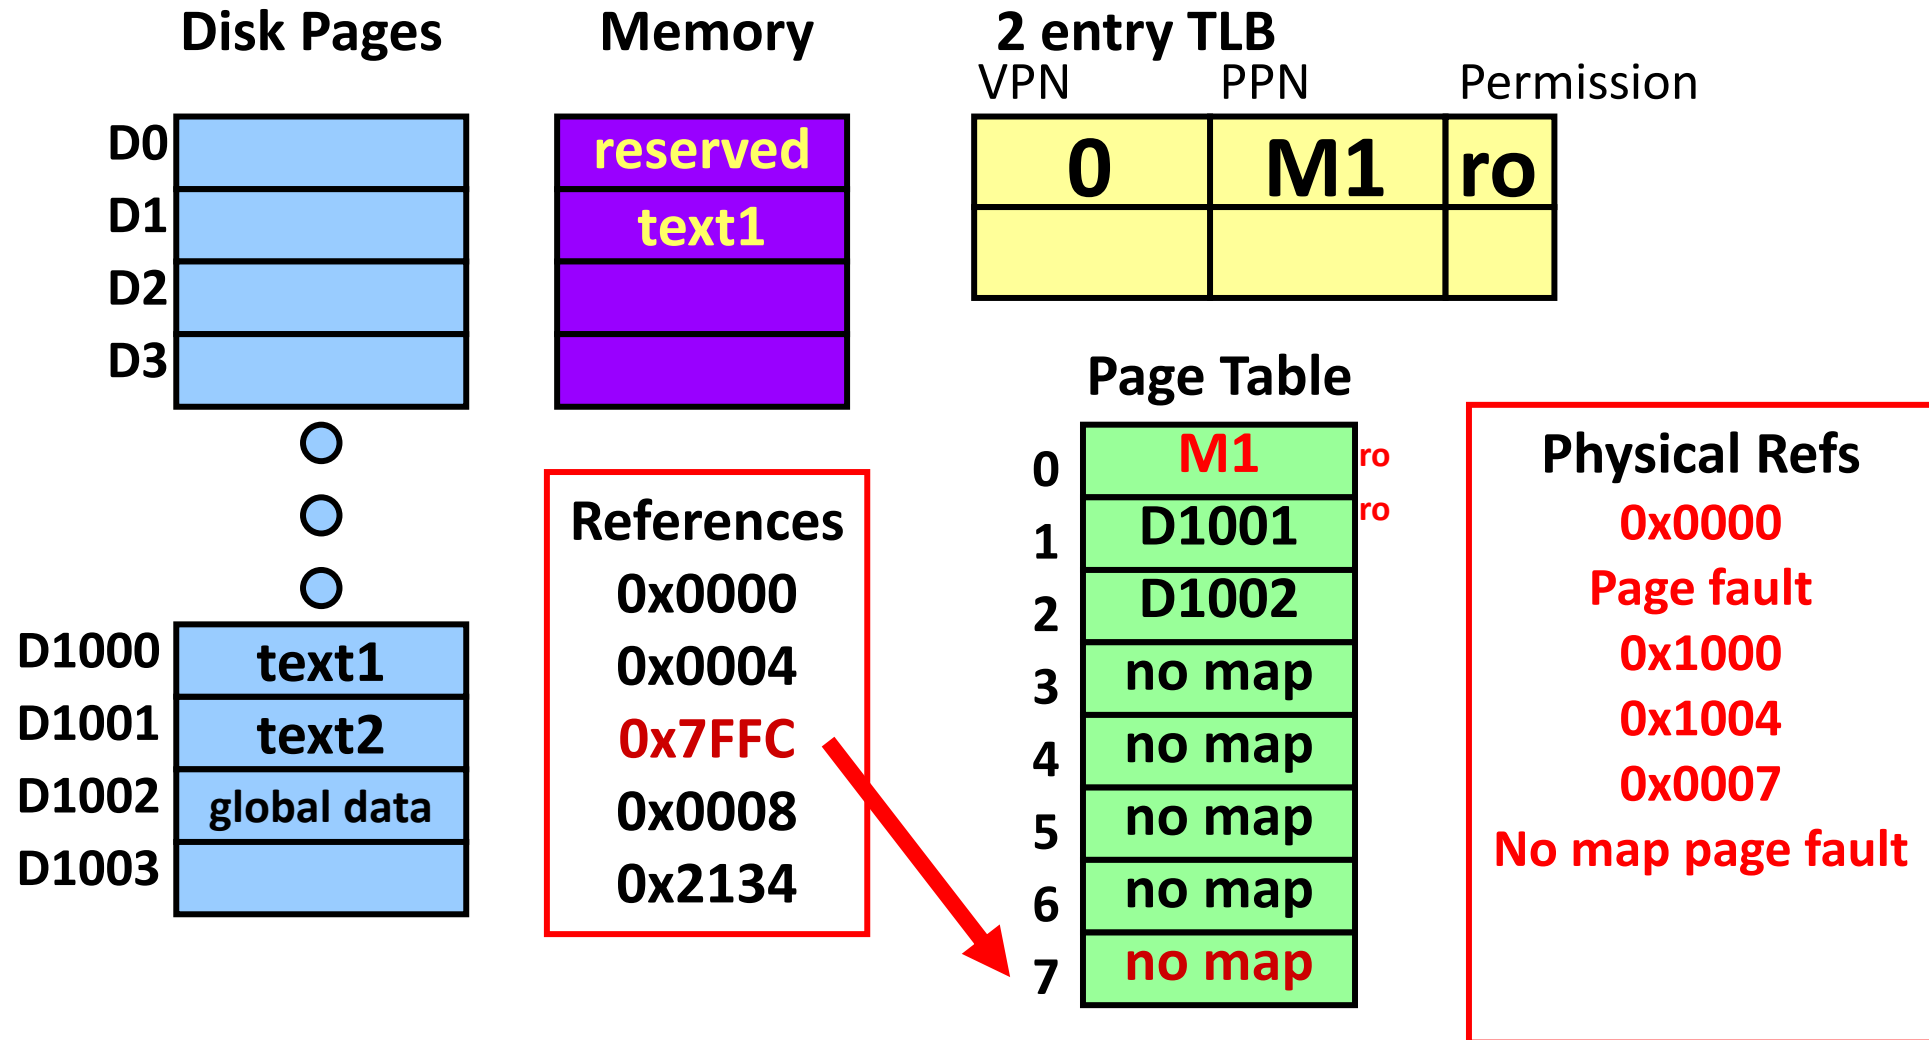
Fetching instruction 0004



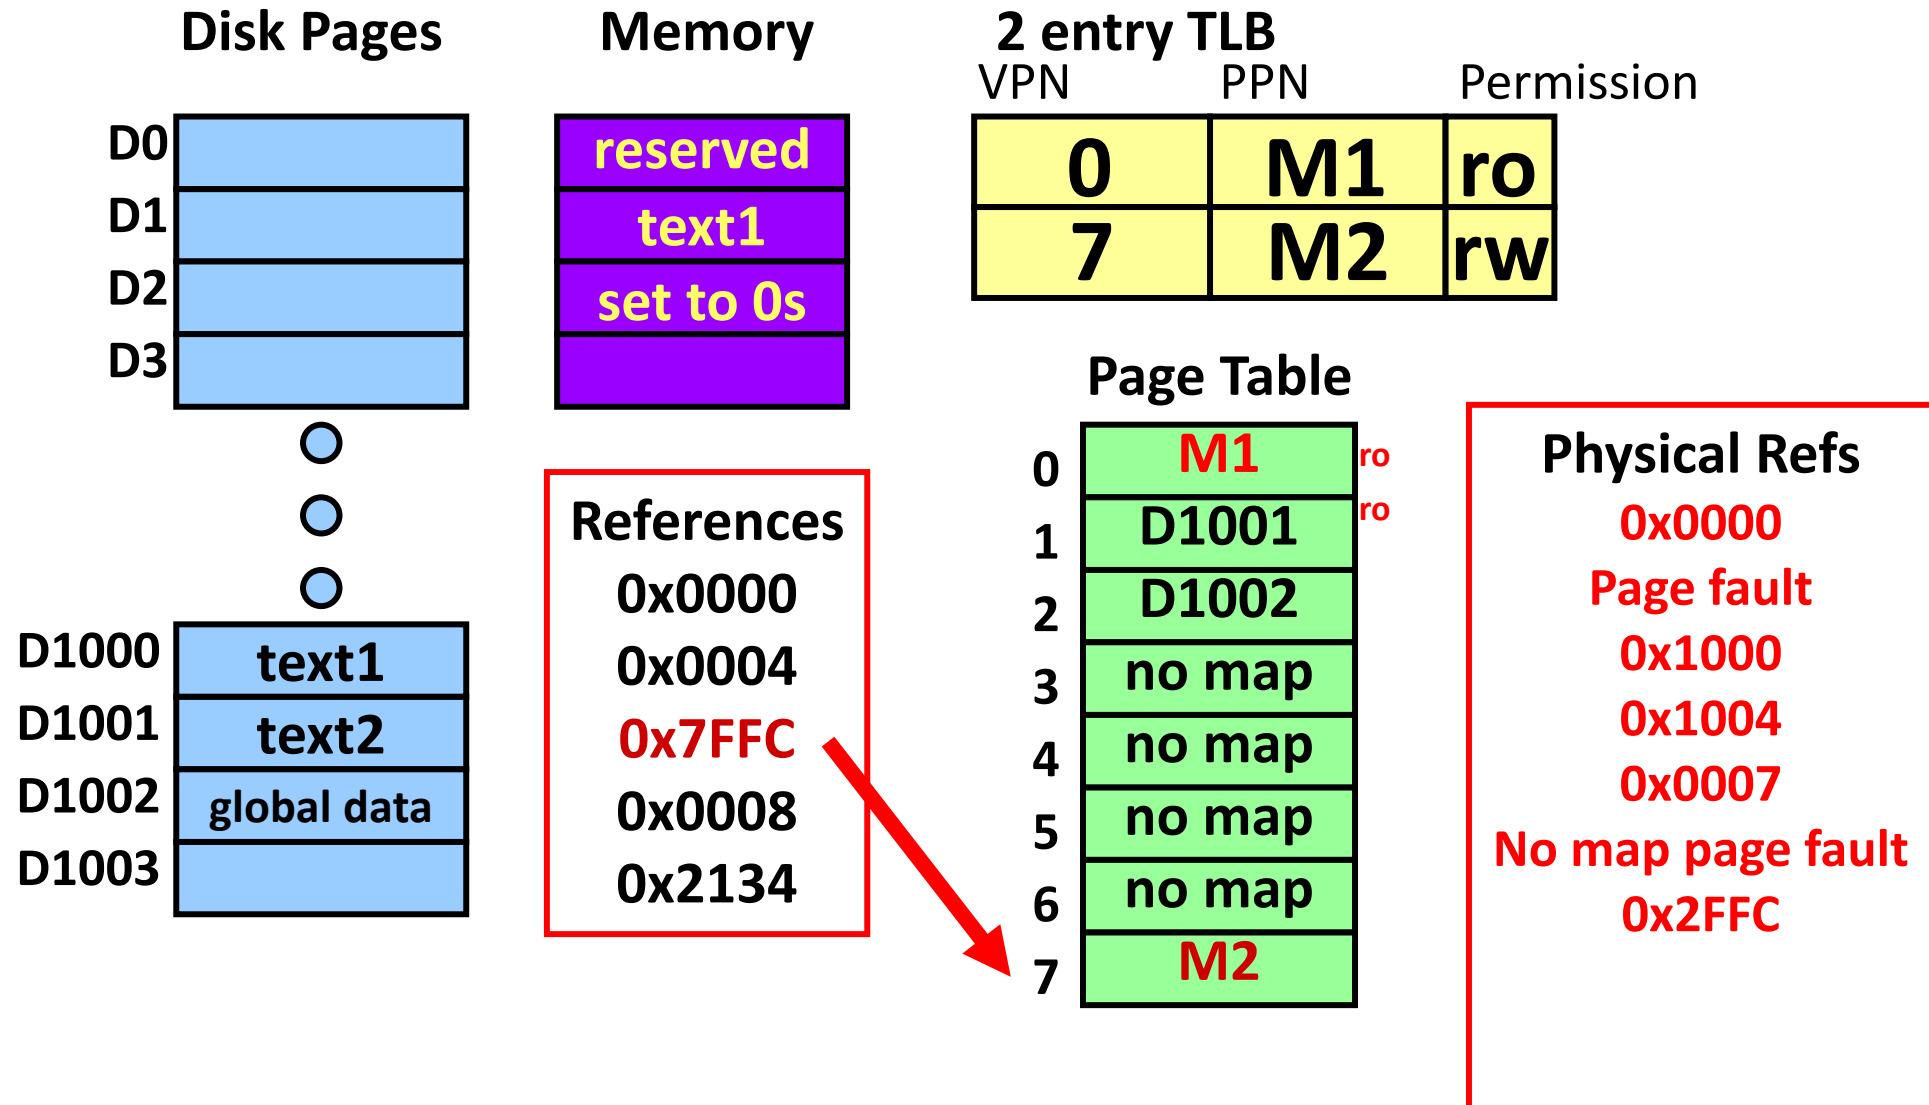
Reference 7FFC



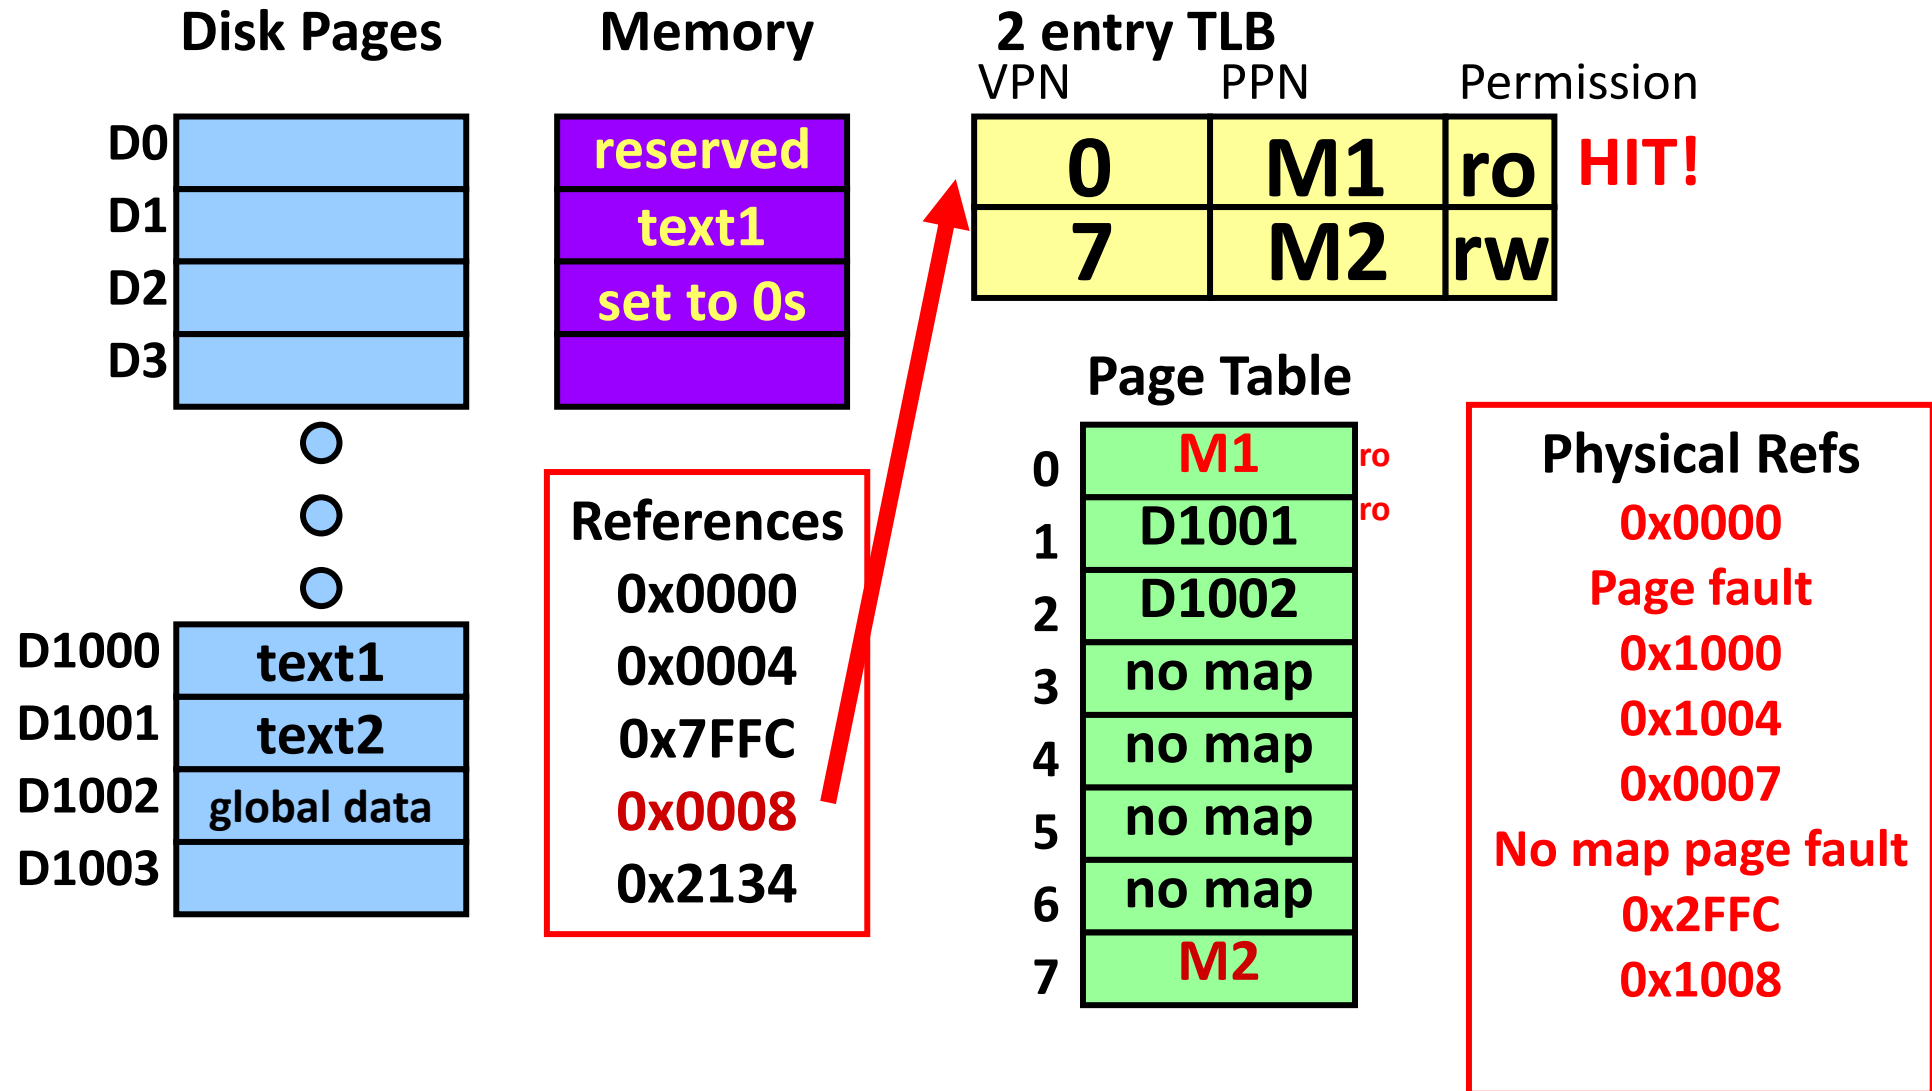
Reference 7FFC



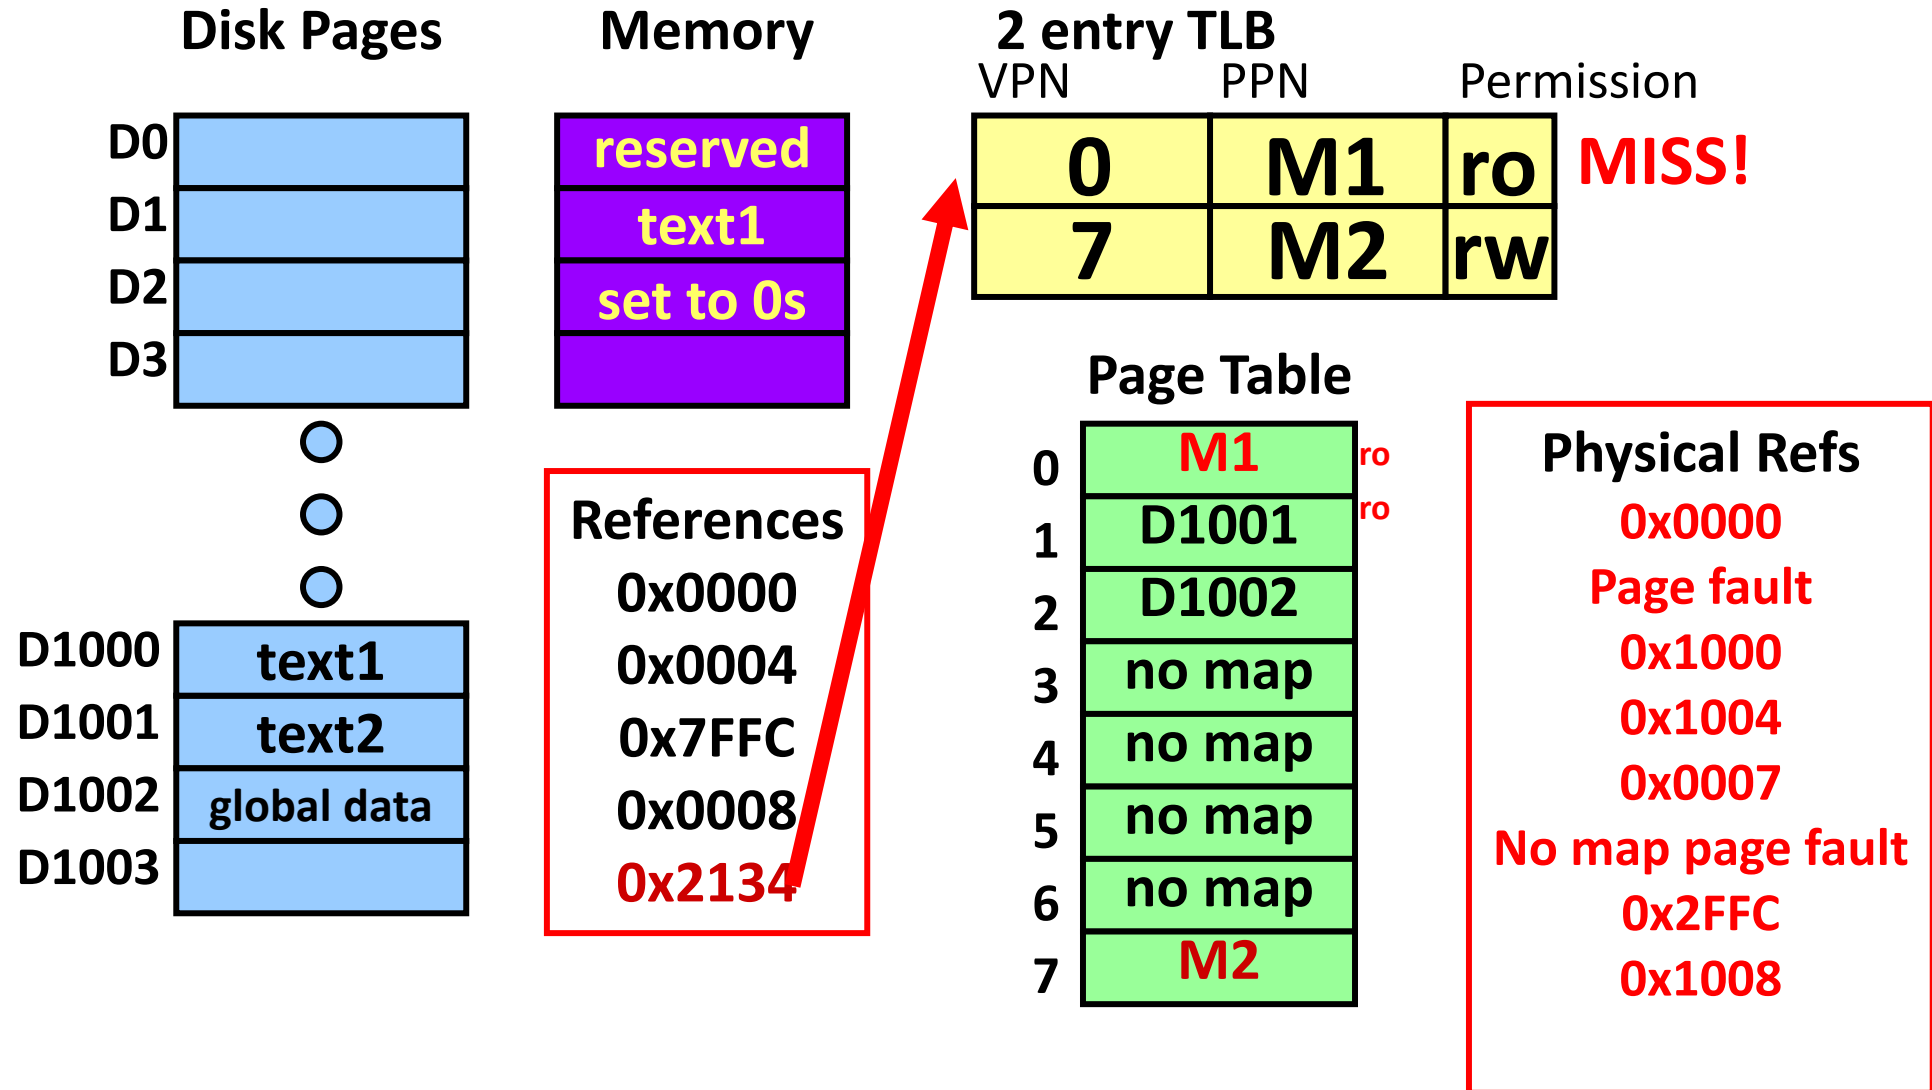
Reference 7FFC



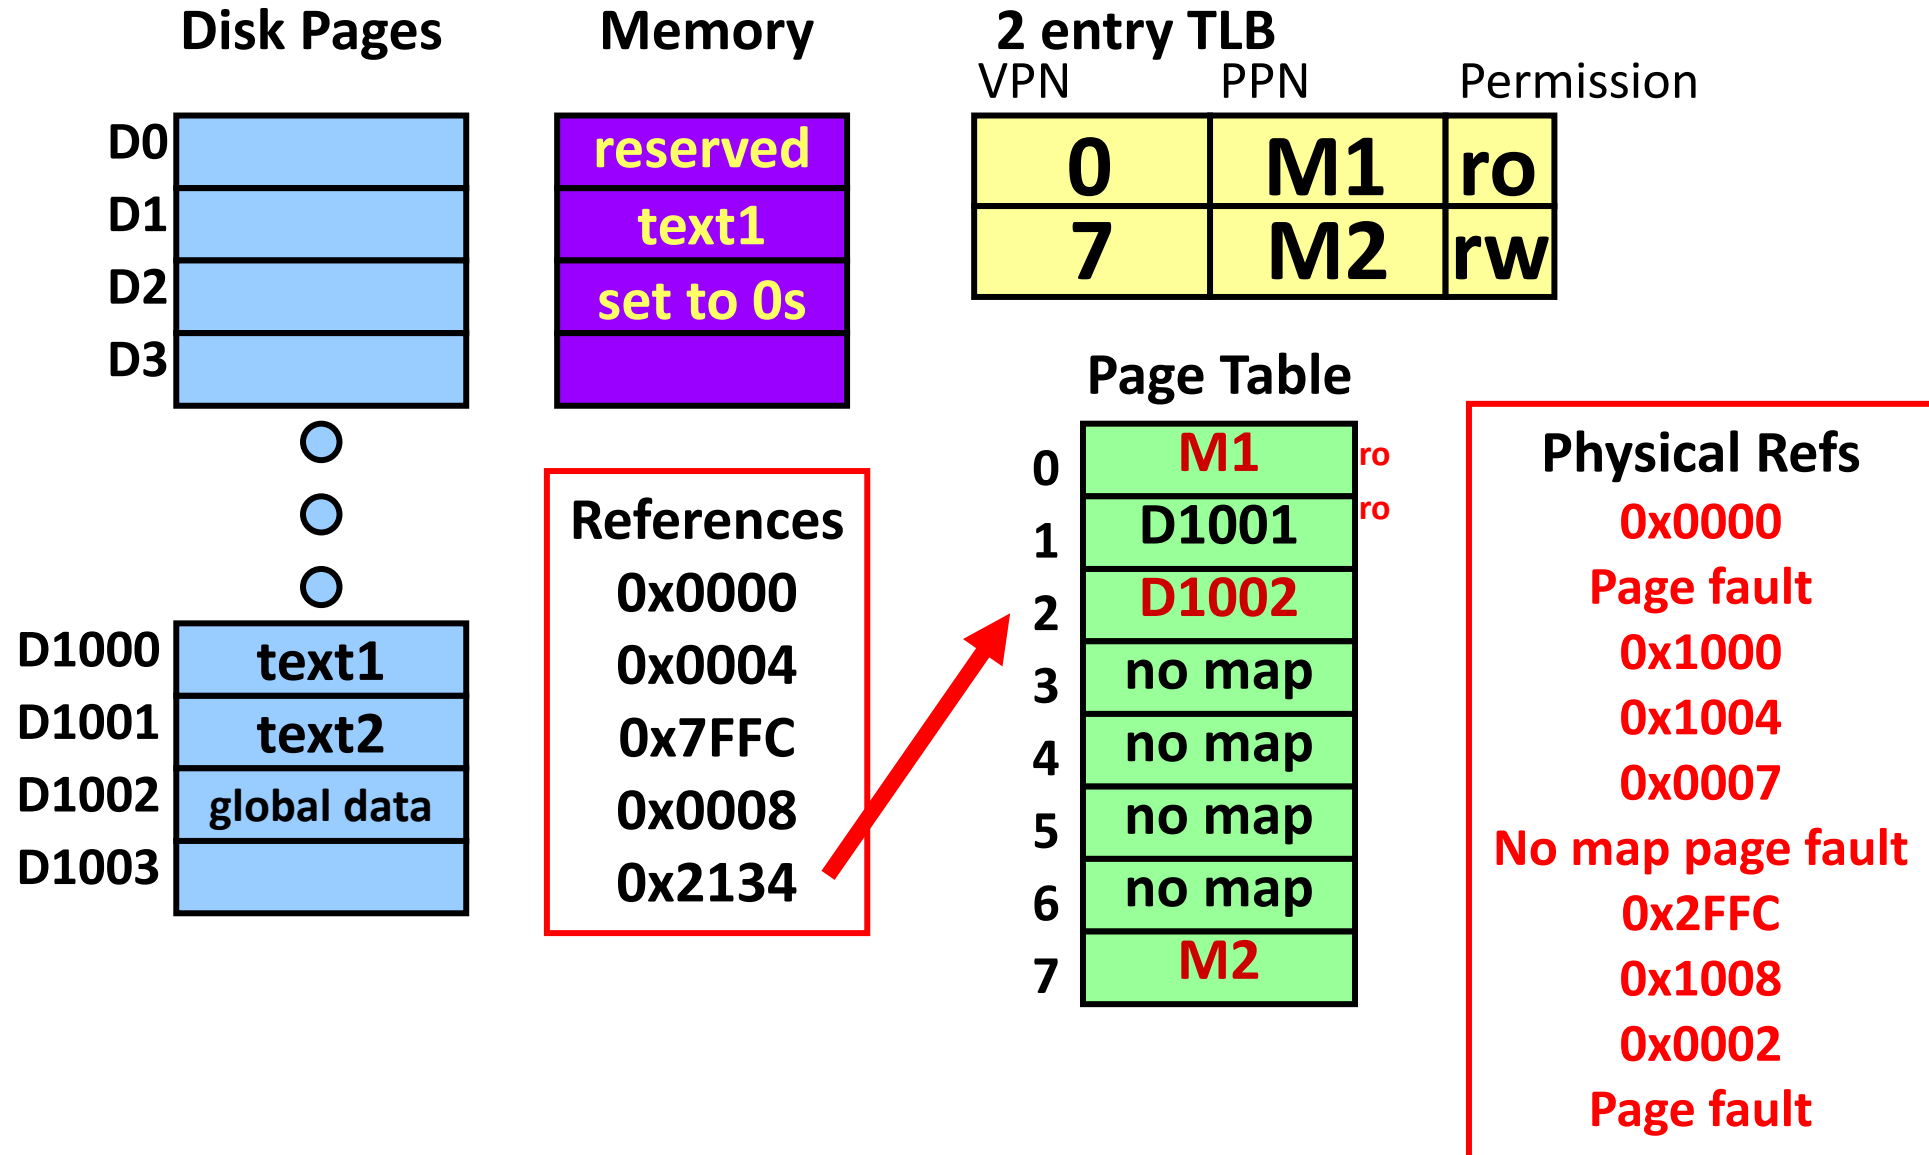
Fetching instruction 0008



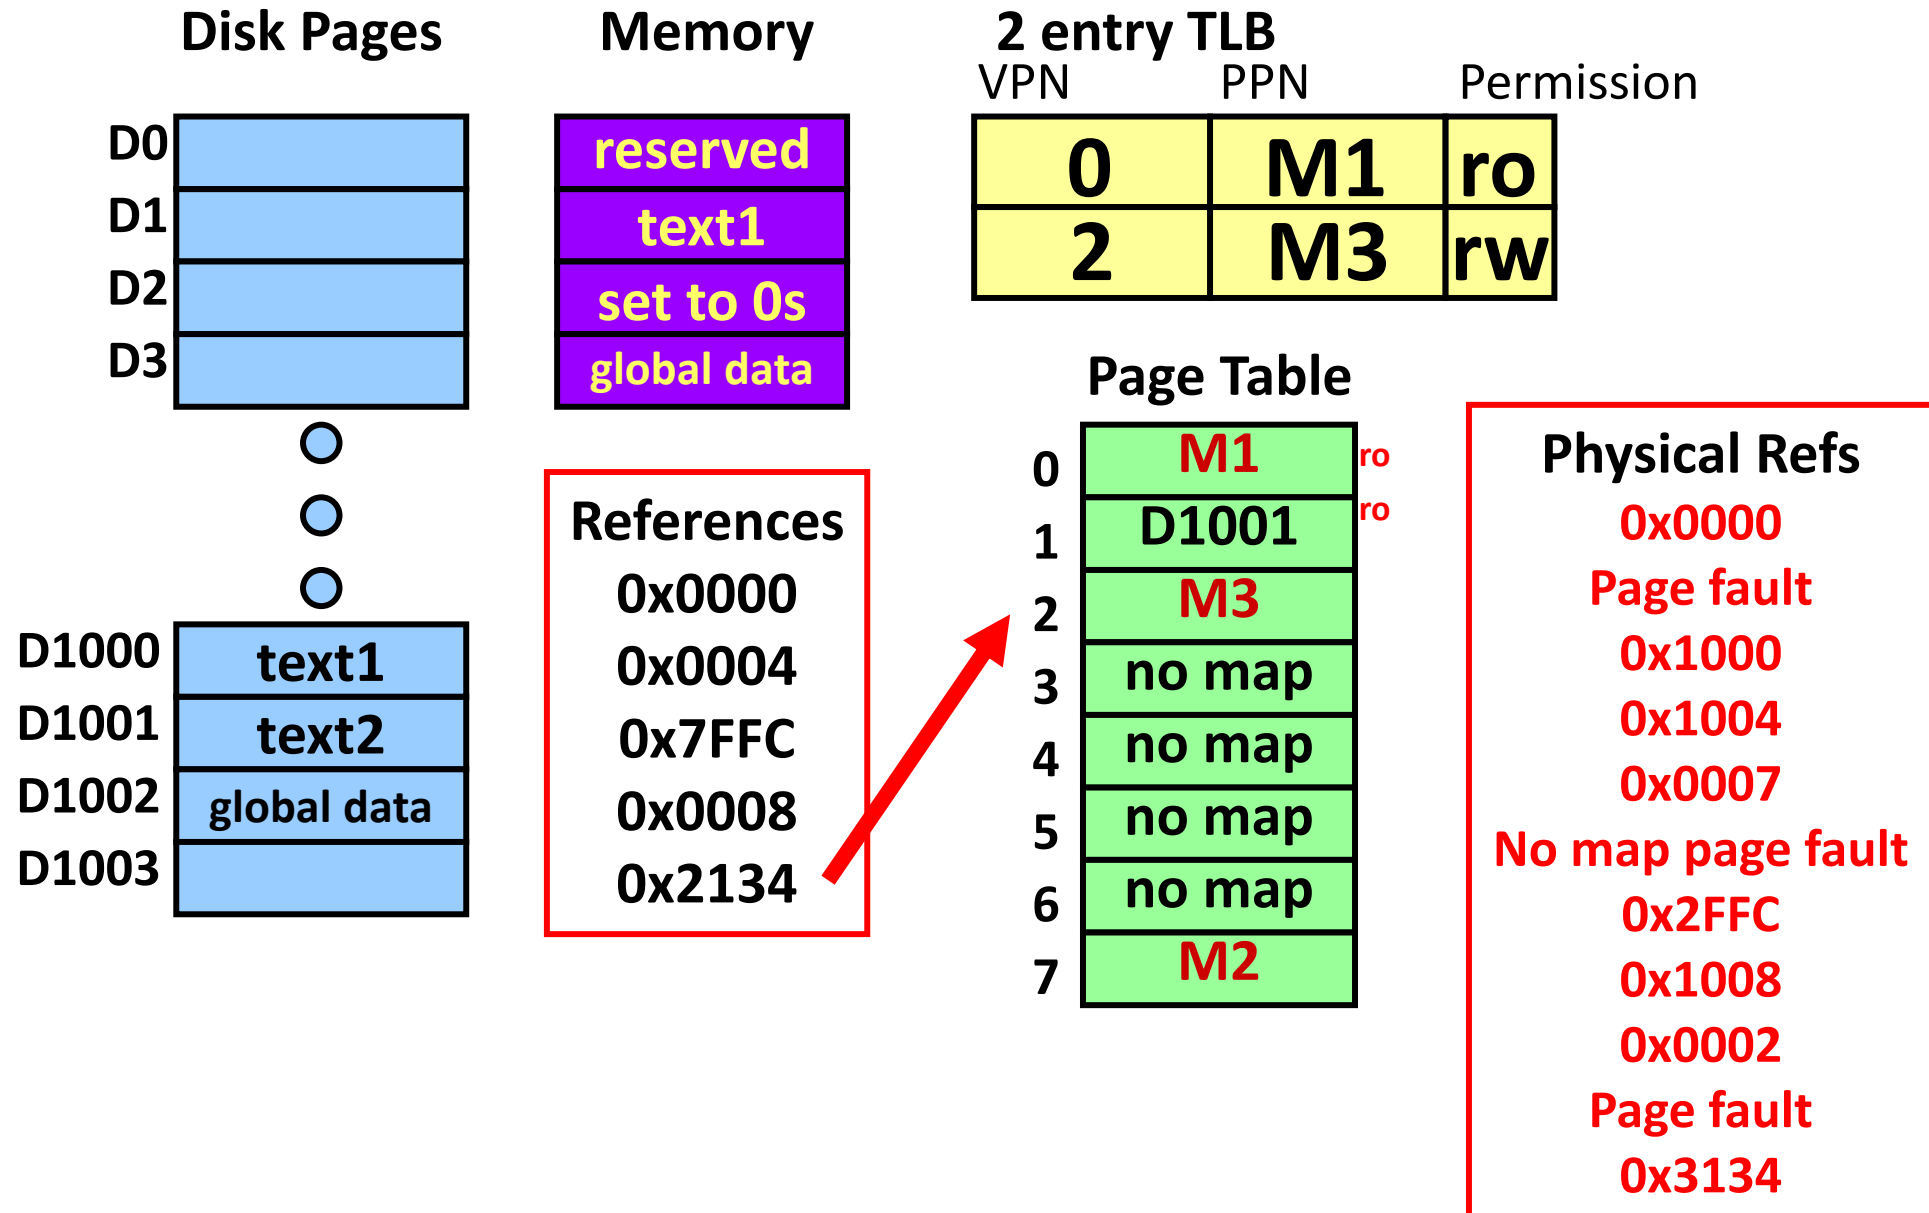
Reference 2134



Reference 2134



Reference 2134



Next topic: Placing Caches in a VM System

- VM systems give us two different addresses:
 - virtual and physical
- Which address should we use to access the data cache?
 - Physical address (after VM translations).
 - We have to wait for the translation; slower.
 - Virtual address (before VM translation).
 - Faster access.
 - More complex.

Poll: Which would be faster to access?

- a) Address cache with virtual address
- b) Address cache with physical address

Cache & VM Organization: Option 1


Physically-addressed Cache



Slower



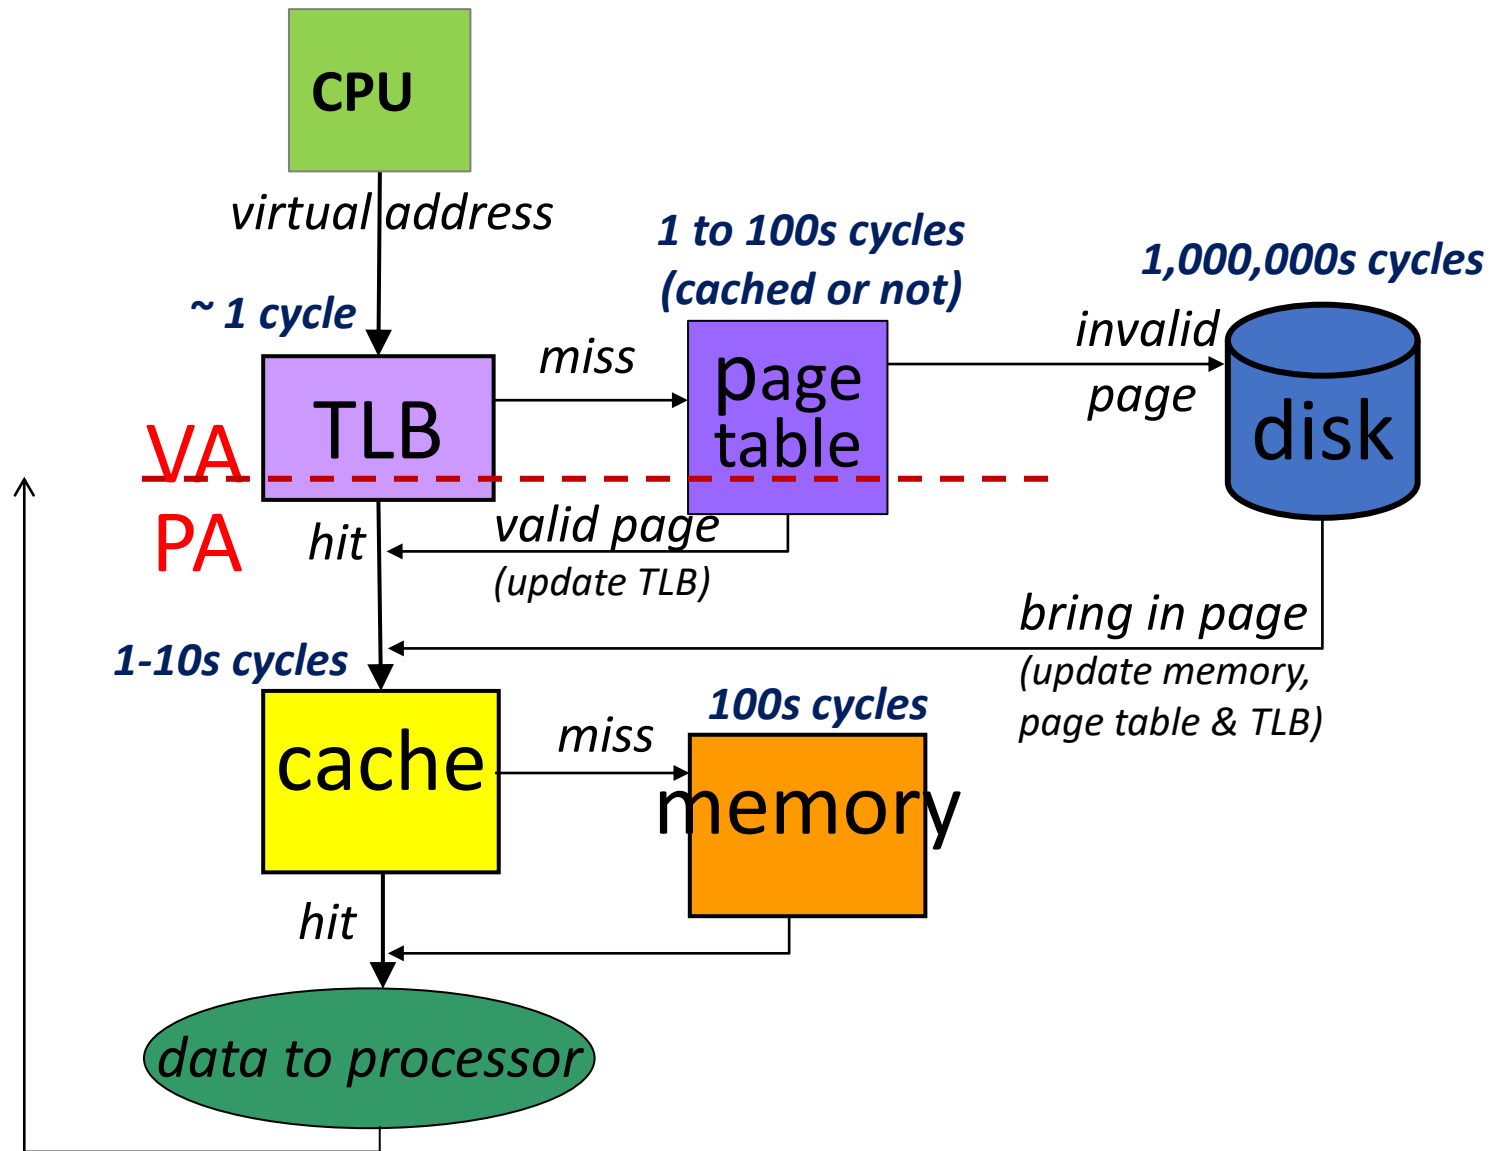
Low complexity



Physically addressed caches

- Perform TLB lookup *before* cache tag comparison
 - Use bits from *physical* address to access cache (tag, set index, and block offset bits)
- Slower access?
 - Tag lookup takes place *after* the TLB lookup
- Simplifies some VM management
 - When switching processes, TLB must be invalidated, but cache OK to stay as is
 - Implications? Might result in fewer cache misses if context switches very common (but they generally are not)


Physically addressed caches: detailed flow



Cache & VM Organization: option 2

Virtually-addressed Cache

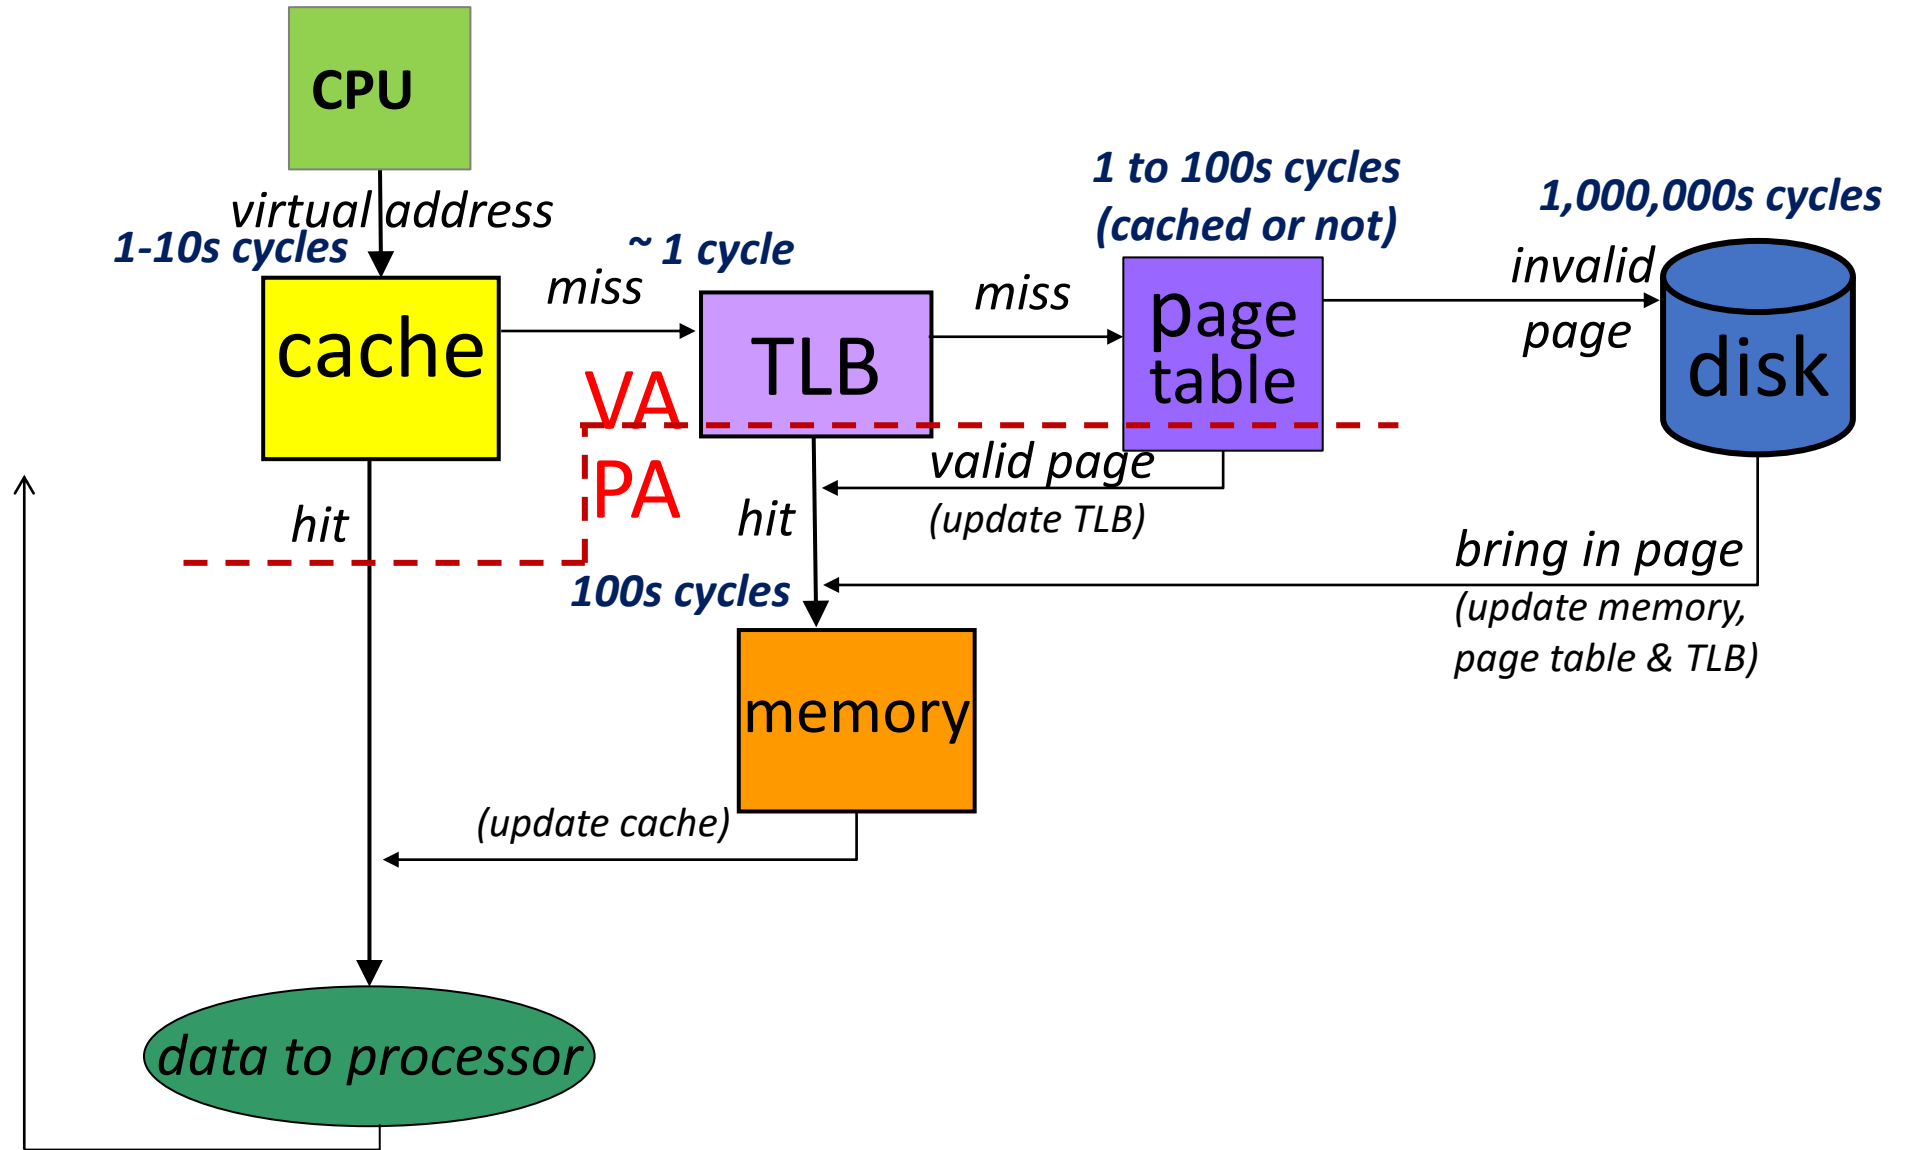
- ✗ High complexity (*aliasing*)
- ✓ Faster



Virtually addressed caches

- Cache uses bits from the virtual address to access cache (tag, set index, and block offset)
- Perform the TLB only if the cache gets a miss.
- Problems:
 - Aliasing: Two processes may refer to the same physical location with different virtual addresses (synonyms)
 - Two processes may have same virtual addresses with different physical addresses (homonyms)
 - When switching processes, TLB must be invalidated, dirty cache blocks must be written back to memory, and cache must be invalidated to solve homonym problem

Virtually addressed caches: detailed flow



OS Support for Virtual Memory

- It must be able to modify the page table register, update page table values, etc.
- To enable the OS to do this, **BUT** not the user program, we have different execution modes for a process.
 - **Executive** (or **supervisor** or **kernel** level) permissions and
 - **User level** permissions.

Next time

- No lecture Thursday
 - Exam review next Tuesday
- Lingering questions / feedback? I'll include an anonymous form at the end of every lecture: <https://bit.ly/3oXr4Ah>



Extra Problem

Exercise using previous multi-level VM

Virtual Address	1 st level	2 nd level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C						
0x001F0C						
0x020F0C						

1st level = 7b	2nd level = 8b	Page offset = 9b
----------------------------------	----------------------------------	-------------------------

Virtual address = 24b

Physical page number = 9b	Page offset = 9b
----------------------------------	-------------------------

Physical address = 18b

Assume memory for page tables are “somewhere else” in memory

Exercise using previous multi-level VM

Virtual Address	1 st level	2 nd level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C						
0x020F0C						

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

Physical address = 18b

Assume memory for page tables are “somewhere else” in memory

Exercise using previous multi-level VM

Virtual Address	1 st level	2 nd level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C	0x00	0x0F	0x10C	Y	0x001	0x0030C
0x020F0C						

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

Physical address = 18b

Assume memory for page tables are “somewhere else” in memory

Exercise using previous multi-level VM

Virtual Address	1 st level	2 nd level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C	0x00	0x0F	0x10C	Y	0x001	0x0030C
0x020F0C	0x01	0x07	0x10C	Y	0x002	0x0050C

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

Physical address = 18b

Assume memory for page tables are “somewhere else” in memory

What's next for You?!

- What classes should you look at if your interested in this kind of stuff?

More Classes?

- EECS 470 – Computer Architecture
- Picks up where we leave off in 370
- Discuss more sophisticated enhancements to processor design
 - Out-of-order execution
 - Multi-core / multi-threaded processors
- Major Design Experience / Capstone
 - Work in teams to design an actual processor (not simulator) over the course of the semester
 - **LOT'S** of work
- Requires EECS 270
 - For Verilog (a Hardware Description Language)



More Classes?

- EECS 373 – Embedded Systems
 - Learn about parts of computer systems that aren't in the processor or cache
 - Input/output
 - Timers
 - Bus interfaces
 - Followed up by 473 (MDE / Capstone)
- EECS 471 – Applied Parallel Programming with GPUs
 - How are graphics processing units different than normal processors?
 - How do we take advantage of their raw power when writing software?

More Classes?

- EECS 427 – VLSI Design
 - Design a processor at the circuit level
 - Actually layout all the transistors
 - Less architecture – more circuit design
- EECS 570 (Parallel Computer Architecture) + 573 (Microarchitecture)
 - Learn about more current research into architecture
 - Requires 470

More Classes?

- EECS 483 (Front-end compilers) + 583 (Back-end compilers)
 - 583 is more relevant to architecture
 - How to design compilers to write efficient assembly?
 - Requires knowledge of hardware optimizations
- EECS 482 – Operating Systems
 - How do moderns systems support multiple threads running simultaneously?
 - How does OS manage virtual memory?
 - How do file systems work?
- EECS 388 – Computer Security
 - How do things like caches "leak" information about a program's data?
 - How can call stacks be tricked into executing arbitrary code?

More Classes?

- EECS 498-001 – Quantum Computing for the Computer Scientist
 - Winter 23
 - How can we redesign the computing stack to take advantage of the bizarre rules of quantum mechanics?
 - Rather than having bits be 0 or 1, have quantum bits be in **superpositions** of 0 and 1
 - Perform massive number of computations simultaneously... sort of
 - Requires 370 and 281

Other questions?!